

Fighting Cognitive Overload: The Case for Feed Control in the Generative AI Era

Introduction: Every piece of content we encounter on social media carries a *cognitive cost*. Scrolling through an endless feed, our brains must parse each post, comment, and notification – even if only to decide it's not worth our attention. This constant micro-decision-making adds up to a significant cognitive load on users. The problem is compounded when we're forced to see content we find irrelevant or low-value. Yet on many platforms, users have limited control over filtering out such noise. For example, LinkedIn allows blocking individual users but **does not allow blocking company pages**, no matter how spammy their posts or comments might be. As UX expert Stéphanie Walter lamented, “*you can't block companies, only people*,” which means LinkedIn users are stuck seeing content from companies that might be flooding threads with AI-generated junk. Every unnecessary item in your feed – whether a random corporate post or a bot's auto-generated comment – taxes your attention for no benefit. In this article, we explore why giving users more control over their feeds is crucial to reduce cognitive overload, and how the rise of generative AI is making this issue far more pressing (while ironically paving the way for better solutions).

The Cognitive Toll of an Unfiltered Feed



Figure: Multiple screens showing an overload of information – a visual metaphor for the barrage of content users face. Each additional post or comment we have to mentally process contributes to cognitive load and potential fatigue.

The human brain has a limited bandwidth for processing information. When a social media feed is cluttered with irrelevant or low-quality content, it forces users to expend mental energy filtering signal from noise. Research on *digital overload* finds that the **constant need to process excessive information can overwhelm our senses and impose a cognitive burden**, making it difficult to concentrate ¹. In the context of social media, being “*bombarded with too much information*” leads our

brains to “struggle to filter out the noise, resulting in a persistent state of distraction” ². In other words, an overstuffed feed not only wastes our time but can also diminish our ability to focus on more important tasks, contributing to stress and burnout ².

Unwanted content in the feed isn’t just harmless fluff – it actively detracts from user experience and well-being. Each irrelevant post or robotic comment is a **micro-interruption**: even if you scroll past it in a second, your brain had to acknowledge and evaluate it. Psychologists note that constant tiny distractions can accumulate into significant mental fatigue and decision overload. This is why *controlling the feed* is more than a matter of preference; it’s about protecting one’s mental clarity. If you could customize your feed to *only* show content you truly care about (and hide or block everything else), you would reduce the cognitive load required to use the platform on a daily basis. *User agency* in feed curation thus becomes key for digital well-being. Unfortunately, today’s mainstream platforms give users very limited agency. Personalized recommendation algorithms try to guess our preferences for us, but **users lack meaningful control over what is recommended to them**, which “reduces their sense of agency and their ability to adapt social media feeds to their own needs and values”. In the absence of robust filtering tools, users are essentially at the mercy of the platform’s algorithm – or the persistence of spammers – when it comes to what appears on their screen.

Why Blocking and Filtering Matter (The LinkedIn Example)

Given the cognitive toll of extraneous content, the ability to **block or filter out sources of unwanted posts** is crucial. On Twitter/X, one can mute keywords or accounts; on Facebook, one can hide posts and unfollow sources. LinkedIn, however, presents a peculiar and frustrating limitation: it lets users block individual people but offers *no way to block content from company pages*. This means if a company page is posting or commenting in ways you find annoying – say, an aggressive recruiter firm commenting on every one of your posts, or a vendor page leaving promotional spam – you’re stuck seeing it unless you report it for spam (which often has no effect). Many LinkedIn users have observed this gap. “*I can report and block users that spam me... but not companies,*” one user wrote, noting the “tyranny” of LinkedIn’s design that rails users onto whatever content companies push out. In practice, this policy “gives spammers an easy way to flood posts with low-quality AI-generated comments,” as one LinkedIn observer noted, because the spam can simply come from company profiles that you cannot banish from your feed.

From the perspective of someone trying to maintain a high-quality feed, LinkedIn’s lack of a “block company” feature is a significant problem. **Every irrelevant comment incurs a cognitive cost** – you have to read a bit of it to realize it’s nonsense, or at least register its presence and get mildly annoyed. When that happens repeatedly, it “pollutes” your feed and distracts from genuine interactions. The **cognitive load** on you and your readers increases with each junk comment that slips in. As the user in our opening scenario described, “*every time you interact with a comment, there’s a cognitive cost to process that information.*” If half the comments on your post are boilerplate gibberish from bots, that’s a lot of mental energy wasted for you and anyone else trying to follow a meaningful conversation.

LinkedIn’s design choice likely isn’t an oversight; cynically, one might argue that high engagement metrics (even from bot comments) benefit the platform’s appearance of activity. Indeed, commentators have pointed out that platforms sometimes tolerate a degree of spam because those “*bogus comments help to fool human users and advertisers into thinking their activity ... is more effective than it really is*”. In LinkedIn’s case, allowing company pages free rein in comments (while users can’t easily block them) tilts the balance away from user control. It highlights exactly why more user-centric controls are needed: *the user, not the platform, should decide what they’re exposed to*, especially when the “*third party has different agendas and isn’t paying the cost*” of that exposure, as our voice memo author put it.

Generative AI: Fueling a New Wave of Spam

The rise of **Generative AI** content has dramatically worsened the signal-to-noise ratio on many social networks. What used to be a trickle of human-written spam has become a flood of AI-generated text and media, because **the cost of producing content has plummeted to near-zero**. Spammers and opportunists can now deploy bots to generate endless comments, posts, or even entire articles in seconds. As one analysis notes, *"if the cost of additional traffic is zero, and a cheap machine can produce traffic, then it's worth pumping out unlimited amounts... even if each individual communication has only a very low chance of profit"*. In other words, generative AI has supercharged the old economy of spam: when you can churn out **limitless junk content at no cost**, why not flood every corner of the network? The result is an avalanche of low-quality posts and comments – **noise** that each user now has to mentally sift through.

Nowhere is this more evident than in comment sections on platforms like LinkedIn. In recent months, many LinkedIn users have noticed replies that *just don't sound human*. For example, you might announce a personal achievement or share an article, and *immediately* get replies like "**Congratulations on your achievement!**" or "**Thanks for sharing your valuable insights!**" from people you've never heard of. The responses are generic, overly positive, and often don't fully align with the tone of your post (e.g. a cheerful *"Exciting times ahead!"* in reply to a somber update about layoffs). These are telltale signs of **AI-generated comments**. As an internet risk expert describes, *"every so often you get a comment so bland and generic you realize it had nothing to do with the post... the user is just harvesting attention by leaving huge numbers of comments everywhere."* Such comments are generated by simple formulas (e.g. take a positive phrase + reword the post title) well within AI's capabilities. The profiles behind these comments often belong to self-described "automation marketers" or fake personas, and they manage to pepper countless posts with auto-comments in a short span – something a human could hardly do manually.

For the average user, encountering these **bot comments** is more than just a mild annoyance. It **dilutes genuine discourse** and again imposes cognitive costs: you have to read an insipid comment (or several) before realizing *"oh, this adds no value - it's basically a form of spam."* The original poster may feel obligated to engage or be polite ("Thank you for your comment...") before noticing there's no real person behind the words. Multiply this by dozens of posts and you see the problem: **AI-driven spam is clogging the channels of communication**, drowning out real voices. *"Selfish behavior on networks... drowns out real dialogue between people,"* as one commentator put it, and if we don't put better safeguards in place, *"they will clog and kill the channels of communication that real people want and need."* In short, generative AI has made spam easier to produce and harder to detect at first glance – a dangerous combination for user experience.

The user in our scenario specifically noted *GenAI-written comments* as a new kind of plague: *"Fundamentally, this is spam - a new version of SEO or engagement farming,"* he argued. It adds no human insight or authentic voice; it's just there to game the algorithms. Indeed, a growing number of professionals are voicing similar complaints. *"These comments are increasingly common on LinkedIn: comments that simply summarize a post with ChatGPT. What is the point of this?"* one observer asked rhetorically. The point, sadly, is often to create the appearance of engagement – to make it seem like the bot or the person running it is active and interested, in hopes of drawing attention to their profile or boosting some metric. But **for everyone else, it's pure noise**. As Eric Priezkalns writes in *Commsrisk*, *"Your time and attention is taken. You receive the output of a machine that mimics meaning but does not understand or care about anything."* This quote perfectly captures the core issue: AI-generated comments simulate a response, but they carry no real intent or soul. They steal a few seconds of your life as you read them, yet **give nothing of value back** – a net loss from a cognitive perspective.

Quality over Quantity: The Rising Value of Human Content

Ironically, the glut of generative content is leading to a renewed appreciation for *authentic, human-generated content*. The more our feeds get inundated by auto-generated text, auto-generated videos, even AI-generated music, the more we start craving a genuine human touch. Savvy users can often *sense* when a comment or post is AI-written – it feels formulaic, context-insensitive, or “too generic.” Over time, this has created a kind of **backlash against AI content** on social media. Recent industry surveys show consumers are increasingly turned off by content that looks impersonal or automated. In fact, as of early 2026, only about **26% of consumers say they prefer generative-AI-created content** over traditional human-created content – *a sharp drop from 60% just a couple of years prior*. The pendulum is swinging back toward authenticity, likely because people are realizing en masse that *much of the new AI content flooding our feeds is “slop”*. It may be grammatically correct and on-topic, but it lacks the nuance, personality, and credibility of a real person’s input.

This dynamic was summarized well by a creator economy expert, who noted that the oversaturation of AI content will “**force people to find real, authentic creators... an oasis from the low-effort, AI-generated content**.” Audiences, he predicts, will increasingly realize “*I have no connection to this*” when they encounter AI-generated posts, “*and they’ll make their way back to authentic, human creations, because that’s what it means to be human – to connect.*” In other words, the very prevalence of soulless content is driving up the *perceived value of genuine content*. It’s a classic market correction: scarcity increases value. As AI makes bland content abundant, truly personal and creative contributions become rarer gems. We’re already seeing platforms respond to this trend. For instance, Instagram’s leadership has talked about boosting “**originality**” in the algorithm and labeling AI-generated posts clearly. LinkedIn itself (perhaps feeling pressure from user complaints) stated in mid-2023 that it would *deprioritize posts written with AI* and favor human-written content. And beyond algorithm tweaks, culturally there’s a celebration of the “*messiness*” of human creation – those imperfect, quirky traits that remind us there’s a person behind the words or image. “*AI can’t replicate the messiness of human creativity,*” as one marketing executive put it. “*We crave that now – we crave imperfection and things that feel ‘off’ in a human way.*” ³ This quote underscores a fascinating point: attributes that might have been seen as *flaws* in content (a less polished sentence, a minor factual error, an emotional tone) can actually make it *more trustworthy* in an age of robotic auto-generated perfection.

For technical and GenAI professionals, this shift has big implications. It suggests that blindly automating content is a dead-end strategy for engagement. If the goal is to build real audience trust or community, authenticity is key. Generative AI can be a powerful tool to assist creators – e.g. helping draft a piece based on your own voice – but the final output needs a human touch. **Otherwise, it risks being lumped in with the tsunami of AI spam that discerning readers are increasingly tuning out.** In short, the more AI content floods our feeds, the more *social capital* accrues to the human, the original, the genuinely engaging. Paradoxically, GenAI is making the content overload problem worse, but in doing so it’s also highlighting exactly what needs to change.

A Push for Better Solutions: Empowering Users and Platforms

The current state of affairs – an overload of content with users having minimal control – is unsustainable. The silver lining is that the pain caused by generative AI overload is *spurring action* from both users and platform designers. **User awareness** is at an all-time high: communities are actively discussing how to spot AI fakes, how to manage digital fatigue, and are pressuring platforms to give them better tools. We’ve seen influencers openly call out LinkedIn’s shortcomings, using terms like “enshittification” to describe the degradation of the user experience. When enough vocal users complain, platforms eventually take note (if only out of fear of losing engagement). On the other hand,

platforms and researchers are starting to experiment with solutions. One promising direction is giving users more direct control over their feeds. For example, the decentralized network BlueSky has begun offering custom algorithms/feeds that users can switch between or tweak. Academic researchers built a system called *Pilot* on BlueSky as a proof-of-concept for user-tailored feed algorithms – providing simple controls to let people exclude or prioritize certain content. In trials, “*users’ sense of agency*” markedly increased, and they were encouraged “*to think more critically about curating their feeds*” when given such tools. This demonstrates that **meaningful feed customization is feasible** and that users benefit from it. It’s likely only a matter of time before mainstream platforms consider similar approaches (or face user exodus to those that do).

In the near term, even without fancy new algorithms, there are some straightforward steps that platforms could implement: **allow blocking of all entity types (including company pages)**, allow keywords or phrases to be muted (e.g. “AI bot” or other telltale spam markers), and invest in **detecting AI-generated comments** to optionally filter or label them. Not all users will want to hide AI content – but the *user should have the choice*. As one analyst put it, “*we must block bad traffic... otherwise it will clog and kill the channels of communication that real people want and need.*” The externality of spam and noise should be mitigated by giving control back to those who suffer from it – the users. On the flip side, platforms that *do* prioritize authenticity can turn it into a selling point. We already see brands and social networks touting “real human connections” in their messaging, implicitly contrasting themselves against a future of bots. If generative AI was 2025’s shiny new toy, 2026 might be the year of the *Authenticity Renaissance* on social media.

Conclusion: The ability to shape one’s own feed – to block, filter, and mute unwanted content – is not just a nice-to-have feature; it’s becoming essential for maintaining *sanity and signal* in the social media landscape. The explosion of generative AI content has poured fuel on the fire of information overload, making the need for user control painfully obvious. But in that challenge lies an opportunity: by recognizing the value of genuine, human-centric content and demanding better filtering tools, we can push social platforms toward a healthier equilibrium. Imagine a feed where *you* set the terms: no AI gibberish, no irrelevant corporate promos, just the updates and discussions that enrich you. Achieving that will require changes in platform policies and perhaps a cultural shift, but the momentum is building. As generative AI shows us what a worst-case, spam-filled feed looks like, we’re also learning exactly how to build a better one. The path forward is clear – give the power back to users, elevate authenticity, and make social media a place for *meaningful* connection rather than cognitive exhaustion.

Sources:

1. Walter, S. (2023). *LinkedIn “entshitification” post on blocking companies*. LinkedIn
2. Woolley, S. (2021). "Why Can't I Block A Company on LinkedIn?" *Medium (Geek Culture)*.
3. Choi, F. & Chandrasekharan, E. (2025). *Designing Usable Controls for Customizable Social Media Feeds*. arXiv preprint
4. Priezkalns, E. (2023). "AI Comment Spam Appears on LinkedIn." *Commsrisk*
5. Resilience Lab. (2023). "Digital Overload: If Your Screen Time Is Out of Hand." *Resiliencelab.us* 1
2
6. Mercante, A. (2026). "After an oversaturation of AI-generated content, authenticity is in high demand." *Digiday* 3

7. Commsrisk Editors (2023). Comments on AI spam and blocking. *Commsrisk*

8. LinkedIn Community Posts (2023). Discussions on bots and spam on LinkedIn

1 2 Digital Overload: Read This If Your Screen Time Is Out of Hand - Cerebral

<https://www.resiliencelab.us/thought-lab/digital-overload>

3 After an oversaturation of AI-generated content, creators' authenticity and 'messiness' are in high demand - Digiday

<https://digiday.com/media/after-an-oversaturation-of-ai-generated-content-creators-authenticity-and-messiness-are-in-high-demand/>