


Supplementary Attack Chains: OpenClaw Agent

Threat Model & Security Implications
of Read-Only Gmail on EC2

CLASSIFICATION: Security Research – Offensive Threat Modelling
DATE: February 2026
REF: COMPANION DOCUMENT 02-B



API Scopes Do Not Constrain OS-Level Capabilities

The Capability Gap


The primary threat model focused on what the Google API allows. This document focuses on what the Agent allows.

OAuth scopes are application-layer constraints. They do not prevent the underlying OS from executing code, opening network sockets, or accessing cloud metadata.

The Constraint Iceberg



Chain 1: Direct Host Compromise via Reverse Shell



A simple prompt injection triggers an outbound connection. Because most security groups allow egress by default, this bypasses the firewall. Gmail access is not even required for this step.


IMPACT RADIUS:

- > **Identity:** Access to IAM roles & OAuth tokens
- > **Cloud Metadata:** Direct access to AWS IMDS
- > **Lateral Movement:** Pivot point to internal VPC

PREREQUISITES:

[Internet] [Exec]

Chain 7: Persistence & Command and Control



CAPABILITIES:

1. **Adaptive Tasking:** Issue sequential commands based on results.
2. **Scalability:** Manage hundreds of instances via web UI.
3. **Persistence:** Durable channel via HTTP polling.

Instead of a one-off command, the agent polls a third-party URL for instructions. This transforms a single exploit into a managed botnet.

Chain 2: Inbox Proxy & Access Brokering



NARRATIVE

The attacker establishes a proxy tunnel through the agent. This allows them to ‘sell’ a live view of the victim’s inbox to third parties (e.g., corporate spies, stalkers).

Traffic originates from the authorised agent, evading geo-blocking and new device alerts.

USER EXPERIENCE

Dashboard

Active Sessions

✉ Session 001 - Active - \$50/hr

✉ Session 002 - Active - \$50/hr

✉ Session 003 - **EXPANDED VIEW** - Active - \$100/hr


✉ **[New]** Subject: RE: Project Alpha Launch - From: CEO

✉ Subject: Q4 Financials - From: CFO

✉ Subject: Meeting Notes - From: Team Lead

Buyer View:
Real-time
browsing
without login
credentials.

Chain 3: The SMTP Bypass



The Myth: "I have read-only scope, so I can't send email."

The Reality: The agent runs on Linux with `smtpdlib` pre-installed. It can connect to any external mail server. It does not need the Gmail API to send email; it just needs the internet.

Context: The State of SMTP in 2026



Legacy Protocol

SMTP has no built-in authentication. Sender identity is a header, not a verified fact.

SPF Limitations

SPF validates the Envelope Sender (invisible to users), not the 'From' header displayed in clients.



DKIM Loopholes

An attacker can sign emails with their **own** valid domain keys. The signature is technically valid, even if the sender is spoofed.

The Result

Authentication is a 'bolt-on' patch. The underlying pipe remains open.


Why DMARC Is Not a Silver Bullet

01.	Display Name Spoofing DMARC does not validate the Display Name. "CEO John Smith" <attacker@evil.com> remains 100% effective against human targets. 	02.	Cousin Domains Exploiting un-policied lookalikes. Example: bigb4nk.com (No Policy) vs. bigbank.com (Reject). bigb4nk.com (No Policy) → Inbox bigbank.com (Reject) → Rejected/Spam
03.	Incomplete Adoption Many sending domains still publish p=none (monitor-only), allowing spoofed mail to land in inboxes. 	04.	Enforcement Failures Major providers occasionally honour 'Allow Lists' or Spam Confidence Levels over DMARC failures to prevent false positives. 

Chain 4: The Closed Loop (Spoof + Reply)

OUTBOUND: Spoof
Agent sends spoofed email
via direct SMTP (Port 25).
Bypasses Gmail API.


CAPTURE: Intercept
Agent reads the reply via
gmail.readonly API before
the human sees it.



INBOUND: Victim Reply
Target replies to the
legitimate address. Email
lands in Victim Inbox.

Impact: The recipient believes they are conversing with the victim. The agent manages the conversation entirely.

Chain 5: HIBP & Credential Stuffing



HIBP Query

Identify breached services




Breach Dump

Retrieve passwords



2FA Defeat

Trigger login.
Service sends
Email 2FA.




Intercept

Agent reads 2FA
code from Inbox.

Any service using email as a second factor (magic links, codes) is fully compromised. The 'Read-Only' agent neutralises the protection by reading the verification code immediately upon arrival.

Chain 6: Account Self-Reset



The final escalation. The agent triggers a password reset flow for the email account itself. By intercepting the recovery link from the inbox, the attacker changes the password.


OUTCOME:

The gmail.readonly scope is rendered irrelevant. The attacker now possesses primary credentials. Status: TOTAL ACCOUNT TAKEOVER.

Consolidated Risk Matrix

ATTACK CHAIN	PREREQUISITES	STRIDE CATEGORY	RISK RATING
Chain 1: Reverse Shell	Internet, Exec	Elevation of Privilege	CRITICAL
Chain 4: Spoof Loop	Internet, Read-Only	Spoofing / Tampering	CRITICAL
Chain 2: Inbox Proxy	Internet	Information Disclosure	High
Chain 6: Account Reset	Read-Only	Tampering	High
Chain 5: HIBP Stuffing	Read-Only	Elevation of Privilege	Medium

The Compounding Effect of Combined Capabilities



None of these components are vulnerabilities in isolation. Linux is secure; the Gmail API is secure; SMTP is a standard protocol. **It is the architectural combination that creates the critical risk.**

Conclusion & Recommendations

The agent's capabilities are defined by its environment, not its API scopes.

Restrict Egress

Default-deny outbound network policies. Whitelist only essential API endpoints for the agent.

Sandboxing

Remove 'exec' capabilities and standard tools (ncat, python, bash) from the production container image.

Zero Trust

Treat 'Read-Only' agents as potential hostile actors. Assume the host is **compromised** and design the network accordingly.