

Thinking in Graphs: Meaning Through Connectivity

Document: issues-fs__thinking-in-graphs

Version: v1.0

Date: 2026-02-05

Status: Draft

Scope: Foundational — this document underpins all other Issues-FS architecture documents

What This Document Is

This is the foundational document for the Issues-FS ecosystem. It captures the core philosophy, architecture, and reasoning model that everything else is built on. Every other document in the library — the architecture overview, the role-based coordination model, the Lexicon, the framework analysis — depends on the ideas here.

The central claim is simple: **everything is a graph, meaning is not declared but discovered through graph relationships, and confidence in that meaning is proportional to how richly connected a node is to other nodes that provide context.**

This is not a metaphor. It is a literal architectural principle that drives how Issues-FS stores data, how it validates assumptions, how it handles ambiguity, and how it coordinates work across boundaries of language, culture, ownership, and intent.

Part 1: Everything Is a Node

Nodes Have No Obligation to Explain Themselves

A node in a graph is just a node. It may carry some local properties — a title, a value, a timestamp — but it has no inherent obligation to declare what it "is," how it should be used, or what it means. A node labelled "Review" is not a Review in any formal sense. It is a node that someone labelled "Review."

This is a deliberate departure from schema-first thinking, where the type system tells you what something is before you encounter it. In a graph-first model, what something "is" emerges from the edges you can trace from it. A node connected to nothing is meaningless — literally. A node connected to a rich web of other nodes is meaningful in proportion to that connectivity.

This is not a weakness. It is the foundational design choice.

The Safe_UInt__Port Example

Consider a concrete example from the osbot-utils Type_Safe framework:

python

```
class Safe_UInt_Port(Safe_UInt):  
    min_value = 0  
    max_value = 65535  
    allow_bool = False  
    allow_none = False
```

Now consider two scenarios in an Issues-FS graph:

Scenario A: A node has a field `port` with the value `8080`, and that field is typed as `Safe_UInt_Port`, which extends `Safe_UInt`, which is part of `osbot-utils v3.63.4`.

The graph looks like:

```
node__server_config  
├── field: port  
│   ├── value: 8080  
│   └── type ──> Safe_UInt_Port  
│       ├── min_value: 0  
│       ├── max_value: 65535  
│       ├── allow_none: False  
│       └── extends ──> Safe_UInt  
│           └── part_of ──> osbot-utils@3.63.4  
...
```

From this graph, we can say with confidence: "This is a network port. It is an unsigned integer between 0 and 65535. It cannot be null. These constraints are enforced at runtime by osbot-utils version 3.63.4." Every claim is backed by a traceable path through the graph.

Scenario B: A node has a field `port` with the value `8080`, typed as `int`.

```
node__server_config  
├── field: port  
│   ├── value: 8080  
│   └── type ──> int (Python built-in)  
...
```

From this graph, all we can say is: "There is a field called 'port' with the value 8080. The name suggests it might be a network port. The type is `int`, which means at runtime this value could be reassigned to -1, to 999999, to a string (in some contexts), or to None. We *think* this is a network port. The developer probably *intended* it to be a port. But the graph cannot confirm it."

The difference is not in the value. Both scenarios have `(8080)`. The difference is in the **connectivity**. Scenario A has edges that lead to nodes which define, constrain, and contextualise the value. Scenario B has almost none. The meaning is identical in the developer's head. It is radically different in the graph.

This Applies to Everything

The port example is about a primitive value, but the principle applies to every concept in the system:

- An issue labelled "Task" is just a node. If it has edges to an issue-type definition that specifies what a Task is, what statuses it can have, what fields it carries — then we know what this Task is. If it doesn't, we have a node that someone called "Task."
- A role called "QA" is just a node. If it links to a `ROLE.md` definition, to a set of issue types it can create, to handoff protocols — then we know what this QA role does. If it doesn't, we have a label.
- A status called "completed" is just a string. If it links to a state machine that defines what transitions lead to "completed" and what transitions lead away from it — then we know what "completed" means in this context. If it doesn't, someone typed the word "completed."

The graph either provides the evidence or it doesn't. The system's job is to report what it can and can't confirm, not to assume meaning where the edges don't support it.

Part 2: Meaning Through Connectivity

Meaning Is Not Declared — It Is Discovered

In a schema-first system, meaning is declared: "This object IS a Task. It HAS these fields. Its status MUST be one of these values." The schema is the authority, and instances conform or fail validation.

In a graph-first system, meaning is discovered: "This node has edges to other nodes. By traversing those edges, we can determine what this node relates to, what constraints apply, and how confidently we can characterise it."

The distinction matters because declared meaning is brittle and local. A schema works perfectly within the system that defined it. The moment you cross a boundary — a different team, a different project, a different culture, a different language — the schema either forces conformity or breaks.

Discovered meaning is resilient and composable. Each context adds its own edges. The more edges, the more meaning. Crossing a boundary doesn't break the graph — it just means some edges don't extend across that boundary, and the meaning becomes thinner (less confident) on the other side. The system can report exactly how thin: "In your context, this node connects to 12 defining nodes. In the other team's context, it connects to 3. Here's what you share, and here's where you diverge."

The Review Example

Consider five different teams, each with a process they call "Review":

Team A (Japanese manufacturing) has a Review node connected to:

Review-A

has_step

assign_inspector

has_step

assign_recorder

has_step

assign_moderator

has_step

prepare_checklist

has_step

conduct_inspection

(requires: physical_presence)

has_step

document_findings

has_step

classify_findings

(severity: critical|major|minor)

has_step

assign_corrective_actions

has_step

verify_corrections

has_step

sign_off

requires

printed_artifacts

requires

formal_ceremony

governed_by

quality_standard_ISO_9001

language

ja

Team B (open source contributor) has a Review node connected to:

Review-B

has_step

look_at_PR

has_step

leave_comment

("LGTM")

language

en

Team C (Frankfurt compliance) has a Review node connected to:

Review-C

has_step

assign_responsible_officer

has_step

conduct_audit

has_step

document_findings

has_step

produce_audit_trail

has_step

assign_corrective_actions

has_step

submit_to_regulator

requires

retention_7_years

governed_by

BaFin_regulation_MaRisk

language

de

Team D (Lagos startup) has a Review node connected to:

Review-D

```
|— has_step —→ record_loom_video
|— has_step —→ post_comments_in_PR
|— has_step —→ approve_or_request_changes
|— requires —→ async_communication
|— language —→ en
```

Team E (São Paulo agency) has a Review node connected to:

Review-E

```
|— has_step —→ client_presentation
|— has_step —→ collect_feedback
|— has_step —→ document_changes_requested
|— has_step —→ revise_deliverable
|— has_step —→ obtain_sign_off
|— requires —→ client_present
|— language —→ pt-BR
```

None of these nodes are "aware" of each other. None of them declare "I am a schema:Review." They are just nodes with edges.

Finding Common Ground Through Graph Analysis

Now suppose these five teams need to collaborate, or we simply want to understand what "Review" means across the ecosystem. We don't ask each team to conform to a shared schema. We traverse the graphs and look for structural overlap.

Step 1: Identify common subgraph patterns.

All five Reviews share some edges, though they use different local labels:

Common edges found across all 5:

```
has_step —→ [some form of examining the work] → 5/5 teams
has_step —→ [some form of expressing a judgment] → 5/5 teams
```

Common edges found across 4/5:

```
has_step —→ [documenting the outcome] → 4/5 teams (not Team B)
has_step —→ [acting on the outcome] → 4/5 teams (not Team B)
```

Common edges found across 3/5:

```
has_step —→ [assigning a responsible person] → 3/5 teams (A, C, E)
has_step —→ [formal sign-off] → 3/5 teams (A, C, E)
```

Step 2: Surface the structural observation.

The graph analysis can now say: "Team B's Review is structurally thinner than the others. It shares only 2 of the 6 common patterns. Specifically, it lacks documentation and actioning — the steps that make a review consequential. This doesn't mean Team B's Review is 'wrong.' It means that if another team depends on Team B's review having consequences, the graph shows that expectation is unsupported."

Step 3: No one had to agree on anything.

This analysis happened without any team changing their process, learning a shared vocabulary, or conforming to a schema. The common ground was *discovered* in the graph, not *imposed* on it.

Compatibility Is a Graph Computation

This redefines what "compatibility" means. Two nodes are not compatible because they share a type declaration. They are compatible to the degree that their subgraphs overlap when traced toward common reference points.

Compatibility is:

- **Not binary** — it's a spectrum (Team A and Team C are more compatible than Team A and Team B)
 - **Not symmetric** — Team A might satisfy all of Team C's requirements, but Team C might not satisfy all of Team A's (if Team A requires physical presence)
 - **Not global** — Teams A and B are incompatible for compliance purposes but perfectly compatible for the narrow purpose of "did someone look at this"
 - **Computable** — by traversing the graph, not by checking type declarations
-

Part 3: The Fractal Principle

Every Scope Defines Its Own Vocabulary

The Issues-FS ecosystem is not one graph. It is graphs of graphs of graphs, at every level of scope:

```
Issues-FS__Lexicon (root)
├── Issues-FS__Service (project)
│   ├── Project-6 (project instance)
│   │   ├── Epic-14 (epic)
│   │   │   ├── Sprint-3 (sprint)
│   │   │   │   └── Task-23 (task)
```

Every level in this hierarchy can introduce its own nodes, edges, types, and definitions. This is the fractal principle: the same structural pattern — nodes with edges, meaning through connectivity — repeats at every

level of zoom.

The root Lexicon provides anchor nodes for well-understood concepts. These are not mandatory schemas. They are reference points that any node at any level can link to (or not).

A project can introduce issue types that only make sense within that project. A project managing a conference might define `Talk_Proposal`, `Speaker_Slot`, `Room_Assignment`. These types don't exist in the root Lexicon because they are domain-specific. They exist as nodes in the project's scope.

An epic within that project might further specialise: `Lightning_Talk_Proposal` as a refinement of `Talk_Proposal`. This is just another node with an edge to its parent concept.

A single task might define its own sub-vocabulary: checkboxes, sub-steps, acceptance criteria that are specific to that one unit of work.

At no level does a child scope need to "register" its vocabulary with a parent scope. The child scope creates nodes and edges in its own graph. If those nodes have edges that connect (eventually, through any number of hops) to anchor nodes in a parent scope, then the parent scope can make sense of them. If they don't, the child scope's vocabulary is locally meaningful but opaque to the parent.

Name Clashes Are Normal

In a fractal system, name clashes are inevitable and non-problematic. `Task-23` exists in `Project-6` and also in `User-Story-12`. These are different nodes. Their labels happen to collide, just as two directories can each contain a `README.md`.

Resolution is structural, not nominal. Each node has a unique identity in the graph (its node ID). The label is a local convenience. The fully-qualified path through the graph disambiguates: `Project-6 → Sprint-3 → Task-23` is a different traversal than `User-Story-12 → Task-23`.

Type-level clashes are more interesting. If `Project-6` defines a `Review` issue type and `Project-7` also defines `Review` but with different fields, the graph contains two `Review` type-definition nodes with different subgraphs. The system doesn't need to resolve this. It needs to *surface* it: "There are two definitions of 'Review' in scope. They share these edges and diverge on these edges." A human or a Librarian agent can then decide: promote the common parts to a higher scope, or keep them intentionally separate.

Incompatibility and Missing Pieces Are Expected

A fractal system with distributed vocabulary will inevitably produce:

- **Gaps:** An issue references an issue type that doesn't have a definition node anywhere in its reachable graph. The system can report: "This issue claims to be of type X, but no definition of X is reachable from this scope."
- **Partial mappings:** An issue type in one scope maps to 3 out of 5 aspects of a reference definition. The system can report: "This type partially matches the reference. Here's what matches and what

doesn't."

- **Contradictions:** Two definitions in overlapping scopes make conflicting claims. The system can report: "Definition A says status 'closed' is terminal. Definition B says 'closed' can transition to 'reopened.' These are in conflict within scope X."

None of these are errors. They are observations that the graph analysis can make, surfacing them for human (or agent) judgment. The system operates with partial, uncertain, and sometimes contradictory information — just like the real world.

Part 4: Anchor Nodes and Reference Graphs

The Semantic Web's Insight (and Mistake)

The Semantic Web community identified the right problem: how do independent parties exchange meaning without agreeing on everything upfront? Their answer — shared ontologies, RDF triples, linked data — was architecturally sound. The implementations (schema.org, SKOS, Dublin Core, PROV-O) produced genuinely useful reference vocabularies.

But the community made a subtle mistake in practice. They ended up attaching meaning *to nodes* rather than deriving meaning *from edges*. When you write:

```
issues:Review
skos:prefLabel "Review"@en
skos:prefLabel "レビュー"@ja
```

You are declaring: "This node IS a Review, and here are its labels in different languages." The meaning is encoded as properties of the node itself. The node becomes a little document that describes itself.


This is schema-first thinking dressed in graph syntax. The node knows what it is. Other nodes are expected to reference it. The authority flows from the definition node outward.

The Graph-First Correction

In a graph-first model, the schema.org `Review` concept is just another node in the graph. It is a very well-connected node — it has edges to many other well-defined concepts (`reviewBody`, `reviewRating`, `itemReviewed`). It is maintained by a well-known authority (schema.org). But it is not special. It is not a schema that other nodes must conform to. It is a reference point that other nodes may link to.

A team's local Review node does not declare "I am a schema:Review." Instead, the graph may contain edges that say:

Review-A



The edges express partial, specific mappings. "Our `document_findings` step is similar to what schema.org calls `reviewBody`." Not "we ARE a schema:Review." The mapping is granular, honest, and traversable. And critically — the mappings might have been added by anyone: the team itself, a Librarian agent, an automated analysis, or a third party doing an integration.

The local Review node remains untouched. It does not carry any schema.org metadata as properties. The connection exists only as edges in the graph, which can be queried, added, removed, or disputed without modifying the original node.

What Anchor Nodes Actually Are

An anchor node is simply a node that:

1. **Is well-connected** — it has many edges to other well-defined concepts
2. **Is well-maintained** — someone (schema.org, W3C, the Issues-FS Lexicon) keeps it current
3. **Is well-known** — many other nodes across many scopes link to it
4. **Has no special authority** — it does not force conformity; it enables discovery

The Issues-FS Lexicon's job is to provide a curated set of anchor nodes for the issues domain. Some of these will be locally defined (Issues-FS-specific concepts like `Handoff`, `Decision`). Some will be linked to external reference nodes (schema.org, SKOS). All of them are just nodes that happen to be well-connected and well-maintained.

A node in a child scope that links to a Lexicon anchor node gains meaning by association. A node that doesn't link to any anchor node is locally meaningful but opaque to cross-scope analysis. Neither state is wrong. They represent different levels of connectivity, and therefore different levels of confidence in meaning.

Part 5: Confidence as a Function of Connectivity

The Confidence Spectrum

Every assertion the system makes about a node sits on a confidence spectrum:

No edges —→ "We know nothing about this node beyond its label"

|

Few local edges —→ "We know some local properties but can't verify
their meaning"

|
|

Edges to typed definitions —→ "We know the structure and constraints
(like Safe_UInt__Port: 0-65535)"

|

Edges to anchor nodes —→ "We can relate this to well-known concepts
and compute compatibility with other nodes"

|

Edges to external references —→ "We can relate this to globally shared
vocabularies (schema.org, ISO standards)"

|

Rich multi-hop connectivity —→ "We have high confidence in what this is,
how it's constrained, and how it relates
to the broader ecosystem"

The system should always be able to answer: "How confident are we about what this node means?" And the answer is always computable from the graph: count the edges, measure the depth of traversal to anchor nodes, assess the authority of the nodes reached.

Honest Uncertainty Is the Default

Most nodes in most graphs will sit in the middle of this spectrum. They have some edges but not enough for full confidence. The system's default posture is honest uncertainty:

- "This field is called `port` and has value `8080`. We *think* it's a network port based on the name. But it's typed as `int`, so we can't enforce constraints."
- "This issue is labelled `Review`. It has edges to three steps. Based on those steps, it partially matches the Issues-FS Lexicon's reference review pattern. But it's missing the `document_outcome` step, so we can't confirm it's a full review process."
- "This role is called `QA`. It has a `ROLE.md` that defines responsibilities. But the `ROLE.md` doesn't link to any issue type definitions, so we can't confirm what issue types this role is authorised to create."

In every case, the system reports what it found (the edges that exist), what it didn't find (the edges that are missing), and what that means for confidence. It never fills in the gaps with assumptions.

Enrichment, Not Enforcement

When the system identifies low-confidence nodes, the response is not validation failure. It is an opportunity

for enrichment: "Adding an edge from this field to `Safe_UInt__Port` would increase confidence from 'we think this is a port' to 'we know this is a port with constraints 0-65535'."

The Librarian role, or automated tooling, or a developer can add that edge. The node itself doesn't change. The graph around it becomes richer, and confidence increases.

This is how the ecosystem improves over time: not by imposing stricter schemas, but by growing the graph. Every edge added is a small increment in the system's understanding of itself.

Part 6: Graphs of Graphs

The Recursive Structure

The Issues-FS ecosystem is not a single flat graph. It is a graph of graphs, recursively:

- **An issue** is a node in a project's graph. But the issue itself may contain a subgraph: sub-tasks, checkboxes, attachments, comments, each with their own nodes and edges.
- **A project** is a node in the ecosystem graph. But the project contains a graph of issues, milestones, sprints, and roles.
- **The Lexicon** is a graph of reference concepts. But each concept (like "Decision") has its own subgraph: fields, constraints, valid transitions, related concepts.
- **A role repo** is a node in the development graph. But it contains a graph of prompts, templates, issue types, and handoff protocols.
- **The Type_Safe framework** is a graph of type definitions. But each type (like `Safe_UInt__Port`) has its own subgraph: constraints, inheritance chain, version history.

At every level, the same principles apply: nodes are just nodes, meaning comes from edges, confidence comes from connectivity depth, and no node is obligated to explain itself.

Cross-Graph References

The most powerful aspect of this model is cross-graph references. When an issue in `Project-6` has a field typed as `Safe_UInt__Port`, it creates an edge that crosses from the Issues-FS graph into the osbot-utils Type_Safe graph. That cross-graph edge is what allows the system to say "we know this is a port" — the meaning comes from a completely different graph, maintained by a different team, with its own versioning and evolution.

Similarly, when a local Review node has edges to `schema.org/Review` subcomponents, those edges cross from the local project graph into the schema.org reference graph. The meaning is enriched by the connection, not by importing or copying the schema.org definitions.

This is what makes the graph model scale across teams, languages, cultures, and agendas. Each graph is sovereign. Cross-graph edges create interoperability without requiring any graph to change its internal structure.

Versioning and Temporal Graphs

Cross-graph edges raise an important question: when `Safe_UInt_Port` changes in a new version of osbot-utils, does the meaning of the port field in Project-6 change?

The answer depends on which version edge the project graph contains. If the edge points to `osbot-utils@3.63.4`, then the meaning is pinned to that version's definition. If the edge points to `osbot-utils@latest`, then the meaning floats with the latest release. Both are valid choices. The graph makes the choice explicit and traceable.

This applies to all cross-graph references. A link to a Lexicon anchor node can be versioned: "this matches the v1.0 definition of Decision" or "this tracks the current definition of Decision." The graph records the choice, and the system can surface the implications: "This node links to Lexicon v1.0's definition of Decision, but Lexicon v2.0 has changed the valid transitions. Your node may be operating on a stale definition."

Part 7: Practical Implications for Issues-FS

What the Lexicon Actually Provides

Given this graph-first philosophy, the Issues-FS Lexicon is not a dictionary or a schema registry. It is:

1. **A curated set of well-connected anchor nodes** for the issues domain (Task, Bug, Decision, Handoff, Review, etc.), each with a rich subgraph of relationships, constraints, and reference edges.
2. **A set of cross-graph edges** to external reference vocabularies (schema.org, SKOS, PROV-O), maintained as first-class graph structure, not as metadata properties on Lexicon nodes.
3. **Graph traversal tools** that can compute: connectivity depth, compatibility between nodes, coverage against reference patterns, and gaps where edges are missing.
4. **Bootstrap definitions** — opinionated but overridable types for common issues-domain concepts. These are conveniences, not constraints.

What Every Repo Can Do

Every repository and every scope in the fractal hierarchy can:

- **Add nodes and edges** in its own graph without coordinating with any other scope
- **Link to Lexicon anchors** for the concepts it wants to connect to the broader ecosystem
- **Define local types** that only make sense in its own context

- **Override or specialise** bootstrap definitions by creating local nodes with different subgraphs
- **Ignore the Lexicon entirely** for concepts that are purely local

What the System Can Observe

Given any set of connected graphs, the system can:

- **Compute confidence** for any node: how well-connected is it to defining nodes?
- **Compute compatibility** between any two nodes: how much do their subgraphs overlap?
- **Identify gaps**: where are there missing edges that would increase understanding?
- **Surface conflicts**: where do two definitions in overlapping scopes contradict each other?
- **Trace provenance**: where did this definition come from? What authority does it carry? What version?
- **Generate honest assessments**: "We know X. We think Y. We cannot confirm Z."

What the System Never Does

The system never:

- **Rejects a node** for lacking type conformity
 - **Forces a vocabulary** on any scope
 - **Assumes meaning** where edges don't support it
 - **Hides uncertainty** behind default assumptions
 - **Requires registration** of local vocabulary with a parent scope
 - **Treats any node as authoritative** by virtue of its source alone — authority is a property of connectivity, provenance, and governance, not of identity
-

Part 8: A Worked Example — End to End

To illustrate the full model, here is a complete scenario that spans the fractal hierarchy, crosses graph boundaries, involves multiple languages and cultures, and demonstrates honest uncertainty.

The Setup

A distributed team is building Issues-FS itself. The team includes:

- A developer in London writing the core library
- A QA engineer in Tokyo writing integration tests
- A compliance officer in Frankfurt reviewing security practices
- An open source contributor in Nairobi proposing a feature

The Graph

The developer creates a task:

Task-42 (in Project: Issues-FS__Service)

```
|— title: "Add rate limiting to API endpoints"
|— type —→ Task (local definition in project scope)
|   |— links_to —→ Lexicon:Task (anchor node)
|— status: "in_progress"
|   |— defined_by —→ Lexicon:State_Machine:Task
|       |— in_progress can transition to: [completed, blocked]
|— assigned_to —→ Role:Dev (in Issues-FS__Dev__Role__Dev)
|   |— role_definition —→ ROLE.md (defines Dev responsibilities)
|— field: max_requests_per_minute
|   |— value: 100
|   |— type —→ Safe_UInt (osbot-utils)
|       |— but NOT Safe_UInt__Rate_Limit (doesn't exist yet)
|— field: port
|   |— value: 8080
|   |— type —→ Safe_UInt__Port (osbot-utils@3.63.4)
|       |— min: 0, max: 65535
|       |— extends —→ Safe_UInt —→ osbot-utils@3.63.4
```

What the system can say about this task:

- "Task-42 is a Task (high confidence: linked to project definition which links to Lexicon anchor)."
- "Its status is 'in_progress,' which is a valid state with defined transitions (confidence from state machine link)."
- "It is assigned to the Dev role, which has a ROLE.md definition (confidence from role definition link)."
- "The `port` field is a network port, 0-65535 (high confidence: linked to Safe_UInt__Port with full constraint chain)."
- "The `max_requests_per_minute` field is an unsigned integer (medium confidence: linked to Safe_UInt), but we have no domain-specific constraints. We know it's non-negative, but we can't say if 100 is reasonable, what the upper bound should be, or whether this represents requests per minute, per second, or per hour. The field name suggests per-minute, but the graph doesn't confirm it."

The QA Engineer in Tokyo

The QA engineer creates a review request — but in their local vocabulary:

レビュー依頼-7 (in Scope: Issues-FS__Dev__Role__QA)

```
├── title: "Task-42のレート制限実装をレビュー"
├── type ──→ レビュー依頼 (local type definition)
│   ├── has_step ──→ テスト計画作成 (create test plan)
│   ├── has_step ──→ テスト実行 (execute tests)
│   ├── has_step ──→ 結果文書化 (document results)
│   ├── has_step ──→ 品質判定 (quality judgment)
│   ├── has_step ──→ 是正措置 (corrective actions)
│   └── links_to ──→ Lexicon:Review_Request (anchor)
│       └── partial_match: 4/5 reference steps
├── linked_to ──→ Task-42
└── language ──→ ja
```

What the system can say:

- "レビュー依頼-7 links to Task-42 (confirmed: direct edge)."
- "Its type links to the Lexicon Review_Request anchor with a 4/5 partial match (high confidence in compatibility)."
- "It includes a corrective actions step (是正措置) that the Lexicon reference doesn't require. This is a local enrichment."
- "The language is Japanese. The Lexicon anchor's labels don't include a Japanese label. The connection is structural (via graph edges), not nominal (via shared label)."

The Compliance Officer in Frankfurt

The compliance officer needs to verify that the review process meets regulatory standards:

Überprüfung-3 (in Scope: Frankfurt Compliance)

```
├── type ──→ Audit_Review (local type)
│   ├── has_step ──→ assign_responsible_officer
│   ├── has_step ──→ document_findings
│   ├── has_step ──→ produce_audit_trail
│   ├── has_step ──→ retention_period: 7 years
│   ├── governed_by ──→ BaFin:MaRisk
│   └── links_to ──→ Lexicon:Review_Request (anchor)
│       └── partial_match: 3/5 reference steps
├── reviews ──→ レビュー依頼-7 ──→ Task-42
└── language ──→ de
```

What the system can say:

- "Überprüfung-3 (Frankfurt) reviews レビュー依頼-7 (Tokyo) which reviews Task-42 (London). The review chain is traceable through the graph."
- "Both review processes link to the Lexicon Review_Request anchor. Their overlap on the reference steps is: both have `document_findings` and `quality_judgment/produce_audit_trail`. The Tokyo process has `corrective_actions` which Frankfurt doesn't require. Frankfurt has `retention_period` and `regulatory_governance` which Tokyo doesn't have."
- "For the narrow purpose of 'was the code reviewed,' both processes are compatible. For the narrow purpose of 'does this meet BaFin requirements,' only Frankfurt's process has the necessary edges."

The Open Source Contributor in Nairobi

The contributor proposes a feature but has no connection to the role system or the Lexicon:

```
feature-proposal-1 (in Scope: GitHub Issues)
├── title: "Add WebSocket support for real-time sync"
├── type ──→ GitHub:enhancement (GitHub's label system)
│       └── (no link to Lexicon)
├── body: "It would be great if..."
└── author ──→ github:contributor_123
```

What the system can say:

- "feature-proposal-1 is a GitHub issue labelled 'enhancement.' We have no link to the Issues-FS type system, so we can't map it to an issue type."
- "If `Issues-FS__Service__GitHub` were to import this issue, it could add an edge: `GitHub:enhancement` → `Lexicon:Feature_Request`. That mapping would give us confidence about what this issue is. Without the mapping, we know it's a GitHub issue, and that's all."

This is honest. The system doesn't pretend the GitHub issue is a Feature_Request just because a human would make that inference. It reports the absence of the connecting edge and notes what adding that edge would enable.

Summary: Core Principles

1. **Everything is a node.** Nodes carry local properties but have no obligation to declare what they are or how they should be used.
2. **Meaning comes from edges.** What a node "is" emerges from the graph relationships that can be traced from it. No edges, no meaning.

3. **Confidence is proportional to connectivity.** The more edges you can trace from a node to well-defined reference points, the more confidently you can characterise that node.
4. **The system is fractal.** Every scope — from the root Lexicon down to a single task — can define its own nodes, edges, types, and vocabulary. The same structural principles apply at every level.
5. **Anchor nodes enable interoperability without enforcing conformity.** The Lexicon provides well-connected reference nodes. Other nodes link to them at whatever granularity makes sense. Partial mapping is normal and expected.
6. **Compatibility is computed, not declared.** Two nodes are compatible to the degree that their subgraphs overlap when traced toward common reference points. Compatibility is a spectrum, not a boolean.
7. **Honest uncertainty is the default.** The system reports what the graph supports and what it doesn't. It never fills gaps with assumptions.
8. **Enrichment, not enforcement.** When confidence is low, the remedy is adding edges, not adding validation rules. The graph grows; it doesn't constrain.
9. **Cross-graph edges are first-class.** The most powerful connections span graphs: from project issues to Type_Safe definitions, from local types to Lexicon anchors, from Lexicon anchors to schema.org references.
10. **No node is aware of how it's used.** The information about what a node means and how it participates in workflows is extracted from the surrounding graph structure, not encoded as properties of the node itself.