


Running GenAI Agentic Teams in Production

A Working Example from February 2026



This is not a theoretical proposal. This is a detailed technical account of how a single human lead orchestrated two AI agent teams to build, ship, and operate a secure, zero-knowledge file transfer service.

AUTHOR: Dinis Cruz
STATUS: Active / Live Production
DATE: 17 February 2026

One Human, Two Teams, Five Days

1

HUMAN LEAD

Directs, Decides, Reviews



18

AGENT ROLES


Research, Code, Security, QA



5

DAYS


13-17 February 2026



75+

DOCUMENTS


950KB Structured Data



Moving beyond chat bots to structured, autonomous agentic workflows.

The Vehicle: SG/Send

A Zero-Knowledge Security Product (send.sgraph.ai)



The Stack

Python Backend,
Vanilla JS Frontend
(No Frameworks),
Web Components.


The Complexity

Requires GDPR
compliance, AppSec
review,
Infrastructure
management, and User
Documentation.

Why This Project?

It is open source, live,
and faces real security
threats. It is not a toy.


The Human-Agent Operating Model



KEY CONSTRAINT:
NO custom orchestration frameworks.
NO agent-to-agent APIs.
Communication happens strictly via files.


The Explorer Team: Innovation & Chaos

Mission: Experiment, Research, Prototype. "It is okay to break things."



The Villager Team: Discipline & Delivery

Mission: Production, Deployment, Uptime. "Breaking things is NOT okay."



Explorers optimise for speed. Villagers optimise for reliability. Features graduate from Explorer to Villager.


Principle #1: Files, Not Conversations

CHAT INTERFACE



Ephemeral. Hard to Audit. Limited Context Window.


FILE SYSTEM (GIT)



Permanent. Versioned. Recoverable.

Agents communicate by reading/writing to the repository.
The file system has no context limit. ``git blame`` traces every decision.


The File System is the Knowledge Base




Immutable File Deployments (IFD):
Documentation is versioned. We do not overwrite history; we commit new versions (v0.1.2 → v0.4.7).

Role definitions are 10–30KB structured text files, not simple prompts.

The Daily Workflow: Voice to Code



Scaling Complexity: Briefing Packs




When tasks are too big for a single prompt, we create Briefing Packs—dedicated context folders containing everything an agent needs.


Method Streams

Code context must reference actual file paths and line numbers. No descriptions; agents use the raw code to ensure high-context execution without hallucination.

Production Velocity

Timeline: 13–17 February 2026 (5 Days)



 **QUALITY NOTE**
Security researchers reviewed code. Agents classified findings on P0-P10 scale.


KEY STATS

18
Roles Created

29
Vulnerabilities

BETA
User Feedback Integrated

What Works: Role Routing & Velocity



Key Points


- **Clarity:** 18 roles creates speed, not confusion. No meetings required.
- **Voice Efficiency:** 10 mins 'rambling' voice note becomes 30KB structured plan. Human thinks, AI formats.
- **Separation:** Two teams (Explorers/Villagers) prevent the 'Speed vs Quality' conflict.

The Friction Points



Context Windows



The Air Gap



Latency




The Future of Work (2026 Outlook)

1 Human + Agents = 10+ Person Team Output

FROM	TO
 Doing	 Directing
 Coding	 Judging
 Writing	 Reviewing

The Quality Gate: The system is designed so the **human's judgment** is the only bottleneck. This is a **feature, not a bug**.




Start Your Agentic Team

The tools exist (**Claude, Git, Voice**). The barrier is discipline.

PRODUCT: `send.sgraph.ai`

SOURCE CODE: `github.com/owasp-sbot`

DOCUMENTATION: `docs.send.sgraph.ai`



What the agents eliminate is not the human's judgment — it's the busywork between the judgment calls.