

Agentics Foundation – Briefing for Consulting & Systems Integration Partners

Purpose

This document outlines how mid-size consulting and systems integration companies can leverage the expertise of the Agentics Foundation community to radically improve their delivery capabilities, internal operations, and client outcomes through agentic AI.

The Agentics Foundation community is composed of practitioners who have been hands-on with AI agents for years — building, experimenting, and shipping real multi-agent systems. This isn't theoretical; these are people with deep operational experience in agentic workflows.

The Strategic Imperative

Generative AI — and agentic AI specifically — is becoming an existential competitive factor for consulting firms. A small team of 5–6 people augmented by 20–30 specialised agents, orchestrated across dynamic workflows, will consistently outperform a traditional team of 20–50 consultants in terms of research depth, deliverable quality, speed, and client communication.

This isn't about replacing humans. It's about solving the perennial problems consulting has always faced: scaling delivery, absorbing client context, producing high-quality outputs consistently, and making senior expertise economically accessible for shorter engagements.

The window is now. The tooling has matured significantly in the last six months. Yesterday would have been the best time to start. Today is good. Tomorrow is materially worse.

Two-Sided Value Proposition

1. **Internal transformation** — These companies (of any size) can use agentic approaches to improve their own processes, research, analysis, and deliverables.
2. **Client-facing services** — Once internally capable, they can offer agentic-powered services to clients who are on the same journey and want to learn.

The clients and the consulting firms are on the same maturity curve — the firms just need to be a step ahead.

Agentic Maturity Model

The community has identified a clear progression of capability levels. The key principle: **start small, then evolve.**

Level 1 — Role-Based Prompting

The simplest form of an "agent": a well-crafted system prompt that tells an LLM to behave as a specific role.

- "You are a researcher / architect / journalist / executive..."
- Still effective — and increasingly so as base models improve.
- Zero infrastructure required. Pure prompt engineering.

Level 2 — Multi-Agent Repos with Claude Code

This is where things get interesting. Using tools like Claude Code, you can spin up sub-agents with defined roles and have them collaborate within a single project.

How it works in practice:

- Create a repository with a `CLAUDE.md` file containing instructions and context.
- Define roles in a `ROLES.md` file (e.g., DevOps, Architect, Librarian, Historian, CISO, DPO, GRC, Journalist, Cartographer, etc.).
- Claude manages the interactions between agents via the conversational interface.
- Start with 2–3 roles, scale to 15–20 as needed — not all agents are active simultaneously; you create sub-teams focused on specific tasks.

Real-world example: The community's own projects run teams of 20 agents across 3 sub-teams covering roles from developer to ambassador to alchemist.

Level 3 — Claude Desktop / Teams

A step up in orchestration: agents can talk to each other more fluidly with better context management. Claude's Teams offering provides a more managed version of Level 2 with improved multi-agent interaction.

Level 4 — Claude Flow (via Agentics Community)

Claude Flow, published by the Agentics community, is a more powerful orchestration layer. It itself has three sub-levels:

Sub-Level	Description
4a	Role-based orchestration similar to Claude Teams, but with richer primitives.
4b	Agents gain access to a broader set of tools and capabilities that Claude Flow ships with — significantly extending what each agent can do.
4c	Full swarm-mode: highly composable, dynamically orchestrated agent environments tackling complex projects with high efficiency. This is the advanced frontier, demonstrated extensively by community members like Reuven Cohen.

The Talent Model — Why Agentics Community Members Are Uniquely Valuable

Traditional consulting staffing doesn't work for agentic delivery. The Agentics Foundation community offers a different model:

Traditional Model	Agentic Model
20–50 consultants on multi-week engagements	5–6 agent-native specialists on focused sprints
Senior talent locked into long commitments	Senior talent contributing high-impact hours, amplified by their own agents
Linear scaling (more people = more cost)	Exponential scaling (more agents = more capability at marginal cost)

Why hire from this community:

- **Agent-native** — They think in terms of multi-agent workflows, not just prompts.
- **High-calibre talent** — Exactly the kind of expertise consulting firms want on engagements.
- **Flexible engagement model** — They don't have time for 5-day consulting blocks, but they can deliver outsized value in a couple of hours, powered by their own agent teams.

Think of it as the **Upwork model for senior agentic consultants**: a talent pool that is currently economically inaccessible to these firms under traditional models, but perfectly viable in an agent-amplified engagement structure.

The human element is not removed — it is amplified. A community member can spin up agents that ask exactly the right questions, conduct the research, and deliver results — then the human brings judgment, context, and credibility to the final output.

Recommended Next Steps

1. **Start a pilot** — Pick one internal project. Stand up a small multi-agent repo (Level 2). Learn by doing.
2. **Engage Agentics community members** — Bring in 1–2 agent-native practitioners for a short sprint to accelerate learning.
3. **Build internal capability** — Use the maturity model above as a roadmap. Move through levels deliberately.
4. **Package for clients** — Once you have internal confidence, begin offering agentic-powered delivery to clients.