


Capturing Data Before Entropy

The Commercial Value of Structured Source Data



Prepared for Senior Leadership & Strategy Teams

**Enterprise data management
is trapped in a costly paradox.**

90%


of enterprise data is unstructured
(documents, PDFs, emails, slides)

Organisations are currently spending billions
on Generative AI, RAG pipelines, and Vector
Databases to 'read' this noise.


The Inversion

Unstructured data is largely
an artificial problem of
our own making.

Data is born structured.
It is rich with context,
relationships, and
meaning until we
intervene.




We break our data, then buy expensive glue to fix it.



Result: We are paying to solve a problem we created.

Publishing structured data into documents introduces 'Data Entropy'.


Data Entropy is the loss of order, context, and provenance that occurs when structured inputs are compressed into a static report.



The Reality Gap: The document is just a static 'photo' of the data. It is no longer connected to the source reality.

The hidden tax of the document-centric workflow.

The Cost of Reconstruction




The Risk: Broken Provenance

-  **Manual Extraction:** Re-keying data into spreadsheets introduces human error.
-  **Traditional ETL:** Brittle scripts break when document formatting changes.
-  **AI Hallucination:** Without original context, models 'guess what the author's models author 'ream the author meant,' breaking explainability.

Without lineage, auditors cannot trace conclusions back to evidence.

Visualising the waste: The Status Quo


The A → B → C → D Lifecycle



The Bottleneck: The B → C conversion is expensive, replicated across every consumer, and introduces errors.

The Strategic Pivot: Distributing the Source.

Optimised Lifecycle (A -> D)



Structured Source Data is a revenue-generating asset.

Optimised Lifecycle (A -> D)



Clean Source Data



Trust & Provenance



**Extraction Cost
Avoided by
Consumer**

£££  Revenue Opportunity

- **Monetise the Ingredients:** Sell API access or Knowledge Graph subscriptions, not just PDF downloads.
- **The Value Proposition:** If extraction costs the consumer £100k, they will gladly pay £30k for the clean source data.
- **Competitive Edge:** In Regulated Industries (Finance, Pharma), full lineage is a premium differentiator.

Buyers pay for certainty, speed, and reduced risk.



Immediate Machine-Readability

Zero time spent on parsing. Machines can ingest JSON-LD or RDF graphs directly.



Eliminating Guesswork

Definitions and relationships are specified in the ontology. No need to 'infer' intent.




Traceability

Context is preserved. Consumers understand not just the 'what' but the 'why' (lineage of insights).

Drastically reduced risk of AI hallucinations.

Unlocking network effects through interoperability.


Silos



Documents do not interoperate.

vs.

Ecosystems



Structured Data can be joined.

New Business Models: Data Marketplaces and Mashing Up.


Key Concept: One-to-many distribution of *knowledge*, rather than one-to-many extraction of *text*.

Implementation requires a shift in tools and governance.

Maturity Model



Stop fixing unstructured data. Start preventing it.



From:
Parsing broken data and
hallucinating context.

To:
Capturing meaning at
the point of creation.

The future of Enterprise AI isn't better extraction; it's provenance-preserving publication.