





**High LLM API costs
are an engineering
deficit.**

The providers have no incentive to help you fix your context infrastructure. It is not malicious; it is just market economics.



The fatal flaw of conversational history: maintaining context means sending the full history every time. Every request carries the dead weight of the past.



Unchecked agents burn tokens blindly. They solve problems by brute force, not by understanding. The result is a fragile system held together with chewing gum.

Most “hallucinations” are just accurate responses to briefs with gaps.



The background features a light blue line-art illustration of a chain with several links. One link in the upper right is broken, with a globe of the Earth positioned at the point of fracture. Several sharp, triangular shards are depicted around the broken link and the globe, suggesting a shattering or breaking apart. The overall style is technical and conceptual.


**The Radical Principle:
Every LLM request
must be one-shot.**

No history. No accumulated context.

The LLM Runtime



The Engineered Pipeline



Continuity is maintained through your engineering, not the LLM.
Curate the context. Strip the irrelevant. Yield 10x-100x reductions in token usage.

Case Study 1: The Naive Approach.


Using the LLM as the runtime to modify HTML structure.

Result: Slow execution, broken tags, and a cost of \$1.50 for every single page transformation.



The Engineered Solution: Use the LLM to write the code, not do the transformation.


The LLM acts as the author, not the runtime.
After the initial generation, the script runs deterministically, forever, for free.





Treat agent management like project management.

Never run all agents simultaneously. Curate inputs through structured markdown briefs. Employ specific agents dedicated entirely to preventing token waste.



Micro-teams




Structured Briefs



Cost-saving Roles
(Librarian & Accountant)

The LLM Maturity Map

Most teams build a prototype, see that it “works”, and stop at Genesis. The high costs you are paying are simply the price of accepting Genesis-level infrastructure.



The Town Planner Playbook. Rigorous infrastructure for relentless optimisation.

1. **Hash-based caching:** Only process content that has changed.

2. **Model routing:** Dynamically assign the cheapest capable model per task.

3. **Skeleton prompts:** Send structure, not raw content.

4. **Pre-computed context:** Use site descriptions instead of raw HTML.

The Virtuous Cycle.

Stop treating every token as an inevitable tax. Build the infrastructure to go from thousands in API debt to predictable, scaled efficiency.

