The GenAI Usage Paradox: Why Experts Should Be the Power Users

A conversation last week crystallised something I've been thinking about for a while.

I was talking to someone who recently started using ChatGPT. They were getting into it, discovering the value bit by bit, getting excited about the parts they understood. I started describing agentic workflows, how I delegate entire workstreams to AI agents, how the agents coordinate across repos and produce structured deliverables. The usual pushback followed: that sounds nice, but I'm not there yet.

Then they said something that stopped me: "I use AI for the things I'm not good at. The areas where I have doubts, where I don't understand something. That's where I go to the LLM."

And there it is. The paradox.

The Wrong Way Round

What this person described is the most common pattern I see. People use LLMs for domains where they lack expertise. They ask the AI to write legal contracts when they're not lawyers. They ask it to debug code when they're not developers. They ask it to draft medical summaries when they're not clinicians.

It feels productive. The output looks real. It has the right structure, the right vocabulary, the right tone. But here's the problem: because they're not experts in that domain, they cannot tell when it's wrong. They can't spot the hallucination in a legal clause. They can't see the subtle bug in the generated code. They can't catch the clinical error in the medical summary.

This is why there are so many horror stories. This is why so many AI-assisted projects don't fulfil their potential. People are using AI as a substitute for expertise they don't have, and they lack the domain knowledge to quality-check the output.

The Real Power Is the Other Way

The real value of agentic AI is not for the things you don't know. It's for the things you know deeply.

When you are an expert in a domain, you know what good looks like. You can delegate with precision. You can review output and immediately spot what's off. You don't need the AI to think for you. You need it to execute for you. To handle the scaffolding, the plumbing, the non-functional requirements, the repetitive structuring that consumes 80% of the time but requires only 20% of the expertise.

I run 7-8 Claude sessions in parallel across different repos. I delegate architecture briefs, code implementation, investor materials, documentation. But I know what the output should look like. I can read a brief and tell in thirty seconds if the architecture is sound. I can review generated code and spot the design flaw. The AI does the heavy lifting. I do the critical thinking, the direction-setting, the quality gate.

That is the pattern that works. Expert plus agentic AI equals scaled expertise.

The Formula One Analogy

Think of it this way. The best racing drivers in the world don't say "I'm a great driver, I don't need a good car." They say "give me the most advanced machine you have, because I can extract every ounce of performance from it."

A Formula One car in the hands of someone who has never driven is dangerous. The same car in the hands of a professional driver breaks records.

The person I was talking to was essentially saying: "I'm a great driver, so I'll just use the AI to help me park. But I don't know how to fly a plane, so I'll let the AI fly for me." That's exactly backwards. The great driver should be using AI to win races, not to park. And nobody should be letting AI fly a plane they can't land themselves. Use your expertise where it's deepest. Let the advanced tooling amplify what you already know. And if you don't have the expertise, don't hand the controls to an AI you can't supervise.

The Paradox

The people most poised to get transformative value from agentic AI are domain experts. They know how to delegate. They know what quality looks like. They can review and iterate at speed. Agentic AI doesn't replace their expertise. It makes their expertise reach heights and execution speeds that were never previously possible.

But these are exactly the people who push back. Some of it is the fear of being replaced. Some of it is the assumption that if they're already good at something, technology can't add much. Some of it is that they haven't seen what agentic workflows actually look like when driven by someone who knows the domain.

Meanwhile, the people who lack domain expertise are the enthusiastic adopters. They use AI for the things they don't understand. And they produce output that looks right but isn't, because they can't tell the difference.

The Opportunity

There are entire professions that have never had good technology. Professions that rely on spreadsheets, emails, and manual processes because nobody ever built proper tools for them. These professionals have deep domain knowledge and zero tooling.

Agentic AI changes this. For the first time, a domain expert can describe what they need in natural language and get structured, executable output. Not a generic tool they have to adapt to. A workflow that adapts to them.

But only if they use it for what they know. Not for what they don't.

The expert who delegates within their domain gets scaled expertise. The non-expert who delegates outside their domain gets confident-sounding nonsense. Same technology, opposite outcomes.

The question isn't whether AI is useful. It's whether the right people are using it for the right things.