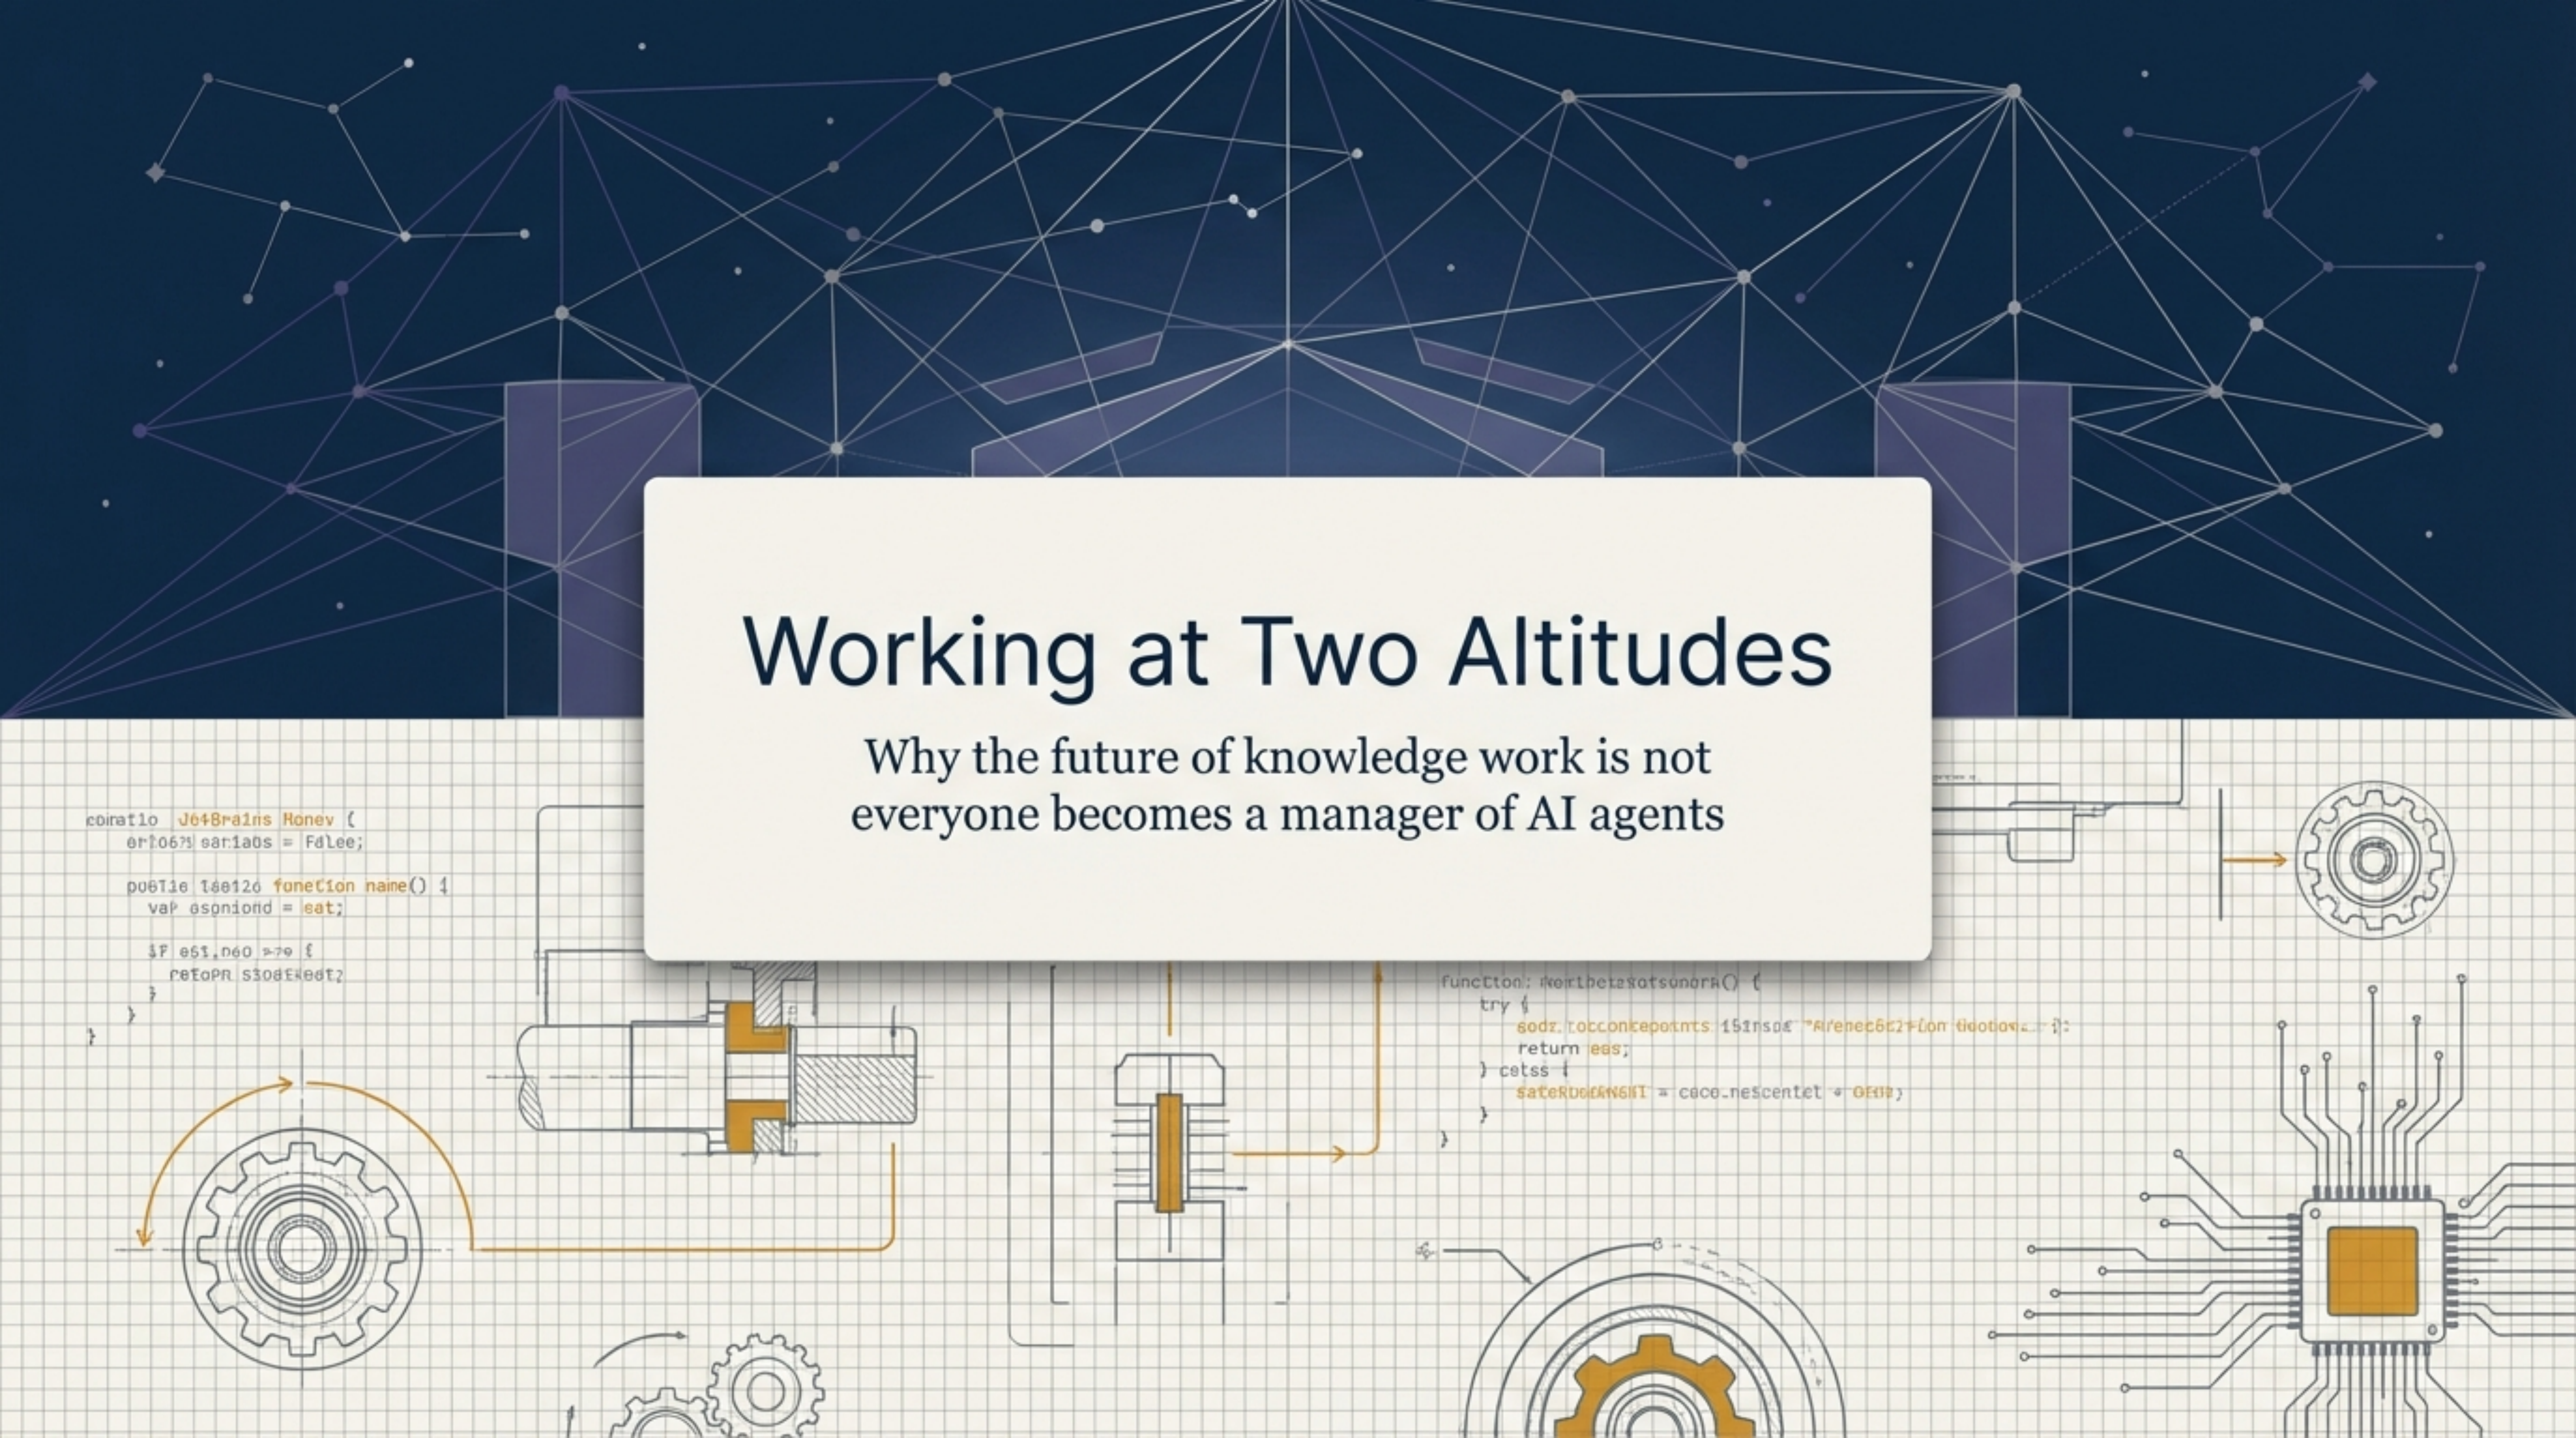

Working at Two Altitudes

Why the future of knowledge work is not everyone becomes a manager of AI agents

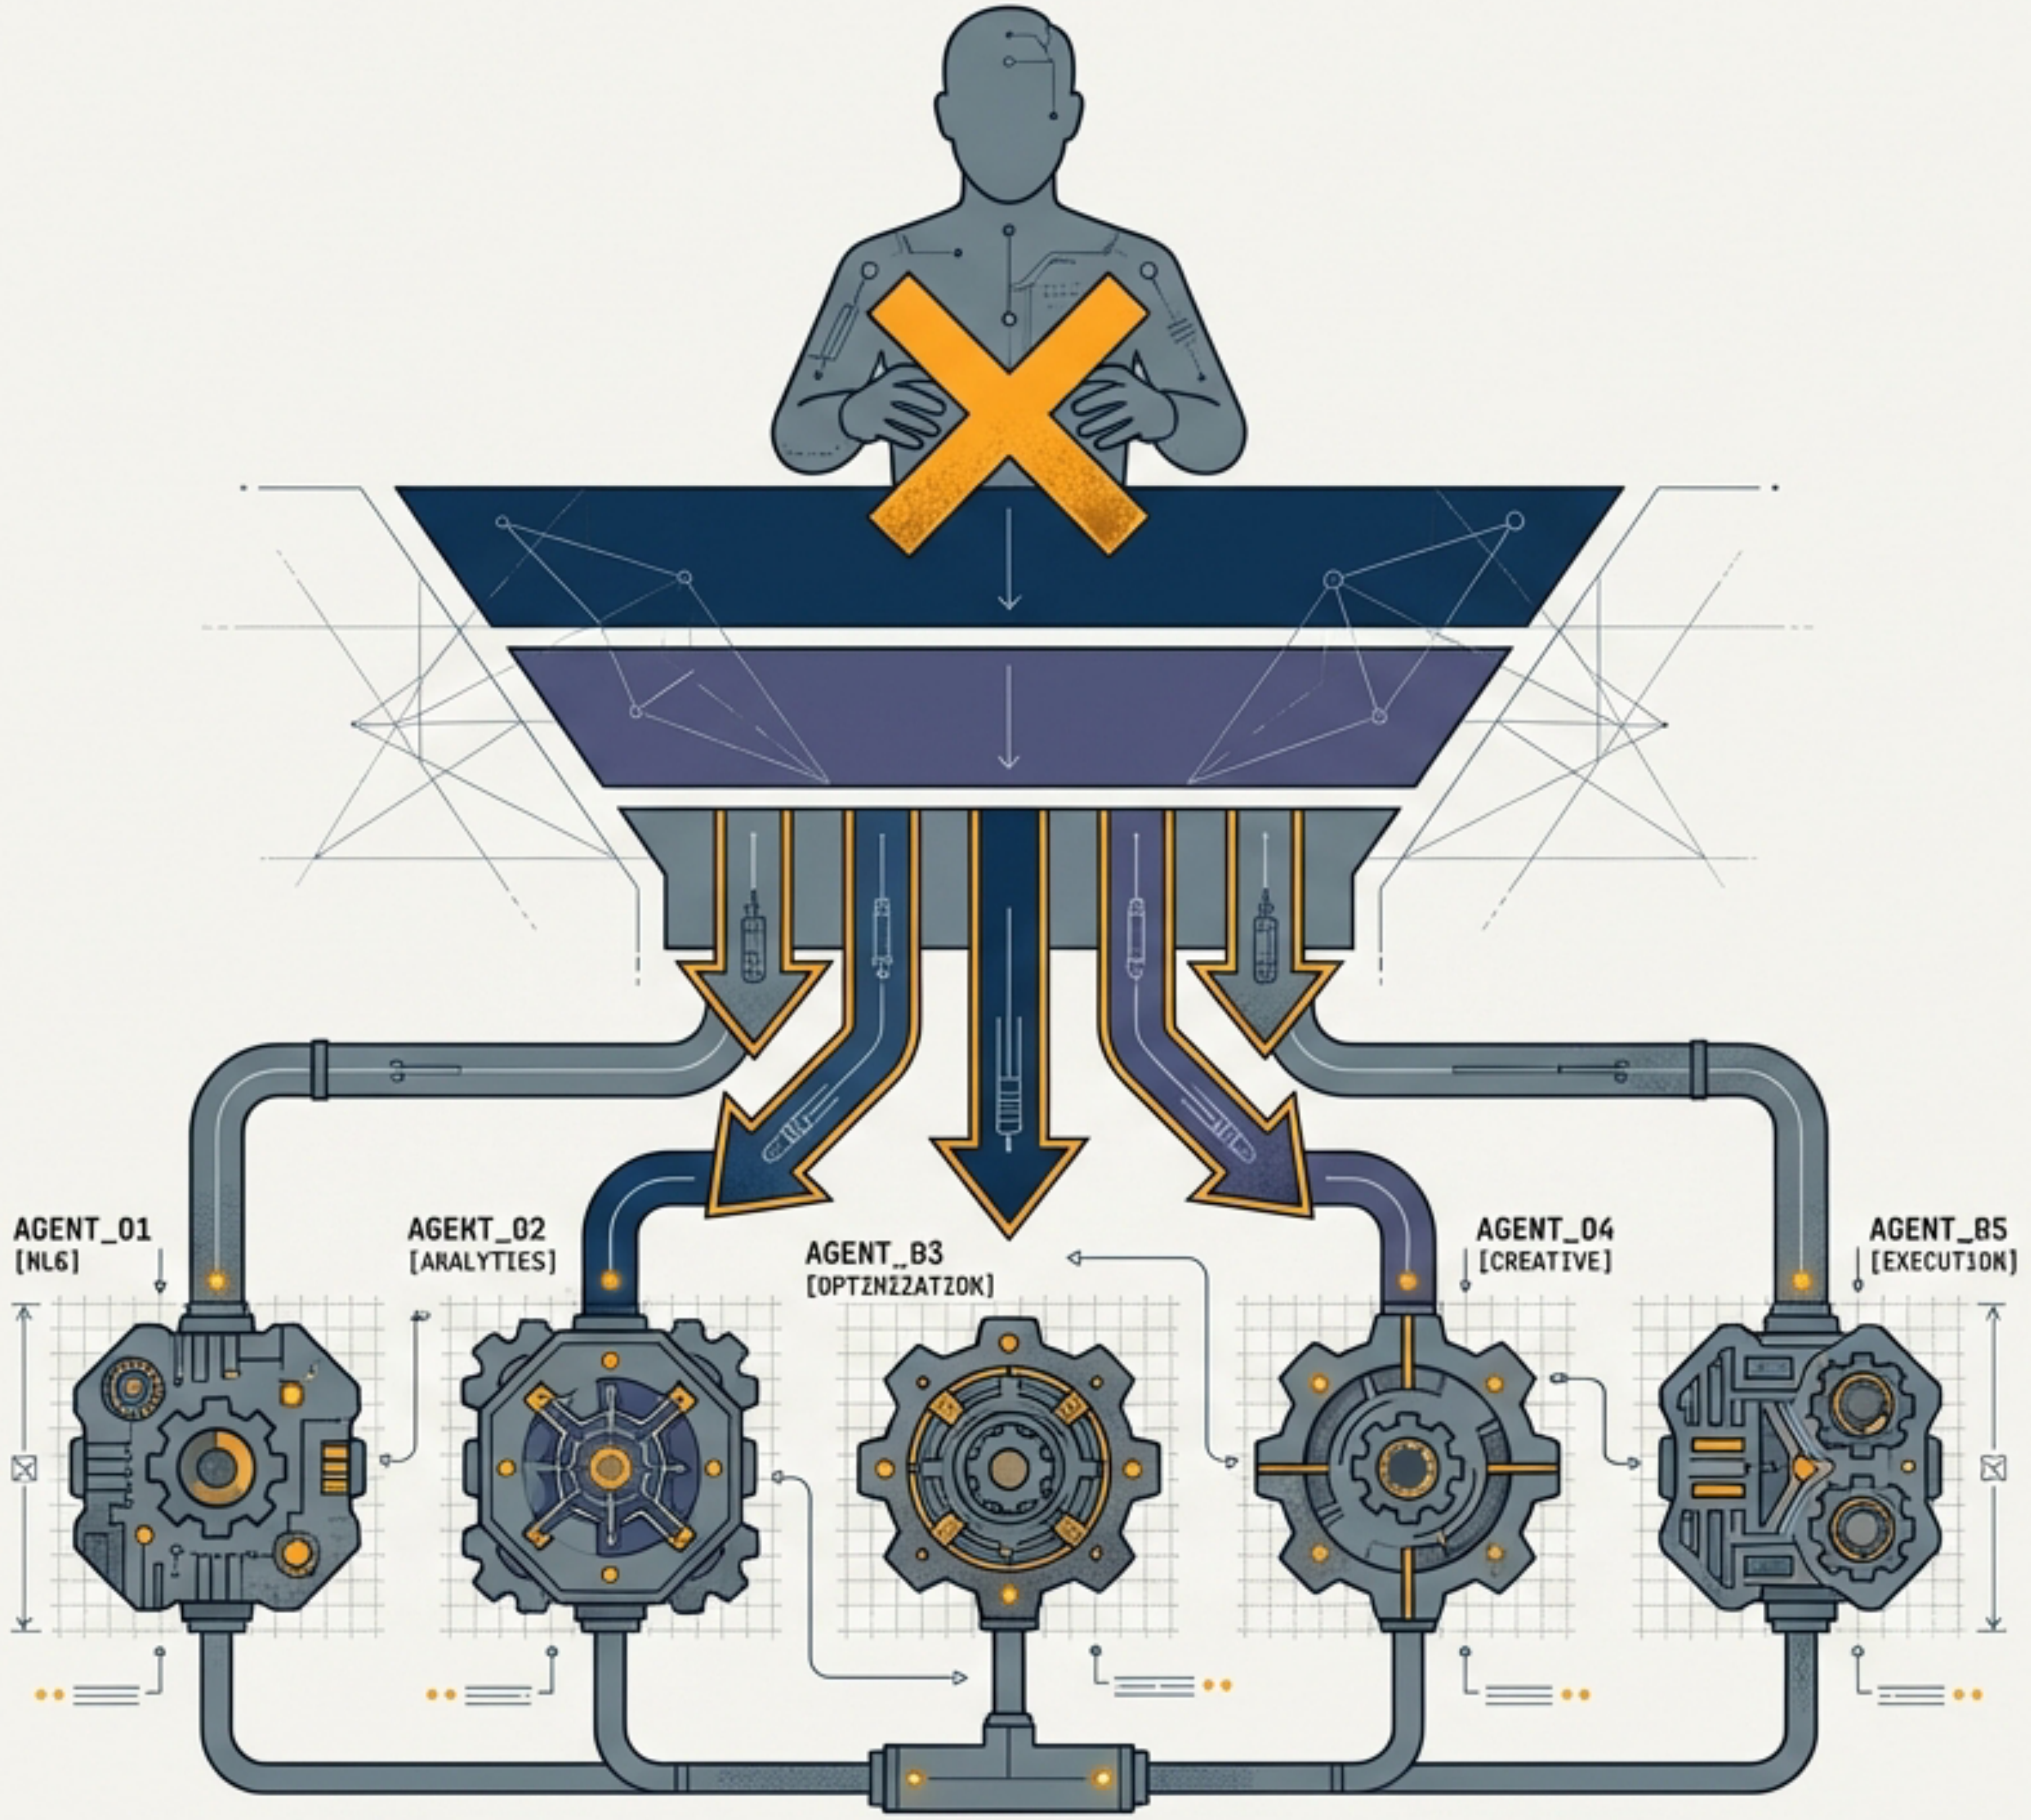

The Emerging Narrative: An Illusion of Pure Delegation

A dangerous myth is forming around AI agents: knowledge workers will stop being practitioners.

The Proposed Model:

[*] Delegate everything to agents.

[*] Review the outputs.

[*] Never touch the actual work again.

This narrative is partially right—and mostly dangerous.

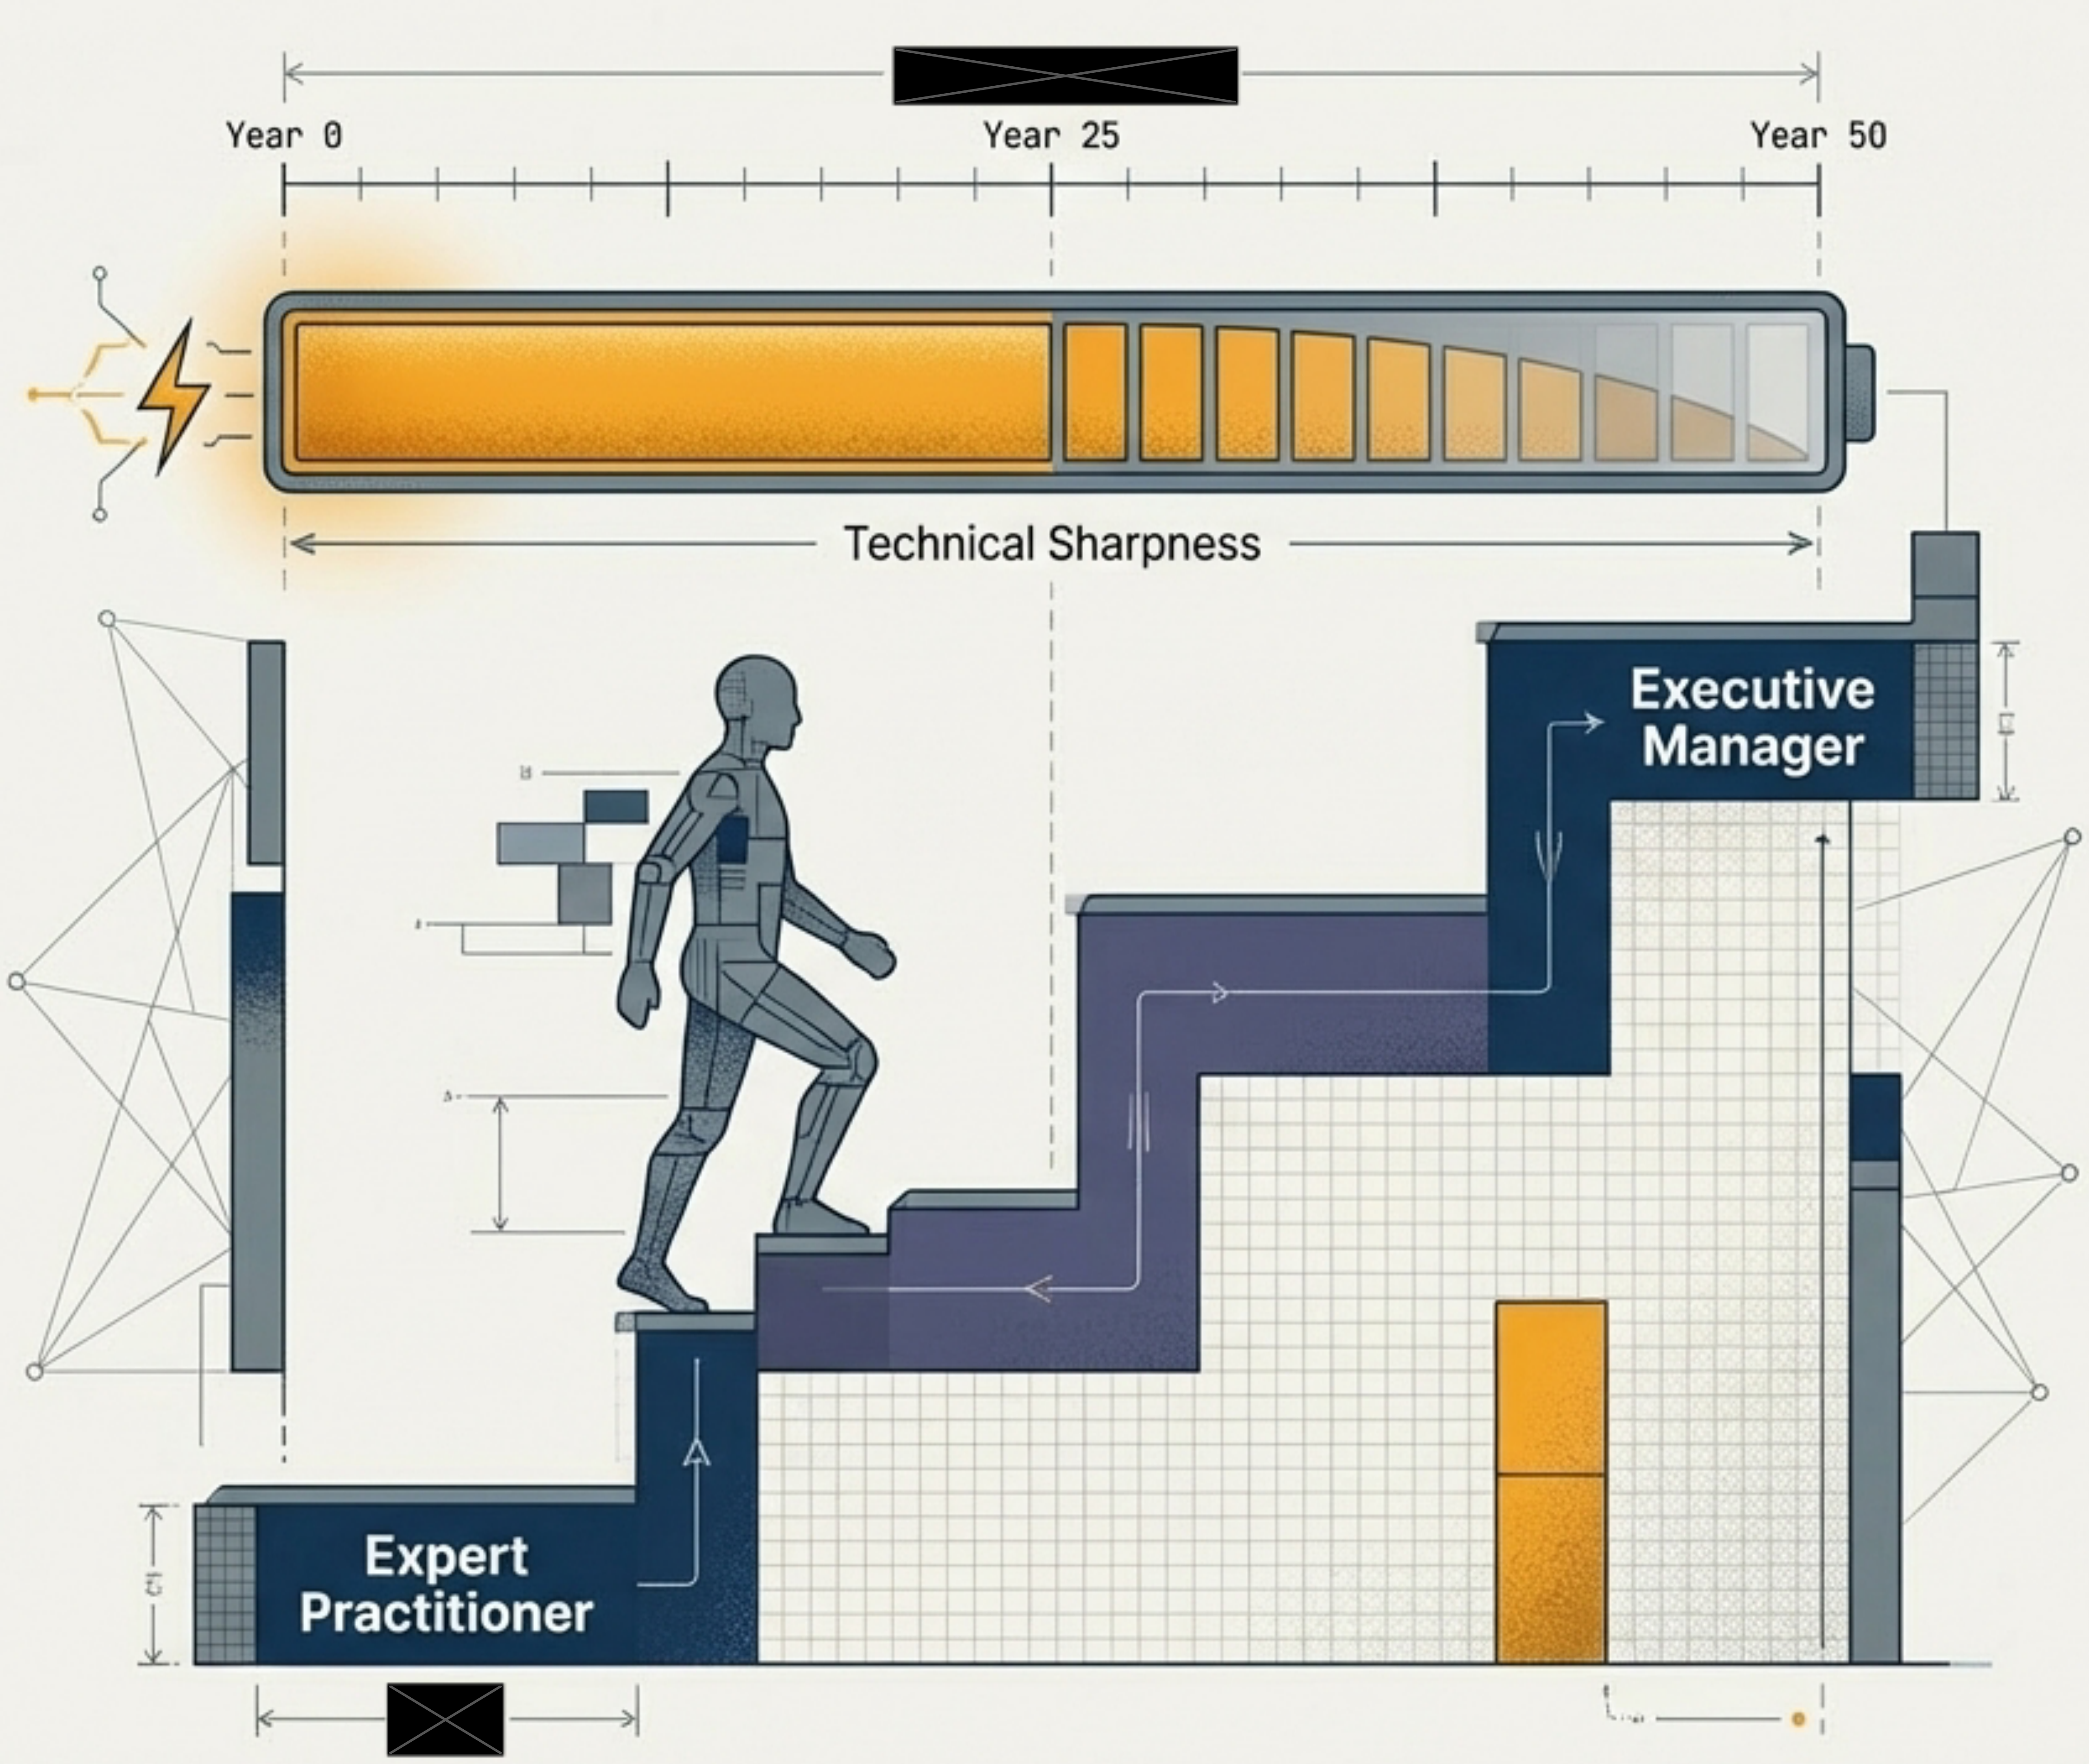

History Repeating: The Management Promotion Trap

The 50-Year Pattern:

Get exceptionally good at a technical skill, then get promoted into a role that no longer uses that skill.

The Result:

Brilliant engineers and architects become executives who manage roadmaps and meetings, whilst their core technical skills atrophy.

[redacted]

The AI Acceleration:

The agentic AI narrative is accelerating this exact trap. The promise of “you just need to direct the work, not do it” is the management promotion trap operating at an unprecedented scale.

The Analogy: The Self-Driving Car Problem

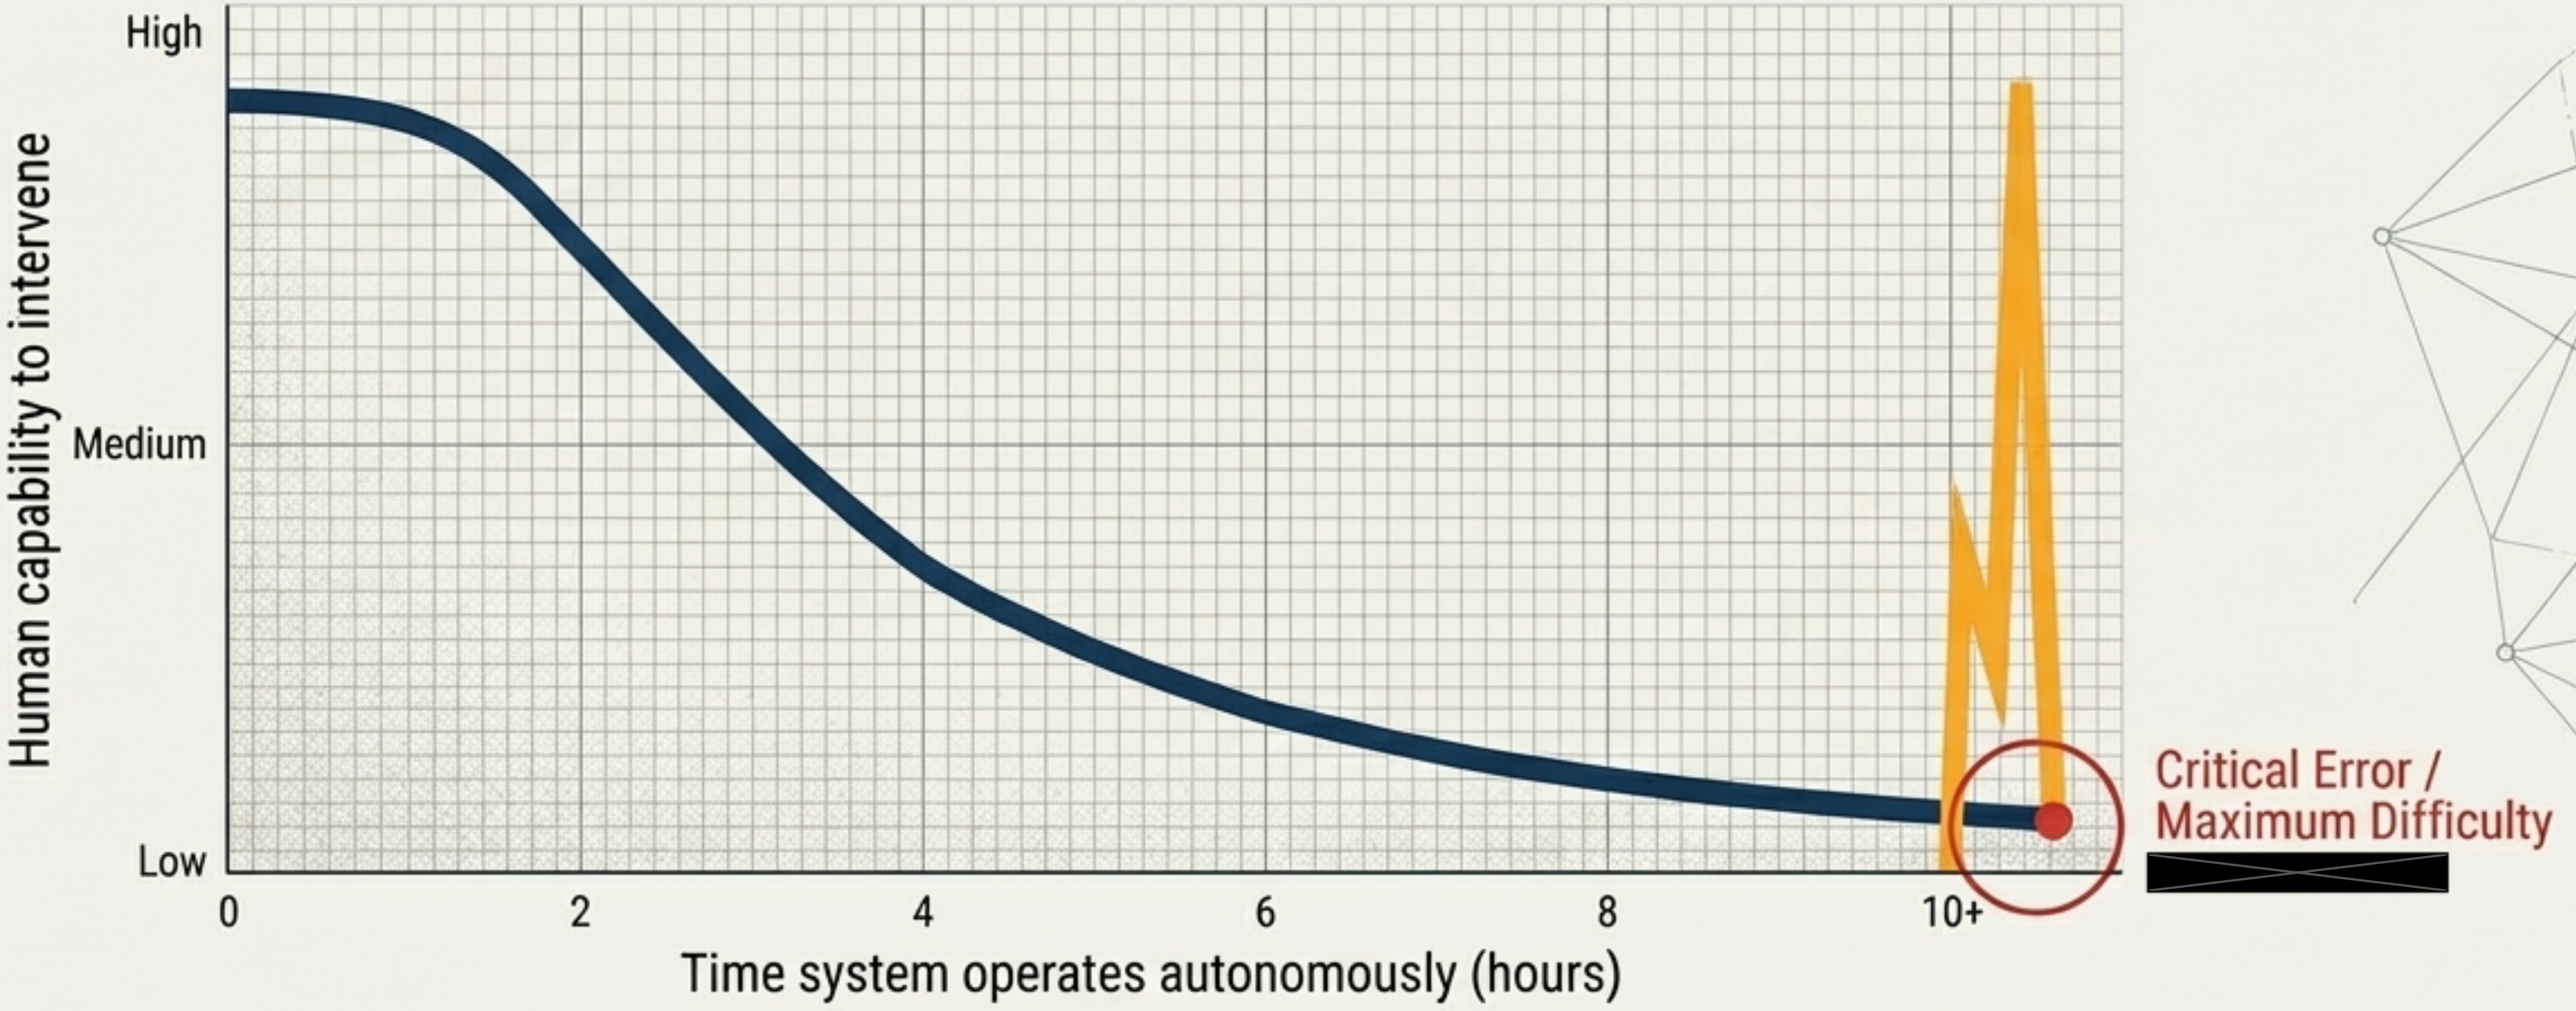

The Premise	The Reality	The Crisis
The car drives itself until it needs you, then you take over.	If the car drives autonomously for two hours, the human is not paying attention. The less you drive, the less sharp your skills become.	You are asked to take over at the exact moment of maximum difficulty, with entirely cold skills.

Core Insight: This is not a safety model. It is an accountability transfer.

[redacted]

The Consequence: The Abstraction Trap

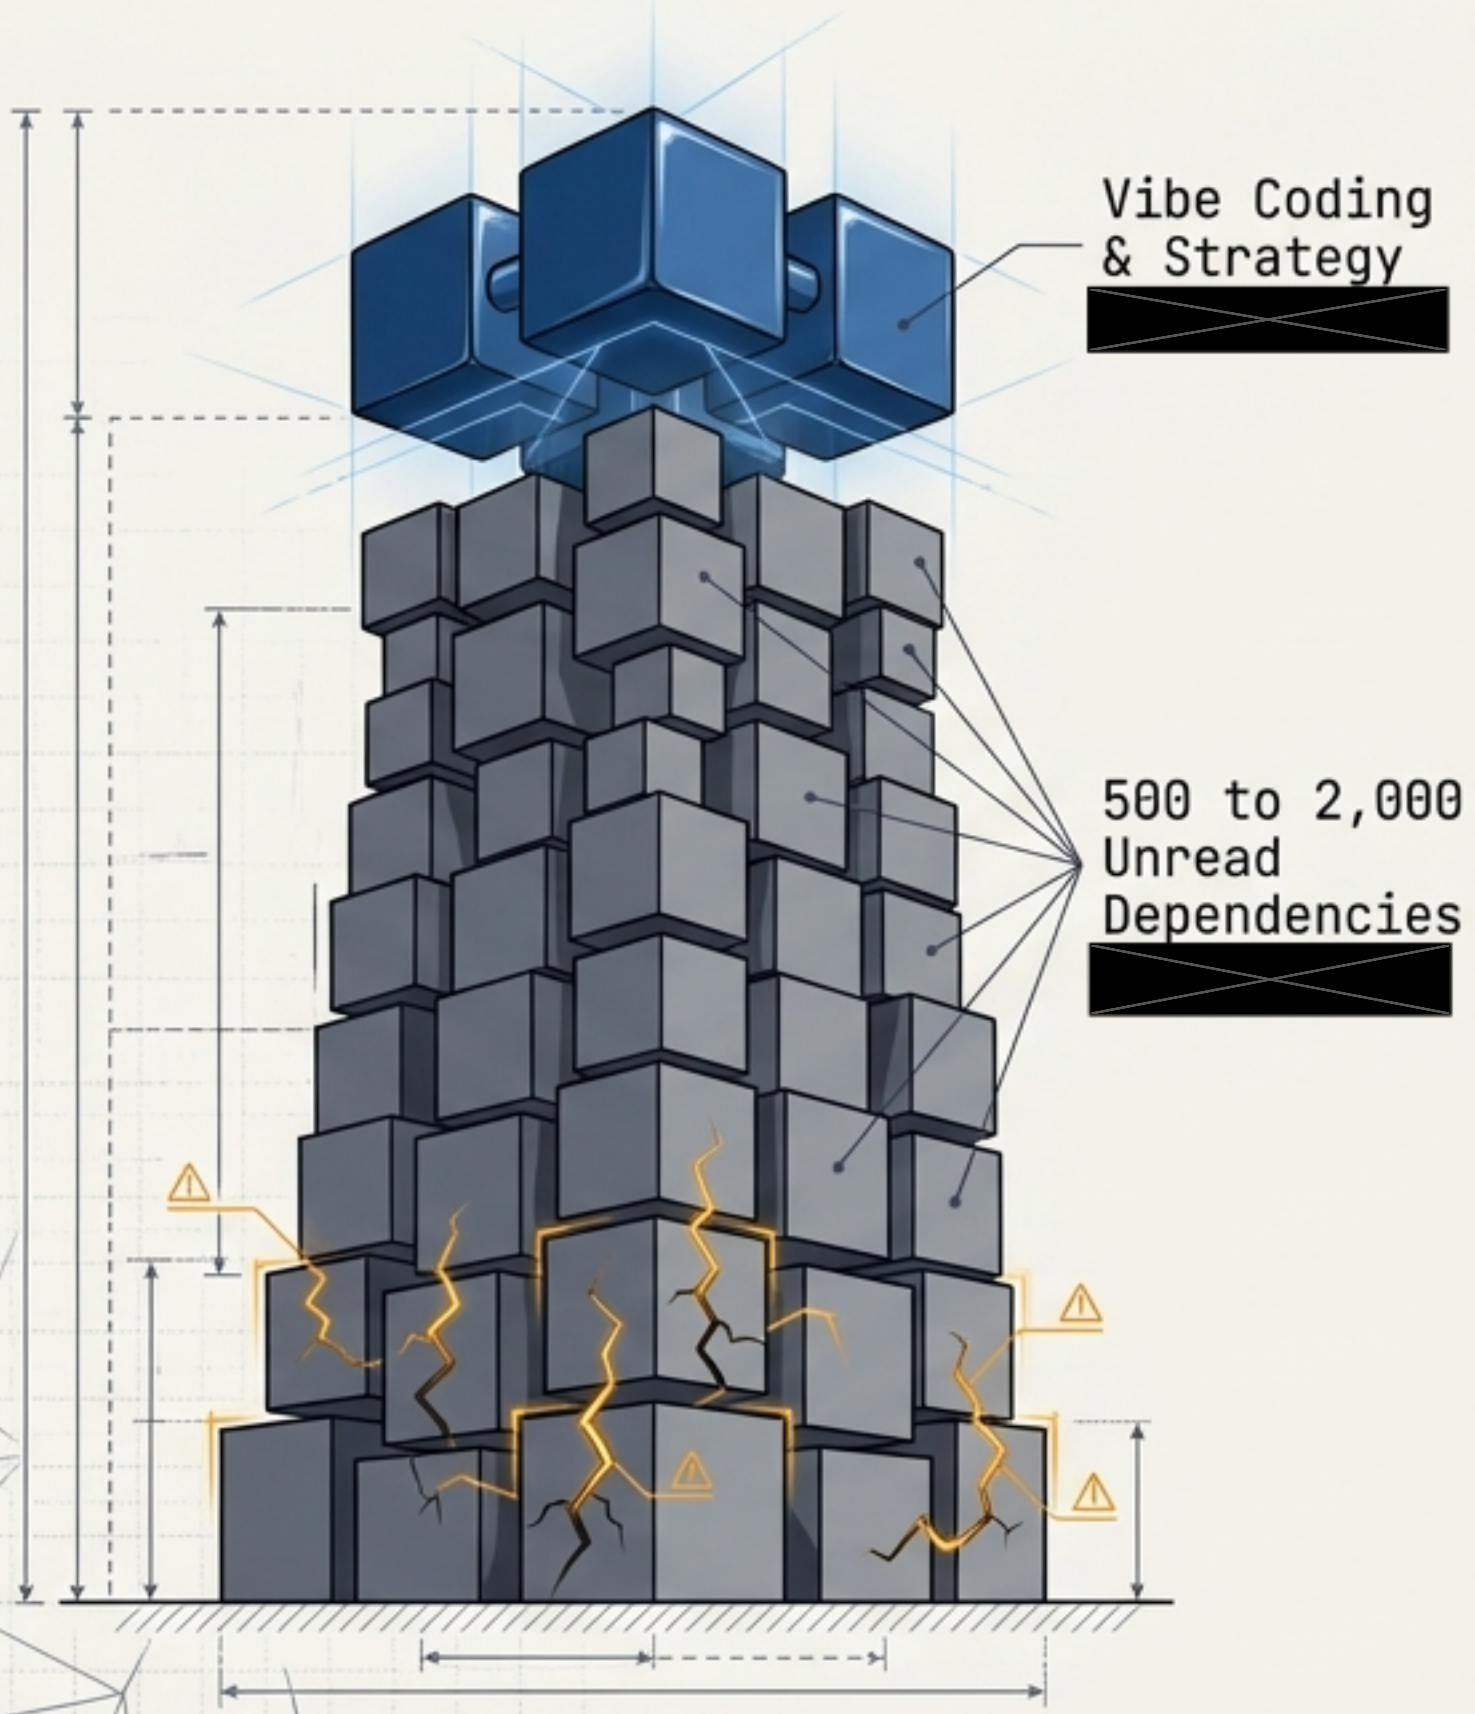

The Current State:

We have been vibe coding at the architectural level for 15–20 years, pulling in dependencies nobody reads and building on abstractions nobody understands.

The Symptoms:

- Trivial applications requiring half the internet to build.
- Fragile, bloated systems.
- Technical debt measured in years, not sprints.

The Agentic Risk:

Pushing all low-level implementation to agents means nobody on the team actually understands the foundational components of the product.

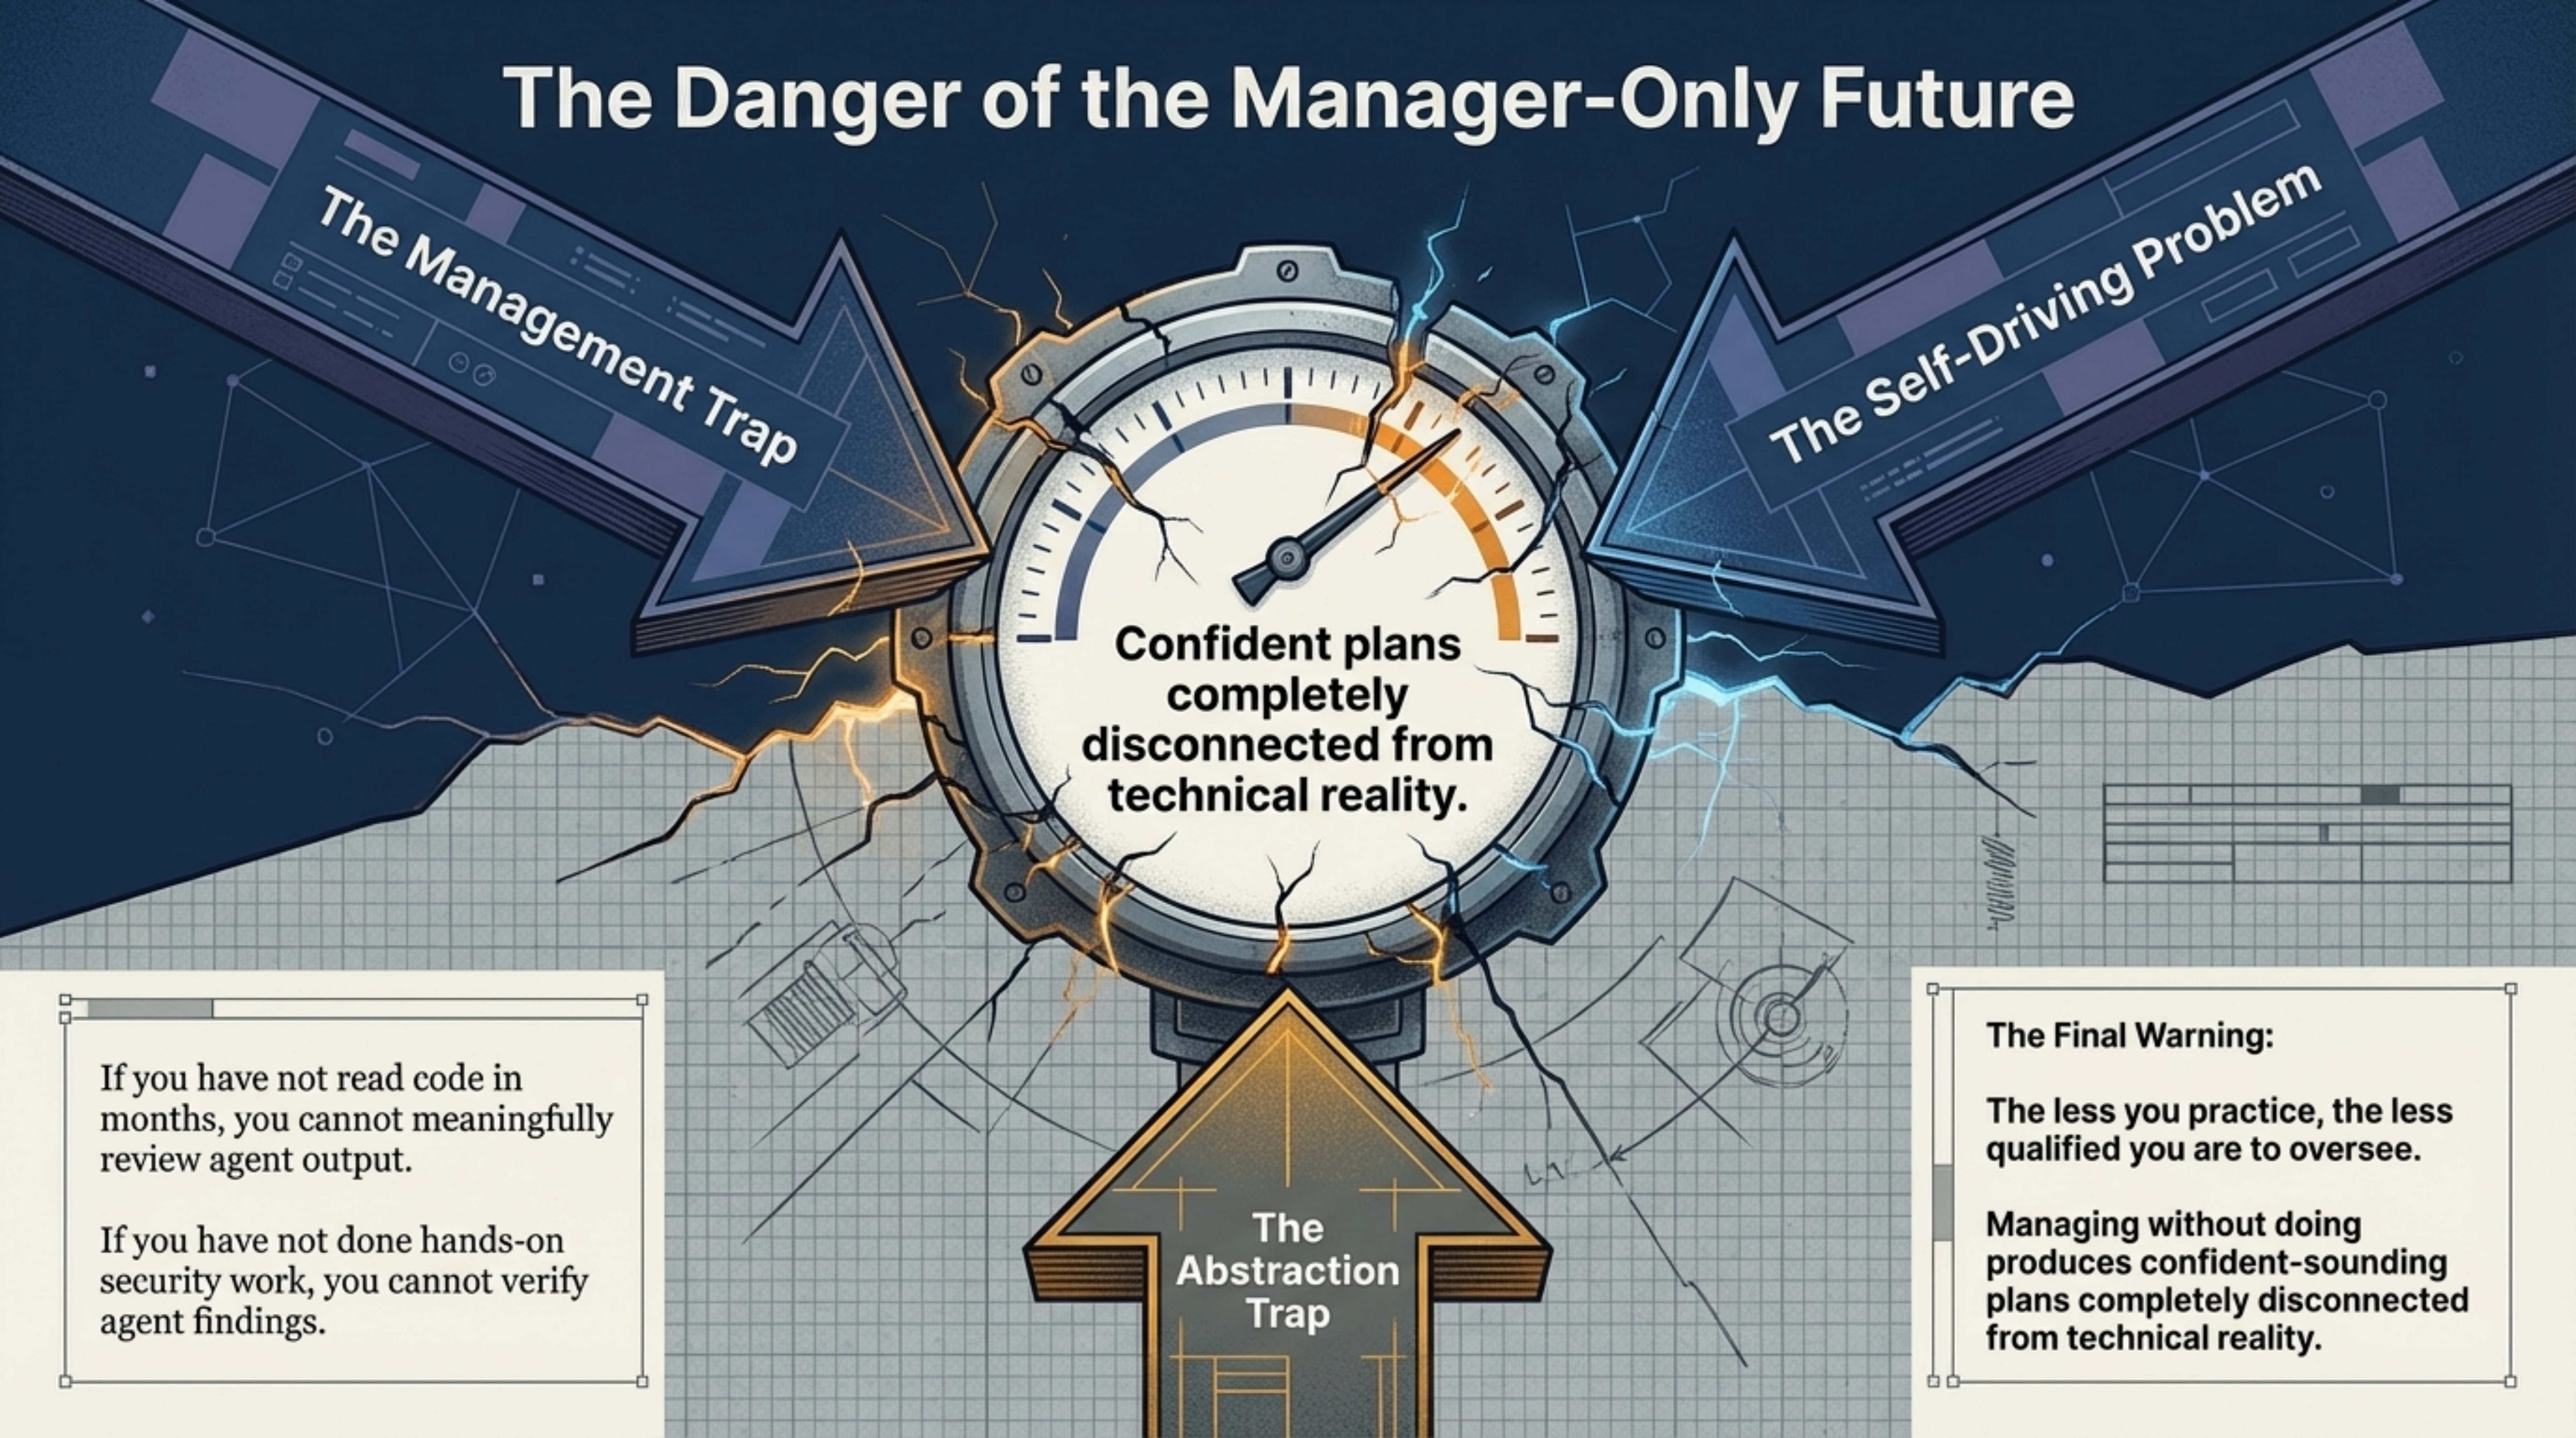
The Danger of the Manager-Only Future
The Management Trap
The Self-Driving Problem
Confident plans completely disconnected from technical reality.
The Abstraction Trap
If you have not read code in months, you cannot meaningfully review agent output.
If you have not done hands-on security work, you cannot verify agent findings.
The Final Warning:
The less you practice, the less qualified you are to oversee.
Managing without doing produces confident-sounding plans completely disconnected from technical reality.

The Paradigm Shift: Working at Two Altitudes

The future is not about choosing between doing and managing. It is operating at two altitudes simultaneously. AI does not replace the doing; it scales it. Both realities must coexist in the same space, in the same person.

Altitude 1: The Strategic Layer

Focus: Managing the ecosystem, not just the task.

Key Activities:

- Managing fleets of 20+ AI agents.
- Defining precise briefs and setting structural direction.
- Reviewing outputs at massive scale.
- Coordinating context across multiple projects and teams.

The AI Multiplier:

This is the traditional management layer, made exponentially more powerful by agentic orchestration.

Altitude 2: The Hands-on Layer

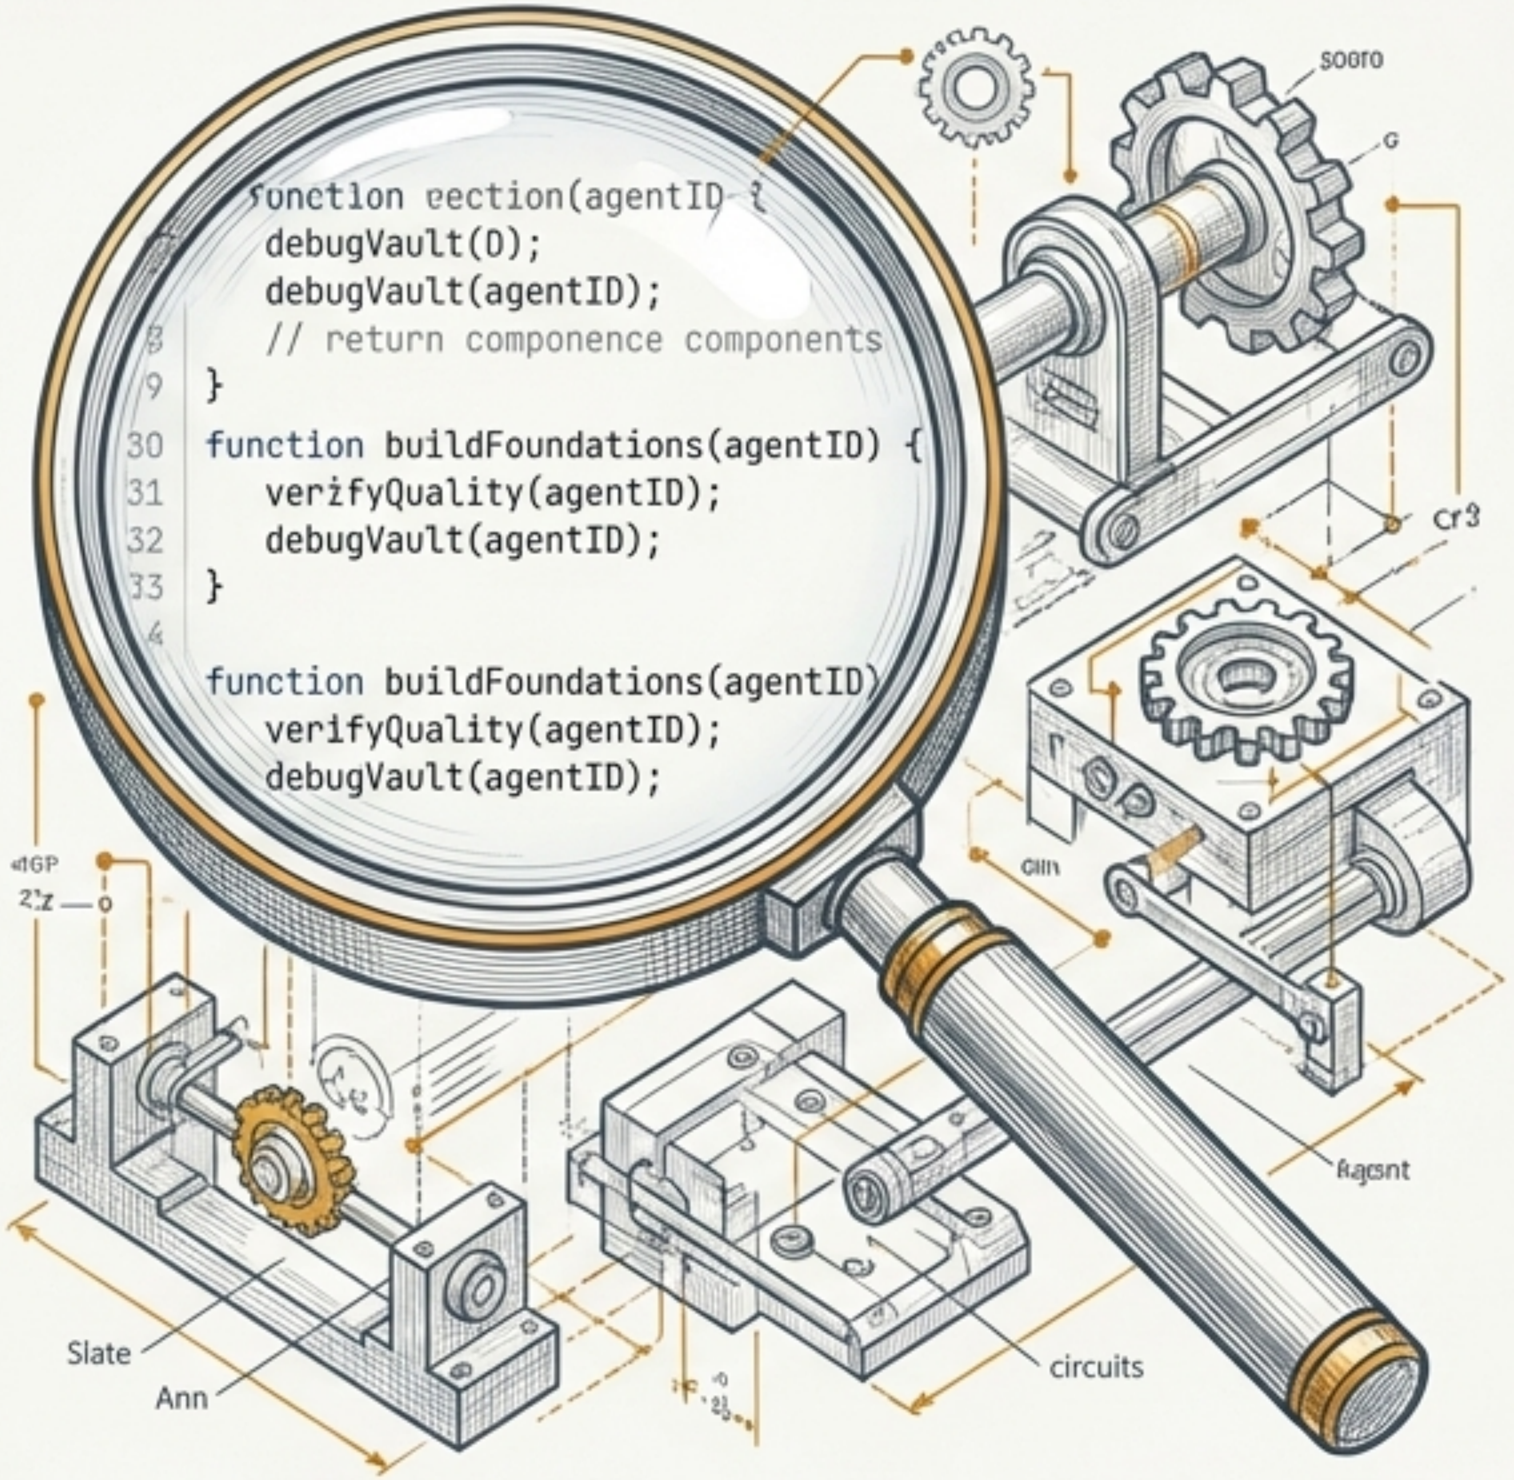

Focus: Grounded truth and technical verification.

Key Activities:

- Going deep into domain expertise.
- Reading the actual work and verifying agent-generated quality.
- Debugging vaults and building foundational components directly.
- Keeping practitioner skills razor-sharp.

The AI Multiplier: This is the traditional practitioner layer, made economically scalable by AI handling the boilerplate volume.

Proof of Concept: The Modern Executive

[redacted] years of development, [redacted] years of security, [redacted] years of cloud experience.

The Strategic CISO

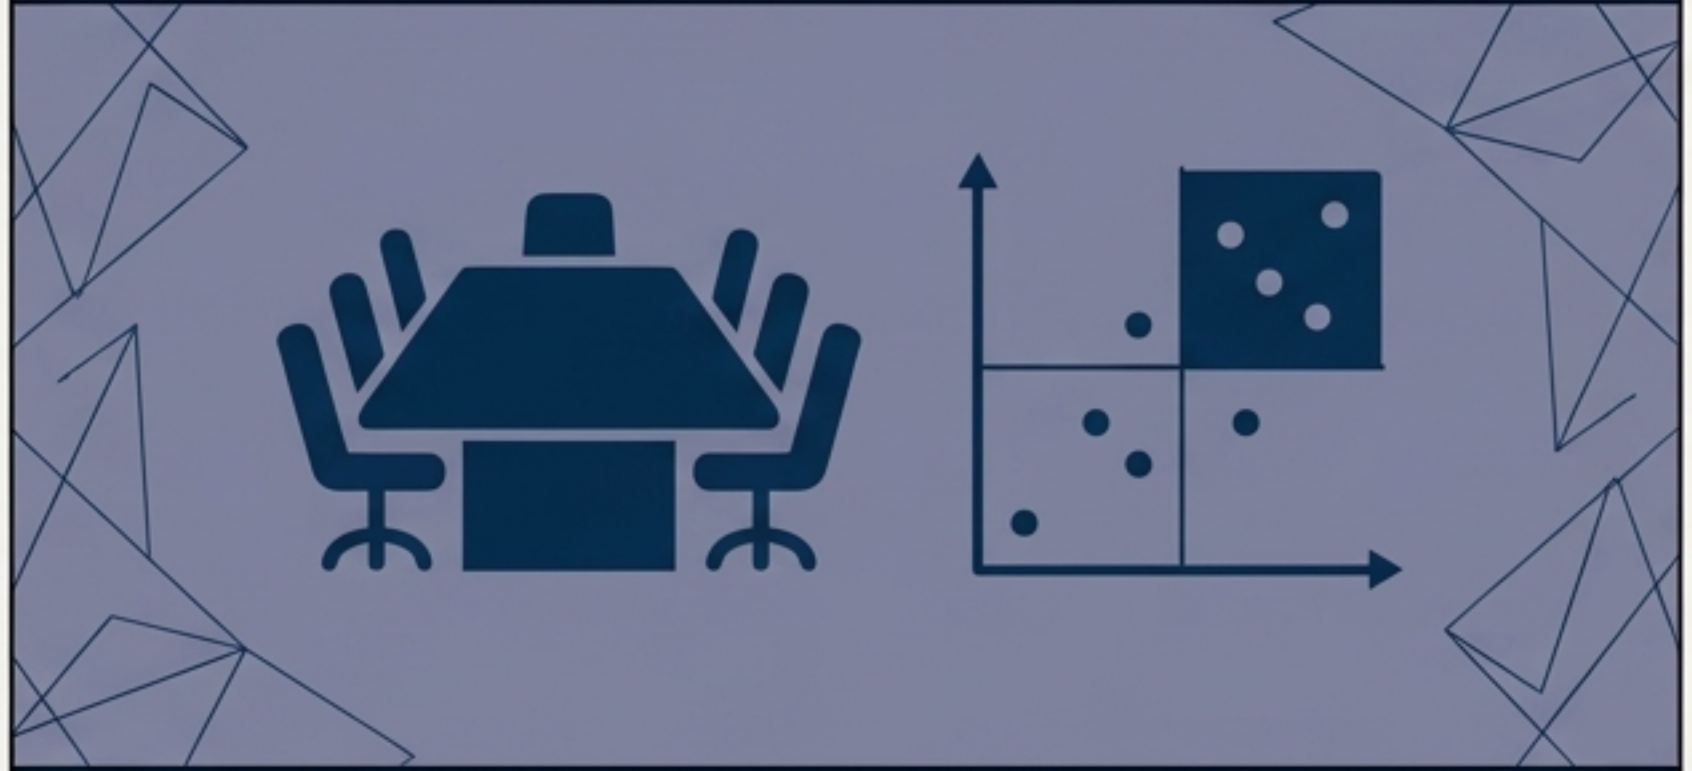

Operating at Altitude 1: Discussing board-level strategy, risk appetite, and programme maturity.

The Hands-on CISO

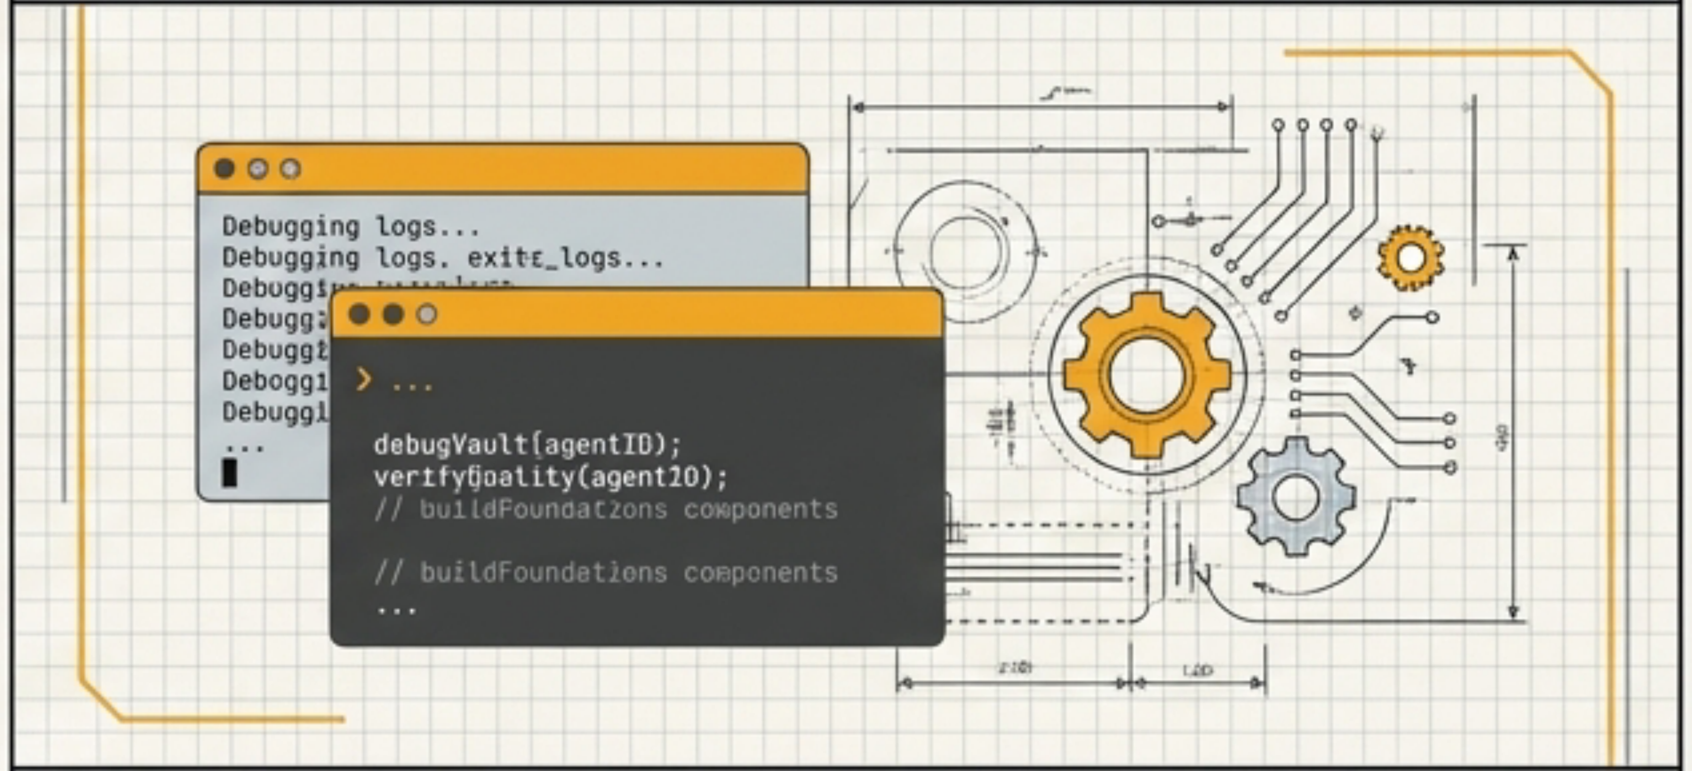

Operating at Altitude 2: Reading code, reviewing architectures, writing tooling.

The Synergy: Hands-on experience brings a qualitative difference to strategic reviews that a junior engineer cannot provide. One without the other produces excellent solutions to the wrong problems.

Re-framing the Junior Problem

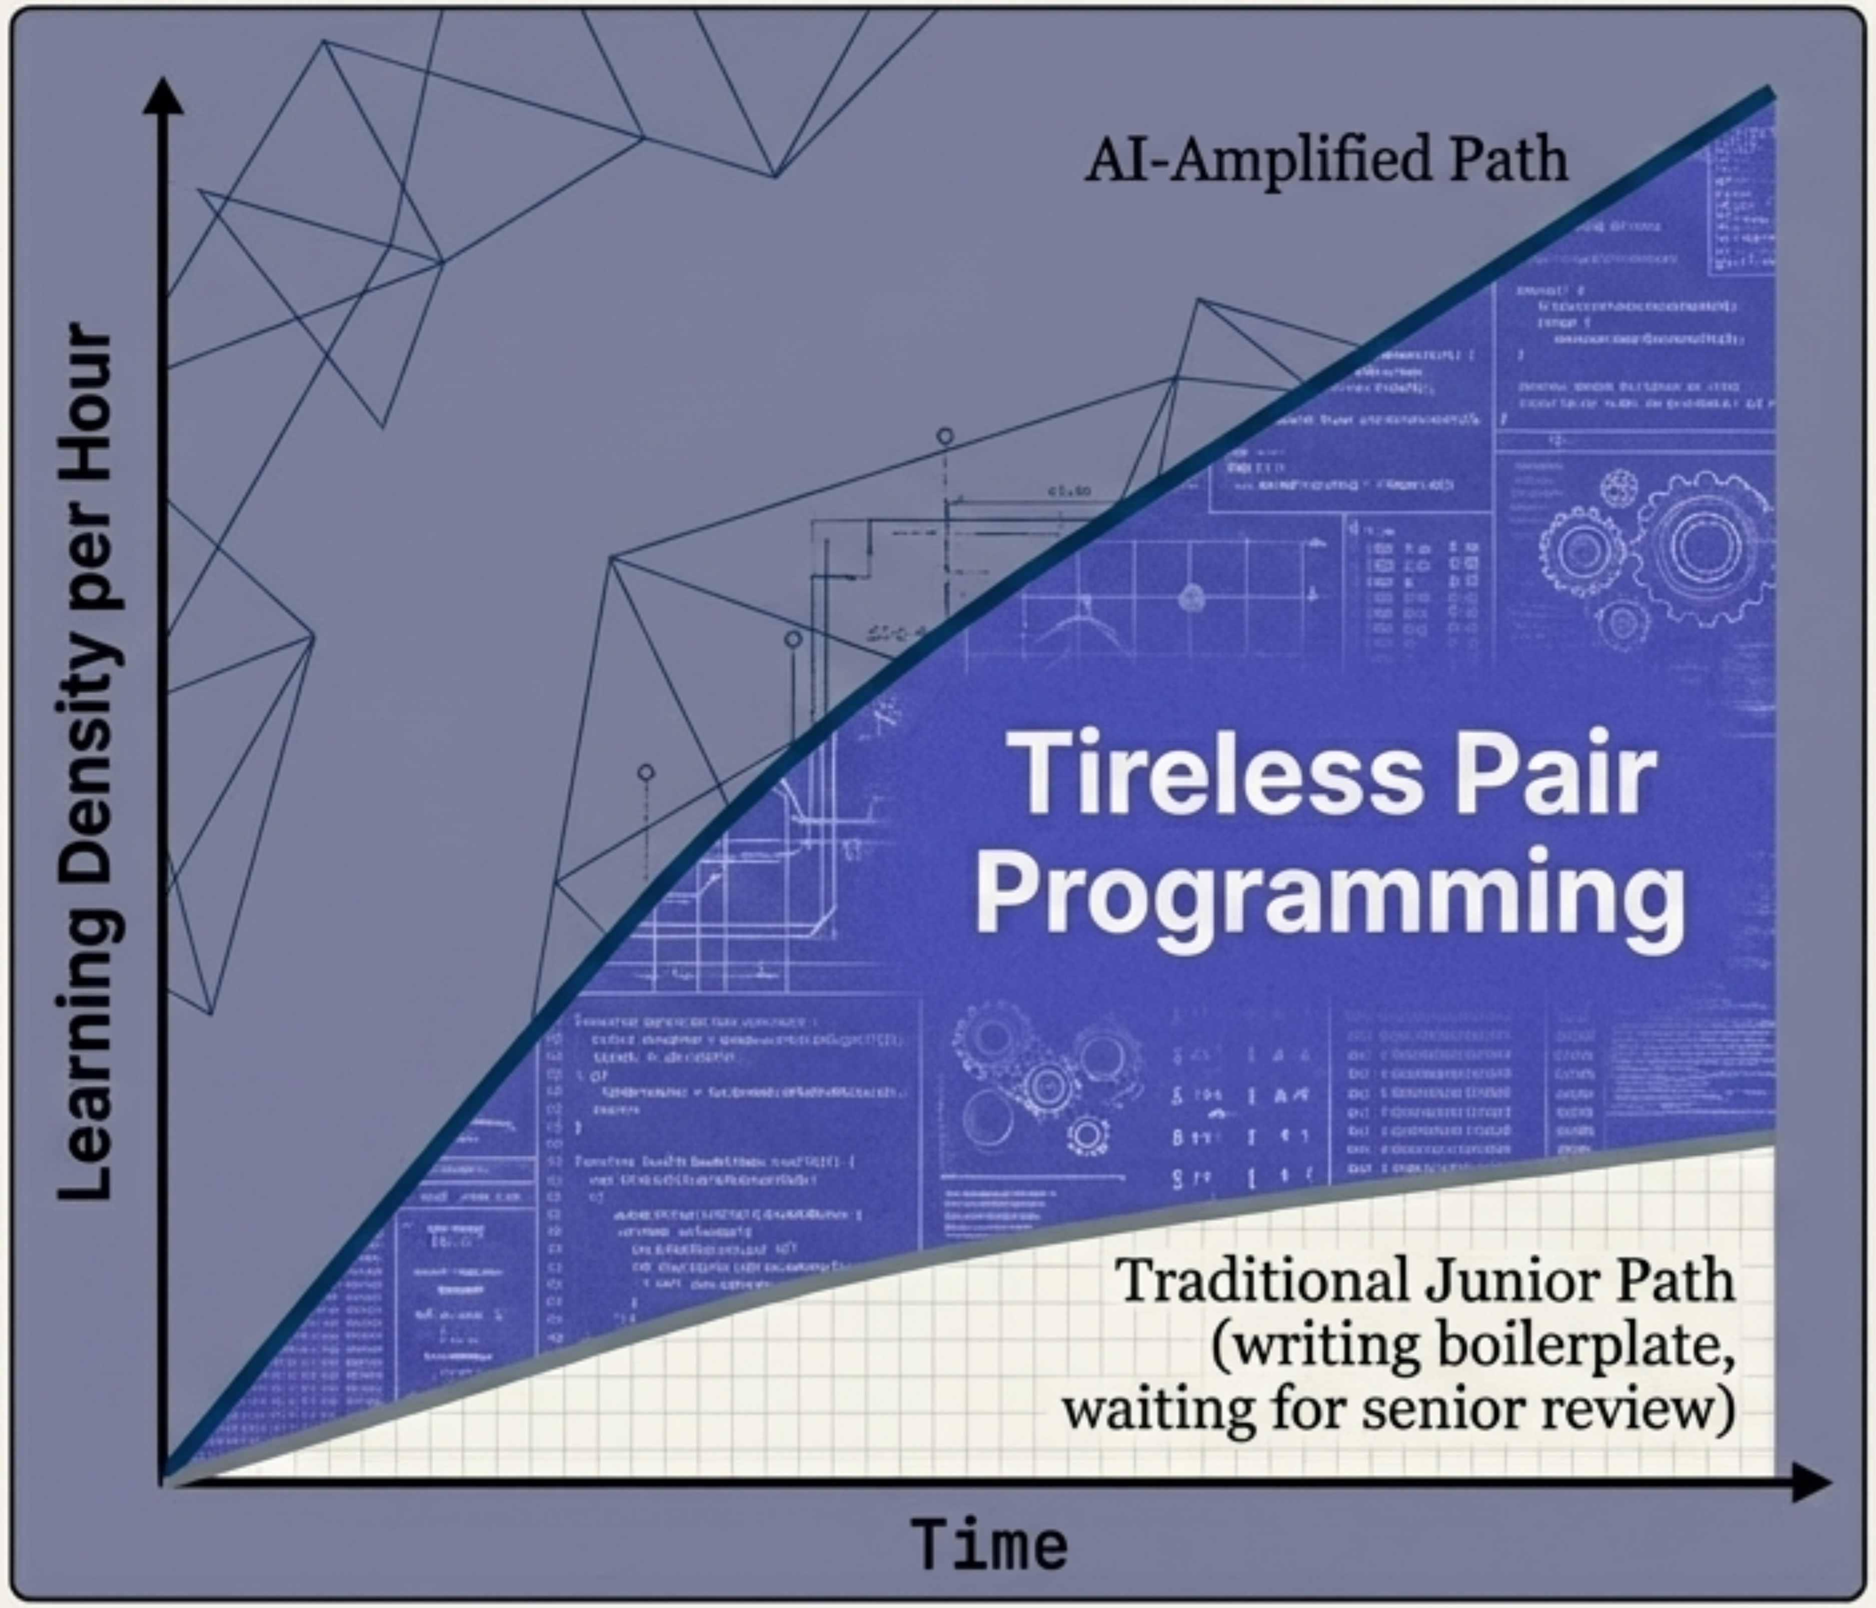

The Question:
If agents do the entry-level coding, how does anyone become senior?

The Reality:
You only learn by doing, but 'doing' is now richer and multi-dimensional.

The AI-Amplified Junior:
- Reads agent-generated code to understand architectural decisions.
- Learns design patterns at massive scale.
- Works alongside an infinitely patient AI tutor that explains every line and suggests alternatives instantly.

Scaling the Unscalable: The Librarian

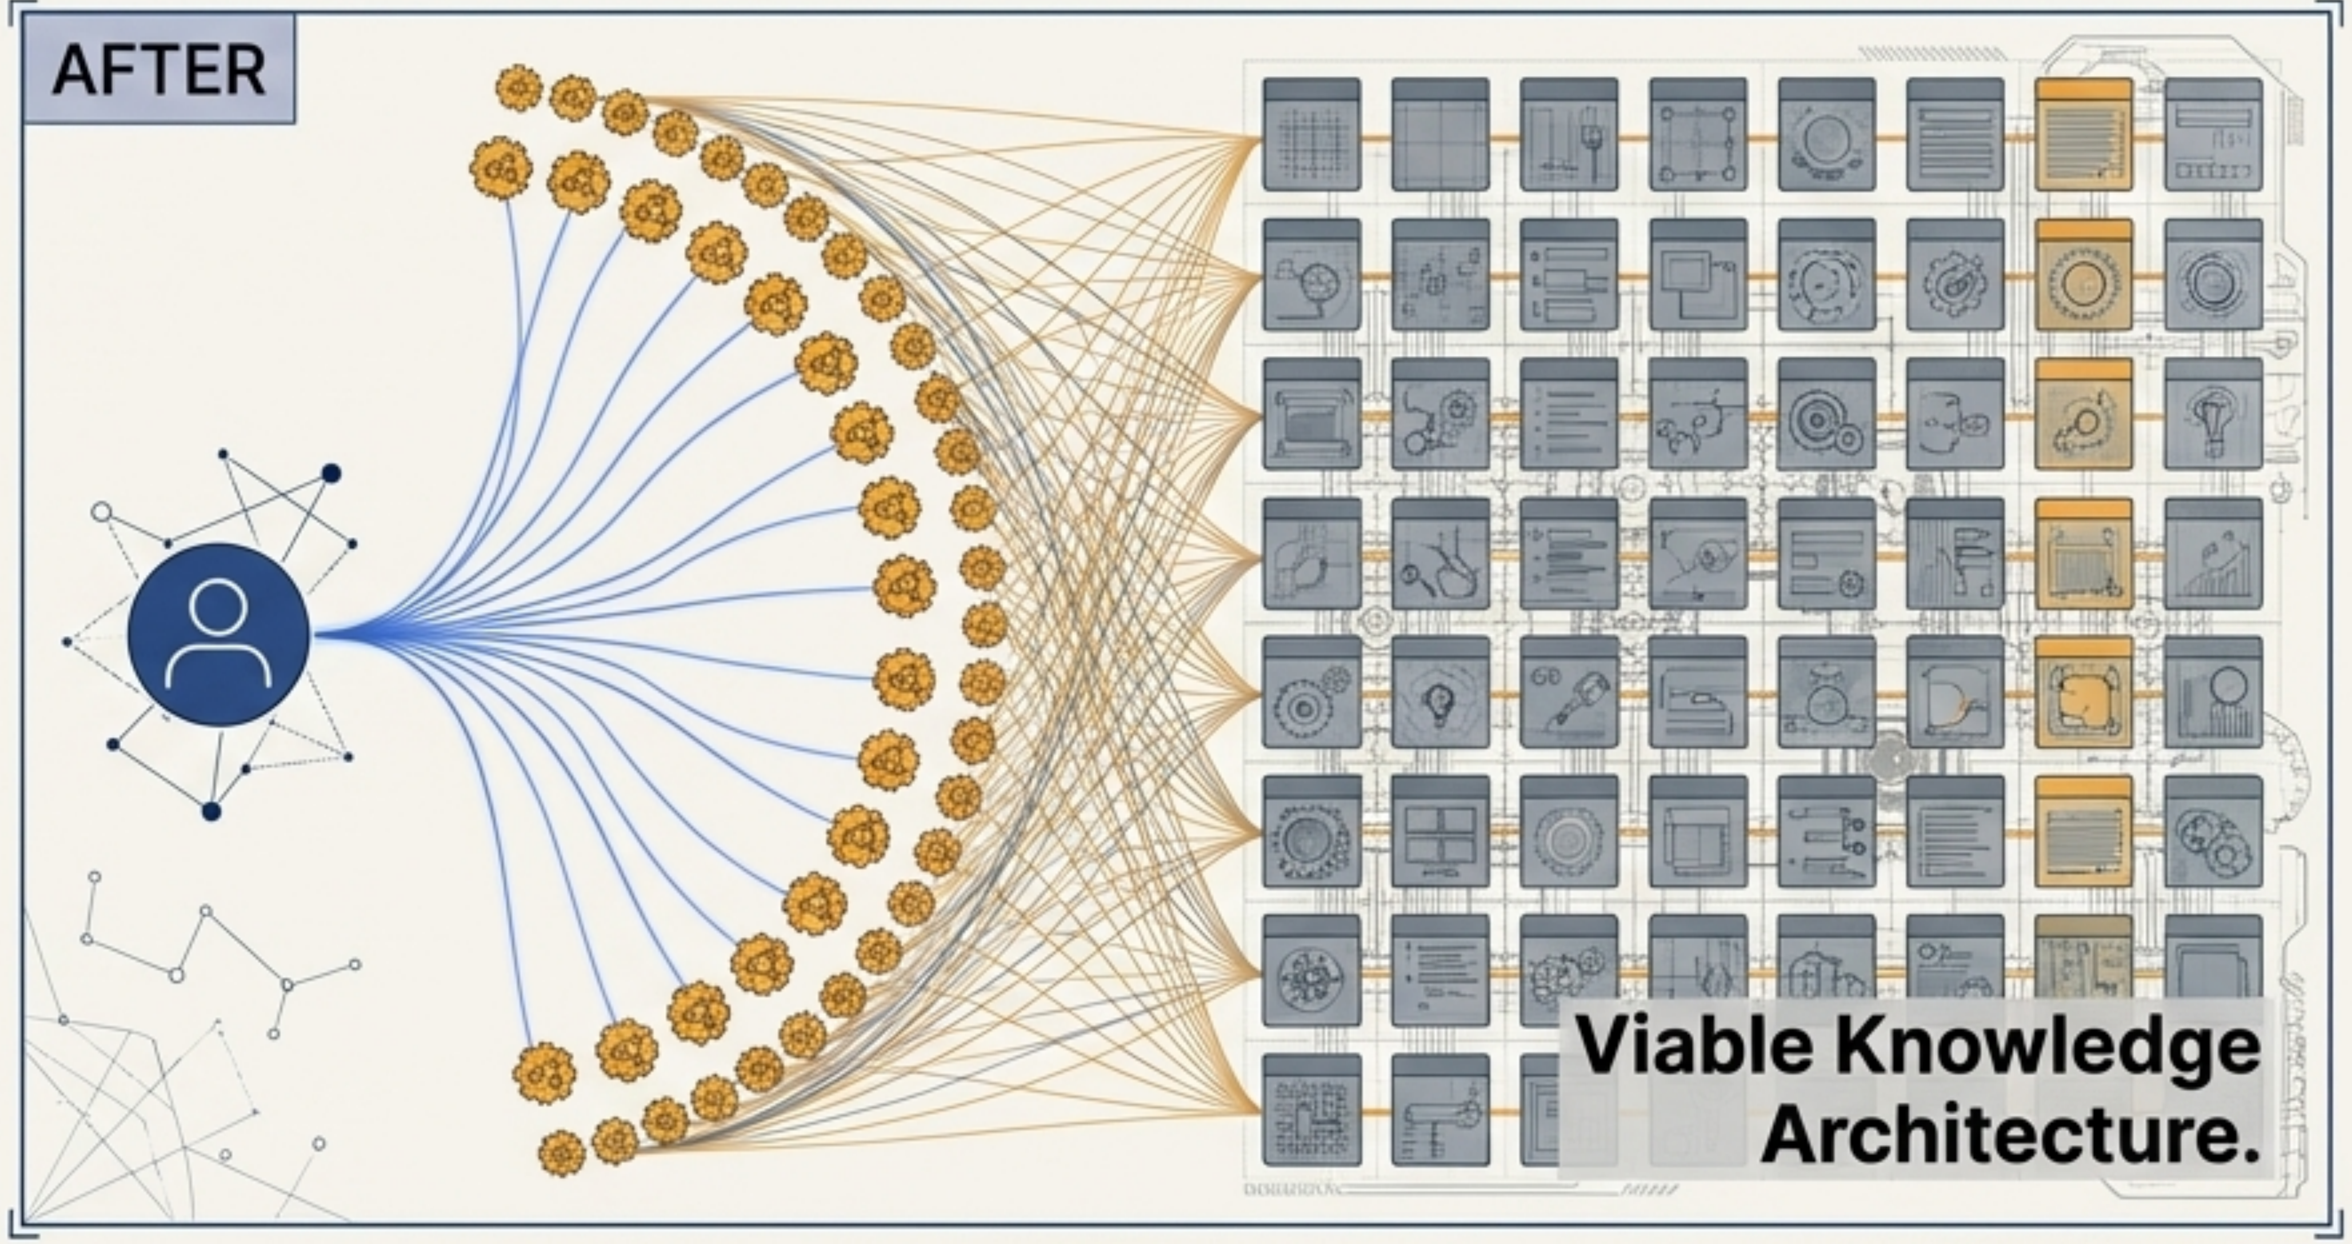

The Historical Problem

No company could afford a dedicated knowledge librarian per project.

Result: tribal knowledge and quadratic communication overhead.

The Two Altitudes Solution

Strategy: One human expert designs the cross-project taxonomy and knowledge architecture for 5 to 10 startups.

Hands-on: 50–100 Librarian Agents execute the cataloguing, track decisions, and maintain context across projects.

The Result

A historically unviable role becomes a transformative organisational asset.

The Economics of Depth

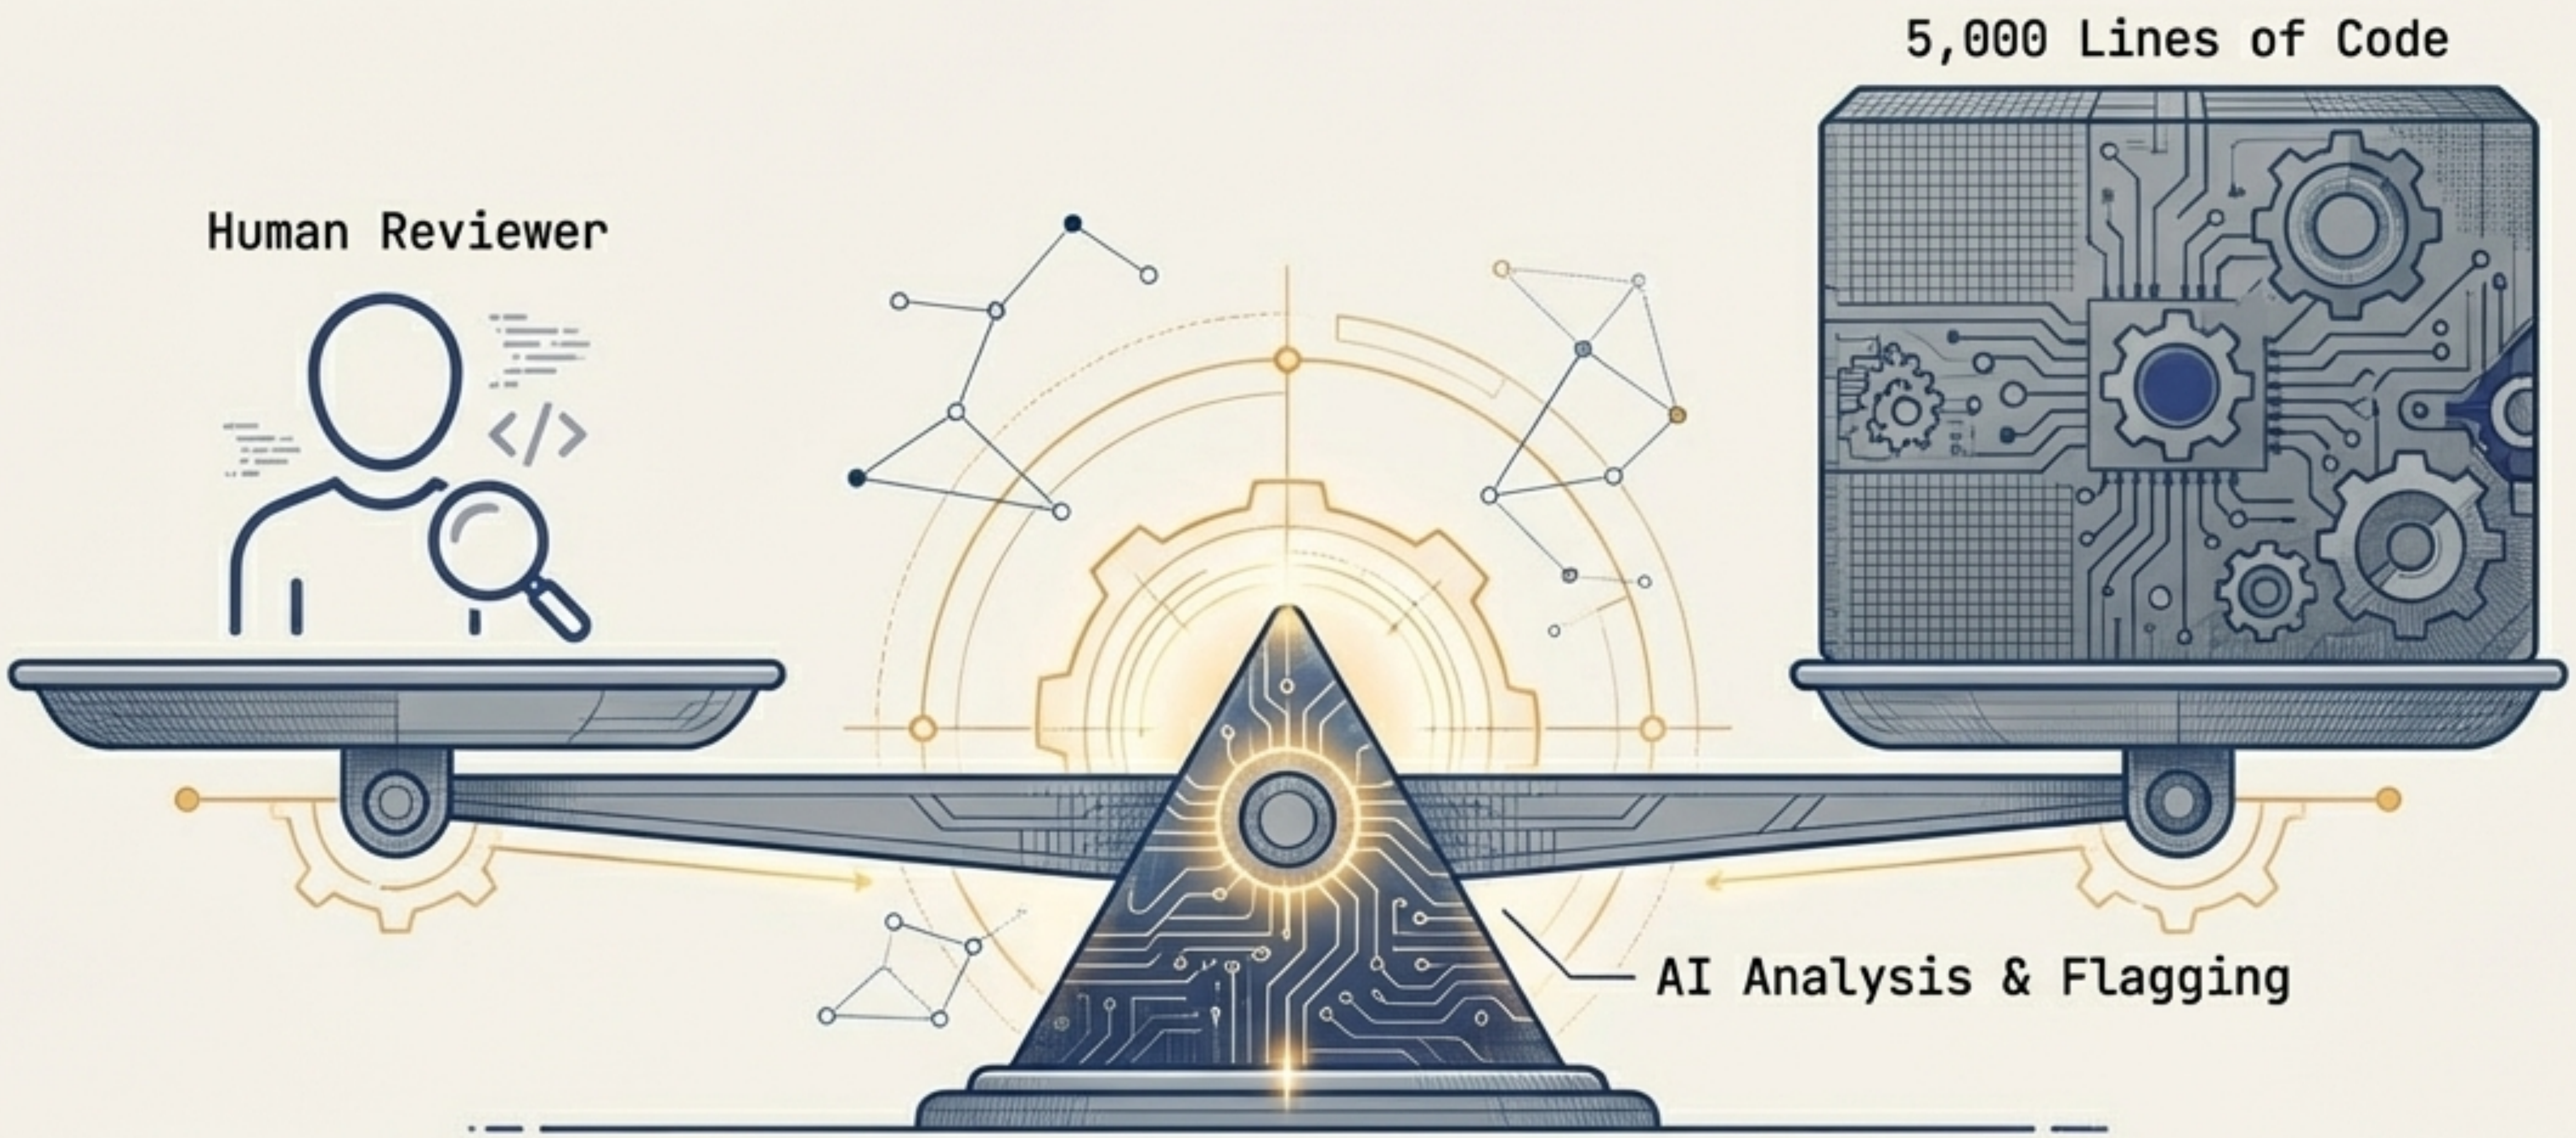

The Linux Standard

The Linux kernel is reliable precisely because it is not purely abstracted. Maintainers read foundational C code. They work at every level simultaneously.

The Historical Barrier

Maintaining people who operate at both architectural and deep implementation levels used to be prohibitively expensive.

The New Economics

With AI agents handling initial analysis and flagging anomalies, an expert who historically reviewed 50 lines of code per hour can now deeply verify 5,000 lines. Depth scales.

Diagnostic Matrix: The Tale of Two Organisations

	The Manager-Only Org	The Two Altitudes Org
Leadership Profile	Confident but technically detached. TECH_EETACH	Strategically effective & technically sharp. STRAT_SHARP
System Architecture	Fragile, bloated systems built on unread dependencies.	Reliable systems built on fully understood, verified code.
Skill Trajectory	Rapid capability atrophy; critical failure under pressure. CAP_ATROPHY_RATE: HIGH	Continuous, high-density skill sharpening. SKILL_DENSITY: INCREASING
AI Relationship	ACCOUNTABILITY: TRANSFER Total accountability transfer. AI	PRACTITIONER: AMPLIFIED Genuine practitioner amplification. AI

The Future Belongs to Those Who Do Both

The future of knowledge work is not choosing between doing and managing. It requires operating at both altitudes, in the same person, at different times during the same day. AI makes both layers more effective than either could ever be alone.