

Phase E Performance Analysis: HTML → MGraph Pipeline

Bottleneck Identification & Optimisation Strategy

DATE: JAN 2026

STATUS: ANALYSIS COMPLETE

SOURCE: Perf__Phase_E__Conversion



Three Critical Findings Define the Pipeline Performance

The Bottleneck



97%

of processing time is consumed by 'Dict->MGraph' conversion.
Optimisation elsewhere is ineffective.

The Win



2.3x

Speedup confirmed by enabling 'fast_create'. Immediate implementation recommended.

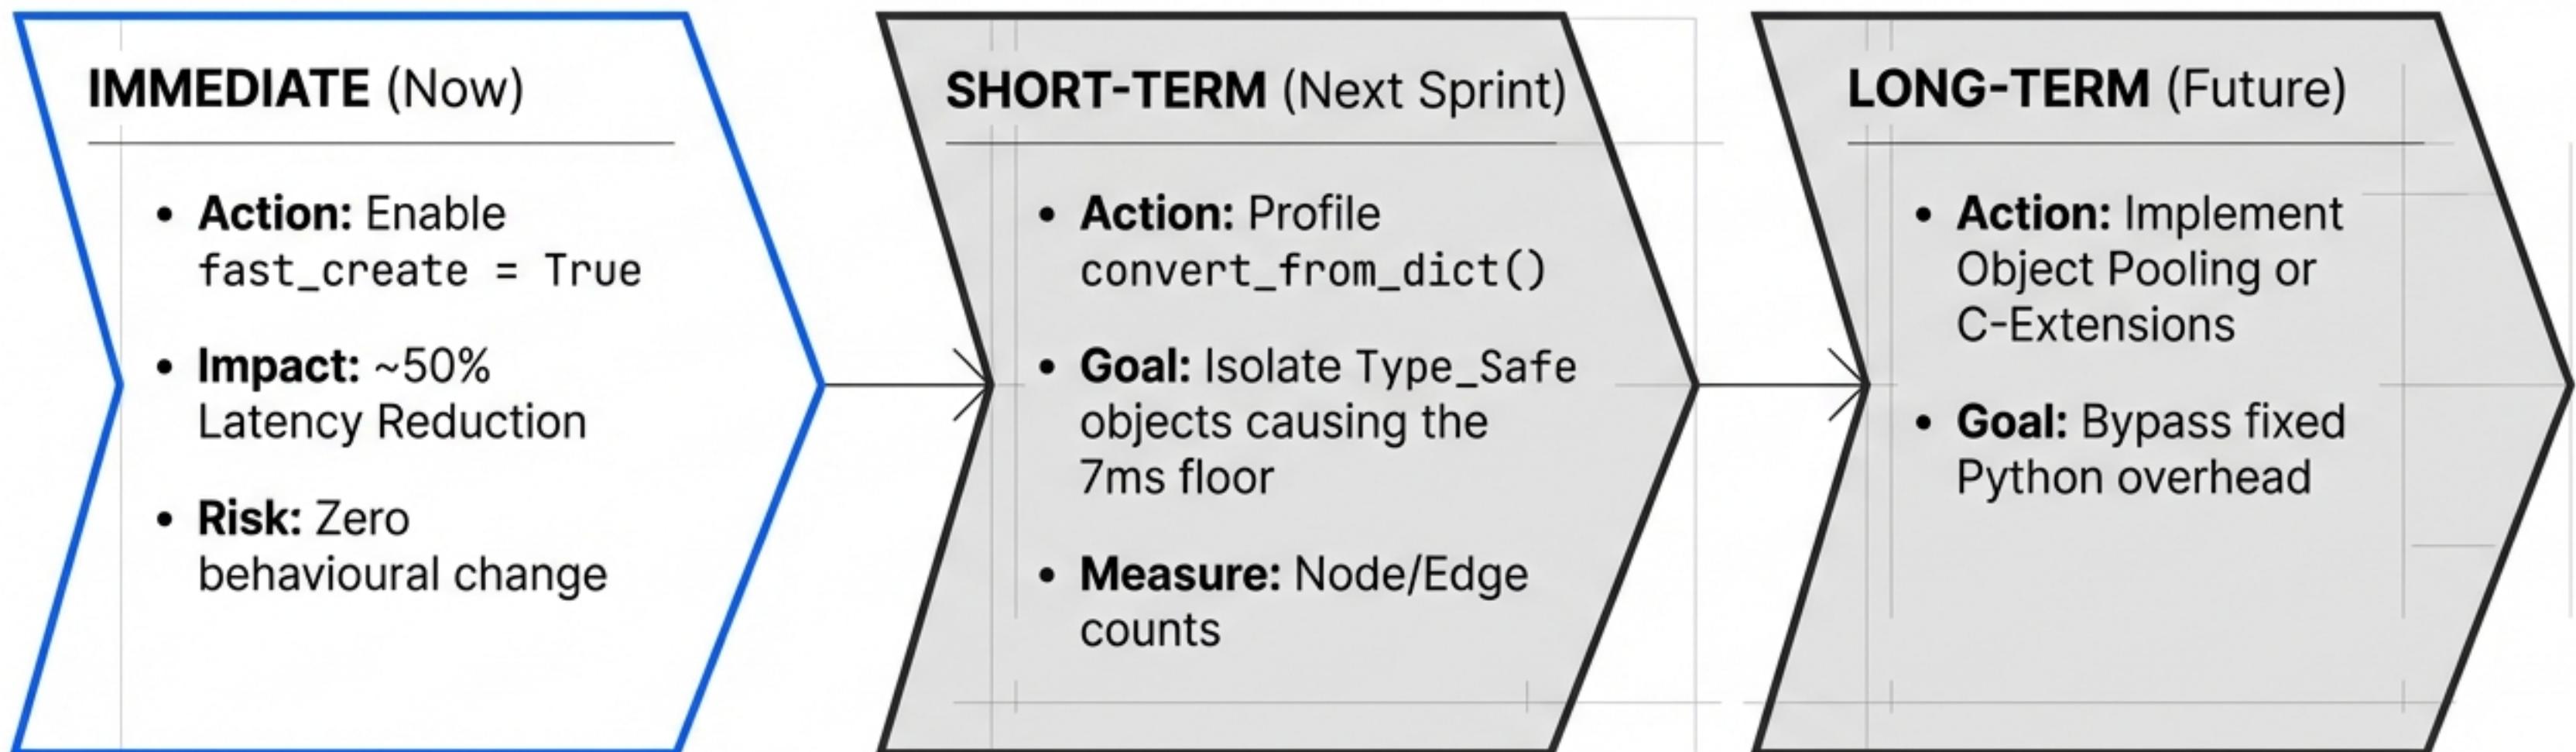
The Mystery



7ms

Fixed architectural baseline exists regardless of input size.
Suggests heavy object infrastructure.

Strategic Roadmap: From Immediate Gains to Architectural Refactoring



Methodology & Integrity

Protocol

Metrics derived from 'measure_fast' protocol using **87 Fibonacci** iterations to ensure stability and eliminate outliers.

Environment

Test machine under normal load. Results consistent across multiple runs.

Test Conditions



1 Node



10 Nodes

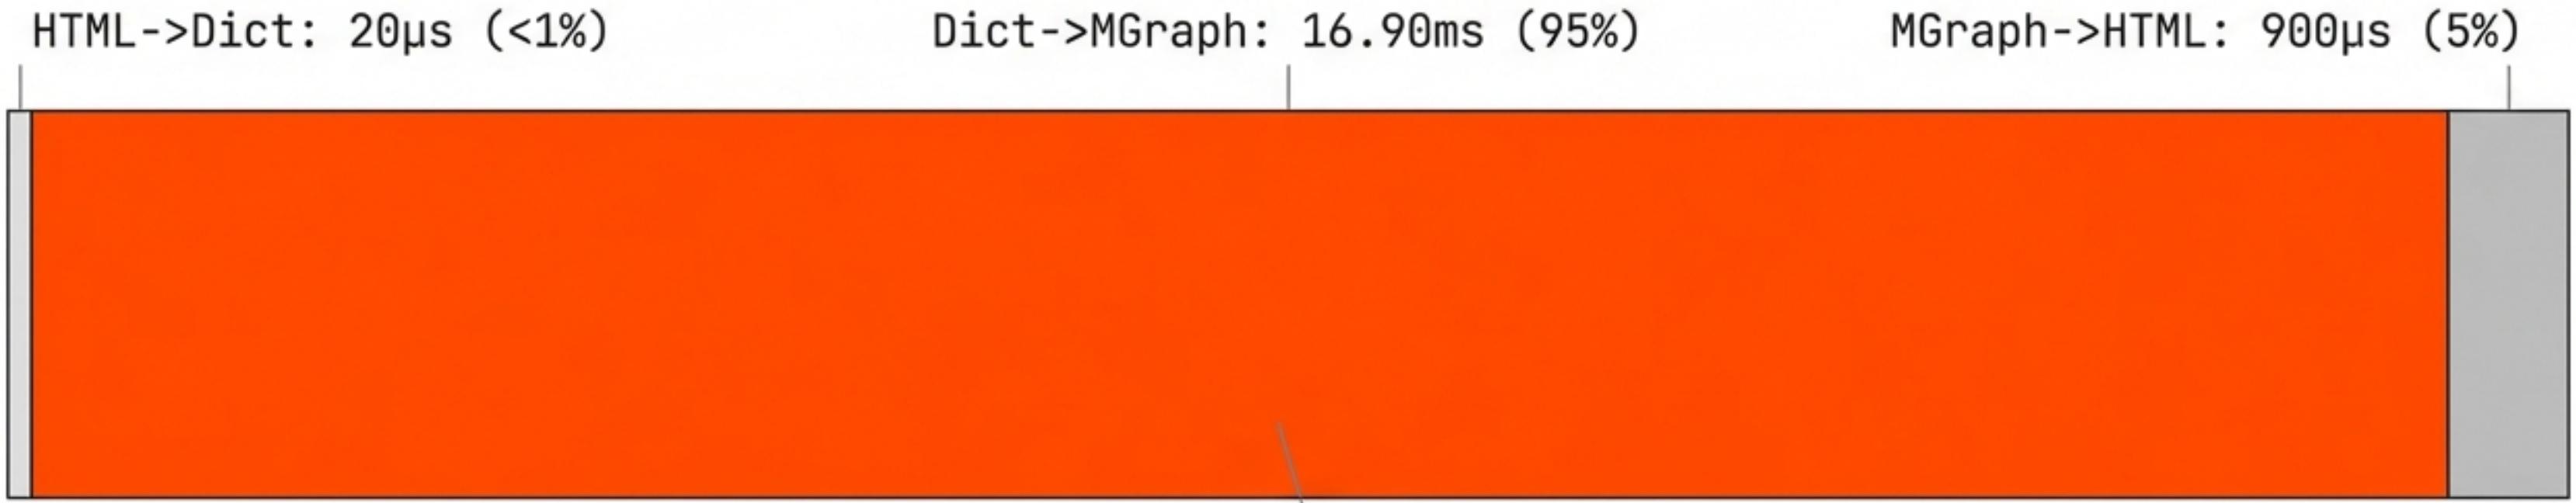


100 Nodes

Generated HTML files compared across:

- Default 'Type_Safe' Mode
- '**fast_create=True**' / '**skip_validation=True**'

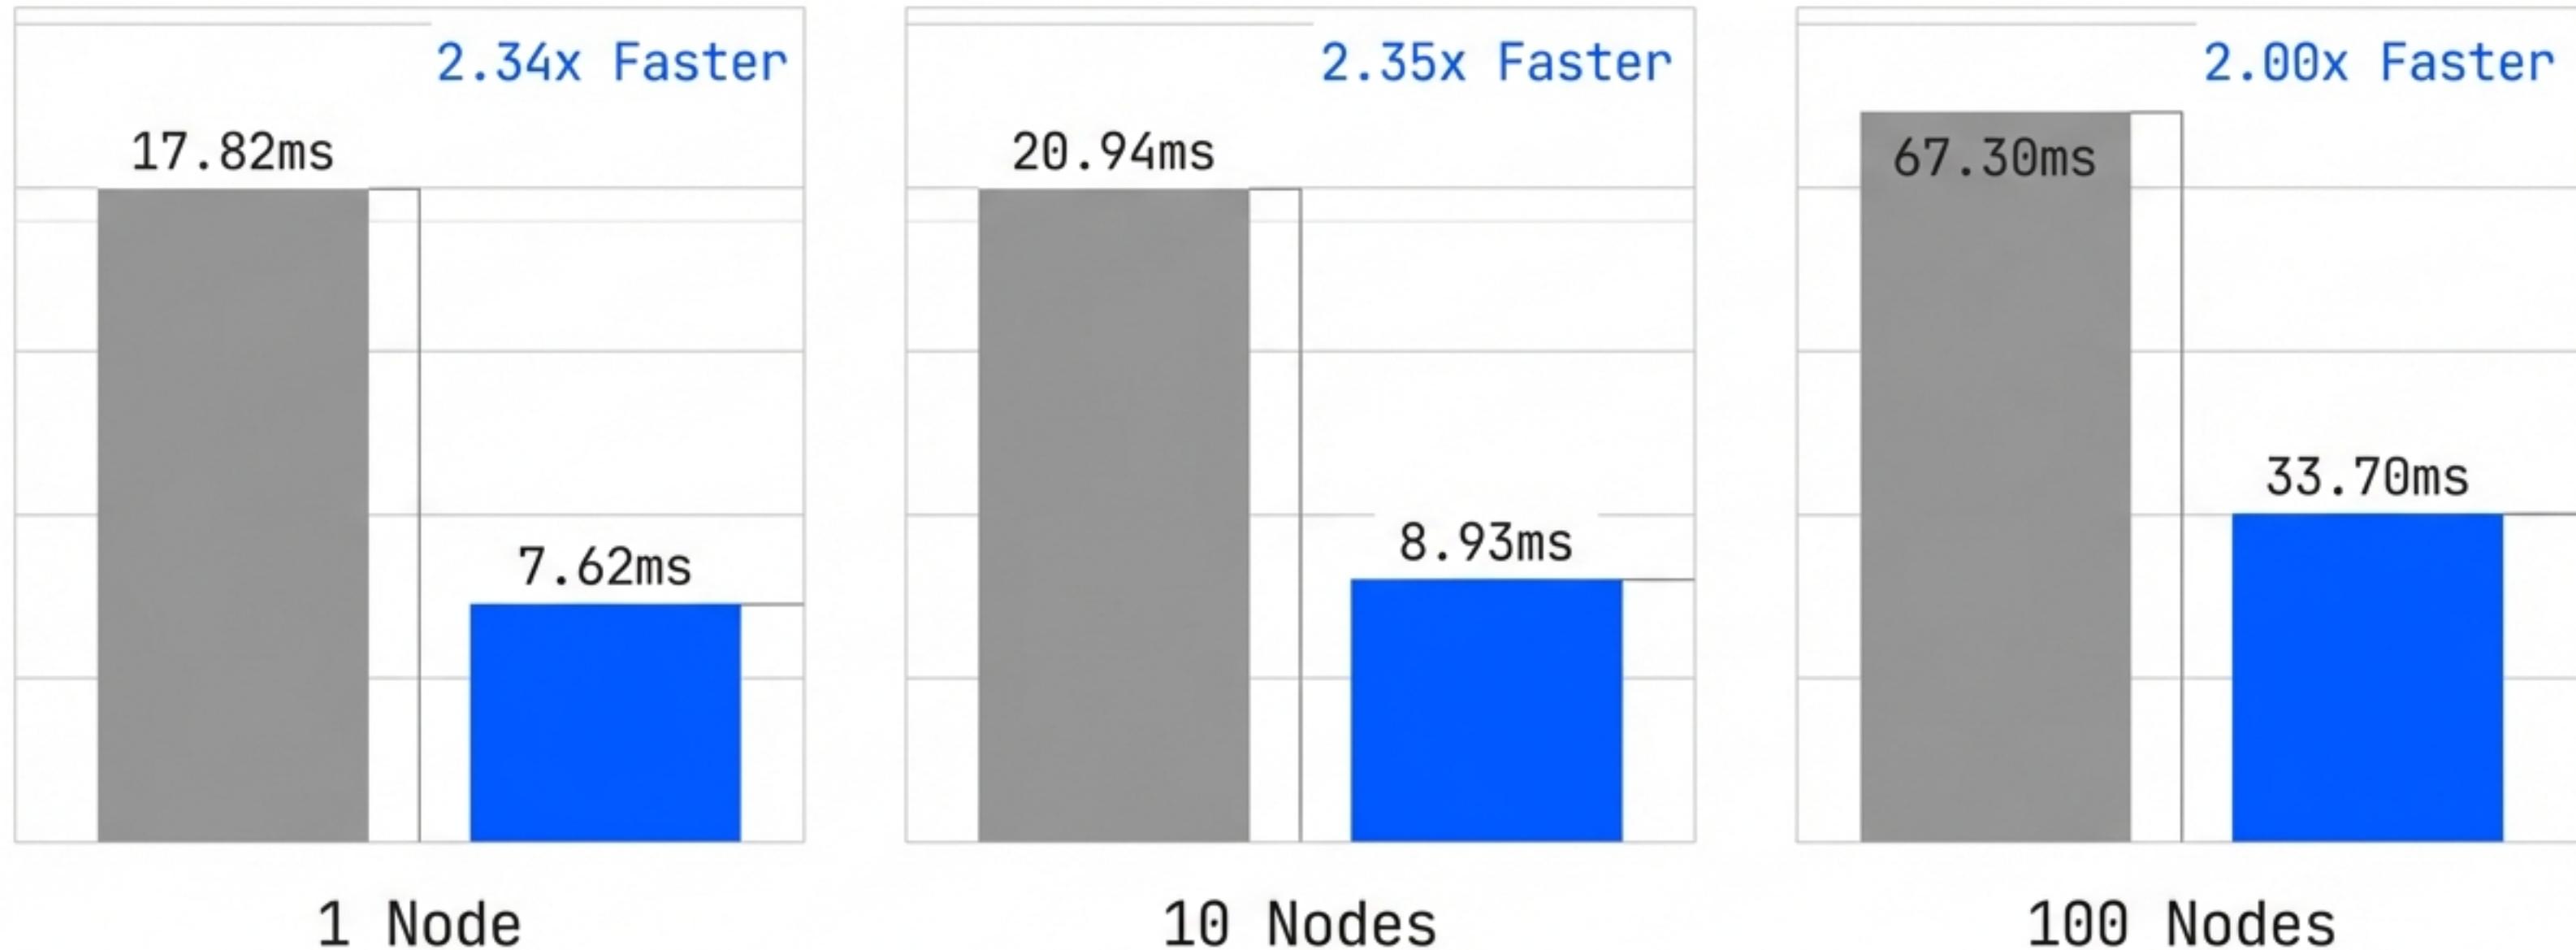
The ‘Dict->MGraph’ Stage Dominates Execution Time



CRITICAL FOCUS AREA

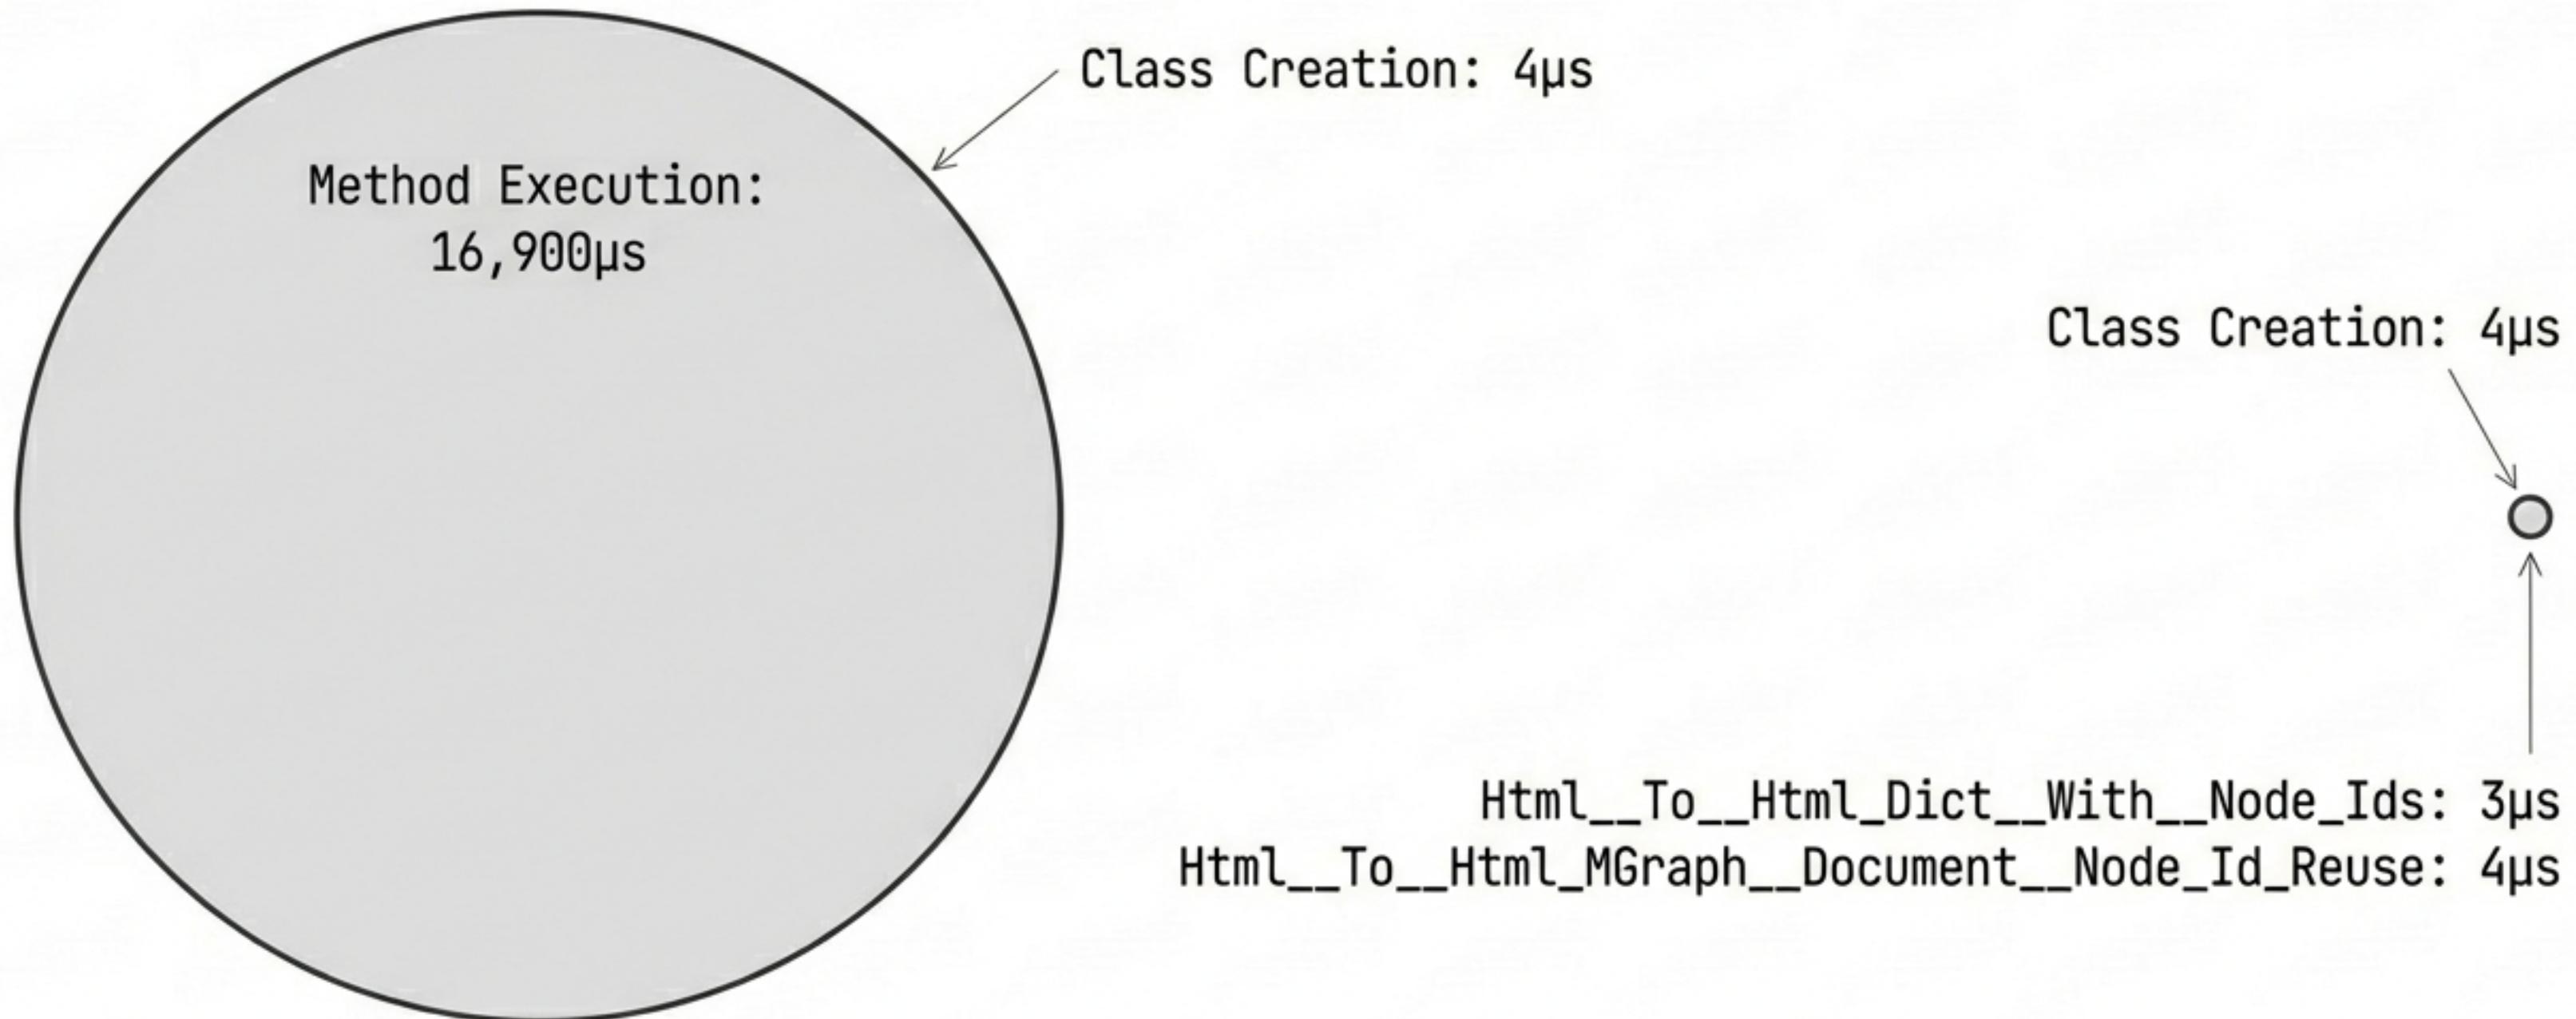
Optimization focus must be exclusive to the “convert_from_dict” method. This 16.90ms block dwarfs the input stage.

'fast_create' Delivers a Consistent 2x Speedup



The speedup is consistent (~50-57% reduction) across all scales.

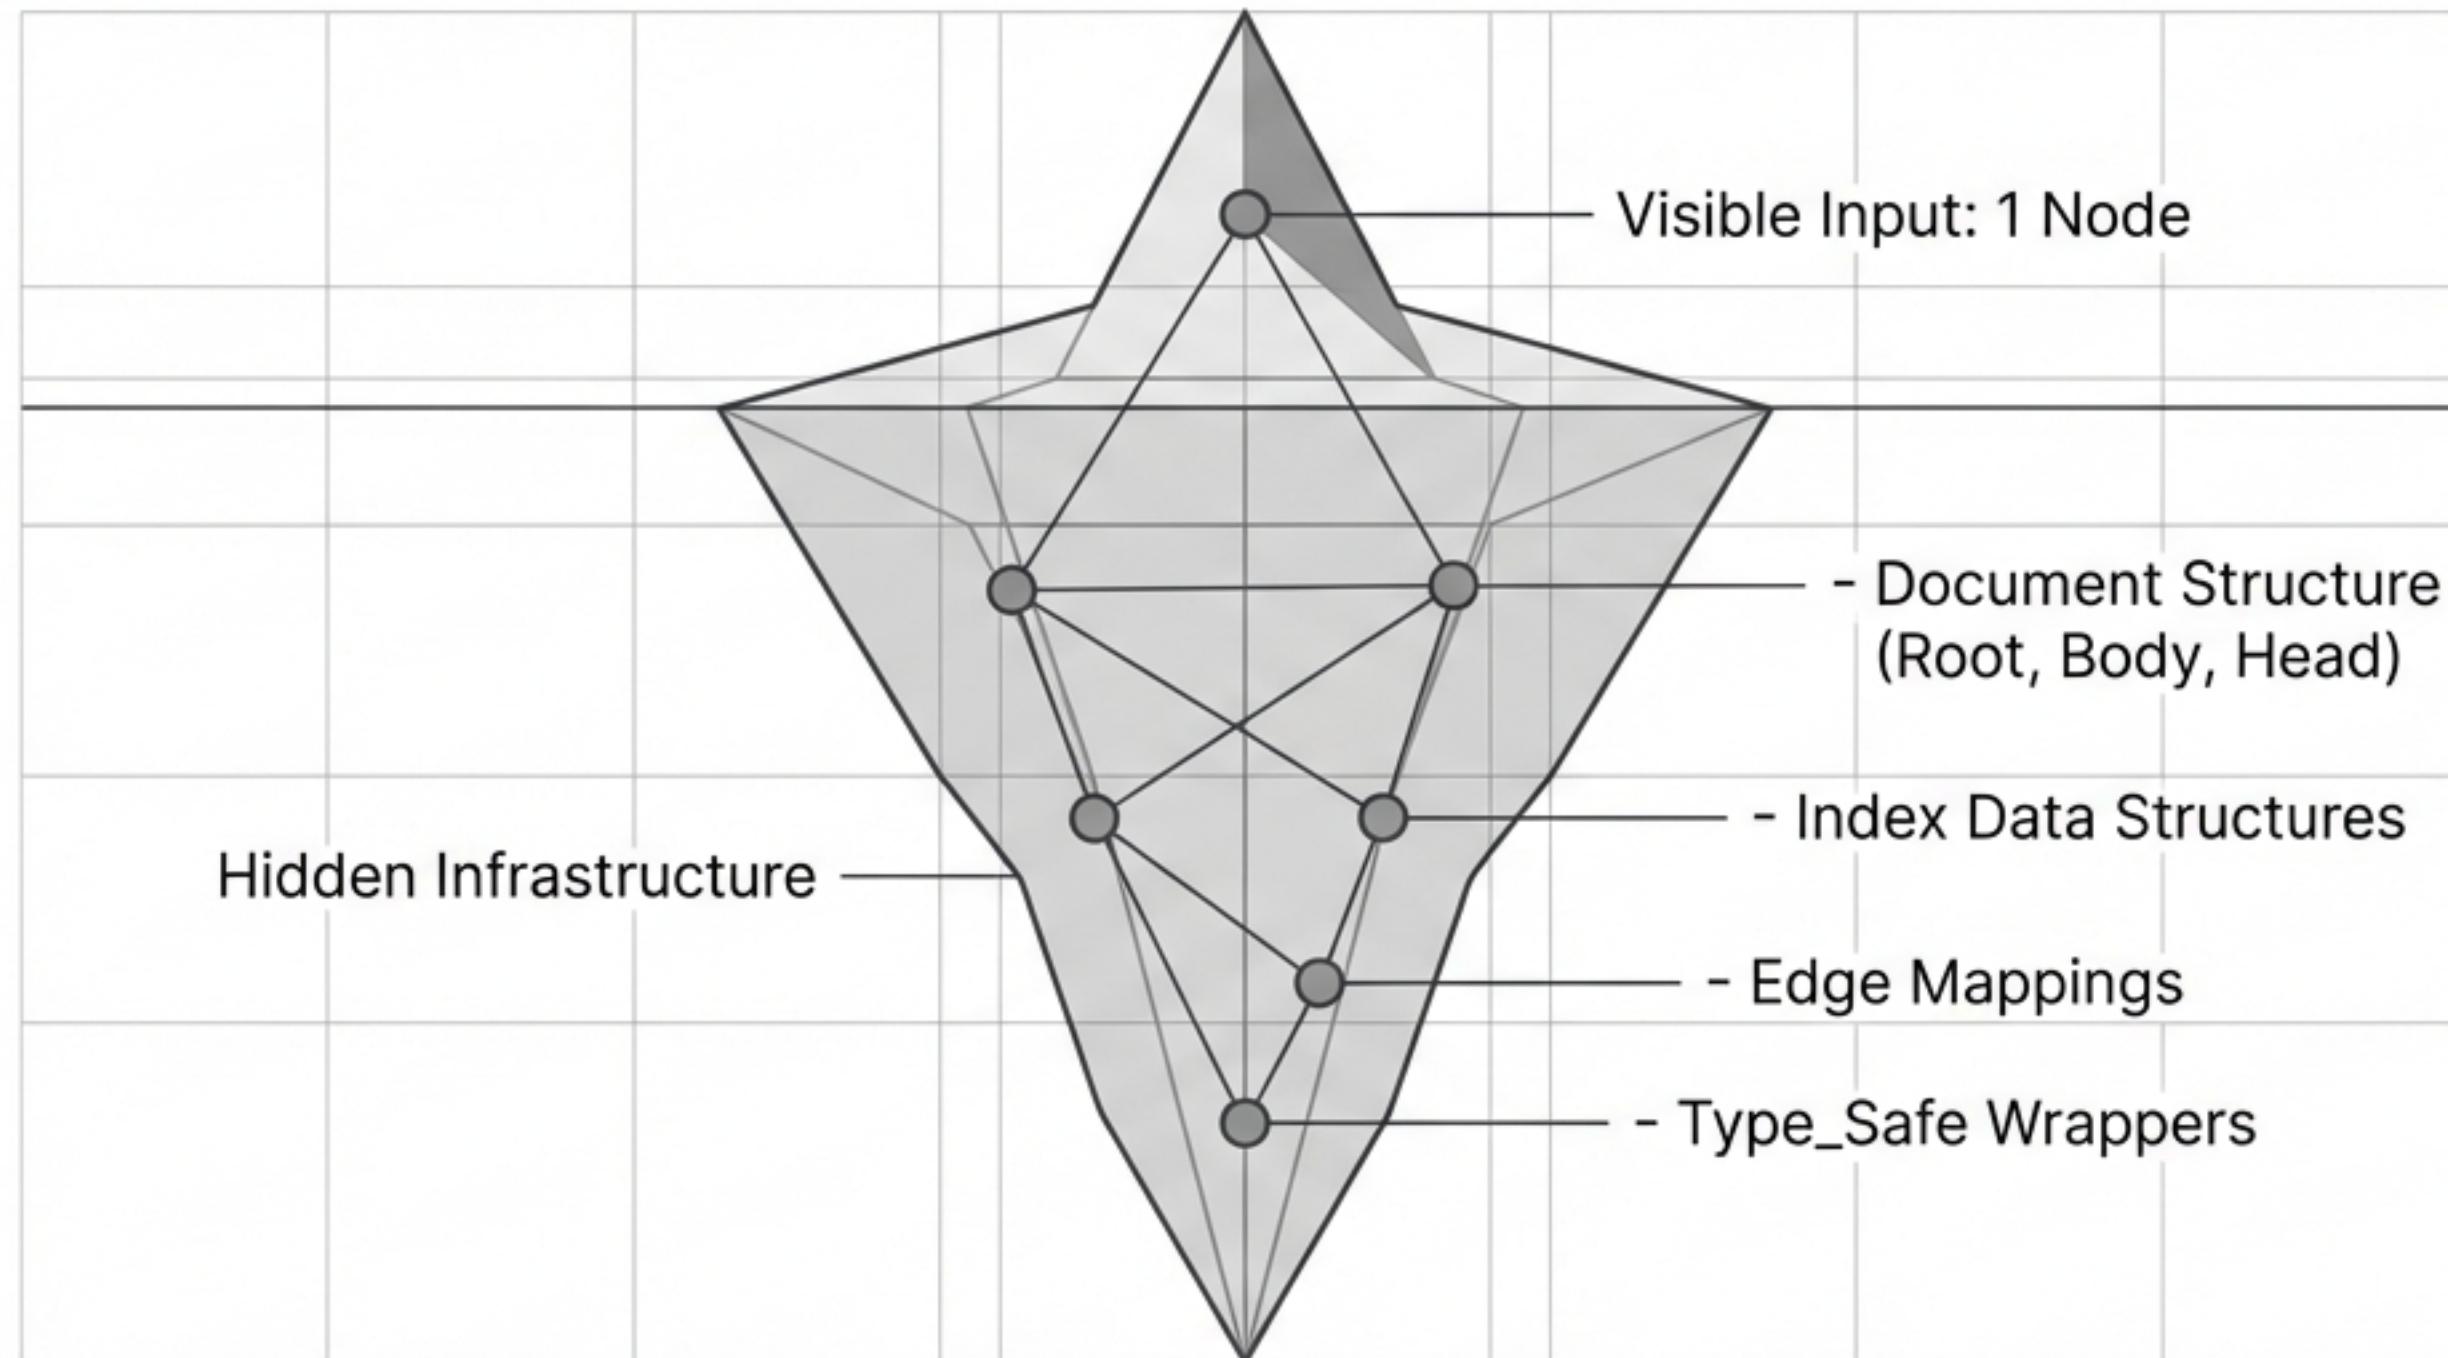
De-risking: Converter Instantiation is Negligible



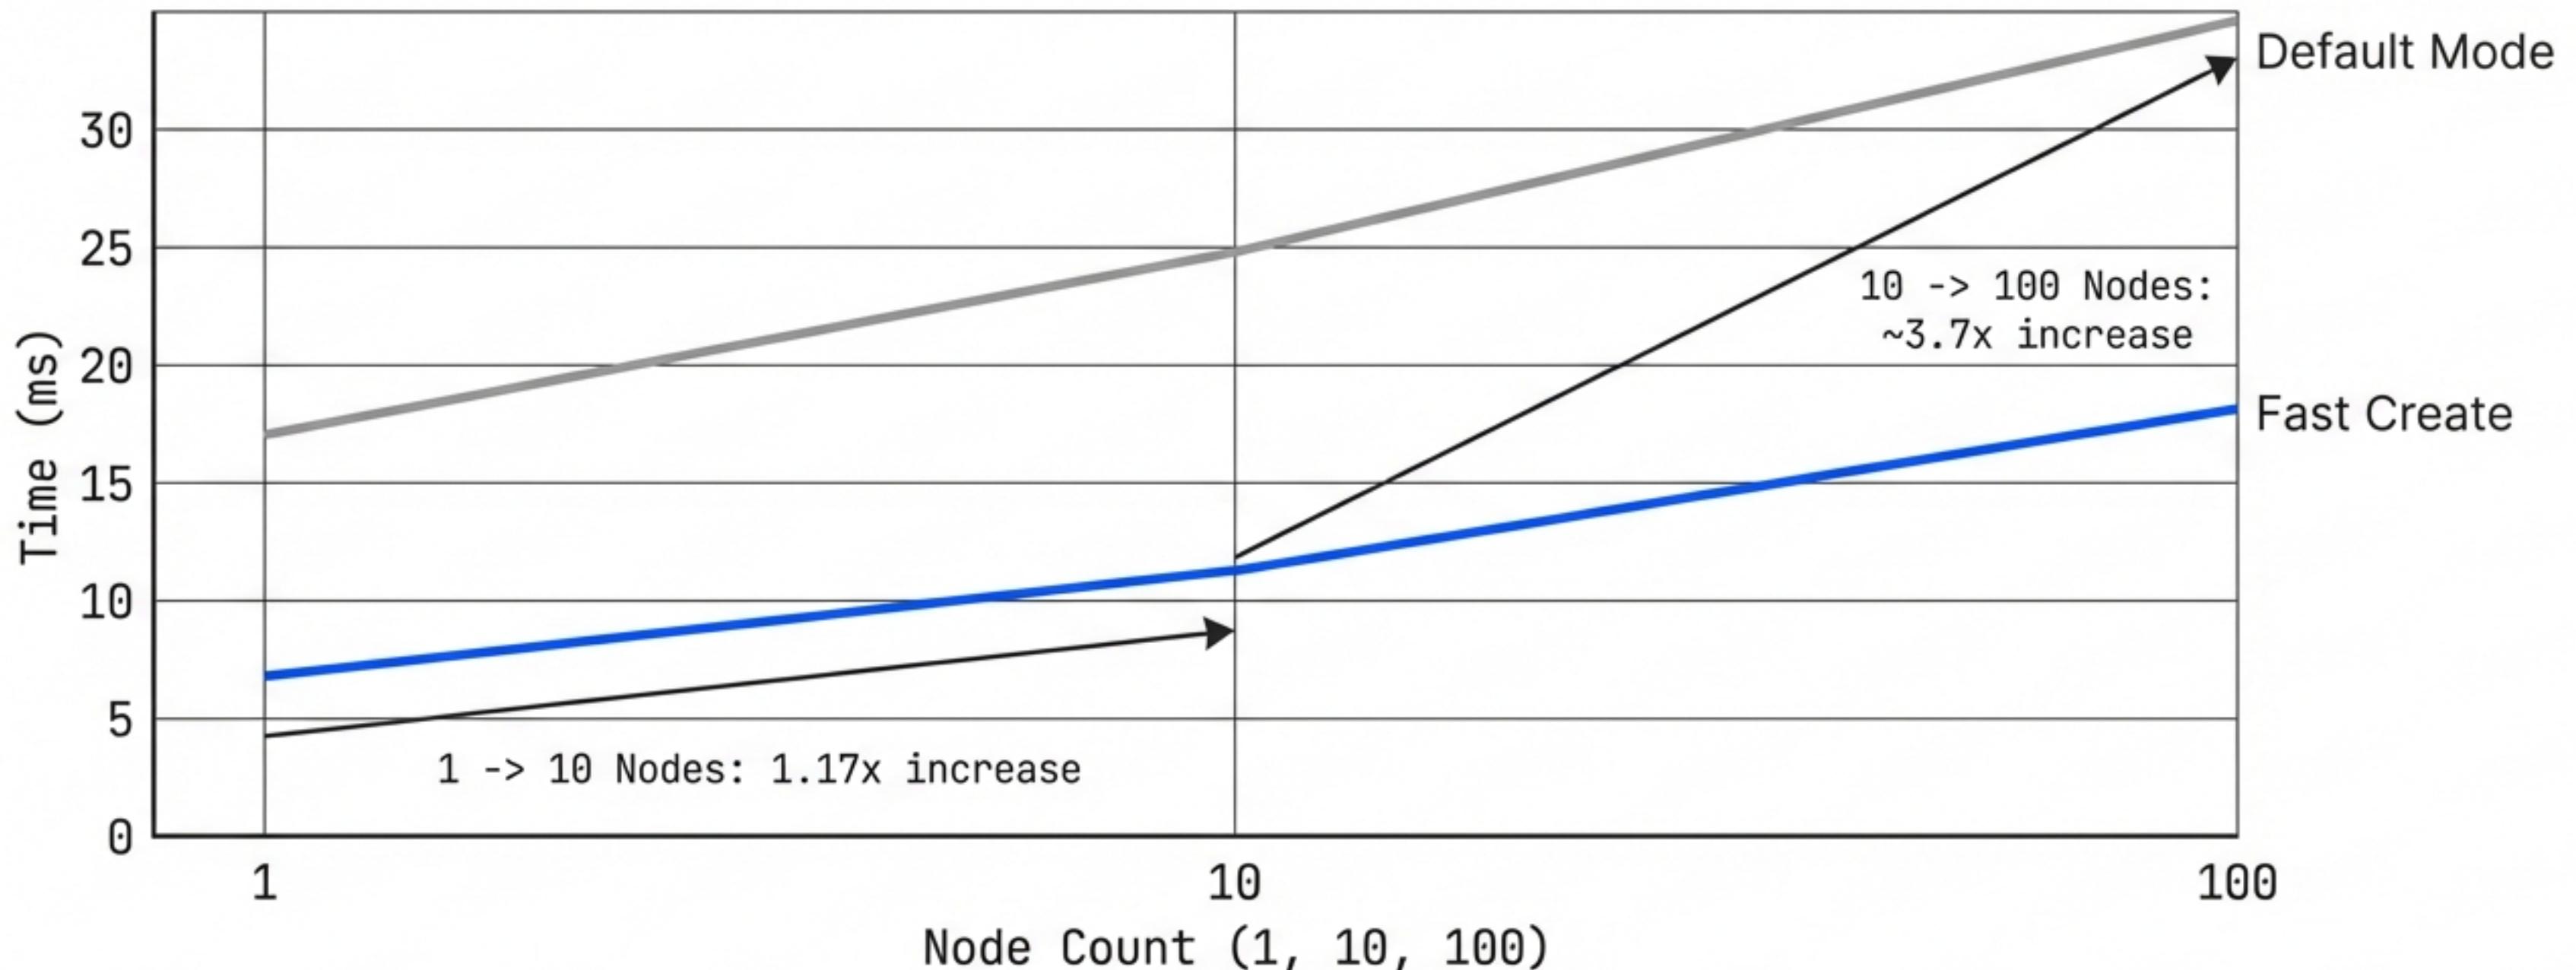
We do **NOT** need to cache converter instances. The cost is entirely within the `.convert()` methods.

The 7ms Baseline Mystery

Even a single <p>Hello World</p> takes 7.40ms. Why?

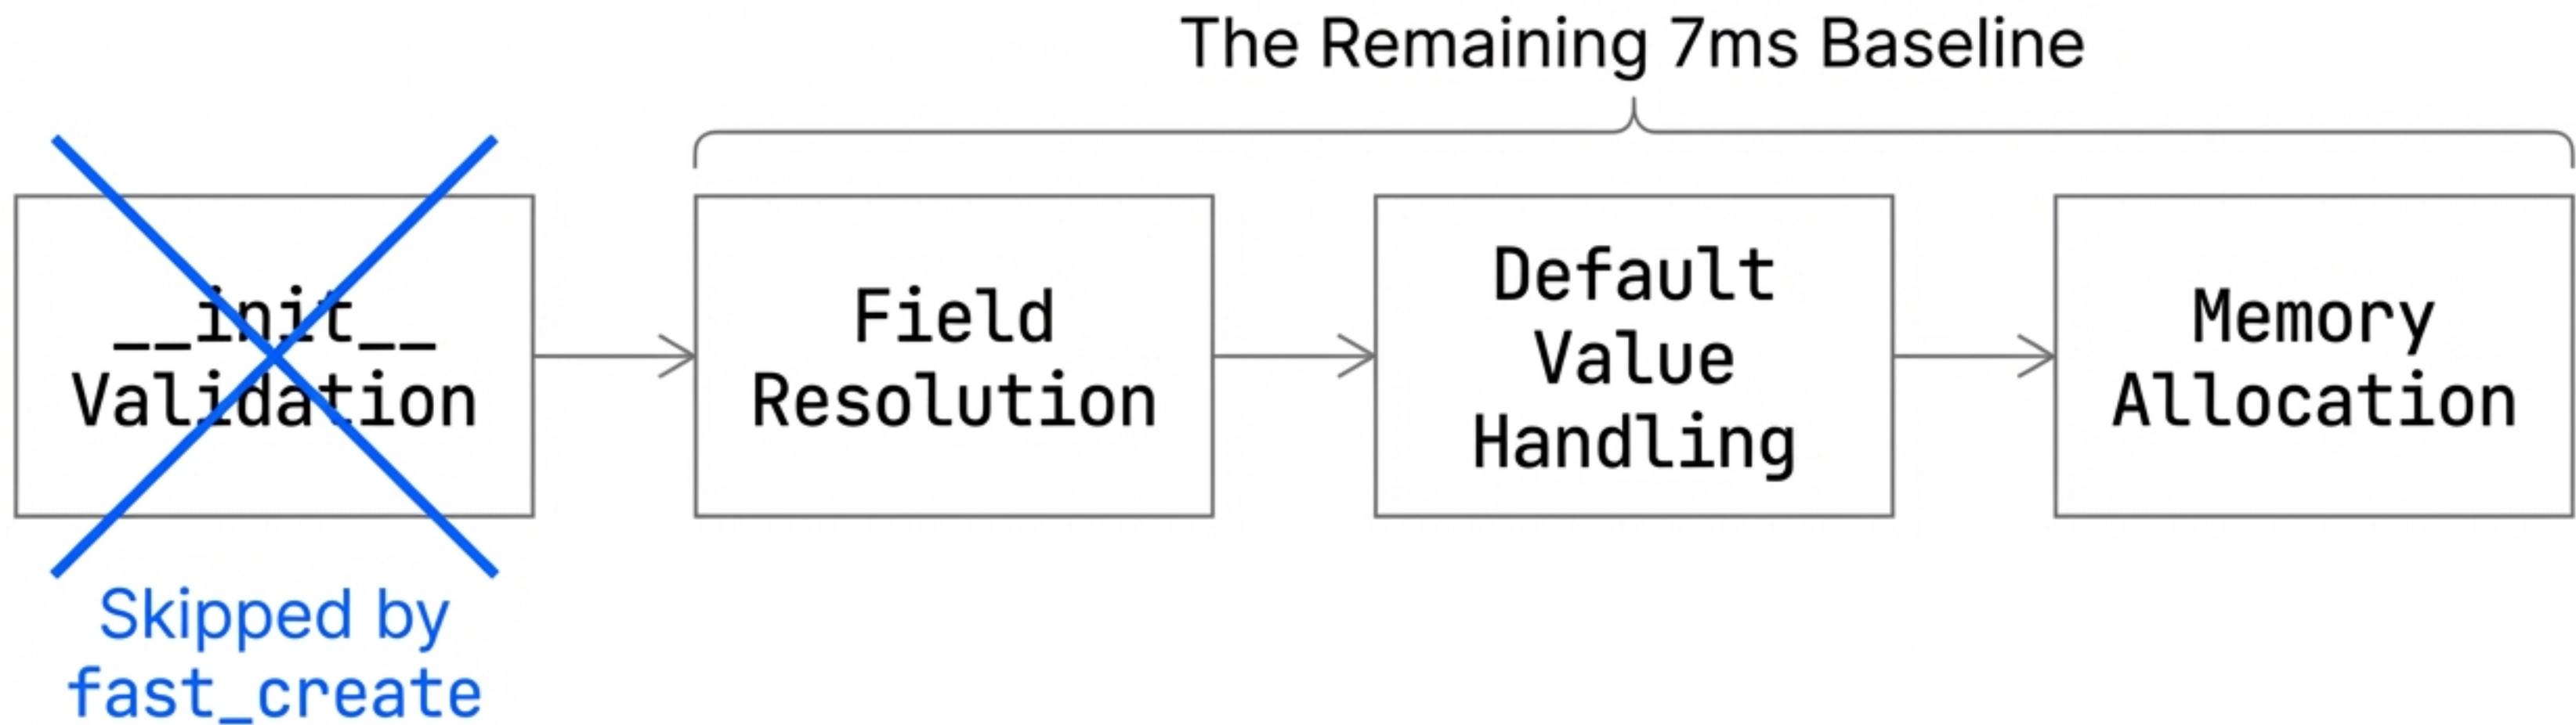


Scaling Behaviour is Linear and Healthy



Conclusion: No $O(n^2)$ behaviour. The algorithm is sound. The issue is the high fixed cost ('the floor'), not the scaling ceiling.

Root Cause: The Cost of “Type_Safe” Overhead



`fast_create` bypasses validation (~56% savings), but allocation costs remain. Further gains require Object Pooling or C Extensions.

Immediate Action: Enable ‘fast_create’

Target: `Html__To__Html_MGraph__Document__Node_Id_Reuse.convert_from_dict()`

Setting: `fast_create = True`

Validation: `skip_validation = True`



Justification: Zero behavioural change. Immediate **50%** latency reduction. Confirmed **2x Speedup**.

Future Strategy: Breaking the 7ms Barrier

Investigation Targets

- **Metric Analysis:** Benchmark **MGraph** creation directly to isolate graph infrastructure.
- **Lazy Initialisation:** Assess if indexes must be built immediately.



Architectural Shifts

- **Schema-Based Creation:** Pre-compute field layouts.
- **Object Pooling:** Reuse node/edge objects to avoid allocation costs.

Summary & Conclusion

We have a sound linear algorithm with a high fixed cost.
Enabling 'fast_create' buys us performance today while we
solve the architectural overhead for tomorrow.

TODAY: Enable 'fast_create' ([2x Speedup](#))

TOMORROW: Optimise internal object allocation

Phase E Analysis Complete. Ready for implementation.