

Phase E Scalability Analysis

Performance Debrief & Data Review

January 2026

Status: ✓ Analysis Complete

Source: Perf__Phase_E__Scalability test suite



Executive Summary: System Health & Key Findings



Scaling Behaviour

O(n) Linear

Verdict

Algorithm is sound; no O(n^2) issues detected.



Optimization Win

1.9x–2.2x

Speedup

fast_create mode delivers ~50% overall improvement consistently.



Primary Bottleneck

Dict→MGraph

78–95% of Time

The critical path for future optimization.

Note: Fixed overhead identified at ~9ms minimum (Architectural).

Methodology & Test Scope

Analysis conducted using 87-iteration Fibonacci sampling (measure_fast) to ensure statistical significance.

Quick Analysis



Scope: 10, 50, 100 nodes

Result: **49.8%** improvement detected

Standard Analysis



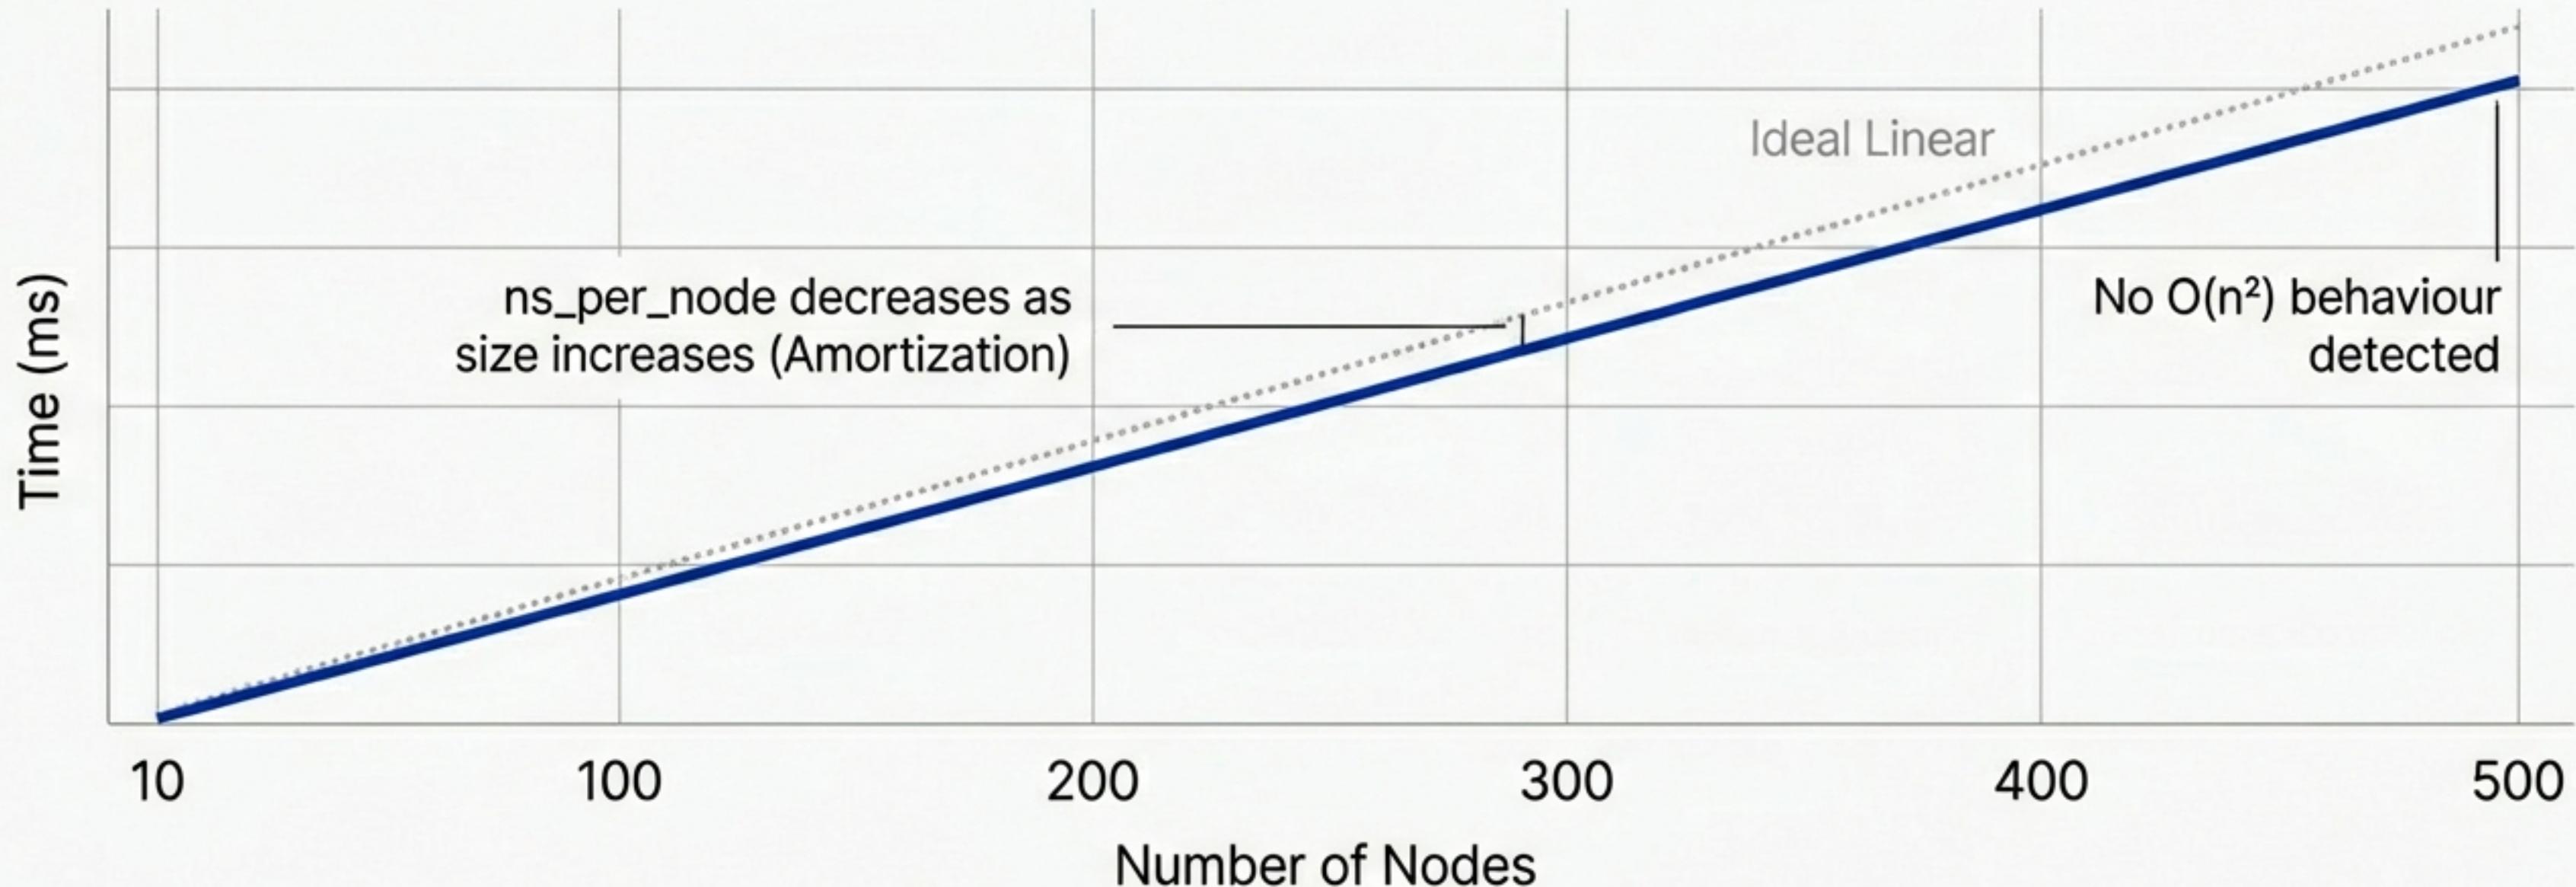
Scope: 10, 50, 100, 500 nodes

Result: **49.2%** improvement detected



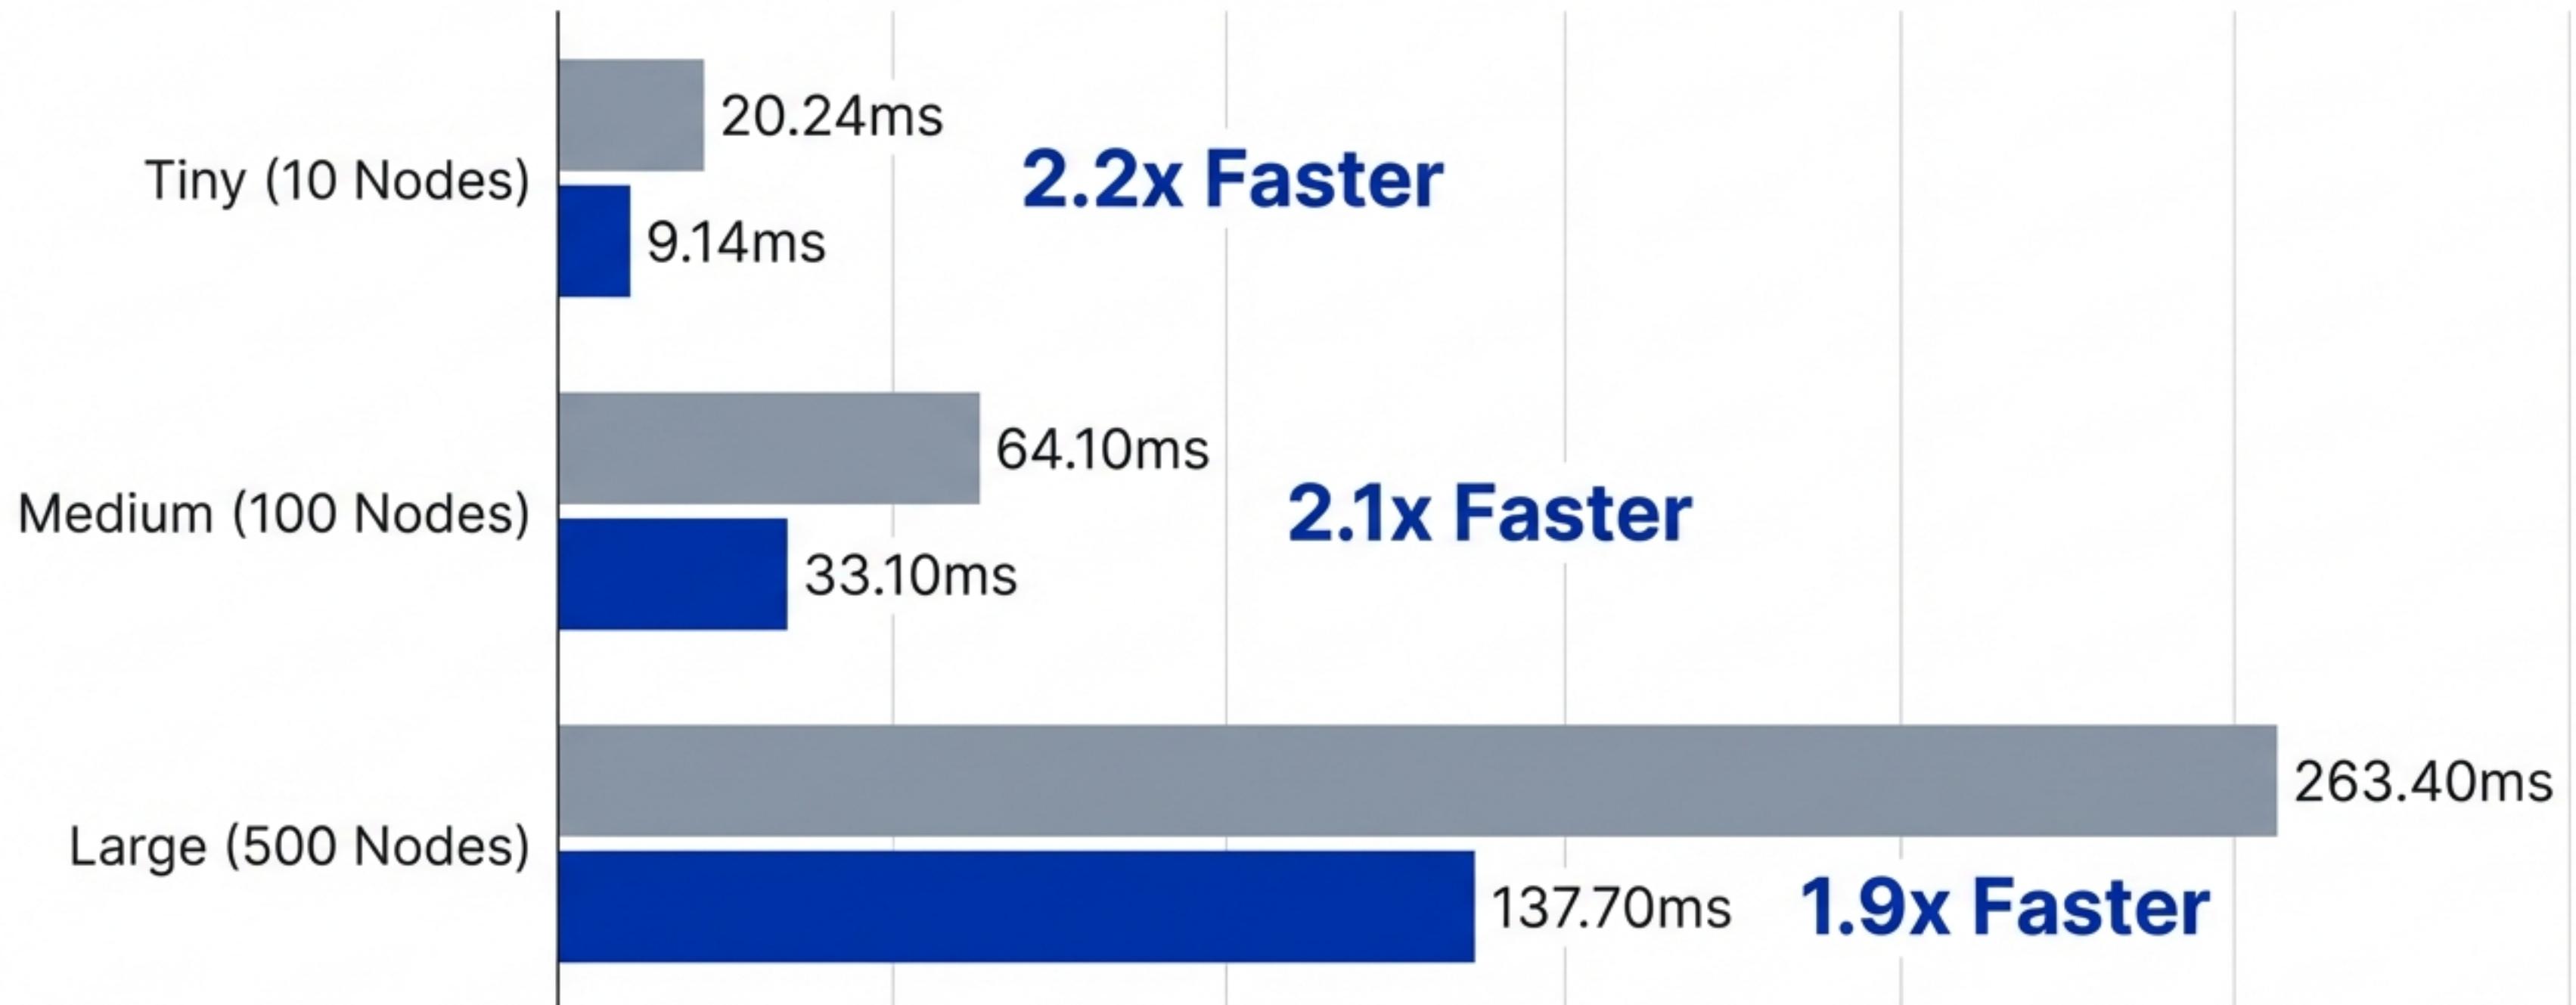
Insight: Both analysis modes produce consistent verdicts. 'Quick' analysis is sufficient for general profiling; 'Standard' is required only when investigating large-scale behaviour.

Core Finding: Healthy O(n) Scalability



The algorithm performs better than pure linear scaling due to amortization of fixed overheads.
The system becomes more efficient per-node as the document size grows.

The Win: Validation of fast_create



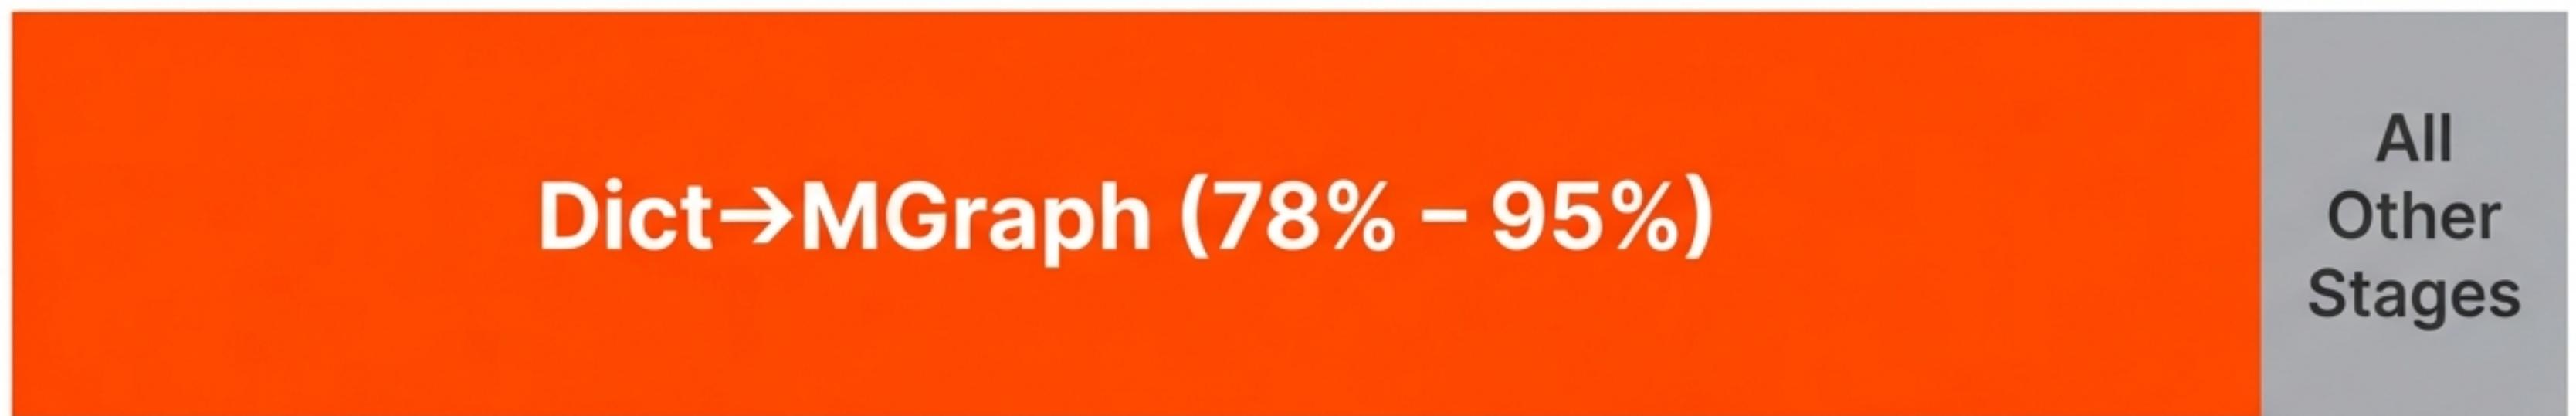
The optimization consistently halves execution time across all tested scales.

Speedup Consistency Across Scales

Size	Nodes	Default Time	Fast Create Time	Speedup
Tiny	10	20.94ms	9.44ms	2.2x
Small	50	40.40ms	19.60ms	2.1x
Medium	100	71.00ms	33.30ms	2.1x
Large	500	266.60ms	140.30ms	1.9x

Analysis: The slight decrease at 500 nodes suggests irreducible complexity in per-node conversion logic, but the gain remains substantial (~50% time reduction).

The Bottleneck is Dict→MGraph



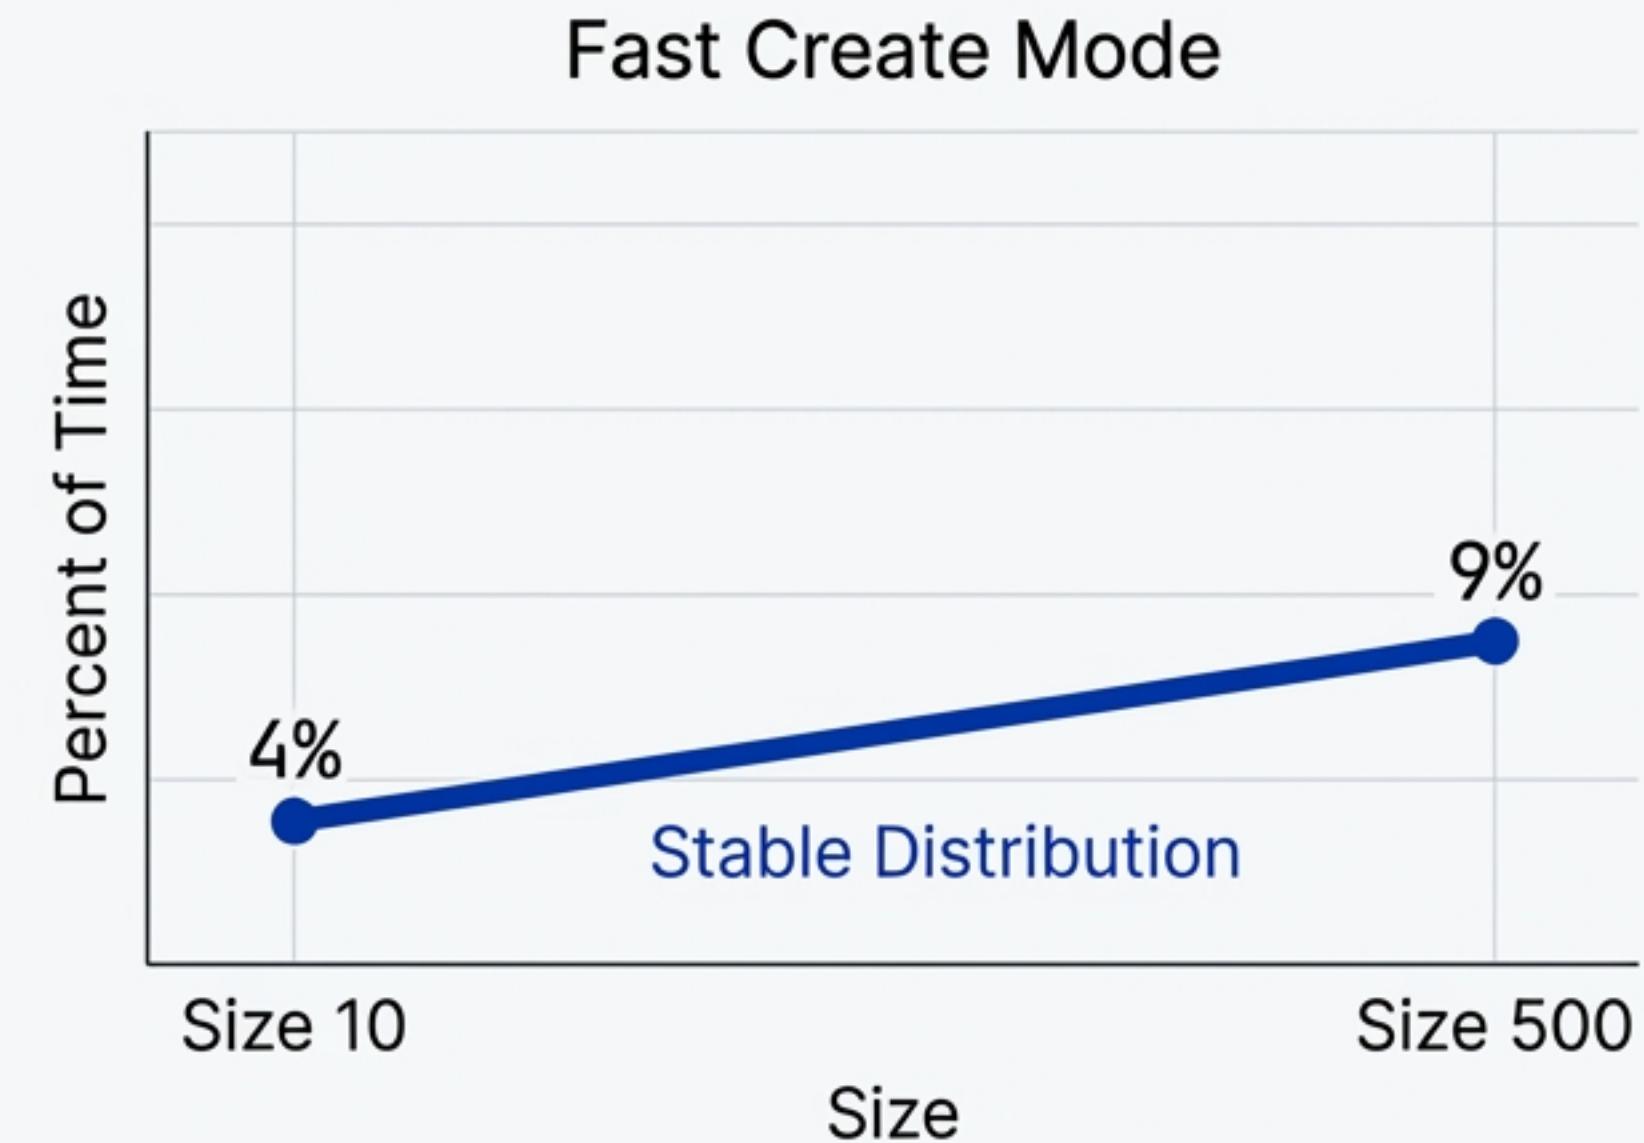
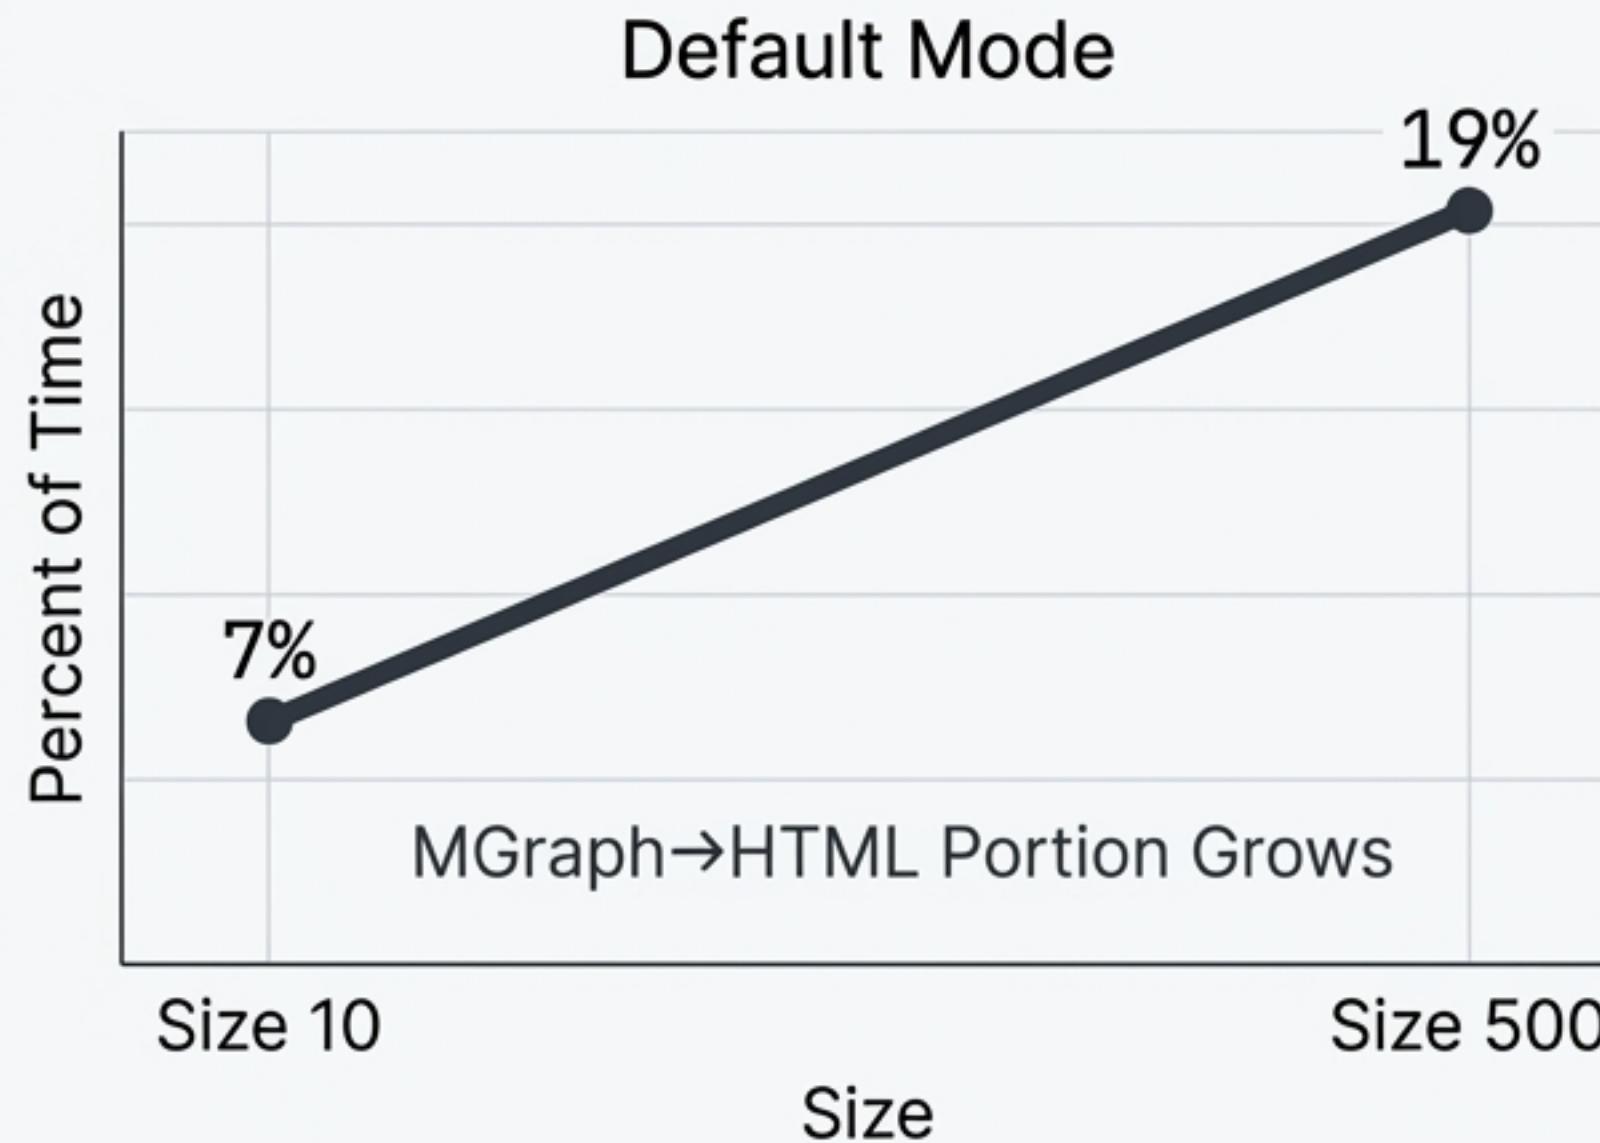
Tiny Scale (10 nodes):
92% of total time

Large Scale (500 nodes):
81% of total time

“This is the only stage worth optimizing further.”

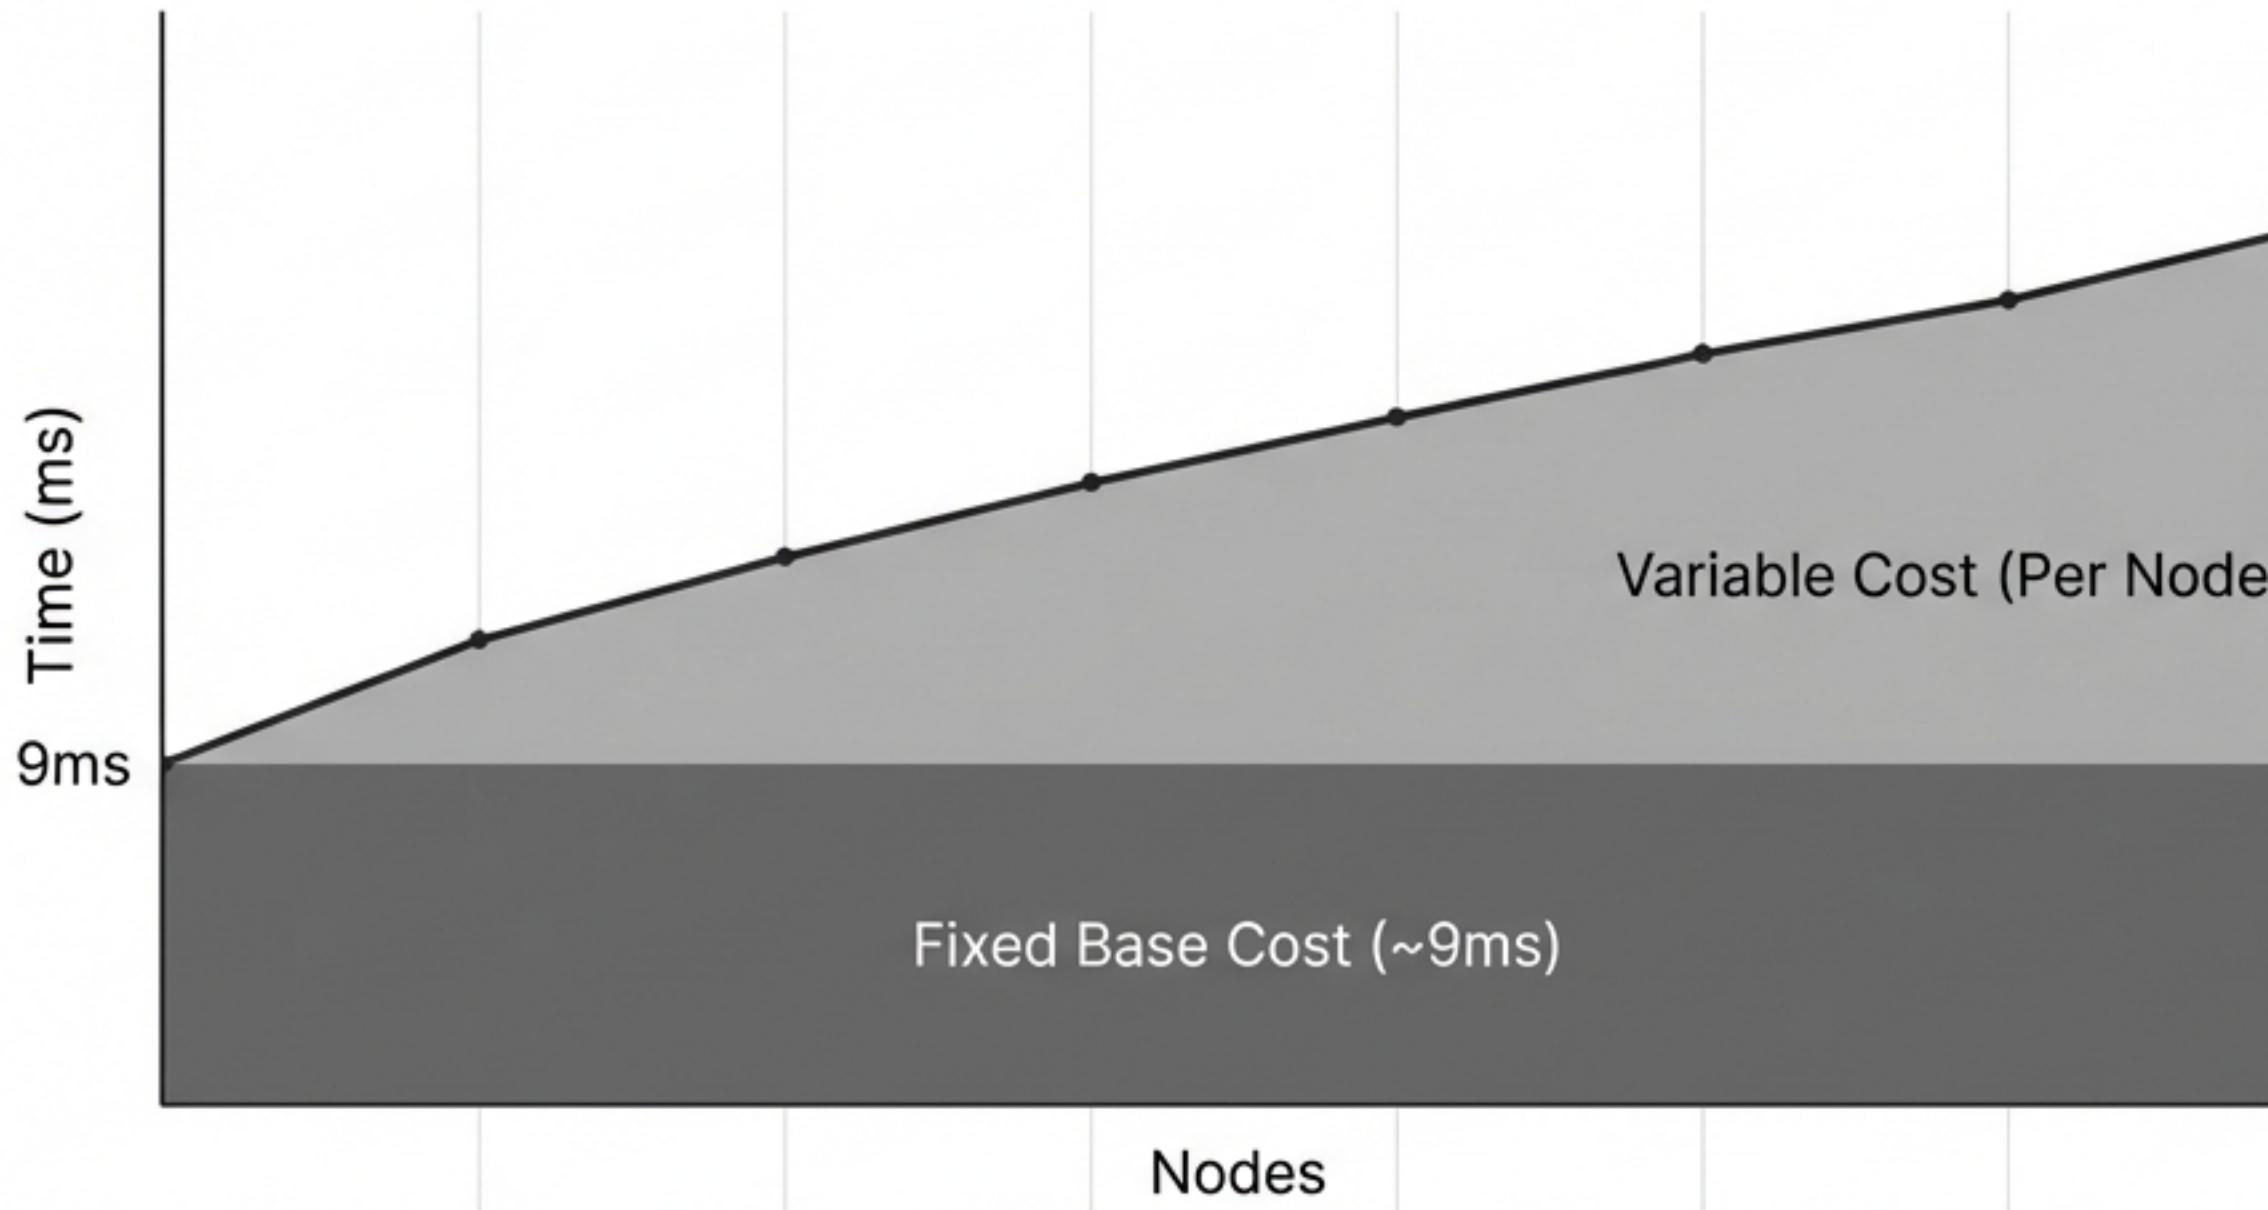
Stage Distribution Shift & Stability

Comparing behaviour of MGraph→HTML stage



fast_create benefits the HTML reconstruction phase more at larger sizes, likely because it handles Type_Safe object creation more efficiently.

The ‘Fixed Overhead’ Floor



Minimum floor: **~9ms**
(even 1-node HTML takes 7ms)

Marginal Cost (10 nodes):
~900µs per node

Marginal Cost (500 nodes):
~275µs per node

The high initial cost is driven by document structure creation, index initialization, and MGraph infrastructure setup.

Per-Node Efficiency (Amortization)

2,024,000 ns

Per Node Cost (Tiny Scale)



526,800 ns

Per Node Cost (Large Scale)

As the node count increases, the fixed overhead is spread out, resulting in a 74% reduction in per-node cost at scale. This confirms the architecture scales efficiently.

Recommendation: Production Deployment

Always enable `fast_create`.

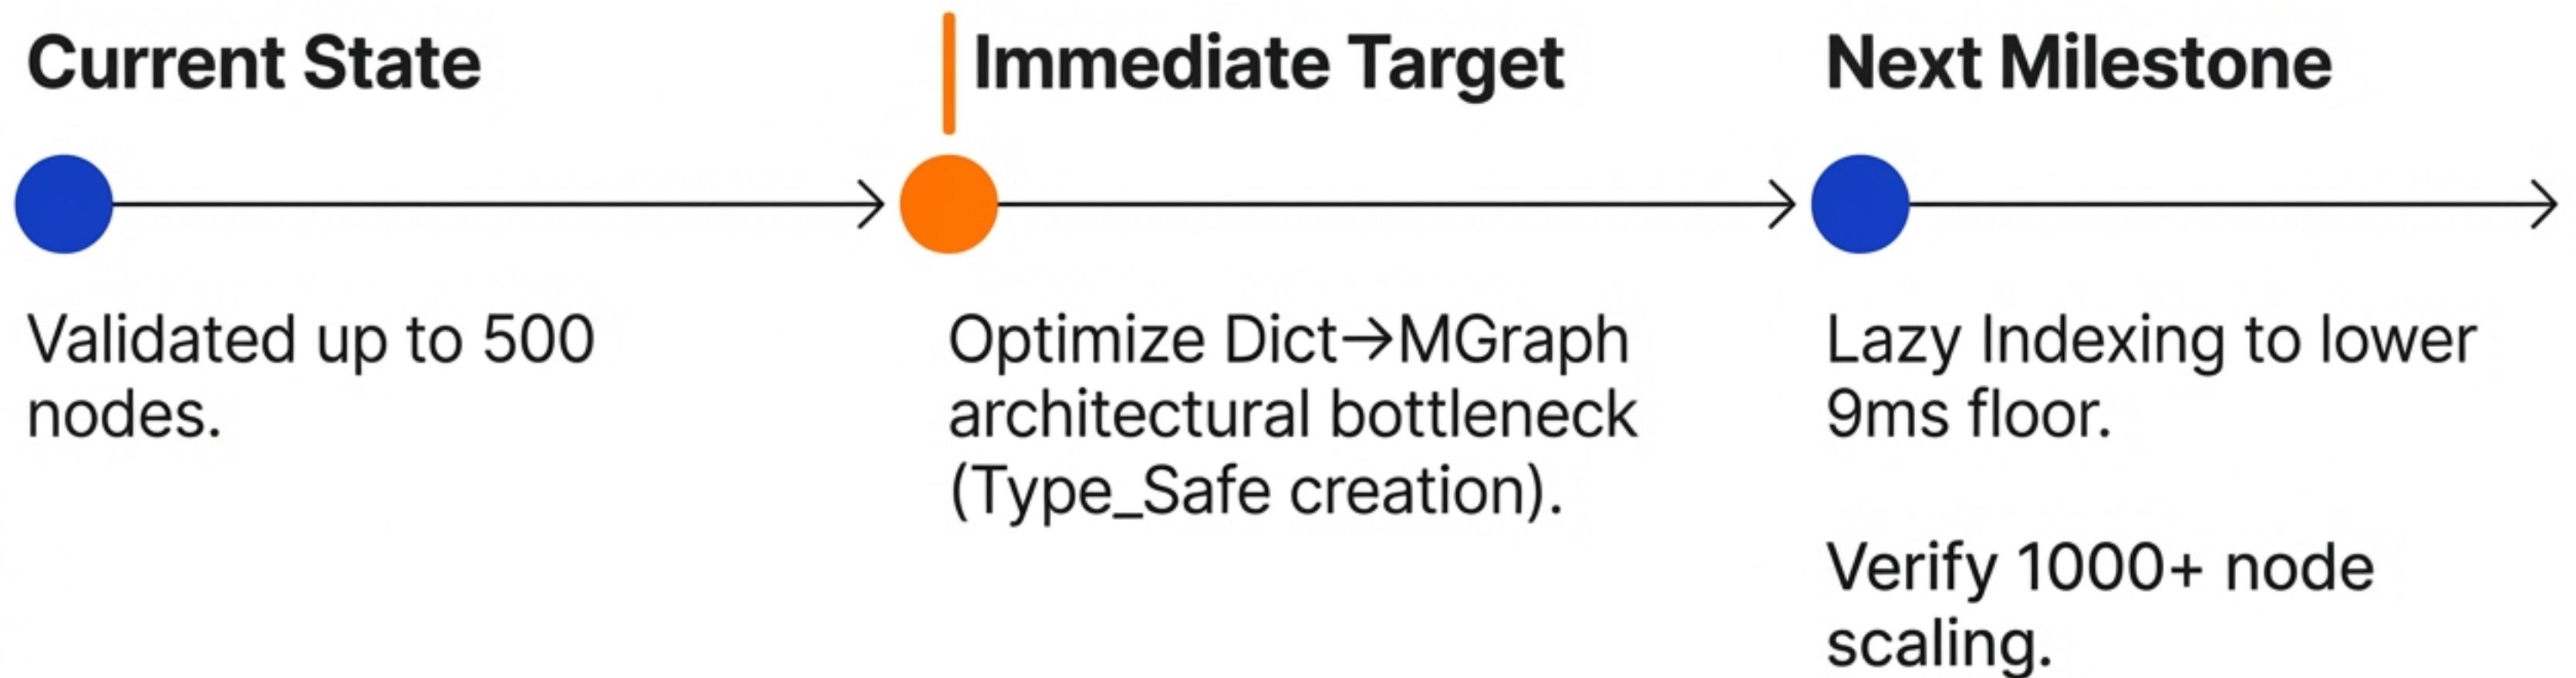
- ✓ Consistent 1.9x–2.2x speedup.
- ✓ No behavioral changes or regressions.
- ✓ Improves stability of MGraph→HTML at scale.

Result: Immediate ~50% performance gain for all users.

Engineering Roadmap & Priorities

Priority	Task	Goal
P0 (Critical)	Profile inside <code>convert_from_dict()</code>	Identify specific object creation costs.
P1	Count MGraph objects created per node	Understand the multiplication factor.
P2	Test lazy index initialization	Reduce the ~9ms fixed overhead.
P3	Benchmark at 1000+ nodes	Verify scaling holds at extra-large sizes.

Future Outlook



Generated Report Artifacts

- 📄 scaling__quick__default.txt
- 📄 scaling__quick__fast_create.txt
- 📄 scaling__standard__default.txt
- 📄 scaling__standard__fast_create.txt
- 📄 scaling__comparison__quick.txt
- 📄 scaling__comparison__standard.txt

All measurements used the `Perf_Phase_E_Scalability` test suite.

Final Verdict

Algorithm is Sound: $O(n)$ linear scaling verified.

Speedup is Real: `fast_create` delivers a consistent 2x improvement.

Target is Defined: Future gains lie exclusively in optimizing `Dict`→`MGraph` object creation.

Phase E Analysis Complete.

