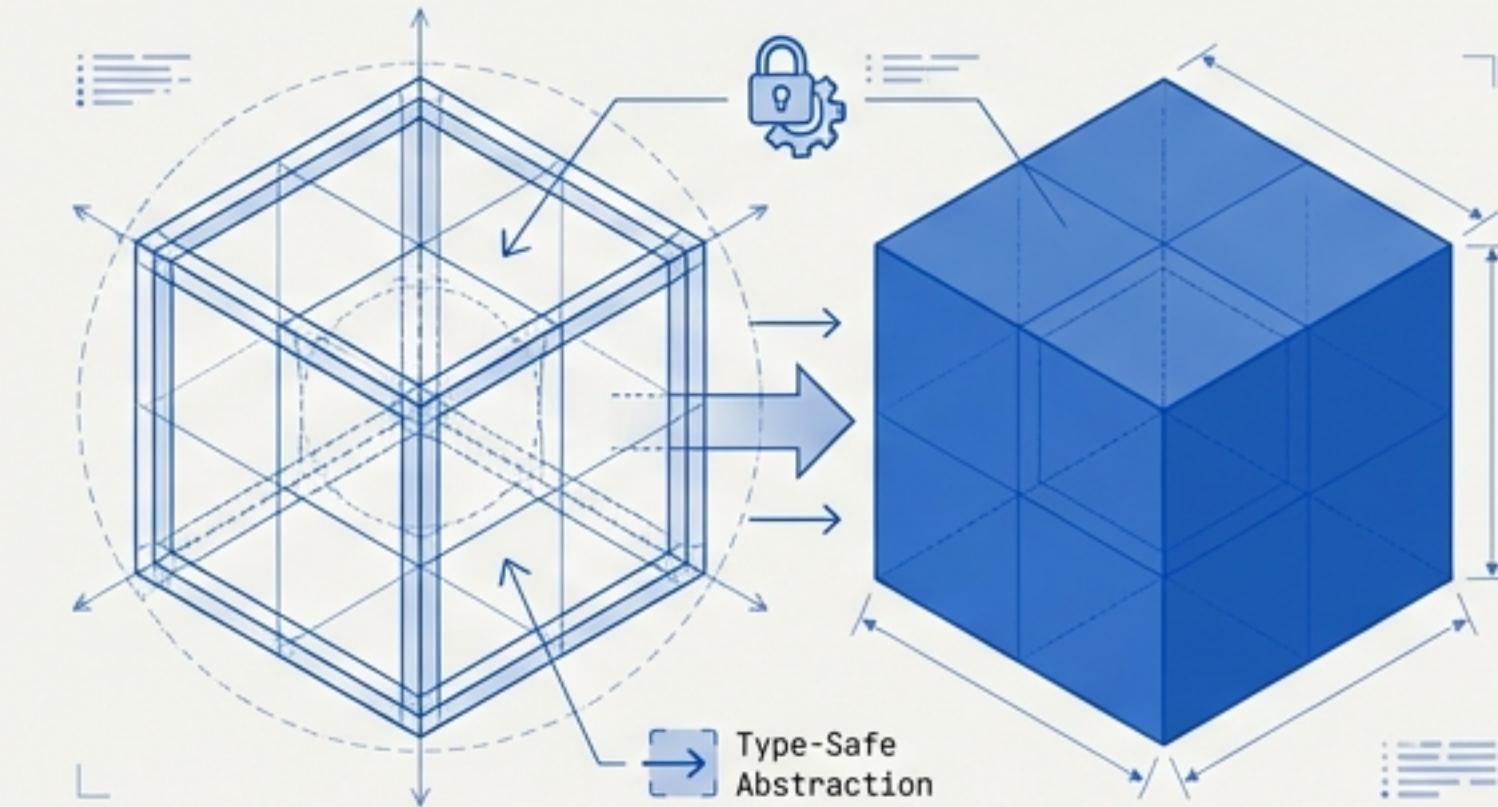


Phase E_3 Debrief: Cache Service Storage Integration

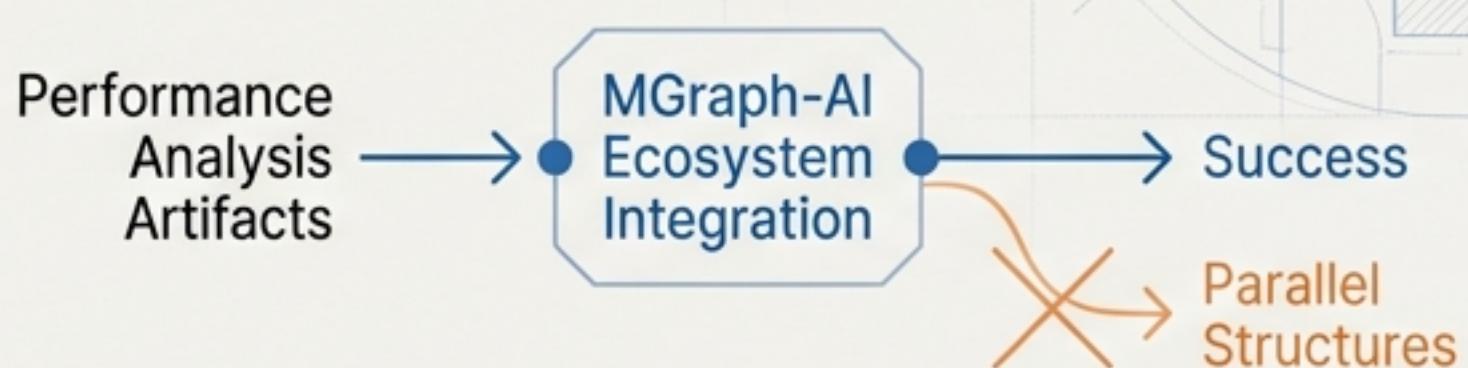
Implementing a Pluggable, Type-Safe
Storage Abstraction Layer for MGraph-AI.

📅	DATE: January 2025
◆	STATUS: Completed
☒	DURATION: Single session with iterative refinement



Executive Summary & Objectives

We aimed to create a flexible storage system for performance analysis artifacts. We succeeded by integrating deeply with the MGraph-AI ecosystem rather than building parallel structures.



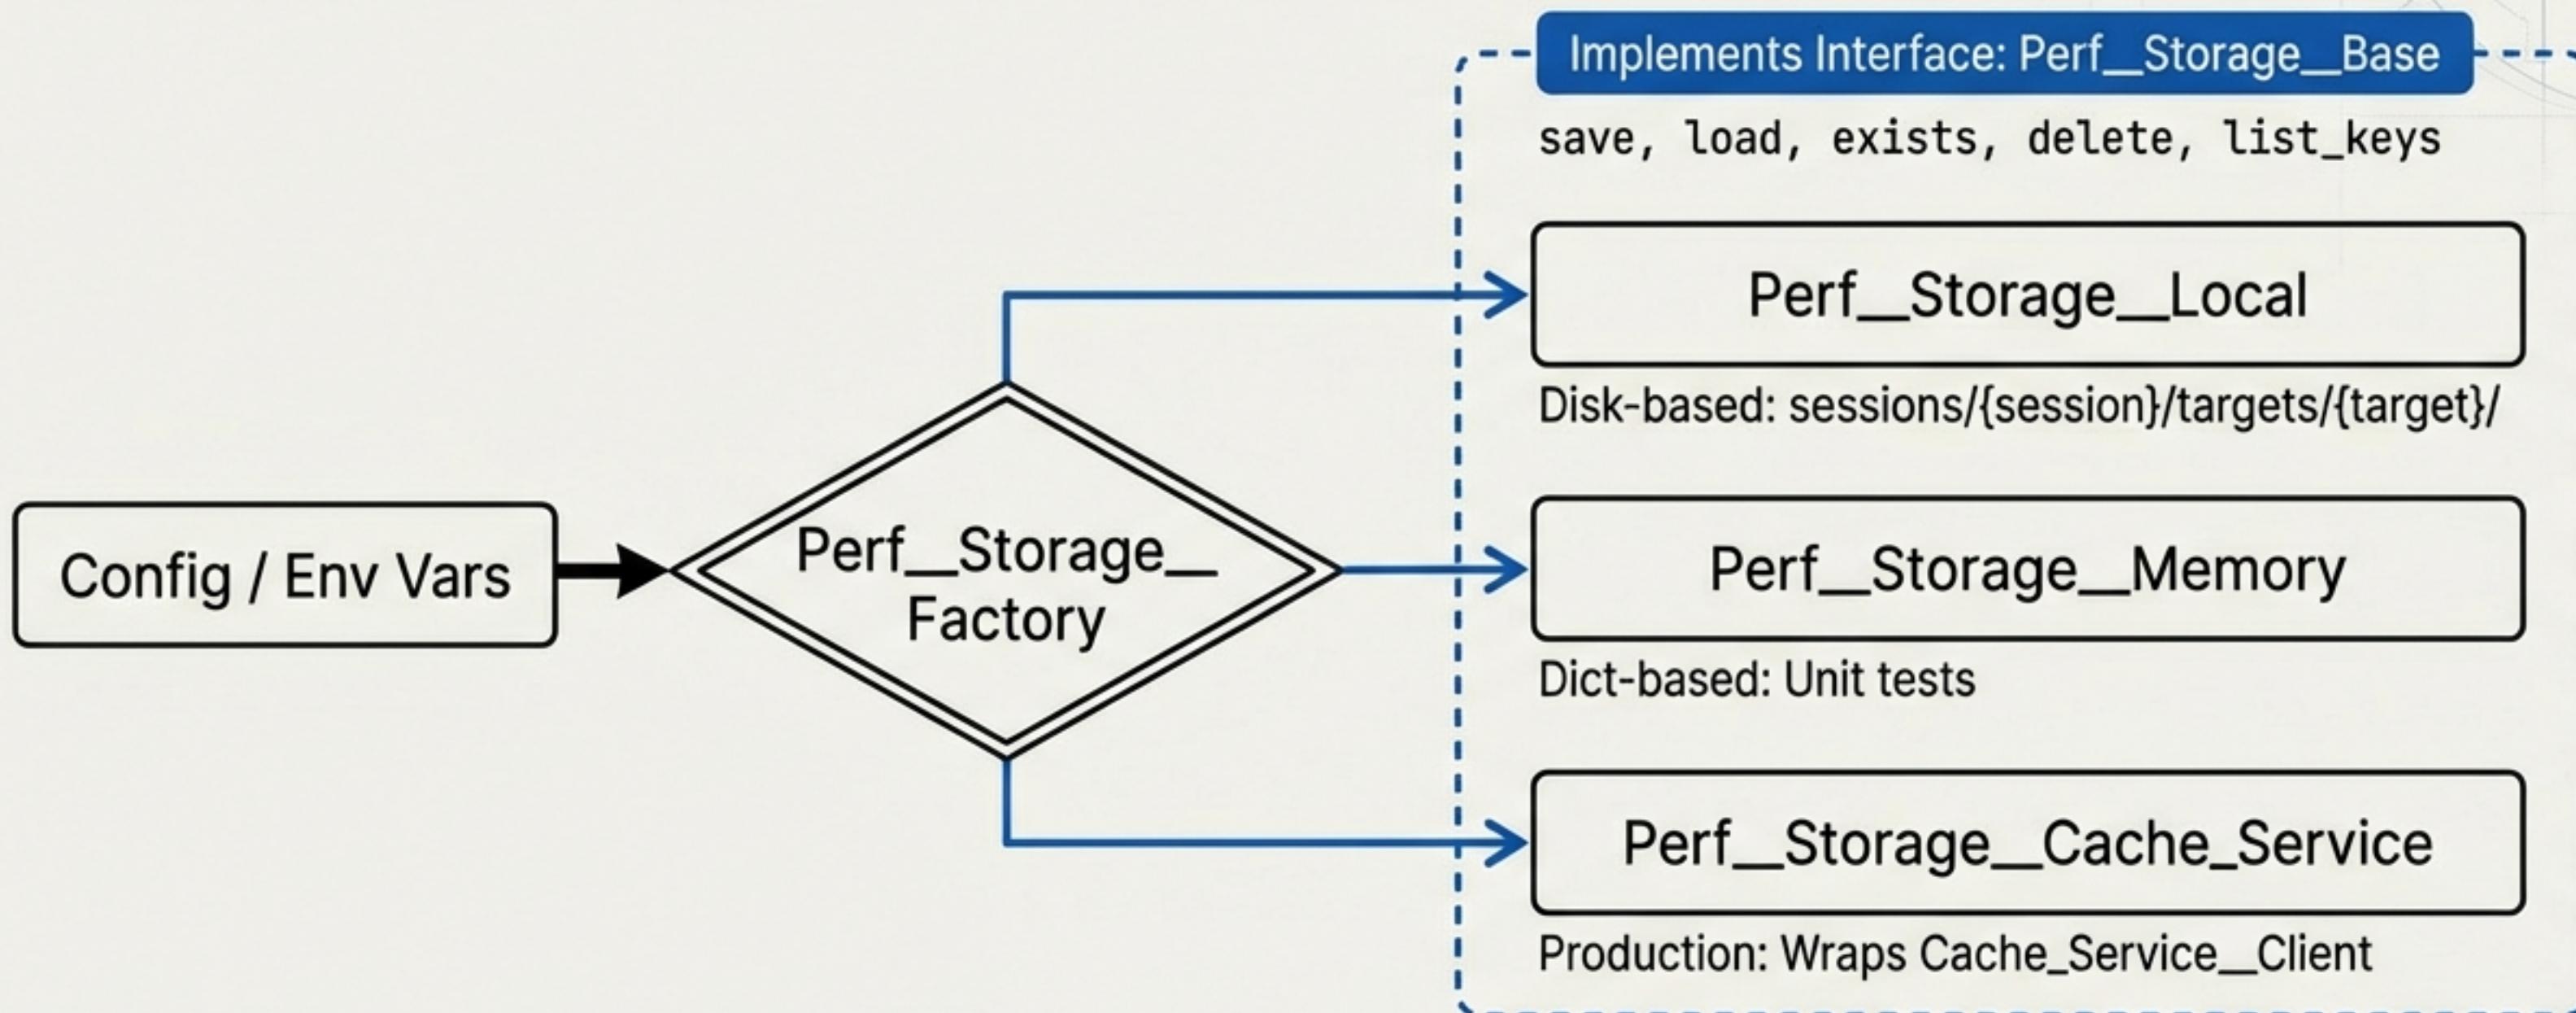
Objectives

- 1. Create pluggable backend system.
- 2. Integrate with MGraph-AI Cache Service (Production).
- 3. Maintain Local/Memory backends (Dev/Test).
- 4. Enforce Type_Safe compliance.

Key Achievements

-  **Architecture:** Three functional backends (Local, Memory, Production).
-  **Code Volume:** 40 Files Delivered (25 source, 15 test).
-  **Headline Win:** Successful integration with the [official MGraph-AI Cache Service client library](#).
-  **Quality:** Zero mock maintenance due to realistic in-memory testing.

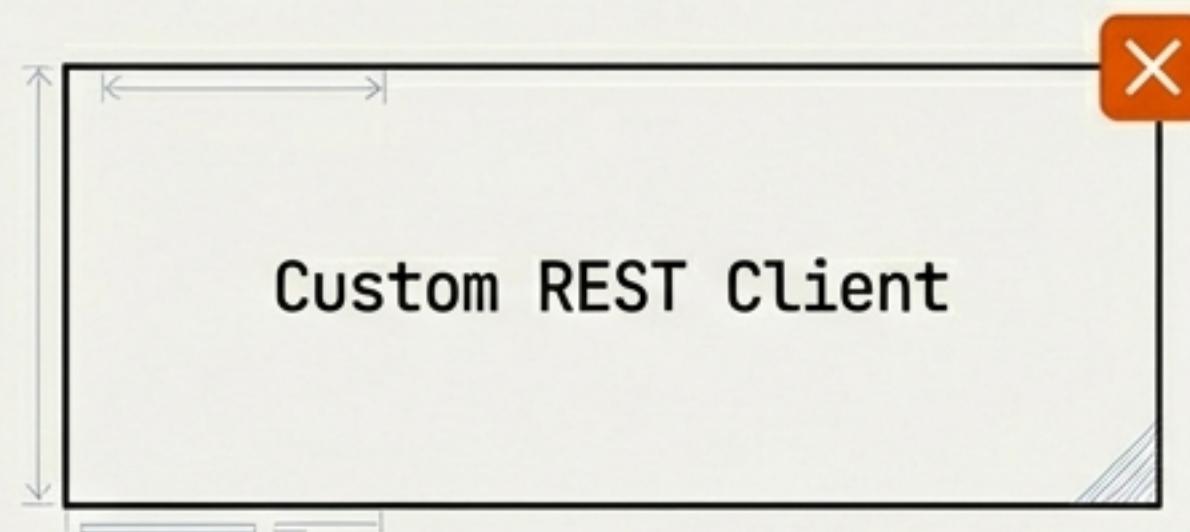
System Architecture Overview



The Factory pattern isolates application logic from storage implementation. The calling code never knows—or cares—where the data lives.

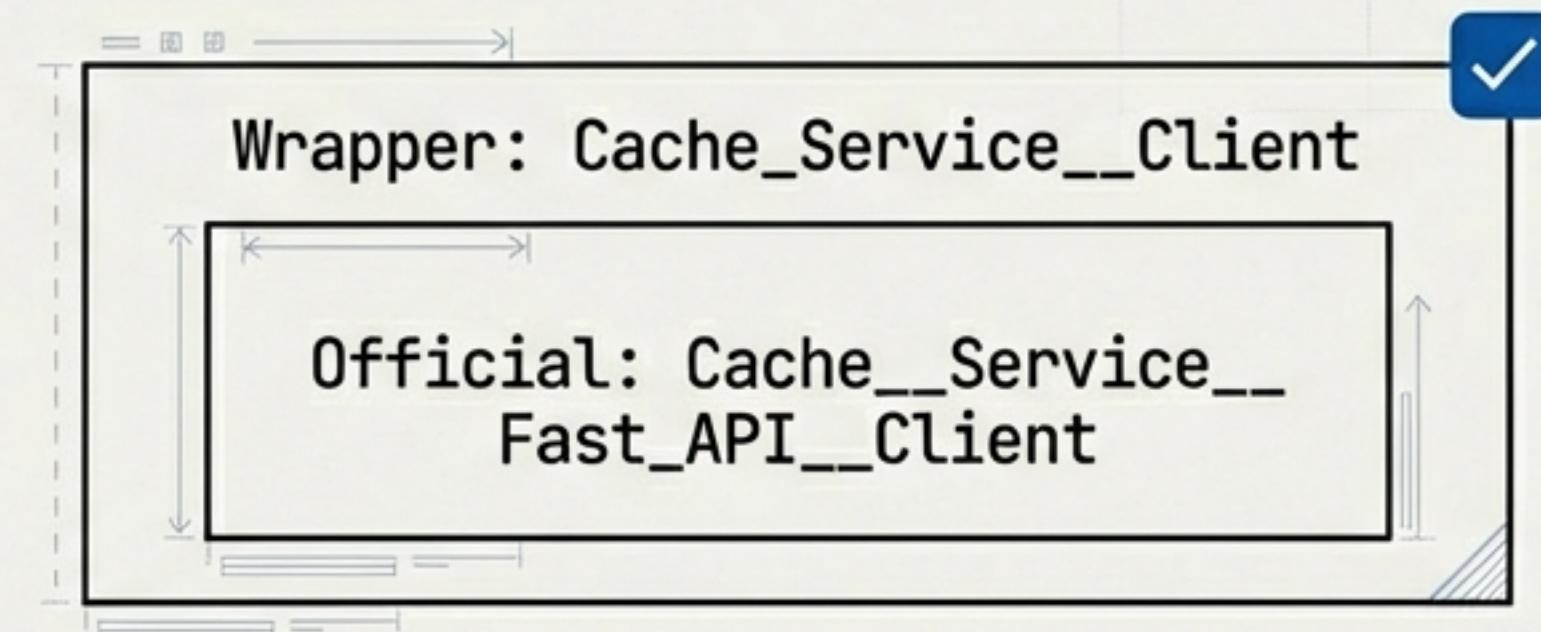
Evolution of the Client Strategy: Wrapper vs. Direct REST

Iteration 1 - Abandoned



- Implementation: Built using 'requests' library.
- **Issues:** High maintenance burden.
- **Issues:** Duplicated authentication logic.
- **Issues:** Brittle mock requirements.

Iteration 2 - Adopted

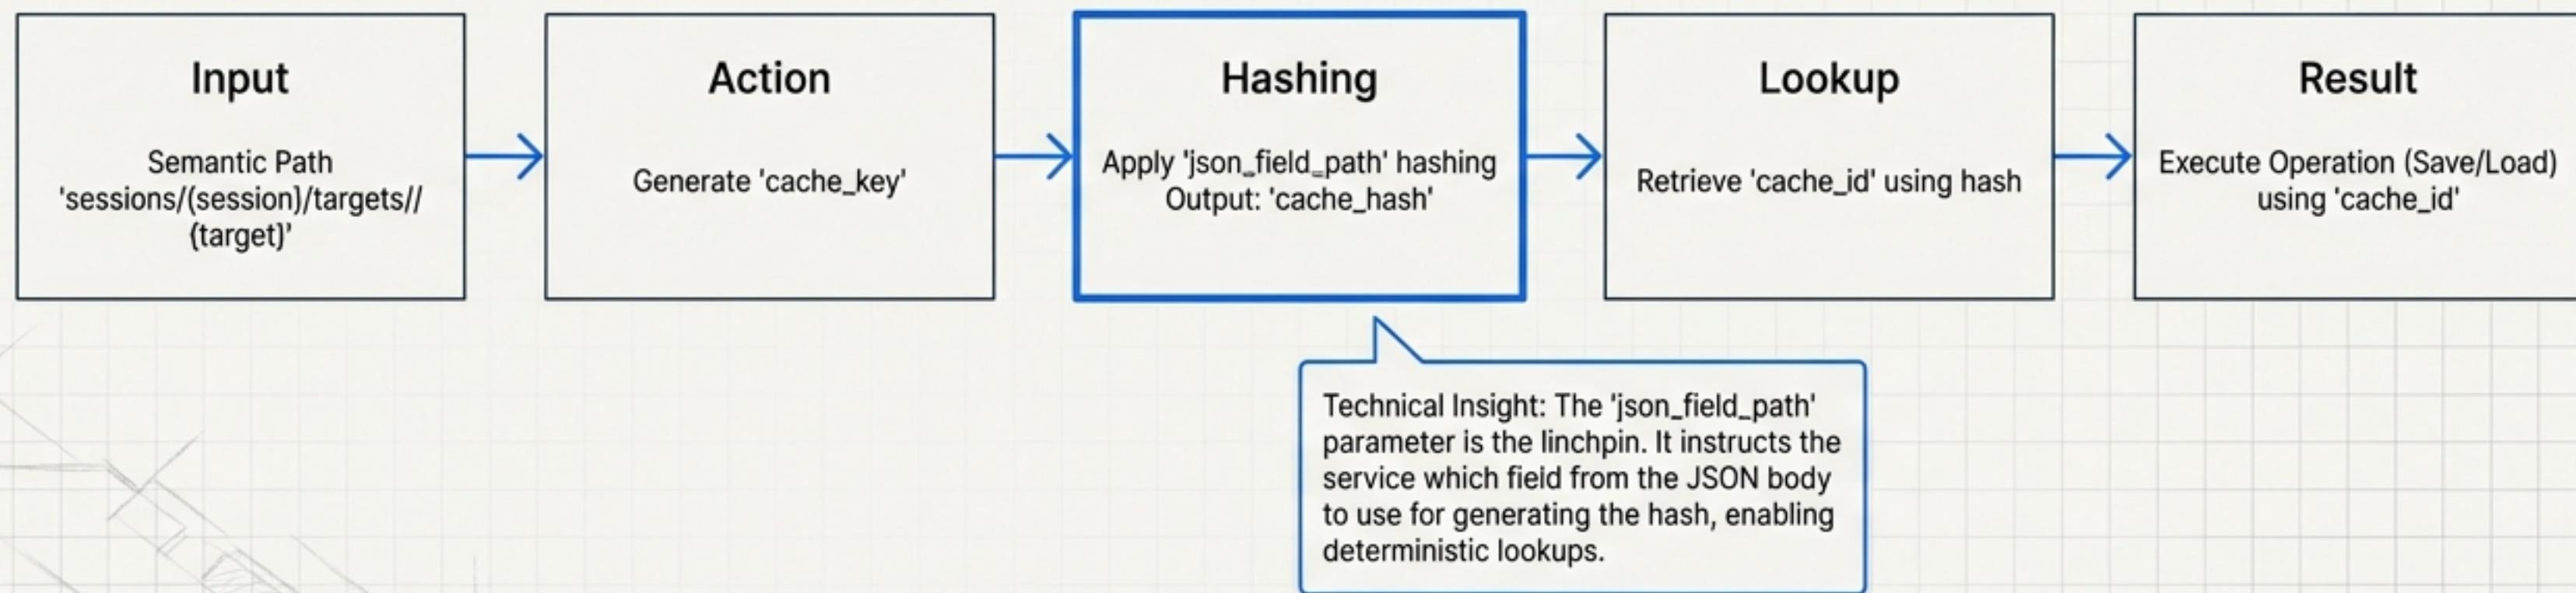


- Implementation: Wraps official client.
- **Benefits:** Consistent API with MGraph-AI services.
- **Benefits:** Built-in auth handling.
- **Benefits:** Type-safe schemas.
- **Benefits:** Support for REMOTE, IN_MEMORY, LOCAL_SERVER modes.

Bottom Line: We chose to leverage existing infrastructure rather than reinventing the wheel.

The Content Addressing Challenge: Hash-Based Lookup

Problem Statement: The Cache Service is content-addressable. It does not natively understand semantic paths like 'sessions/session_1/targets/target_A'.



Quality Assurance: Shifting from Mocks to In-Memory

Metric	Initial Approach (Mocking)	Final Approach (In-Memory Service)
Implementation	Used "unittest.mock"	Real "Cache_Service" in memory mode
Reliability	Brittle tests, low confidence	<input checked="" type="checkbox"/>
Maintenance	High burden	Zero mock maintenance
Performance	N/A	Startup time ~100ms

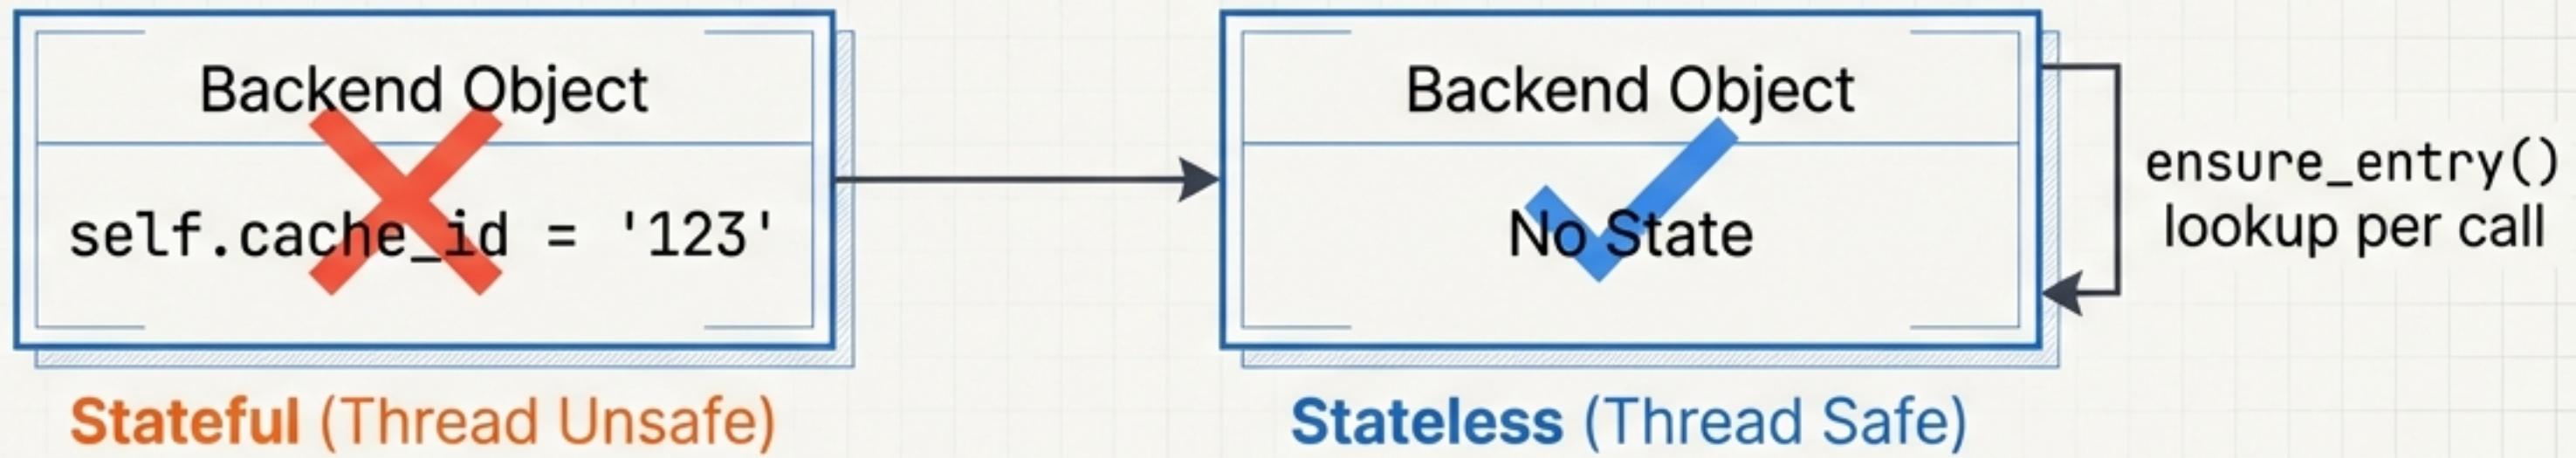
“Testing against the real implementation provides higher confidence with zero mock maintenance.”

Key Benefits

- Catches real API contract issues.
- Eliminates mock drift.
- Ensures integration behavior validity.

Backend Design – Statelessness & Configuration

Stateless Backend Design



Benefit: Predictability. The caller controls the caching behavior.

Configuration Simplification

Initial: Explicit URLs in config files.



Final: Configuration handled by client library via environment variables.
AUTH__TARGET_SERVER__CACHE_SERVICE__KEY_VALUE

Benefit: Leverages standard env vars; removed redundant logic in 'Perf_Storage_Config'.

Retrospective: Validated Patterns vs. Discarded Approaches

What Worked

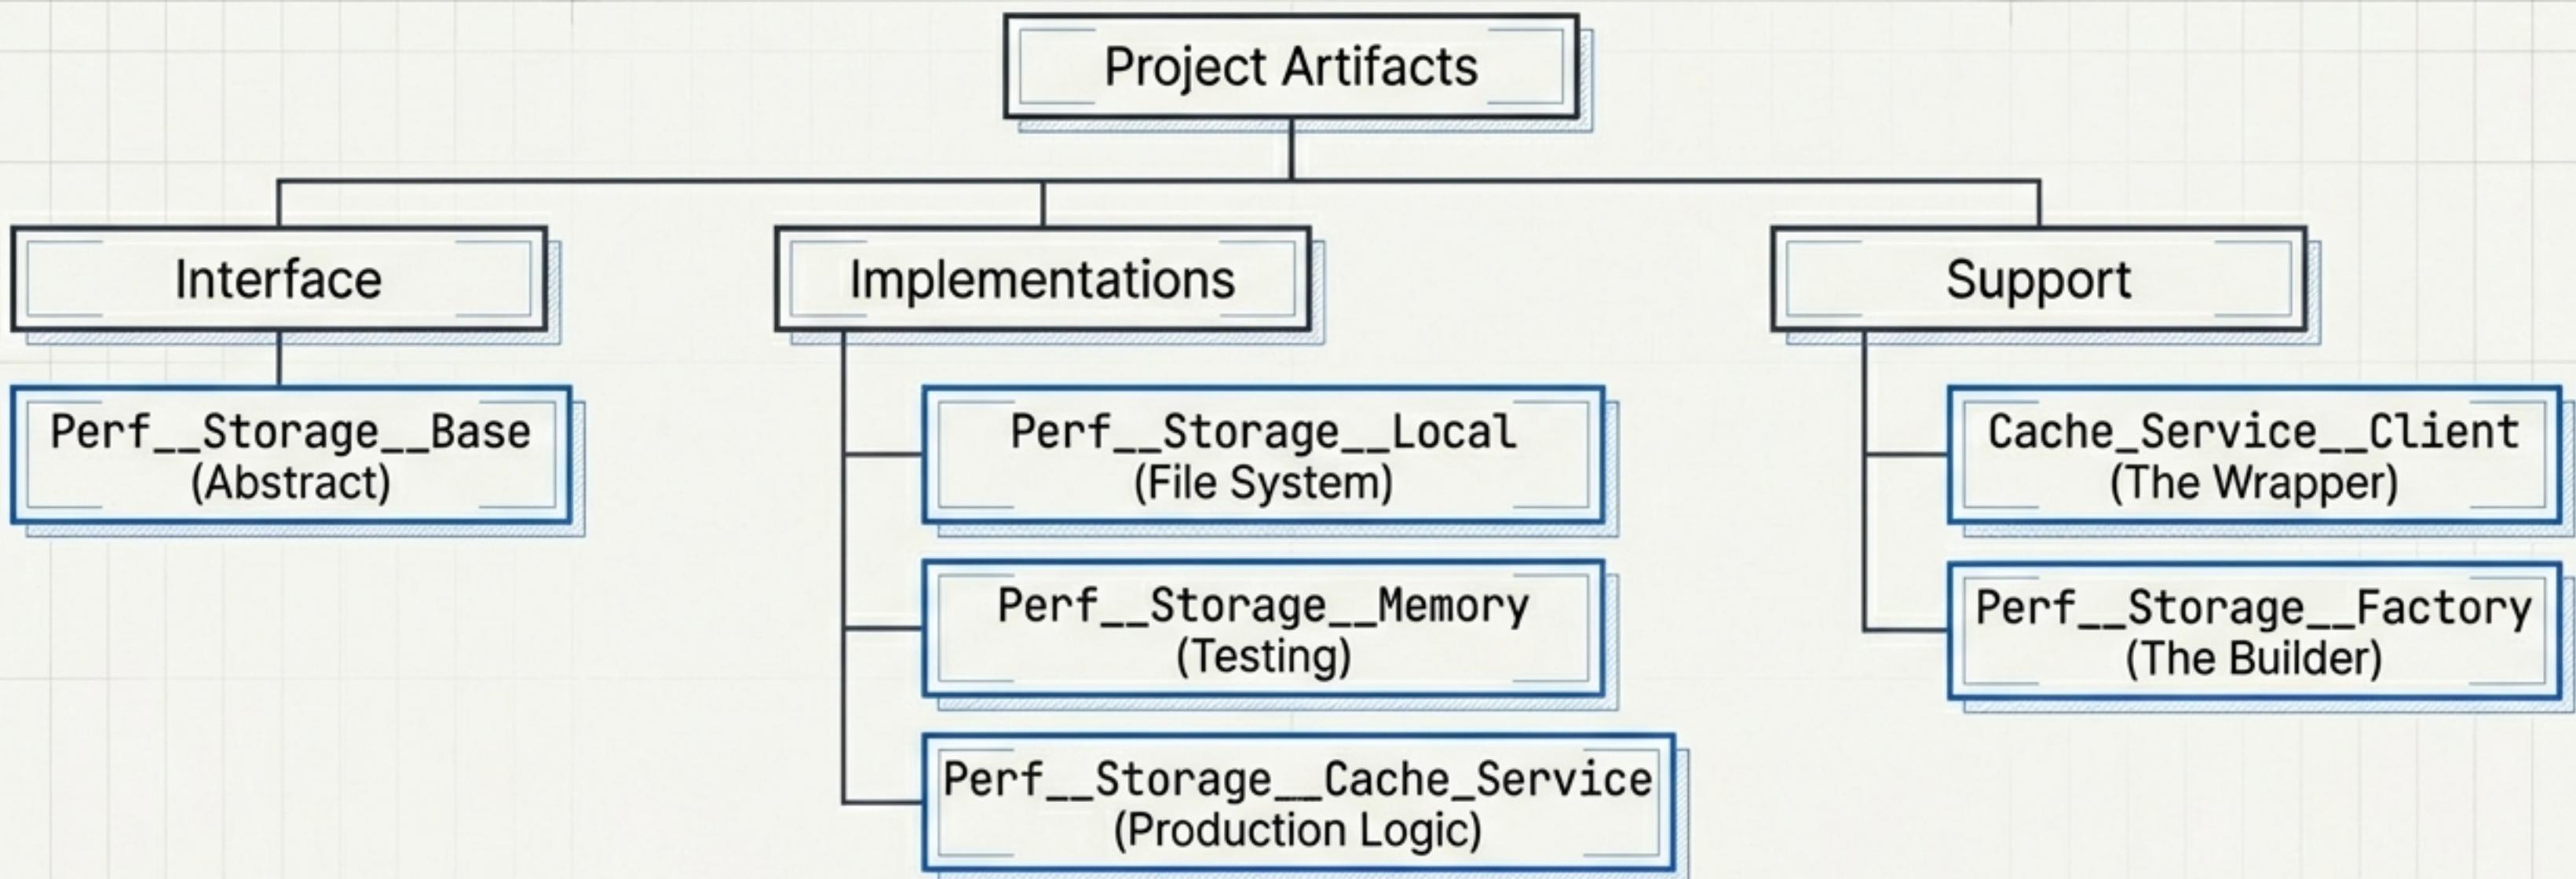
- ✓ **Hierarchical Child Data:** Fits natural artifacts structure (results/reports).
- ✓ **Official Client Wrapper:** Reduces code debt and maintenance.
- ✓ **Factory Pattern:** Clean dependency injection via config.
- ✓ **In-Memory Testing:** High confidence, low latency.

What Didn't Work

- ✗ **Direct REST Calls:** Replaced to avoid duplicating client library features.
- ✗ **Mock-Heavy Tests:** Replaced due to brittleness.
- ✗ **Stored 'cache_id':** Removed to ensure thread safety.
- ✗ **Separate 'setup()' method:** Removed in favour of passing dependencies via constructor.

Delivered Artifacts & Component Breakdown

40 Total Python Files
(25 Source, 15 Test)

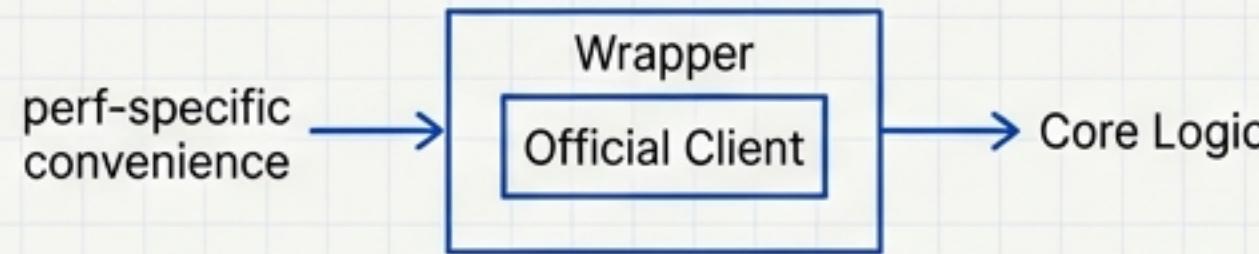


Schema Updates: 'Schema_Perf_Entry.py' (Added `cache_key`), 'Schema_Perf_Storage_Config.py' (Cleaned).

Technical Decision Matrix

Decision: Wrapper vs. Direct Client

Rationale: Add perf-specific convenience while delegating core logic to the official implementation.



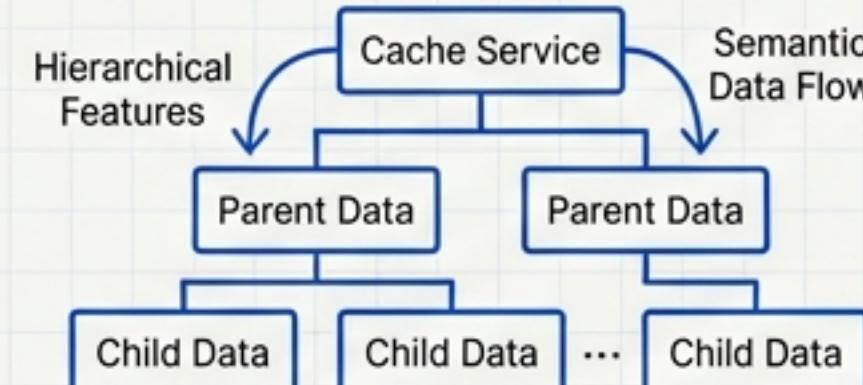
Decision: Constructor Arguments

Rationale: Pass Session/Target defaults explicitly to avoid hidden state changes.



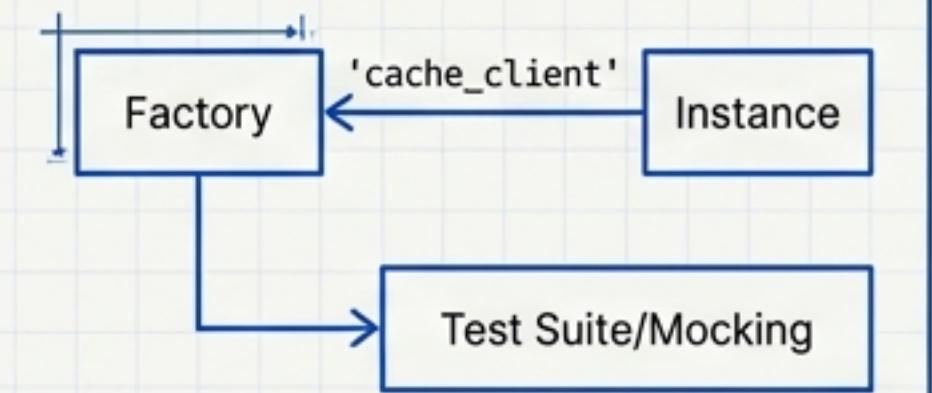
Decision: Child Data Pattern

Rationale: Leverages cache service's native hierarchical features to keep data semantic.



Decision: Dependency Injection

Rationale: Factory accepts 'cache_client' to support easier testing and mocking scenarios.



Roadmap: Phase E_4 Candidates

1. List Keys Implementation

Use the cache service 'refs' API to list child data files.

2. Batch Operations

Mechanism to store/retrieve multiple files efficiently.

3. Cache Invalidation

Logic to delete an entry and all its children.

4. Compression

Handling for large string content.

5. Metrics

Tracking storage operation latencies.

Conclusion: Production-Ready Storage Abstraction

Phase E_3 has successfully bridged the gap between local analysis needs and scalable cloud storage.

Core Assets Delivered



Hash-Based Lookup

The key pattern unlocking semantic storage on the Cache Service.



Type_Safe Wrapper

Ensuring compliance and reducing runtime errors.



In-Memory Confidence

A testing strategy that mimics production without the latency.

Final Status: Ready for deployment. Extensible for Phase E_4.