



# Semantic Text Service: Unlocking Knowledge from Unstructured Data

## Introduction

Semantic Text is a new service concept built around **semantic knowledge graphs** – data structures that represent information as interconnected entities and relationships, capturing the meaning of content. The core idea is to convert unstructured text (documents, chats, articles, etc.) into a structured **knowledge graph** that can be more easily searched, analyzed, and used by AI systems. Unlike traditional top-down knowledge management approaches that impose a single taxonomy or ontology, Semantic Text takes a **bottom-up, flexible approach**. It can work with multiple ontologies and taxonomies simultaneously – essentially an “ontology of ontologies” strategy – and then connect the dots between them dynamically. This avoids the need to rigidly agree on one master schema upfront. As one industry source notes, using open standards and allowing both top-down *and* bottom-up modeling lets organizations iterate and extend their semantic models as needed while avoiding vendor lock-in <sup>1</sup>. The result is a knowledge graph that organically grows from the data and use-cases, rather than from a pre-fixed hierarchy.

This approach addresses a key reason past enterprise knowledge graph projects struggled to scale: historically, building and maintaining semantic graphs required extensive manual effort to define ontologies and curate data. Those efforts often stalled because the schemas were inflexible and the feedback loops (to refine the graph based on usage) were weak. Semantic Text aims to change that by **automating knowledge graph construction** with AI and creating continuous feedback loops. Modern large language models (LLMs) can read text and propose connections, while graph databases and algorithms can integrate new information over time. In essence, Semantic Text compresses raw data into a semantic graph representation – distilling *meaning* and *connections* from the noise. This “meaning first” approach can make generative AI and agentic workflows far more effective, by giving them a rich, factual substrate of knowledge to work with.

## What Is a Semantic Knowledge Graph Service?

A **semantic knowledge graph service** like Semantic Text provides on-demand conversion of text into a structured graph format (nodes and edges with semantic labels). In practice, this means you can feed the service a body of text – for example, an article, an email, a transcript – and it will return a **graph** of the key entities, concepts, and the relationships between them. The graph is essentially a compressed interpretation of the source data, focusing on who/what is involved and how things relate. This graph can then be stored, queried, or further processed. Crucially, the service can integrate with existing ontologies/taxonomies if provided: it will map the text onto those predefined categories when possible, and even highlight where new concepts don’t fit the existing schema (suggesting ontology updates or extensions). If no ontology is provided, the service can induce one by identifying emergent clusters of concepts. This aligns with the idea that an effective knowledge graph can incorporate one or many ontologies and make that knowledge available to applications <sup>2</sup>.

Under the hood, Semantic Text uses a pipeline somewhat analogous to ETL (Extract, Transform, Load) but adapted for semantics – often referred to as “**LETS**” (**L**oad, **E**xtract, **T**ransform, **S**ave) in our architecture. The process works as follows:

- **Load:** Ingest the raw data (text documents, transcripts, etc.), possibly from multiple sources (files, RSS feeds, APIs).
- **Extract:** Use NLP and AI (including LLMs and tools like AWS Comprehend) to extract key elements: entities (people, places, things), topics, dates, sentiments, and candidate relationships stated in the text. This step pulls out the raw facts and mentions.
- **Transform:** Convert the extracted elements into a structured **semantic graph** format. This involves assigning semantic types from ontologies (e.g. recognizing that “ACME Corp” is a Company, “Jane Doe” is a Person) and linking entities with relationships (e.g. Jane Doe **works for** ACME Corp). The graph may also infer additional connections or normalize different terms referring to the same entity. This step dramatically reduces noise by focusing only on meaningful connections, essentially compressing the source into an ontology-backed representation. (Any irrelevant or unsupported data from the source is dropped in the transformation.)
- **Save:** Store the results – both the raw extracted data and the final knowledge graph – in a repository or database. The graph can be stored in a graph database or a serialized format (JSON-LD, RDF, etc.) for retrieval. Saving both raw and processed forms allows traceability (you can trace a node in the graph back to the source text) which is critical for trust and debugging

3 .

Once the data is in this semantic graph form, it becomes far more powerful than the original text. Instead of a blob of text, you now have **machine-readable knowledge**. For example, searching becomes easier (“find all documents where a Person works for Company X and mentioned Project Y”), analytics become possible (e.g. count relationships, find central entities), and content can be matched to user interests or policies. This is why knowledge graphs are seen as turning data into knowledge for AI applications 4 5 . The structured context they provide can dramatically improve AI accuracy and consistency – a knowledge graph grounds AI in facts, reducing problems like hallucinations, and in turn AI can help expand and maintain the graph 5 .

## The Semantic Text Service Capabilities

Semantic Text is envisioned as a **serverless, on-demand service** that provides this knowledge graph creation and utilization capability as a flexible API. Some of the key features and capabilities include:

- **On-Demand Text-to-Graph Conversion:** Provide an API endpoint where users send in text (or documents) and get back a semantic graph representation. This could be delivered as a JSON structure encoding nodes and edges (for developers to consume directly), or even visual graph formats for easy interpretation. The goal is an out-of-the-box service that *just works* – you don’t need to train a model or prepare data; the service understands general text and produces a meaningful graph.
- **Ontology-Aware Classification:** Users can supply their own taxonomy or ontology definitions to the service. The Semantic Text engine will then use those definitions to classify and tag the incoming text. For example, if a finance company has an ontology of risk categories, the service can map phrases from a report into that ontology (e.g. flagging mentions of “inflation” or “credit default” under the appropriate risk categories). If multiple ontologies are provided, the service can cross-reference them – effectively linking taxonomies of taxonomies. This enables *translation* between different vocabularies. (Imagine mapping an engineering-focused taxonomy to a sales-focused taxonomy by finding how concepts in the text relate to both.) Support for user-defined

ontologies ensures the graphs align with the customer's domain language. At the same time, the system doesn't force a single global schema; it's flexible to handle many perspectives in parallel, which is a noted benefit of semantic graph approaches [2](#) [1](#).

- **Dynamic Expansion and Learning:** The service continuously "learns" from new data. Each time content is processed, the knowledge graph grows and refines. If the AI extracts a relationship that doesn't fit the current ontology, that can prompt either an ontology update or be stored as a novel insight. Because the pipeline retains raw data references, human experts or automated validators can give feedback – confirming correct links or flagging mistakes – and these adjustments can feed back into the extraction algorithms. This feedback loop means the knowledge graphs remain high-quality and up-to-date. (Recent research in fact shows that using LLMs to generate graphs and then iteratively refining them can yield dense, useful knowledge graphs from plain text [6](#) [7](#). Semantic Text intends to leverage such techniques under the hood.)
- **Integration with Agentic AI Workflows:** One motivating use-case for Semantic Text is to support **AI agents** – autonomous software agents that perform tasks using AI. These agents (for example, an AI assistant that reads incoming emails and takes actions) benefit greatly from having a structured understanding of data. Semantic Text can feed agents with a knowledge graph representation of their context, so the agent can reason over a connected worldview rather than just raw text. For instance, instead of an agent blindly scanning an email for keywords, it could use a graph that tells it "Email X is from Person Y who is a client about Project Z which has Status=Delayed" – allowing much smarter and context-aware actions. By bridging unstructured inputs into structured knowledge, the service becomes the connective tissue for agentic workflows.
- **Scalable, Serverless Architecture:** The service is designed to be cloud-native and serverless. That means it can scale automatically with demand – processing large document batches or high-volume streams in parallel when needed, but also scaling down to zero when idle (incurring minimal cost). Each customer's data can be handled in isolation (tenant isolation at the cell or container level), providing security. Because it's serverless, organizations don't need to install or maintain complex infrastructure; they simply call the API when they need it. The underlying open-source stack has been optimized for cost-efficiency and scalability, proven out by earlier projects. Compute and storage scale with usage and can be tied directly to customer consumption, which keeps the service cost-effective [8](#) [9](#).
- **Open-Source Core (No Lock-In):** A pivotal aspect is that the core technology powering Semantic Text is open-source. The algorithms for extraction, transformation, and graph management, and the supporting frameworks (serverless functions, etc.) are part of an open-source project. This means users have transparency into how their data is processed and the option to self-host if they desire. Unlike some commercial knowledge graph offerings that are proprietary black boxes, Semantic Text prioritizes openness. Clients are not locked in – they can adopt the service knowing their semantic data could be exported or the solution could be run in their own cloud environment if needed. This open model accelerates innovation and trust: as improvements are made by a community, they can quickly benefit all users. It's a differentiator against closed systems and encourages industry collaboration on schema standards rather than each vendor reinventing schemas. (Schema definitions or ontologies developed can optionally be shared in an **open marketplace**, see below.)

In summary, the Semantic Text service offers a pipeline to **inject understanding** into your data. By producing a knowledge graph layer, it enables a host of intelligent functions on top of previously unstructured information. We next explore those use cases and the value they unlock.

## Use Cases and Applications

Because Semantic Text essentially acts as a foundational layer for making sense of data, it has a wide range of use cases across industries. Below are some primary applications and examples of how this service can add value:

- **Personalized Content Feeds:** One immediate use is curating personalized news or information feeds. By generating a semantic graph of articles and matching them to a user's interest graph, the service can power highly tailored content recommendations. For example, in cybersecurity news, a service like MyFeeds.ai (an early implementation of this concept) built graphs for both incoming articles and a user's persona, then matched them to deliver uncannily relevant daily briefs [10](#) [11](#). Instead of keyword filtering, the graph approach "connects the dots" – if a CISO cares about "cloud security" and "zero-day exploits," the system will surface articles that semantically intersect those topics even if specific keywords differ. This goes beyond what RSS or keyword alerts do. It saves users time and ensures they don't miss critical information. Many professionals drowning in information could benefit from such AI-curated feeds.
- **Intelligent Content Filtering and Moderation:** Another use case is real-time filtering of content (web pages, documents, messages) to enforce policies or preferences. Traditional filters use simple keywords or blacklists, which often fail to catch nuances – or conversely, over-block because they lack context. By contrast, a semantic graph of a web page can understand the context of content. For instance, instead of blocking every page that contains the word "weapon," a semantic filter can discern whether a page is actually describing violent content or maybe talking about "weaponizing AI" as a metaphor [12](#). One of our affiliated projects prototyped exactly this: a **Semantic Content Filter** that acts as a smart proxy, building a knowledge graph of each page on-the-fly and then applying user-defined rules at a semantic level (e.g. "allow news site *but remove political news content*") [13](#) [14](#). Semantic Text could provide the core analysis engine for such filters. Enterprises could use it to enforce compliance (e.g. no leaks of personal data in documents), parents could use it for safety (block truly adult content while allowing benign uses of certain words), etc. The key is the filter makes decisions with an **understanding** of the content, not just pattern matching.
- **Enterprise Knowledge Integration and Search:** Organizations often have data silos – documents, databases, emails – that are not connected, making search and analytics difficult. Building an **enterprise knowledge graph** with Semantic Text can unify these silos. The service can ingest everything from corporate wikis to support tickets, turning each into a slice of a giant interconnected graph. This allows, for example, a query like "show me all projects involving Alice that relate to Machine Learning in the past year" to be answered in seconds, even if that information was scattered across various reports and emails. Tech companies like Stardog emphasize that knowledge graphs naturally link diverse data and can capture the context often lost in relational databases [15](#) [16](#). By integrating data sources in a graph (with each source's schema mapped into a common semantic layer), companies gain a **360° view of their knowledge**. This improves enterprise search – employees can find answers faster – and supports advanced analytics like fraud detection (linking people, accounts, transactions) or R&D knowledge management (connecting research papers, experiments, patents in a knowledge network). Semantic Text could be offered as a managed service to build and maintain such

enterprise graphs, complementing existing data warehouse or data lake investments with a semantic layer.

- **AI-Powered Assistants and Agents:** As mentioned, one exciting application is feeding **agentic AI workflows**. Consider an AI assistant that helps prepare business reports. It might use Semantic Text to parse raw data (like technical reports, spreadsheets, incident logs) into a structured form, then reason about it or even generate summaries for a given audience. In our ecosystem, *The Cyber Boardroom* tool does something similar – translating technical cybersecurity information into board-level language by applying persona-based transformations on top of a semantic representation <sup>17</sup> <sup>18</sup>. The knowledge graph here serves as an intermediary between the raw facts and the final narrative, ensuring that the assistant's outputs stay grounded in the source data. More generally, any multi-step AI **agent** (for customer support, workflow automation, etc.) can use knowledge graphs as a memory and context. Agents can query the graph to recall facts ("Has this issue happened before? Who worked on it?") and update the graph with new events as they execute tasks. This approach yields an auditable memory – one can inspect the graph to see what the agent has assumed or learned. The Semantic Text service could become the go-to knowledge store for such AI agents, which is a direction even other vendors are exploring (for example, Squirro's platform combines knowledge graphs with retrieval-augmented generation to enable "auditable agentic AI" workflows <sup>19</sup> <sup>20</sup> ).
- **Domain Schema Development and Data Classification:** Semantic Text can also be used in a consulting or tool-support capacity for **developing ontologies and taxonomies** in the first place. Many industries (healthcare, finance, legal, etc.) are now keen to model their domain knowledge formally (ontologies for clinical terms, taxonomies of risk types, etc.) but lack easy ways to do so. The service can assist by analyzing a corpus of domain documents and suggesting a draft ontology – essentially surfacing the key entities and how they tend to relate, which domain experts can then refine. It's even possible to have **ontologies of ontologies** – for example, mapping a healthcare ontology to an insurance ontology where they overlap (patient -> policyholder, treatment -> claim) to enable cross-domain data interoperability. The service's ability to ingest one taxonomy and map text to it, then do the same with another and find linkages, becomes very useful here. Companies could leverage this to build better data catalogs, glossaries, and knowledge bases. Some enterprise semantic platforms already hint at AI-assisted ontology building (metaphactory's tool, for instance, uses an AI "modeling agent" to recommend classes or reuse opportunities in ontology design <sup>21</sup> <sup>22</sup>). Semantic Text would provide a more automated, bottom-up complement to that – learning the ontology from data first, which can then be curated. There is also a potential marketplace opportunity: **industry-specific knowledge graph schemas** (and even pretrained models) could be developed and offered to others as starting points. For example, a "Cybersecurity Threat Ontology" or a "Financial Regulatory Taxonomy" could be sold as a package, which the Semantic Text service can readily apply to a client's data. These schema packages could be open-source or under specific licenses, but either way, the value to customers is accelerating their setup of a useful knowledge graph.
- **Enhanced Analytics and Decision Support:** Once data is in a semantic graph, advanced analysis is enabled. Graph analytics algorithms can find important nodes (e.g. the most influential customer in a network), detect communities or clusters in data, and even perform link prediction (guessing connections that might exist). For decision makers, having a semantic layer means they can ask high-level questions. For example, a manager might query, "What factors are linked with project delays in our company's history?" and because the knowledge graph connects projects with people, tools, incidents, etc., an analysis could reveal patterns (perhaps many delayed projects involve a specific department or supplier). This is something traditional BI tools struggle with because the data relationships are not explicit. By unlocking these insights,

Semantic Text helps organizations move towards **deterministic, explainable AI** in decision support. It's no coincidence that knowledge graphs are sometimes called *enterprise brains* – they integrate data into a form that is closer to human-like understanding, which can then drive smarter recommendations.

These are just a few representative use cases. Across the board, the common thread is that Semantic Text provides a **semantic backbone** that can power various applications – from content personalization and compliance filtering to enterprise AI and knowledge management. Whenever making sense of complex, unstructured information is the challenge, a semantic graph approach can be part of the solution.

## Service Delivery and Business Model

How would Semantic Text be delivered to customers, and what is the business model behind it? Being conceived in a **serverless-first, "as-a-service"** manner, the offering is very flexible. Here we outline the potential deployment models and revenue streams:

- **Cloud API Service (Multi-tenant):** The primary offering is a multi-tenant cloud service – essentially an API endpoint or set of endpoints that developers or data teams can call on demand. For example, there might be an endpoint `/generateGraph` to which a client POSTs a document (or a URL or text content), and the service returns the semantic graph JSON. Another endpoint might allow querying or simple analytics on the graph (though in many cases clients might download and use the graph in their own systems). This base service would be **usage-priced** – for instance, charging per number of tokens/words processed, or per API call, or per volume of data. This usage-based model aligns cost with actual value consumption and is easily understandable (similar to how one pays for cloud OCR or NLP API usage). The goal is to make it as frictionless as possible: a developer could sign up and within minutes start converting texts to knowledge graphs for a few cents. Given the serverless backend, costs on our side scale with usage, and idle capacity doesn't significantly incur cost <sup>8</sup>, enabling competitive pricing. The API would include security isolation so each client's data is kept private and separate (each request can operate in a separate container or sandbox). **Target users** for this model include startups, SMEs, or individual developers/researchers who need this capability in their projects without wanting to manage infrastructure.
- **Dedicated Deployments (Single-tenant):** For larger enterprise customers or those with heightened privacy requirements, Semantic Text could be offered as a dedicated instance – possibly even deployed within the customer's own cloud environment. Thanks to the open-source, infrastructure-as-code nature of our stack, we can deploy the entire pipeline into an isolated AWS account or VPC for a client, or onto their on-premise servers (if they use AWS Outpost, Azure Stack, etc.). This essentially gives them a private Semantic Text service with no co-mingling of data with others. Naturally, this would come at a premium cost (to cover dedicated resources and management overhead), likely on a subscription basis (e.g. annual license or hosting fee). We could even structure tiers: for instance, a **Standard Cloud** tier (shared environment, data encrypted in multi-tenant stores), a **Dedicated Cloud** tier (your own cluster managed by us in our cloud), and an **On-Premise** tier (we set it up in your environment, possibly with a support contract). The key selling point here is flexibility and trust: enterprises can adopt the solution without fear of lock-in or unauthorized data access, because if needed they have the option to run it themselves. This is a contrast to some vendors that only offer a SaaS black box. Our architecture – being serverless and containerized – makes these options feasible with minimal changes.

- **Customization and Consulting Services:** While the core service is productized, there is business opportunity in helping customers get the most out of it through customization. This includes developing **custom ontologies or schemas** for clients (as discussed in use cases). For example, a pharma company might engage us to create a semantic model for their research data. We could either do this as a professional service or develop it and license the resulting schema/package. Additionally, integrating the Semantic Text service into client workflows might require custom adapters or connectors (e.g. hooking it up to their specific database, or a custom UI for internal users to view the graphs). We can offer integration services through a network of experts (the concept of dynamic, on-demand teams comes into play – we could bring in specialists on contract for a given project, avoiding fixed overhead). Revenue here would be project-based or through a **premium support** model. The advantage of an open-core approach is that third-party integrators or even the clients themselves can also do customizations, but our deep know-how means we could be the preferred experts if needed.
- **Marketplace for Schemas and Models:** As the user base grows, there could be a marketplace or library of crowd-sourced (or partner-sourced) ontologies, taxonomies, and even pretrained extraction models. Some could be free and open, others could be paid. For instance, a consulting firm might develop the definitive “Retail Product Knowledge Graph Schema” and offer it (possibly with data samples) on the marketplace. Clients could one-click deploy that into their Semantic Text instance to bootstrap their solution. We could take a revenue share for any paid schemas sold through our platform. This ecosystem not only creates additional revenue streams but also accelerates adoption (clients in a given industry will find relevant starting points easily). Open-source community contributions would be encouraged for broad use schemas (which also helps establish Semantic Text as a standard in the long run).
- **Usage Analytics and Ongoing Value:** Another aspect of the model is demonstrating value through analytics. Because Semantic Text can track how content is categorized and linked, we can provide dashboards to enterprise customers about their data (e.g. “this month you ingested 10,000 documents and discovered 2,000 unique entities across 5 ontologies; here are the top trending topics...”). These insights can help justify the spend and also identify new opportunities (perhaps the client realizes they should integrate another data source, or invest in curating a certain part of the ontology that appears frequently). Some of these analytics features could be part of a premium tier.

In terms of **pricing strategy**, a combination of usage-based pricing (for the core API) and tiered subscription (for enterprise features or dedicated instances) is likely. The market for enterprise knowledge graph services is growing rapidly – valued around \$1.2 billion in 2024 and projected to reach ~\$1.5 billion in 2025 <sup>23</sup> (with further growth to several billions in the years beyond). This indicates strong willingness of organizations to invest in solutions that help make sense of their data. With an open-core, cloud-flexible offering, Semantic Text can position itself competitively: lower barrier to entry than heavy enterprise software, but with the ability to scale up to enterprise needs.

## Competitive Landscape

It's important to situate Semantic Text in the context of other players and technologies enabling semantic data management and AI. Several categories of competitors or adjacent solutions exist:

- **Traditional Graph Database Platforms:** Companies like Neo4j, AWS (Neptune), and open source projects (TigerGraph, ArangoDB, etc.) provide the graph database infrastructure. However, by themselves they are just storage/query engines – they don't semantically interpret

raw text. Neo4j, for example, has promoted using its graph DB in combination with custom pipelines or LLMs to extract knowledge, but the user must build that. Some have published how-to guides (e.g. using GPT-4 to parse text and then storing results in Neo4j)<sup>24</sup>, which actually validates our concept but isn't a turnkey service. In contrast, Semantic Text is a higher-level solution: it would likely use a graph database under the hood, but the value is in the *automated extraction and ontology mapping*, not just storage.

- **Enterprise Knowledge Graph & Semantic Suite Vendors:** This includes tools such as **Stardog**, **PoolParty Semantic Suite**, **TopQuadrant's TopBraid**, **Metaphacts (metaphactory)**, and **Franz AllegroGraph**, among others. These vendors have been around in the semantic technology space, typically offering software to build and manage knowledge graphs for enterprises. For example, PoolParty helps organizations manage taxonomies and then auto-classify documents to build an enterprise knowledge graph<sup>25</sup>. Stardog offers a knowledge graph platform that emphasizes data integration and reasoning, and it has highlighted how knowledge graphs paired with AI can improve search and recommendations<sup>5</sup>. Metaphacts focuses on collaborative ontology modeling with some AI assistance<sup>21</sup>. **How Semantic Text differentiates:** First, many of these solutions are largely enterprise software (on-prem or heavy SaaS) that require significant setup and expertise – whereas Semantic Text aims to be a light, serverless service accessible via API. Second, our open-source core and open standards focus means lower risk of lock-in and more extensibility (metaphactory also touts use of W3C standards to avoid lock-in<sup>1</sup>, so we share that philosophy). Third, we emphasize the dynamic, bottom-up creation of graphs from text using LLMs. The legacy players started more from an ontology-management perspective (top-down). We meet in the middle: we can do top-down (if you have an ontology) or bottom-up (discover one from data) or both. In terms of cost, we intend to offer pay-as-you-go cloud pricing, whereas many enterprise vendors have large upfront license fees or require buying a whole platform. That could be disruptive by significantly lowering the entry cost for advanced semantic capabilities.
- **GenAI/LLM-Oriented Startups:** With the rise of ChatGPT and enterprise LLM adoption, a new wave of startups is combining LLMs with knowledge graphs. One example is **Squirro**, which markets an “Enterprise AI platform” using knowledge graphs plus retrieval-augmented generation (what they call GraphRAG)<sup>20</sup>. Squirro’s approach involves defining ontologies for a domain and automatically tagging and classifying incoming content to populate the knowledge graph<sup>26</sup> <sup>27</sup> – very much the problem space we’re in. This validation of the approach is good news: it means customers are recognizing that LLMs need a knowledge graph backbone for accuracy and context. Our competitive edge here again comes down to openness and flexibility. Many of these platforms (including Squirro) are proprietary. Semantic Text could carve out a niche as the open, transparent alternative that can plug into existing open-source ecosystems. Also, our design as a microservice/API can allow others to integrate us into their pipelines. For instance, an enterprise could use Semantic Text alongside an OpenAI service or a LangChain pipeline – it’s not a monolithic platform, it’s a focused service. This composability is attractive to developers who want best-of-breed components.
- **In-House and DIY Solutions:** Some large organizations with AI expertise are building their own semantic extractors and internal knowledge graphs (often using open-source NLP libraries and custom code). The question for them will be: continue maintaining a bespoke system or use a service like Semantic Text? If we ensure the service is developer-friendly, reliable, and cost-effective, it may be an easy choice to switch. Additionally, because our core is open source, even the DIY folks could contribute or adapt it – which is fine, as it grows the ecosystem and might convert them to customers for hosted solutions later (this open-core model has worked well for companies like Red Hat in the past). Our aim is to become a de-facto standard toolkit for

semantic graph workflows, whether someone uses the hosted service or runs the libraries themselves.

To summarize the landscape: **the concept of semantic knowledge graphs is gaining renewed importance** in the age of AI. Many competitors acknowledge that a knowledge graph can greatly improve AI's accuracy and enterprise data integration. The market is currently a mix of older semantic tech companies and newer AI startups. Semantic Text can position itself by bridging those worlds – offering the robustness and standards compliance of semantic web tech, with the agility and AI-native approach of the new generation. As an open service, it can also potentially partner with others (for example, integrating with a vector database or an BI dashboard tool) rather than trying to own the entire stack.

The market size for these solutions is on an upward trajectory as noted, and importantly, **the trend** is that organizations want **trusted, explainable AI**. Knowledge graphs provide a layer of determinism and auditability on top of black-box AI <sup>19</sup> <sub>5</sub>. That macro trend benefits any player in this domain. We will just need to execute well and stay innovative (for instance, keeping up with research like GraphRAG, integrating new LLM advances, etc., which our R&D will actively do).

## Conclusion and Next Steps

Semantic Text as a service represents a convergence of open-source innovation and practical enterprise needs. By turning unstructured data into semantic knowledge graphs, it enables a level of understanding and interoperability that was previously hard to attain without large specialized teams. The time is ripe for this approach: organizations are swimming in data and investing in AI, but they struggle with AI that *truly knows* their domain. A semantic graph fills that gap by serving as a **single source of truth for knowledge**, one that both humans and machines can traverse for insight.

The epiphany driving Semantic Text is that **compressing data into a graph of meanings** unlocks myriad capabilities – from smarter search to AI agents that can reason – and that this can be achieved not by forcing one worldview, but by connecting many worldviews. We don't throw away existing taxonomies; we link them. We don't replace human expertise; we augment it with machine-curated suggestions. We also don't ask customers to rip out their data pipelines; we provide a modular service that slips into their workflow and starts adding value immediately.

In developing this service, the next steps would be: 1. **Refine the Core Algorithms** – likely by leveraging state-of-the-art LLM extraction techniques (such as the KGGen approach in research) to improve the quality of the generated graphs <sup>28</sup> <sub>7</sub>. Ensuring high precision and recall in entity/relation extraction is crucial (so that the graphs are accurate mirrors of the source content). 2. **Hardening and Scaling** – turning the current prototype (which already processes feeds and documents for our internal projects) into a robust public API. This involves adding security layers, monitoring, caching mechanisms for high throughput (we learned from the content filter project the need for caching and speed when analyzing content on the fly). 3. **User Interface & Integrations** – though many will use the API, having a simple UI (even if just for demonstration) where a user can drop in a piece of text and see the resulting knowledge graph visualization, would greatly aid in evangelizing the service. Also, building integrations (for example, a plugin for common platforms like Slack or Confluence, which sends content to Semantic Text and returns annotations) could drive adoption. 4. **Community and Schema Library** – kick-starting the open-source community around this, perhaps by releasing some ready-made ontologies and inviting domain experts to contribute enhancements. The more community adoption, the more organic growth for the service's knowledge base.

In conclusion, Semantic Text aims to become the go-to “**brain**” in the cloud that any application or organization can tap into to make sense of textual data. With the combination of an open-core philosophy, a flexible semantic approach, and alignment with the needs of modern AI (explainability, reliability 29 5 ), we believe this service can catalyze a new wave of intelligent applications. The vision is a future where instead of drowning in documents and chats, companies and individuals will have a rich knowledge graph at their fingertips – answering questions, driving workflows, and continually learning. Semantic Text is the service to build that future on.

---

1 21 22 Semantic Knowledge Modeling

<https://metaphacts.com/solutions/semantic-knowledge-modeling>

2 4 5 15 16 What is a Knowledge Graph | Stardog

<https://www.stardog.com/knowledge-graph/>

3 8 9 10 11 12 13 14 17 18 Dinis Cruz's Multi-Startup Strategy\_ Open-Source Innovation Across Four Synergistic Ventures.pdf

[file:///file\\_0000000a86c71f49f64dfb7ff917724](file:///file_0000000a86c71f49f64dfb7ff917724)

6 7 28 KGGen: Extracting Knowledge Graphs from Plain Text with Language Models

<https://arxiv.org/html/2502.09956v1>

19 20 26 27 Knowledge Graphs For Deterministic AI

<https://squirro.com/knowledge-graphs>

23 Enterprise Knowledge Graph Market Overview And Analysis Report 2025

<https://www.thebusinessresearchcompany.com/report/enterprise-knowledge-graph-global-market-report>

24 How to Convert Unstructured Text to Knowledge Graphs Using LLMs

<https://neo4j.com/blog/developer/unstructured-text-to-knowledge-graph/>

25 PoolParty Semantic Suite - Your Complete Semantic Platform

<https://www.poolparty.biz/>

29 Why New AI Tools Like ChatGPT Need Knowledge Graphs | TopQuadrant

<https://www.topquadrant.com/resources/why-new-ai-tools-like-chatgpt-need-knowledge-graphs/>