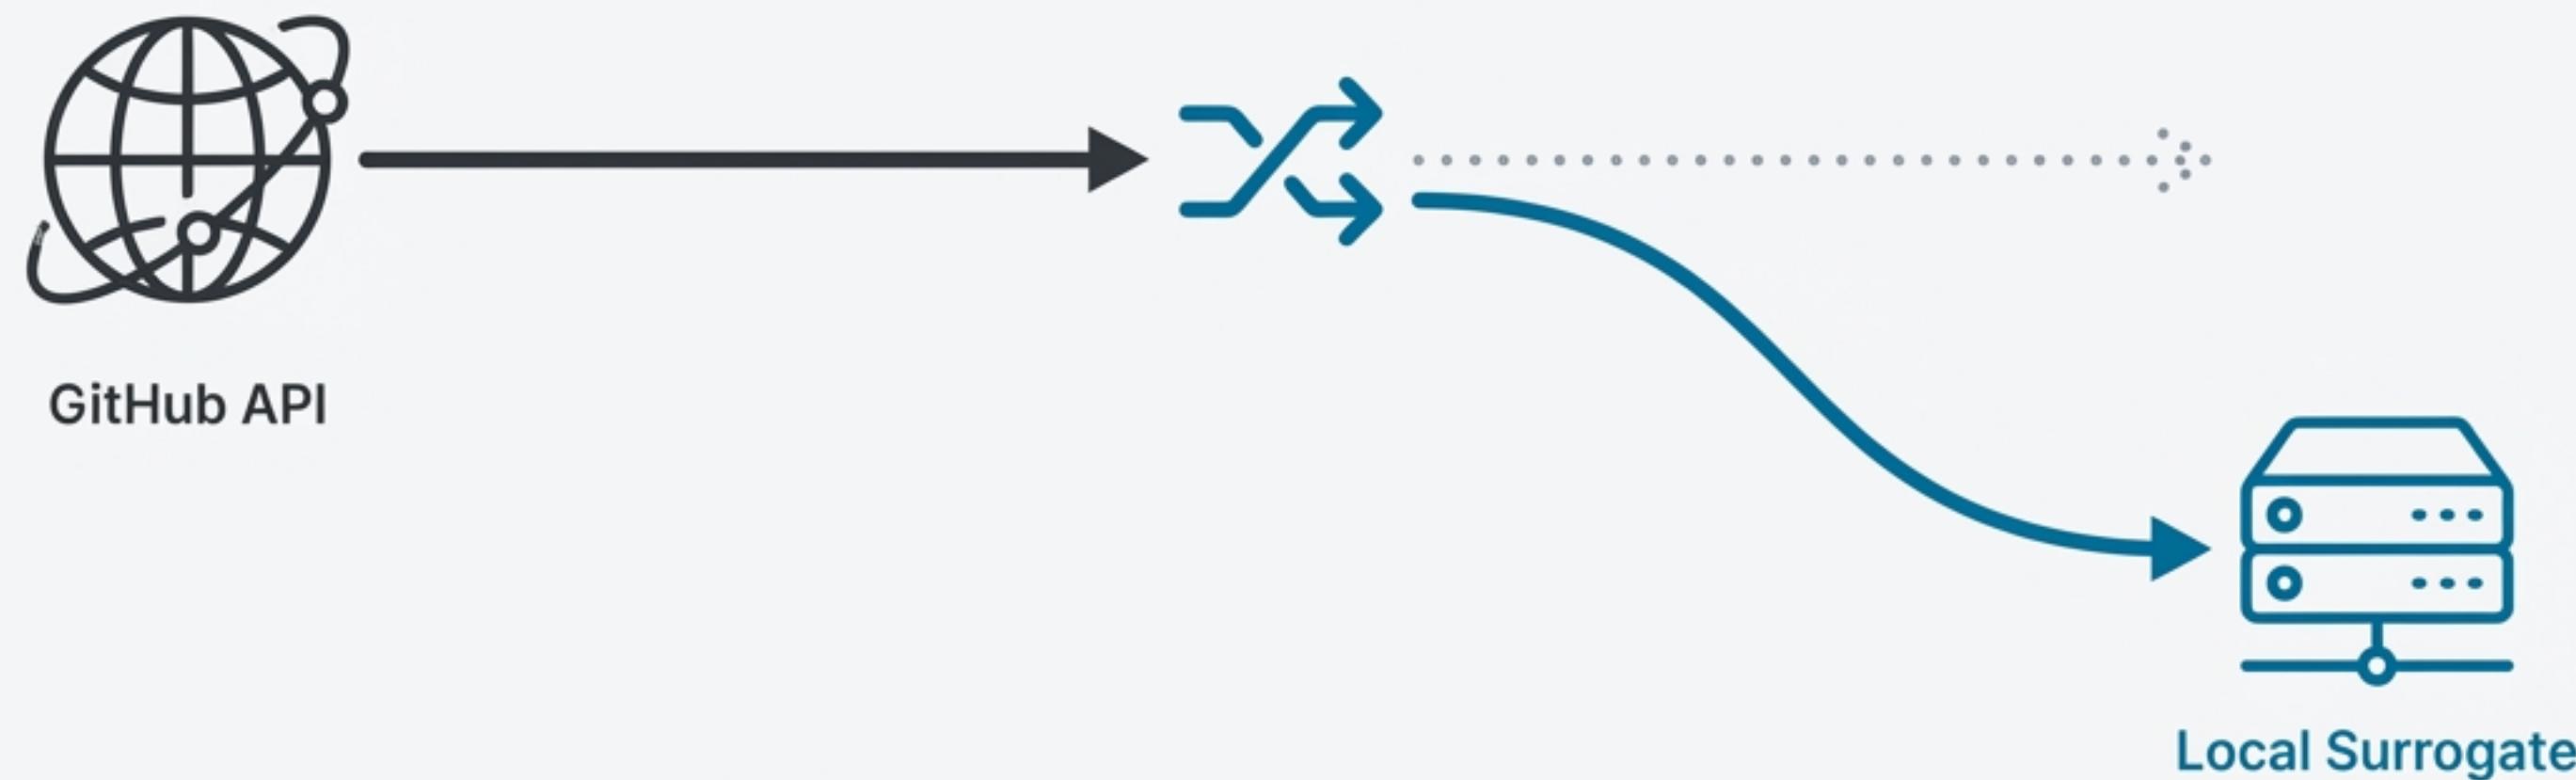


The GitHub API Surrogate

A Case Study in Transforming Our Test Suite



We made our integration tests 20x faster.

BEFORE

20s



AFTER

1.08s



20x Speedup

By replacing live GitHub API calls with a local, in-memory surrogate, we reduced the runtime for 553 tests from over 20 seconds to just 1 second, with zero changes to the core test logic.

Testing against the live GitHub API was slow, flaky, and complex.



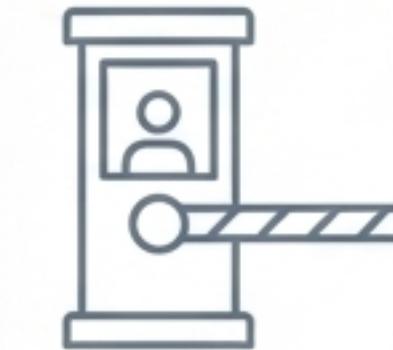
Speed

Each HTTP request takes 200-800ms, making test suites painfully slow.



Authentication

Required valid Personal Access Tokens (PATs), complicating CI/CD pipelines.



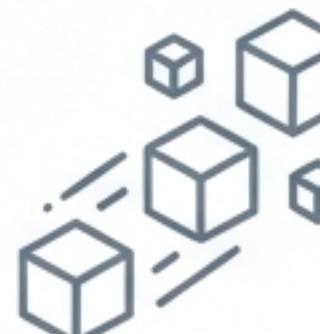
Rate Limits

Hitting the 5,000 requests/hour limit could block development and CI runs.



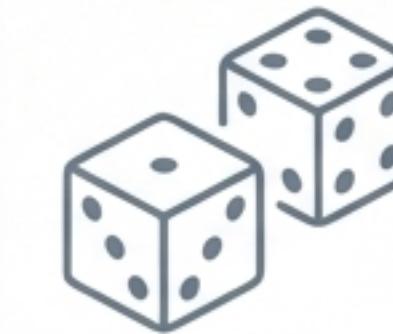
Network Dependency

Tests failed when offline or during GitHub service degradation.



Side Effects

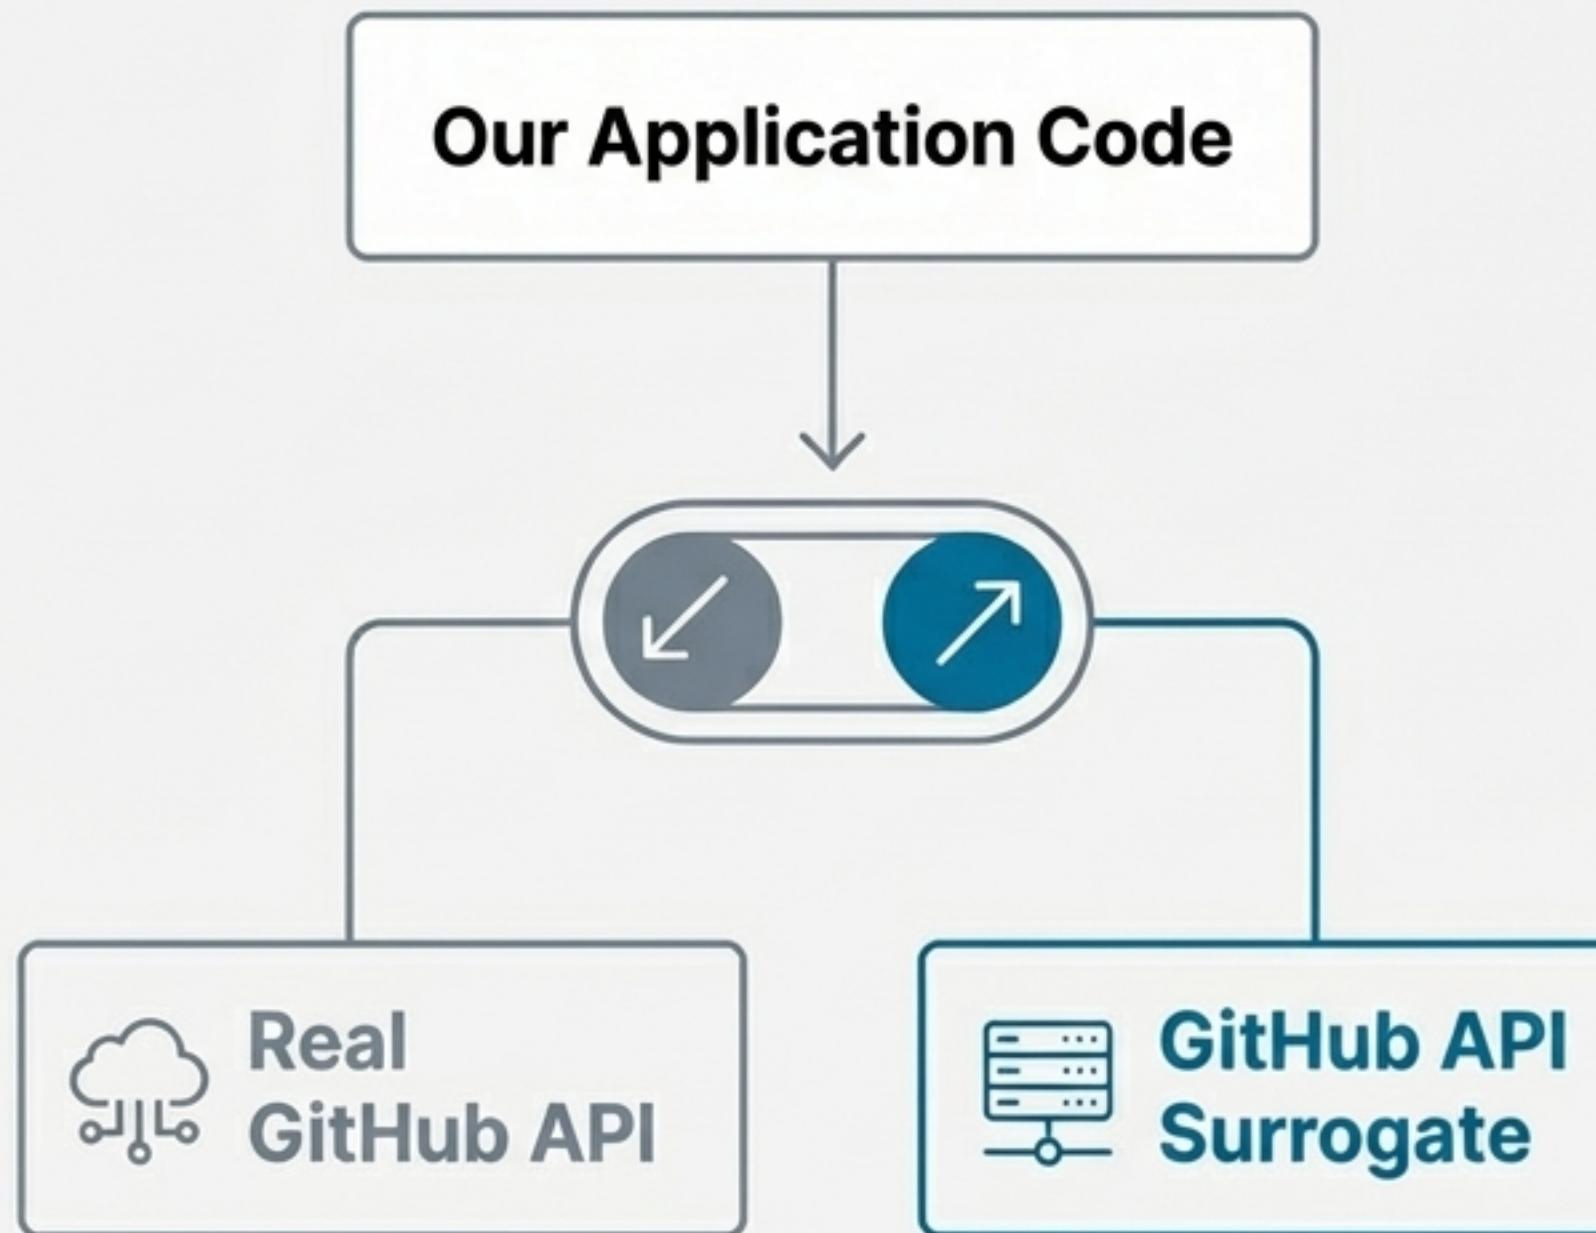
Tests created real resources that required complex cleanup routines.



Non-Determinism

Variable network latency and API responses caused flaky, unreliable tests.

Our solution: A local, in-memory surrogate of the GitHub API.



- A **FastAPI application** that mirrors the GitHub REST API endpoints our code uses.
- Runs **entirely in-memory** with no network calls or disk I/O.
- Maintains its own state for repositories, secrets, and organisations for the duration of the test run.
- Returns responses that precisely match GitHub's real API format, including error conditions.

The transformation is absolute: from network-dependent to self-contained.

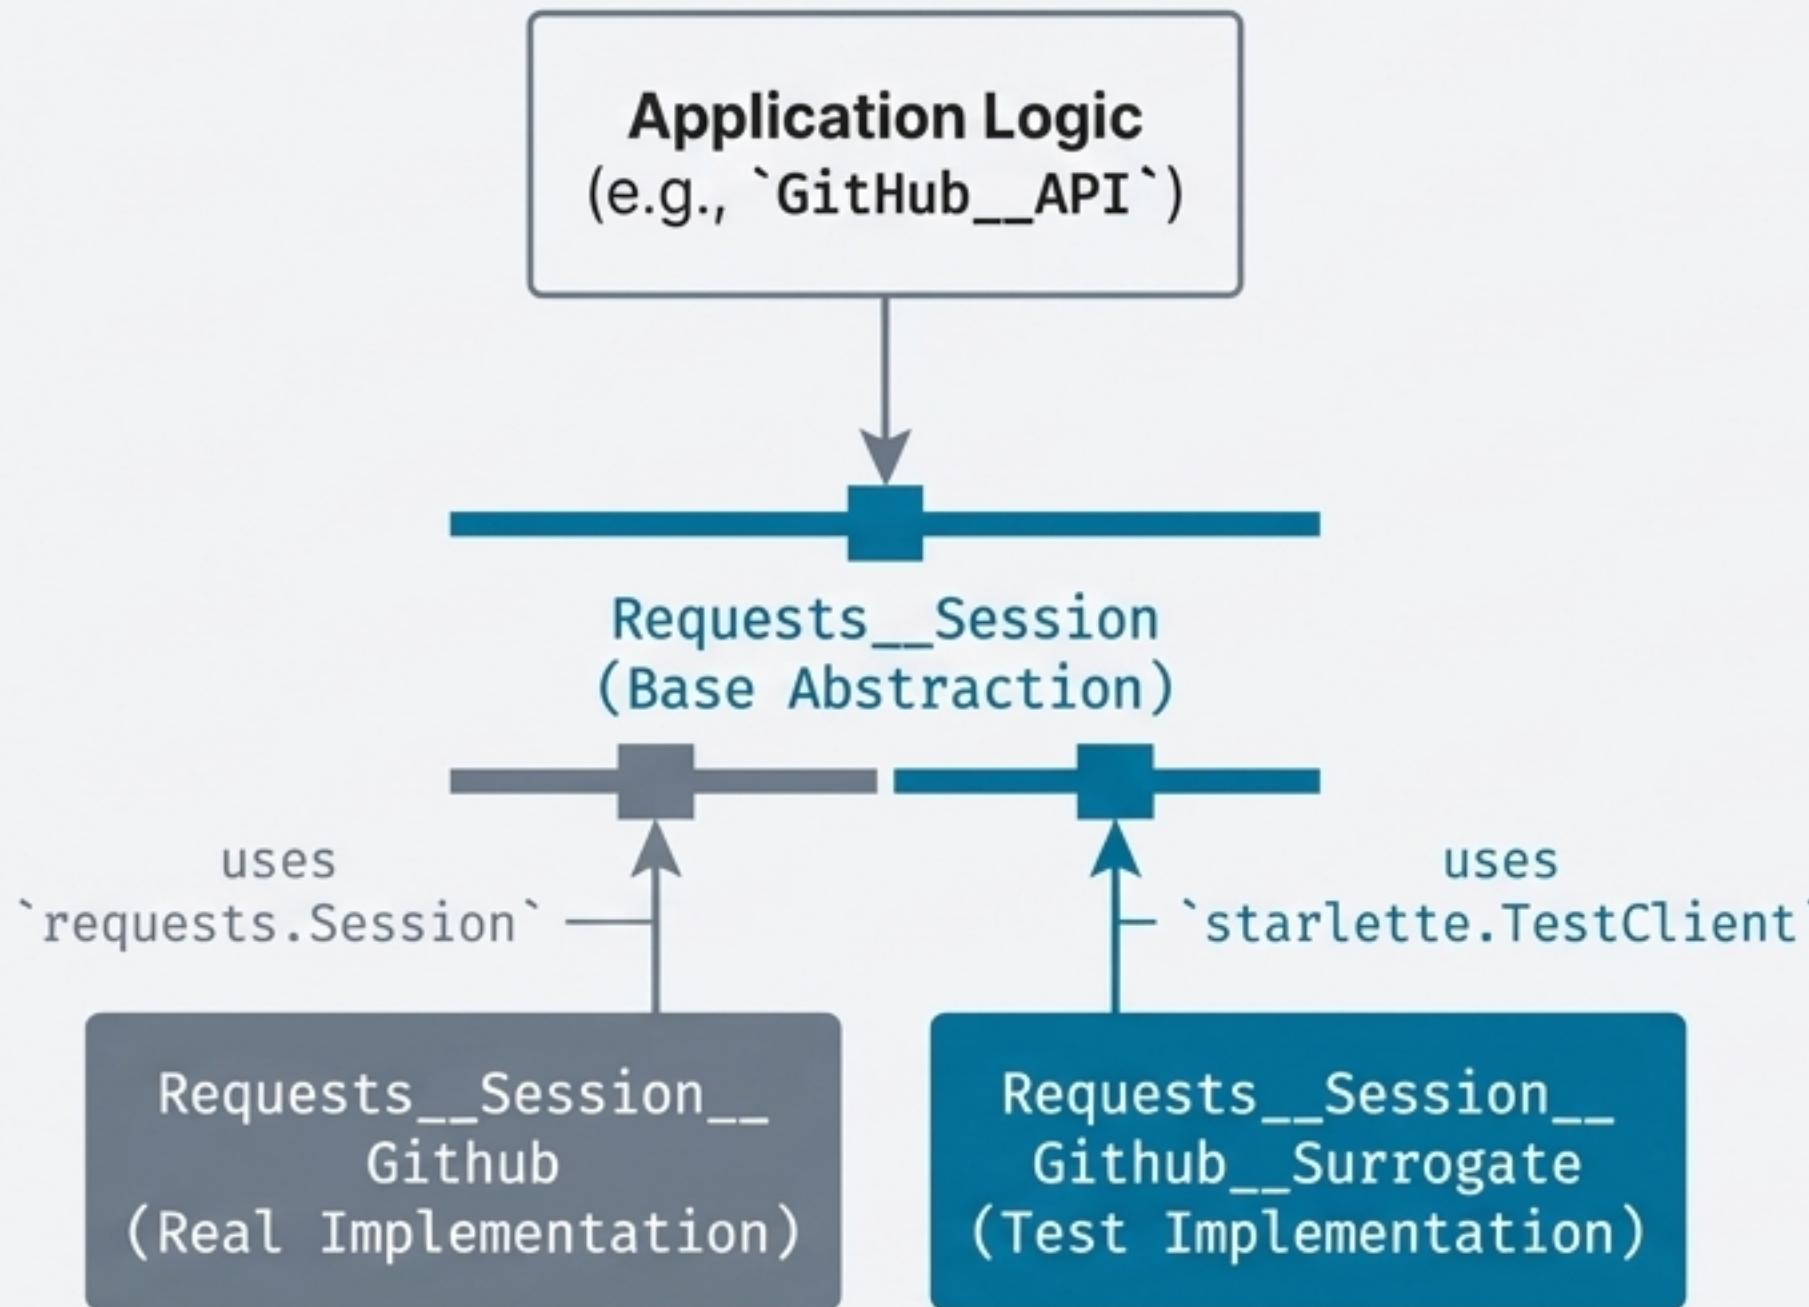
Metric	Before (Real API)	After (Surrogate)
Total test time	~20 seconds	1.08 seconds
Network required	Yes (✗)	No (✓)
GitHub PAT required	Yes (✗)	No (✓)
Rate limit concerns	Yes (✗)	No (✓)
Test reliability	Non-deterministic	100% Deterministic

The most I/O-heavy tests saw speedups of over 200x.

Test Class	Before (Real API)	After (Surrogate)	Speedup
test_GitHub_Secrets	18.5 sec	89 ms	208x
test_Routes_GitHub_Secrets_Env	6.3 sec	52 ms	121x
test_Routes_GitHub_Secrets_Org	6.5 sec	58 ms	112x
test_Routes_GitHub_Secrets_Repo	5.3 sec	64 ms	83x

*553 tests passed using the surrogate. 5 tests were skipped for endpoints not yet implemented.

The key was a clean abstraction for the HTTP session.



We introduced a minimal
'Requests_Session' interface.

Production code uses the real
implementation that makes network
calls.

In tests, we transparently swap in the
surrogate implementation that routes
calls to the local FastAPI app.

The application code is completely
unaware of the difference.

The surrogate implements the exact API surface our application needs.

Key Components

GitHub__API__Surrogate: The FastAPI application that mimics the API.

GitHub__API__Surrogate__Test_Context: A simple context manager for easy test setup and teardown.

In-Memory State: Manages test data (repos, secrets, etc.) for isolated runs.

Implemented Endpoints

Endpoint	Methods
'/user'	GET
'/rate_limit'	GET
'/repos/{...}/actions/secrets/public-key'	GET
'/repos/{...}/actions/secrets/{name}'	GET PUT DELETE
'/repos/{...}/environments/{...}/secrets'	GET PUT DELETE
'/orgs/{...}/actions/secrets'	GET PUT DELETE

Migrating a test requires only a few lines of setup code.

Before

```
1 # Before
2 class TestSomething(TestCase):
3     def setUp(self):
4         # Complex setup with env vars
5         load_dotenv()
6         self.pat = get_env("GITHUB_PAT")
7         self.api = GitHubAPI(self.pat)
8         # Cleanup code in tearDown
```

After

```
1 # After
2 class TestSomething(TestCase):
3     @classmethod
4     def setUpClass(cls):
5         cls.surrogate_context = GitHub__API__S
6         cls.api = cls.surrogate_context.github
7         # Add test data programmatically
8         cls.surrogate_context.add_repo(...)
9
10    @classmethod
11    def tearDownClass(cls):
12        cls.surrogate_context.teardown()
```

The core test logic and assertions remain completely unchanged.

The impact on the developer workflow is immediate.



Run the full suite in 1 second

Get instant feedback before every commit. No more skipping tests locally.



Work completely offline

Code and test on a plane, in a cafe, or during a network outage.



No more secret management

Eliminate `*.env` files and the need to provision PATs for local development.



Fully deterministic tests

The same test code produces the same result, every single time.



Easier debugging

Step directly from your test code into the surrogate's Python code to inspect state.

Our CI/CD pipeline is now faster, more secure, and more reliable.



Faster Builds

Test execution time in CI is reduced by **20x**.



Improved Security

No need to store and manage GitHub PATs as secrets in the CI environment.



Eliminated Flakiness

Network-related test failures are a thing of the past.



No Rate Limiting

Run as many parallel builds as needed without fear of being blocked by the GitHub API.



Perfectly Isolated

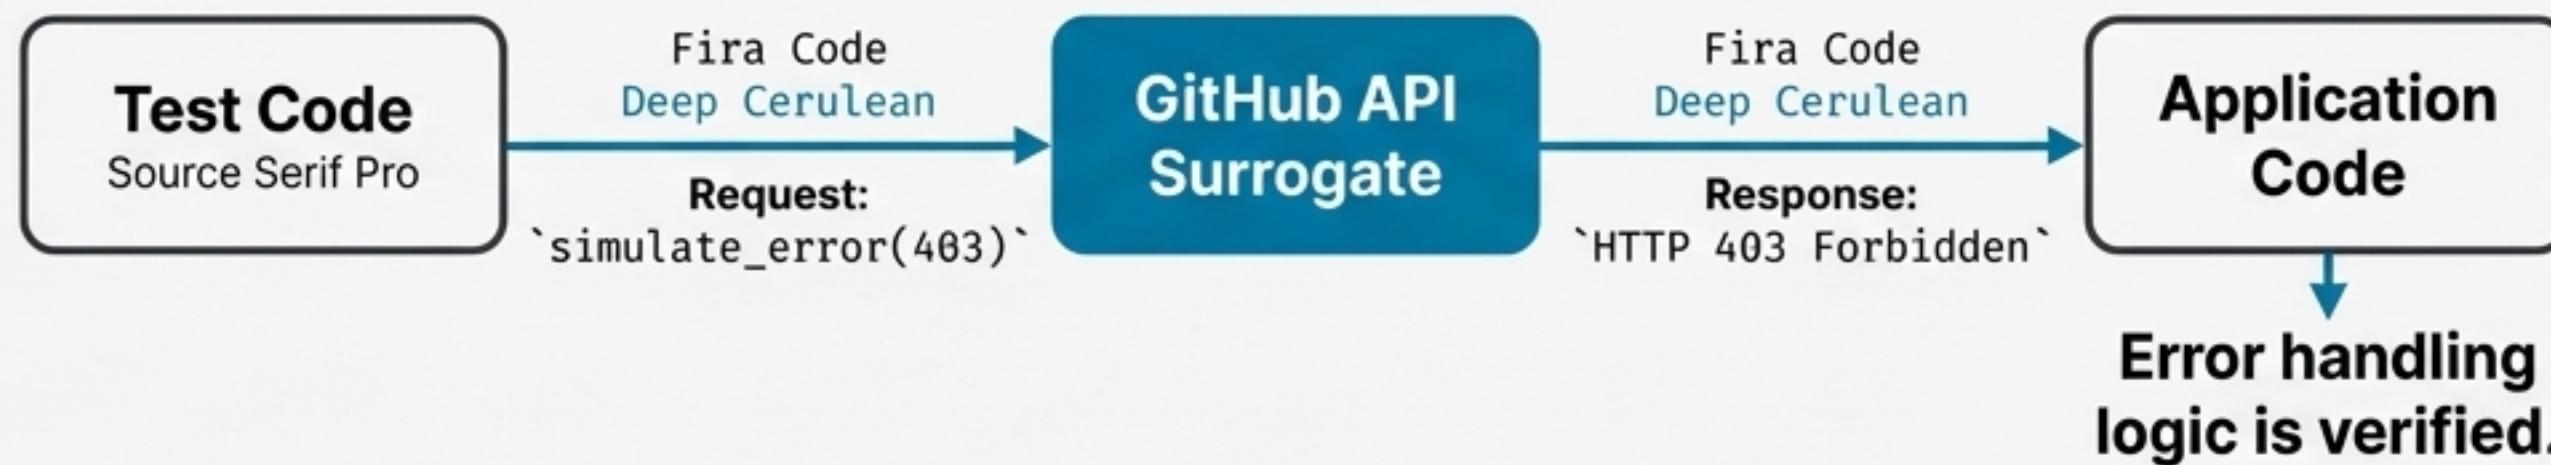
Each parallel test run gets its own isolated, in-memory state.

We can now easily test edge cases and error handling.

The surrogate allows us to simulate failure conditions that are difficult or impossible to create with the live API.

 **Authentication Failures:** What happens when a token is invalid (401 Unauthorized) or lacks permissions (403 Forbidden)?

 **Rate Limiting:** Verify that our application handles 429 Too Many Requests gracefully.



 **Missing Resources:** How does the code handle a 404 Not Found response for a repository or secret?

 **Invalid Inputs:** Test server-side validation errors (422 Unprocessable Entity).

A good abstraction is the ultimate enabler for testing.

By introducing the `Requests__Session` interface, we created a “seam” where production and test implementations can be swapped transparently. The production code has no idea it isn’t talking to the real GitHub API.

Total Implementation Effort: ~4 hours

Ongoing Maintenance: Minimal (add new routes as needed)

Benefit: Immediate and compounding with every single test run by every developer.

This pattern transformed our testing experience from slow and flaky to fast, reliable, and self-contained.