



Quality by Evolution: How Good Design Emerges in Software

Introduction: Quality, Simplicity, and Emergent Properties

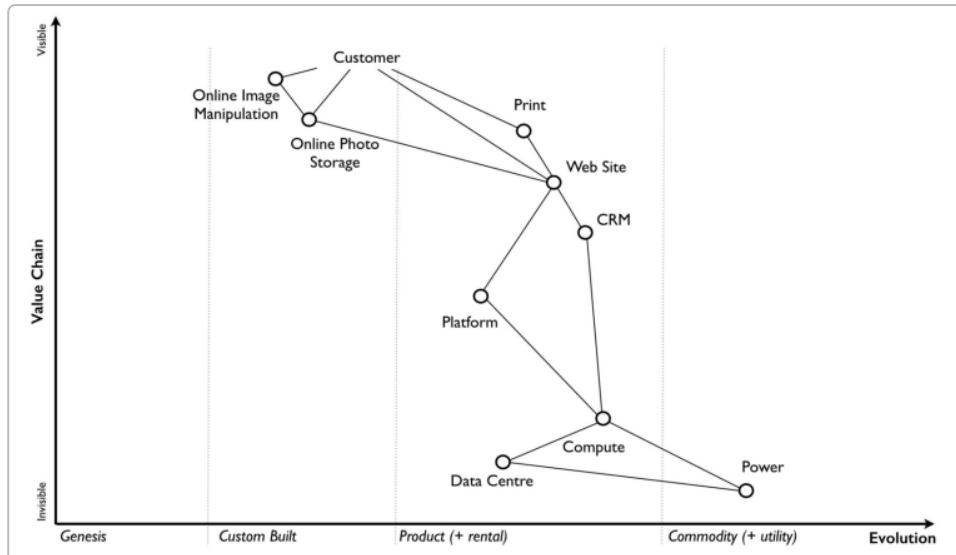
In software development, **quality** and **simplicity** are often seen as hallmarks of good design. But what exactly do we mean by these terms, and how do they relate? Is *simplicity* a subset of *quality*, or vice versa? In practice, high quality software tends to achieve a form of simplicity – not necessarily a simplistic form with few features, but an *elegant* simplicity where everything included serves a purpose and nothing is extraneous. As the famed designer Antoine de Saint-Exupéry put it: “*A designer knows he has achieved perfection not when there is nothing left to add, but when there is nothing left to take away.*” ¹. In other words, a *simple* design often emerges as the end result of a relentless focus on quality and refinement. Quality in this context isn’t merely about avoiding bugs; it encompasses clear structure, maintainability, intuitiveness for the user, and yes, an absence of needless complexity.

Emergent property is a key concept here. *Quality*, especially in complex systems like software, is often an **emergent property** – something that *evolves and crystallizes over time* rather than being fully present from the start. We typically don’t achieve perfect quality in one go; instead, we iterate and improve. In agile software practices, this is well-recognized: “*The best architectures, requirements, and designs emerge from self-organizing teams.*” ². Rather than attempting a complete, “perfect” design upfront, teams start with a workable solution and continuously refine it. Through feedback, refactoring, and iterative enhancements, a high-quality design *emerges*.

This paper explores the idea that **quality is an evolution** – a journey rather than a one-time event – and how simplicity and good design emerge from iterative development. We’ll use analogies like Wardley Mapping’s evolutionary stages, examine why *initially messy prototypes can masquerade as finished products*, and discuss how iterative refactoring and user feedback lead to truly robust, simple, and invisible designs.

Wardley Maps and the Evolution of Software

To frame the evolutionary nature of quality, it helps to borrow the lens of **Wardley Mapping**, a strategic approach introduced by Simon Wardley for visualizing how components evolve ³. Wardley Maps classify the maturity of components along an axis from **Genesis** (novel, chaotic, custom-built solutions) to **Commodity** (standardized, stable utilities) ³ ⁴. In between lie stages often labeled **Custom-Built** (bespoke but starting to solidify) and **Product** (commercially off-the-shelf or standardized offerings) ⁴. Wardley also describes corresponding team personas: **Pioneers** (or **explorers**) excel at the Genesis stage creating new ideas, **Settlers** (the builders or **villagers**) refine these into reliable products, and **Town Planners** industrialize and commoditize solutions for widespread use ⁵.



An example Wardley Map of a software value chain. Components (nodes) evolve from left (Genesis & Custom-built, novel and experimental) to right (Product and Commodity, standardized and stable). The x-axis shows Evolution (stage of maturity) and the y-axis shows visibility in the user's value chain. In software, different parts of a system may sit at different stages of evolution simultaneously. ⁴

In the context of software **code and architecture**, we can think of a brand-new project or feature as starting in a *Genesis* phase – the code is highly customized, likely rough around the edges, and full of unknowns. Over time, with work and refinement, parts of it should move toward *Product* (well-structured, repeatable solutions) and possibly *Commodity* (entirely standardized components or services). The *evolution is not uniform*: one module might become very stable and reusable (commodity-like) while another part is still experimental.

A peculiar property of software is that **it can be at one evolutionary stage internally while appearing to be at another stage externally**. In other words, a software system that is essentially a fragile “*Genesis*” prototype *can present itself like a polished commodity product*. For example, consider a scenario many teams know too well: under pressure to meet a tight deadline, a development team hurriedly cobbles together a prototype-quality implementation with lots of technical debt. Yet, because it has a slick UI and “it runs,” it gets deployed to users as if it were a finished product. To the end-user (and often to management), the software “*looks good* and *seems solid* – *it feels like a mature product or utility*. However, under the hood the foundation is brittle, the architecture ad-hoc, and scalability or maintainability is in doubt. This illusion – **a prototype masquerading as a product** – is a common pitfall in our industry. In fact, Foote and Yoder’s classic “Big Ball of Mud” study observes that many throw-away prototypes have a way of becoming permanent: “*Once the prototype is done... As the time nears to demonstrate the prototype, the temptation to load it with impressive but inefficient functionality can be hard to resist. Sometimes this strategy is too successful – the client, rather than funding the next phase of the project, may slate the prototype itself for release.*” ⁶. In other words, a **“permanent prototype”** is born – a system that *behaves like a product but lacks the quality of a truly evolved product*.

Why does this happen? One reason is that software, unlike physical materials, doesn’t visibly decay or collapse when its internals are shoddy – at least not immediately. A hastily built house would quickly show defects or be unsafe, but a hastily written codebase might run fine for a while, concealing the accumulating **technical debt**. It’s only later, when we try to extend or maintain that code, that the lack of quality becomes painfully apparent (bugs, high change costs, poor performance, etc.). This is why **Wardley Mapping and technical debt** are often discussed together – mapping can highlight if you’re treating something as a “Product/Commodity” that is actually still essentially “Custom/Genesis” under

the hood ⁷ ⁸. A system stuck in a lower evolutionary stage but forced to operate at higher scale or expectations will incur heavy maintenance costs or even fail unexpectedly.

The key lesson is that *true quality requires guiding a system through these evolutionary stages properly*. If we skip the hard work of evolution (e.g. polishing a prototype into a robust product), we end up with a fragile architecture. Thus, quality **design and architecture is something that must emerge through deliberate evolution of the software** – you can't simply slap a “quality” label on a prototype and call it a day. It takes an ongoing investment to move a solution from the exploratory genesis phase into a hardened, simple, commodity-like service.

Emergent Design Through Iteration and Refactoring

How, then, do we *successfully evolve* a software system to high quality? The answer lies in an iterative, feedback-driven development process – one where **continuous refactoring and improvement** are built into the routine. High quality code *starts messy* (as an experiment) and incrementally becomes clean and simple through relentless refinement. In agile methodologies and *Extreme Programming (XP)*, this is a core philosophy: developers implement functionality in the simplest way that works, then refactor mercilessly to improve the design once tests are passing ⁹ ¹⁰. Over time, these small improvements accumulate into a well-architected system. As Martin Fowler quips, “*Any fool can write code that a computer can understand. Good programmers write code that humans can understand.*” Continuous refactoring is how we turn the former into the latter – by restructuring code without changing its behavior to improve clarity and reduce complexity ¹¹.

Importantly, **emergent design is not the absence of design or architecture** – it’s the idea that the *best* architecture unfolds through constant adaptation. You still need skill and intent, but you allow the design to **grow organically** rather than attempting a perfect blueprint upfront. Kent Beck described this approach as: “*Start with something small, then add to it, keep adding function one piece at a time ... Eventually, you may end up with a framework, but you won’t have built one – you’ll have evolved one.*” ¹². In other words, *architecture is continuously redesigned and simplified as the system grows*, rather than set in stone from day one. Given sufficient refactoring discipline, a clean architecture and “*simple design*” naturally emerge from the process ¹³ ¹⁴.

This counters the misconception that developers always “over-engineer” or that upfront architecture is the only way to get design. In a healthy engineering culture that values quality, developers focus on doing the right thing for the current needs (guided by principles like YAGNI – “You Ain’t Gonna Need It”) and then improving the code as new needs arise. Over-engineering is less likely because any complexity that doesn’t prove useful can be refactored out. A well-known Agile principle states: “*Simplicity – the art of maximizing the amount of work not done – is essential.*” ¹⁵. Teams practicing emergent design live by this: implement only what’s needed now, but **continually clean up** and refine the code so that it remains as simple as possible for the functionality it must deliver.

Consider how a piece of code might evolve in this paradigm: you might start with a single monolithic method to get something working (because you’re in “genesis” and still figuring it out). As the concept proves itself, you refactor that method into a class with clearer responsibilities. Later, as features grow, that class might spawn a small cluster of classes or modules, each handling a part of the task (in line with single-responsibility principle). Some of those modules might generalize into reusable libraries or services – effectively becoming commodity components for your team. This *constant factoring and abstraction* means that **over time, the codebase actually gets more modular and easier to work with, even as it grows**. In fact, mature codebases often enable *faster* development of new features than when the system was small and chaotic, precisely because the heavy-lifting has been refined into

reliable, reusable parts. A well-evolved system avoids the classic nightmare of “the more code, the slower progress”; instead, adding new capabilities is faster because the groundwork (clean abstractions, infrastructure, automated tests, etc.) is laid. As one Agile expert notes, *“by continuously improving the design of code, we make it easier and easier to work with... if you get into the hygienic habit of refactoring continuously, you’ll find that it is easier to extend and maintain code.”*¹⁶.

Refactoring – the act of improving internal code structure without changing external behavior – is the engine of emergent quality. It should be frequent and proactive, not an afterthought. Teams that treat refactoring as an occasional, big event (“let’s spend two sprints cleaning up all our debt later”) often fail to ever get around to it, or find the system has grown too brittle by then. Instead, the best practice is to integrate small refactorings into daily work: *“Refactor as you go”* is the mantra¹⁷. This could mean, for each new story or feature, you first tidy up the relevant code to make adding the feature easier, then implement the feature, then clean up again after the change. By doing this, you **pay the cost of quality incrementally**, rather than letting “interest” on technical debt accumulate to crisis levels¹⁸¹⁹.

Of course, there’s a balance to strike. As code can *always* be improved, developers must use judgment on when to stop iterating on a particular design. We need to be mindful of **diminishing returns** – the point at which further refactoring yields minimal benefit. Effective teams refactor with an economic mindset: focus on what’s most problematic or highest value to clean up, and don’t polish parts of the code that aren’t actually causing pain or likely to be revisited soon²⁰²¹. One guideline is to “fix the code you touch”: when adding or changing a feature, clean the areas of the codebase you are actively working in (and that therefore deliver user value). Avoid speculative “maybe we’ll need this cleaned someday” efforts. By aligning refactoring with delivering user stories, you ensure effort is spent on areas that impact value, and you naturally limit refactoring scope to what’s needed¹⁷²². If a section of code is never used or touched, its ugliness isn’t a priority (with the caveat that truly dead code should be removed). As one refactoring principle states: *“If you were in charge of a city budget, you wouldn’t fix roads nobody uses. Limit the refactorings on first-time visits. You must know when to stop due to the diminishing returns.”*²¹. Timeboxing refactoring tasks can help – e.g., “we’ll spend one day improving this module” – to enforce a sense of when enough is enough²³. The goal is sustainable improvement: always leave the code better than you found it, but also always deliver progress on user-facing functionality. When practiced well, this approach leads to a virtuous cycle where each iteration yields both new value and a cleaner system than before. Over multiple iterations, **a robust, well-designed architecture “emerges” as a result**¹³²⁴.

It is worth noting that when a truly **good design emerges, it tends to stick**. There is sometimes a concern, *“but couldn’t you just keep refactoring forever?”* In theory yes, but in practice a well-factored design reaches a point where any further change makes it objectively worse. You naturally stop refactoring when the design **feels “just right”** – when *adding or removing anything would degrade the clarity or functionality*. Good modular designs exhibit a kind of **stable equilibrium**: once responsibilities are in the right places and abstractions are clean, that part of the system often remains unchanged for long periods because it simply works. Developers have no reason to change it until requirements change, and even then the extension points are usually clear. This is the “gravity” the user alluded to – a positive inertia that sets in when you hit upon a really solid solution. A famous measure of good design is exactly that: nothing further to remove, nothing additional needed¹. When teams reach this point on a feature or module, that is a signal to move on and tackle the next challenge.

Invisible Quality: The Hallmarks of Good Design

One remarkable aspect of a high-quality, evolved design is that **it often becomes invisible** – in the sense that it *just works* so intuitively that neither users nor developers are acutely aware of it day-to-day.

As Apple's Steve Jobs noted, the mark of good UX design is that it "doesn't scream for attention – it just works, making things feel natural, intuitive, and easy to use" ²⁵. When software is well-designed, users can accomplish their goals without frustration, and they may not even stop to think about *why* the experience is so smooth – it feels obvious. In contrast, poor design is *very* visible: the user stumbles, tasks take too long, the interface is confusing or cluttered, and the frustration is palpable. Thus, **good design tends to be transparent**: "Good design is invisible — when done right, users don't notice it because everything feels natural, intuitive, and easy." ²⁶. This is true not only in the user interface but also in code architecture (though "users" of code are developers). A clean API, for example, is "*invisible*" in that developers can use it without reading a tome of documentation; its design is evident from how it behaves.

One interesting point raised is that **simplicity in design is not the same as minimalism for its own sake**. It's a common mistake to equate a simplistic, sparse interface with a truly simple (usable) experience. In fact, "*false simplicity*" can harm usability ²⁷. For instance, consider an application that hides all its functionality behind a single menu or button in the name of a clean look – it might have a very minimalist appearance, but users could struggle to discover features or perform multi-step tasks that were one-click in a more explicit interface. As a design article notes, "*we want our designs to be usable and intuitive, but that can often be mistaken for clean or minimal visual designs. [Such] simplicity can be misleading and could be harming the usability of your designs.*" ²⁷. In other words, **intuitive does not always mean fewer buttons**; it means the *right* buttons at the right time. Sometimes, a UI with 10 obvious options is far easier for a user than a UI with 1 hidden option that tries to infer what the user wants. Good design is about *intuition and flow*, not just visual austerity. If a complex task requires multiple controls, presenting them upfront can be more *simple* for the user than a clean slate that forces guesswork. As evidence, usability studies have shown, for example, that hiding navigation menus behind a single "hamburger" icon significantly cuts usage and slows users down, compared to exposing at least the top-level options directly ²⁸ ²⁹. In a real sense, **simplicity is about efficiency and clarity**, not the sheer number of elements on screen.

What does this mean for software quality? It implies that a *quality outcome often has an elegant simplicity*, but achieving that may require adding or refining features in a way that initially looks "more complex" on the surface. We should not confuse *simplistic* design (which might drop essential capabilities) with *simple-to-use* design (which might actually have rich functionality, presented logically). The emergent quality process – with tight user feedback loops – helps get this right. Through iterations, designers and developers can test different approaches (maybe adding that 50-button panel in a beta, or conversely simplifying into a wizard) and find which truly feels more intuitive to the end user. **User feedback is the compass** for simplicity. If the user finds the new version more confusing or less efficient, the design isn't truly simpler from a quality standpoint. On the other hand, if users adapt quickly to a new design and then *hate going back* to the old version, it's a strong sign the new design is a success (the old now feels "really bad" by comparison, as the user in the transcript noted).

This underscores the importance of involving real users throughout the evolution of a product. The sooner and more frequently you test with users, the quicker you can converge on a design that *feels right*. Studies have found that **continuous feedback loops enable faster, smarter design iteration**, reducing the need for costly redesign later. One report notes that integrating user input early can cut rework by up to 25% ³⁰. Rather than guessing what might be simple or high quality, teams that iterate with feedback turn design into a data-informed exercise, where each cycle is a learning opportunity ³¹. In practical terms, this might mean releasing incremental improvements to a subset of users (A/B tests, beta features) or doing regular usability testing sessions each sprint. The development team and the users form a tight feedback loop, co-evolving the product. Under this model, **quality truly emerges as an outcome of many small corrections and enhancements guided by user reactions**. It aligns perfectly with the idea of emergent design – not only is the code being refactored, but the *user*

experience itself is being constantly “refactored” based on feedback until it reaches that invisible state of just working.

Hitting the Sweet Spot: When is Design “Good Enough”?

Because evolution is continuous, one might wonder: when do we know a software product has attained *quality*? In practice, there are a few telltale signs of a mature, well-evolved design:

- **Diminishing Returns:** As mentioned earlier, the team notices that further refactoring or adding of features in a certain area yields little improvement. The glaring problems have been solved; what remains might be micro-optimizations or aesthetic tweaks that hardly impact user satisfaction or robustness. This is often a signal to shift focus elsewhere. It’s important to periodically assess cost vs benefit: *not every part of the system needs gold-plating once it’s “good enough”*. Hitting diminishing returns on improvements means the design is approaching an optimal state for current requirements ²⁰ ²¹.
- **User Delight and Intuition:** Users (whether end-users or internal developers using an API) are largely happy, and new users can get up to speed quickly. Support tickets or complaints drop off for that feature. If you *removed* the new design and gave users the old one, they would complain – a sign that you’ve really delivered value. As the transcript insightfully noted, when design B is truly better than A, going back to A feels *awful* to users. In contrast, if users don’t notice a difference or prefer the old way, then B was not actually an improvement. So a high-quality solution is often one that, once introduced, becomes the *new normal* that people wouldn’t want to live without. In other words, quality design often has **irreversibility** – it’s hard to imagine going back to a clunkier design after experiencing the better one.
- **Stable Architecture:** Internally, the code exhibits a **high signal-to-noise ratio**. It’s easy to understand and change by any competent developer on the team. There are few if any areas labeled “here be dragons” in the codebase. New features can be added without major rewrites because the modular structure supports extension. Also, the code has **low defect rates** – changes don’t break unrelated parts, indicating a good separation of concerns. When bugs do arise, they can be quickly pinpointed and fixed, which reflects clarity in design. All these are signs that the architecture has evolved to a point of *healthy equilibrium*. It’s worth noting that even when a design is “done” for now, it may need to evolve again if new requirements come in – but a quality architecture makes that evolution relatively painless, whereas a poor one makes any change dangerous.
- **Alignment with Constraints:** A good design also means it’s appropriate for its context – the performance is adequate, it scales as needed, and it’s within budget. Sometimes, over-engineering can produce a “technically elegant” solution that overshoots what was needed (say, ultra-low latency code for a problem where it didn’t matter). A sign of quality is that the solution meets all the explicit requirements and constraints without excessive complexity. It’s *elegant in its context* – not too much, not too little. This is the art of hitting the point of diminishing returns and stopping there. The iterative process naturally encourages this, because teams are continuously evaluating real needs and adjusting accordingly, rather than implementing speculative features “just in case.”

When these conditions are met, we can say the design is **“good enough” for its current purpose** and likely of high quality. At that point, the team’s focus often shifts more toward *maintenance mode* on that component (monitoring, minor tweaks) and heavier innovation moves to other areas or new products.

Simon Wardley's model would say such a component has become a *commodity/utility*: reliable, predictable, and not needing constant attention ⁴.

Crucially, reaching this stage is not the end of innovation; it frees up resources to tackle the next *Genesis*-stage idea. This is how mature software organizations can continuously innovate: they commoditize parts of their platform (making them robust and low-effort), which allows them to devote creative energy to new experimental features. They have Pioneers working on new ideas while Town Planners keep the core stable – and settlers bridging between the two ⁵. It's a continuous cycle of emergent quality at different scales.

Conclusion: Quality as a Journey, Not a Destination

In summary, **quality in software is achieved through evolution, not revolution**. A brilliant design is rarely conceived fully-formed at project inception; instead, it *emerges* from iterative exploration, continuous refactoring, and tight feedback loops with users. We start with something rudimentary (and often ugly under the covers), then, like a craftsman honing a piece of wood, we shave off rough edges, reinforce weak joints, and polish the result step by step. The process is messy and nonlinear – much like Wardley Mapping's evolution, there may be loops and step-changes – but over time, order and simplicity emerge from the initial chaos.

It's important to foster an environment that **encourages this evolutionary approach**. That means budgeting time for refactoring and improvement, not just churning out features. It means trusting developers to make architectural changes when needed (with appropriate safety nets like tests in place) so the codebase can improve iteratively. It means involving users early and often, so that "quality" is defined by real-world success, not just theoretical purity. When teams operate this way, they inherently avoid many of the classic failure modes: they won't end up with a big ball of mud that's impossible to maintain (because they're constantly untangling it), and they won't deliver a product users hate (because they validated ideas with users and evolved the UX). Instead, they produce systems that are **robust, intuitive, and adaptable**.

Finally, this emergent perspective reframes the relationship between *quality* and *simplicity*. Rather than asking if one is a subset of the other, we see that **simplicity is both a cause and effect of quality**. By striving for quality (through good practices), we tend to simplify the design (removing unnecessary complexity). And by valuing a simple, clear solution, we increase the overall quality (making the product easier to use and maintain). They reinforce each other. The end state of a high-quality evolution is often a beautifully simple solution – one that might even make people say, "Wow, that seems obvious, why didn't we always do it this way?" That reaction is a testament to all the unseen complexity that was whittled away through careful, emergent effort. In the world of software, achieving that kind of simplicity is hard work – but it is the work that defines quality.

Sources:

- Agile Manifesto Principles – emergent design ²
- Wardley Mapping concepts (Genesis to Commodity; Pioneers, Settlers, Town Planners) ⁴ ⁵
- Foote & Yoder, *Big Ball of Mud*: on prototypes becoming production ("Throwaway Code") ⁶
- Fowler/Beck, emergent architecture via continuous refactoring ¹² ¹³
- Luís Soares, *Guiding Principles for Refactoring*: on focusing refactoring and stopping when returns diminish ²¹
- Tom Kenny, *False Simplicity in UI Design*: simplicity vs minimalism in usability ²⁷
- MightyFine Design, *Good Design is Invisible*: intuitive design is "invisible" to users ²⁶

- UserTesting Blog: feedback loops reduce rework (user input = quality UX) 30
-

1 Antoine de Saint-Exupery - A designer knows he has...

https://www.brainyquote.com/quotes/antoine_de_saintexupery_121910

2 10 11 13 15 18 19 24 Iterating Toward Legacy - Scrum's Achilles Heel | Agile Alliance

<https://agilealliance.org/iterating-toward-legacy-scums-achilles-heel/>

3 Wardley Mapping: A Strategic Tool for Addressing Technical Debt | by Mark Craddock | Context Engineering | Medium

<https://medium.com/prompt-engineering/wardley-mapping-a-strategic-tool-for-addressing-technical-debt-84ddfb46a99>

4 Landscape

<https://learnwardleymapping.com/landscape/>

5 How to organise your teams | Simon Wardley | 32 comments

https://www.linkedin.com/posts/simonwardley_how-to-organise-your-teams-activity-7288507899289518080-jevD

6 The Big Ball of Mud and Other Architectural Disasters

<https://blog.codinghorror.com/the-big-ball-of-mud-and-other-architectural-disasters/>

7 8 abusedbits.com: Wardley Mapping Technical Debt of a System

<http://www.abusedbits.com/2018/11/wardley-mapping-technical-debt-of-system.html>

9 16 17 20 21 22 23 Guiding Principles for Refactoring | by Luís Soares | CodeX | Medium

<https://medium.com/codex/guiding-principles-for-refactoring-668f13d06374>

12 14 A Primer on Emergent Design

<https://www.pmi.org/disciplined-agile/a-primer-on-emergent-design>

25 26 Good Design Is Invisible: The Brilliance of Excellent Design

<https://mightyfinedesign.co/good-design-is-invisible/>

27 28 29 Tom Kenny Design | False Simplicity in UI Design: Simple is not always better

<https://tomkenny.design/articles/false-simplicity>

30 31 How feedback accelerates design iteration and reduces rework

<https://www.usertesting.com/blog/feedback-drives-design-iteration>