

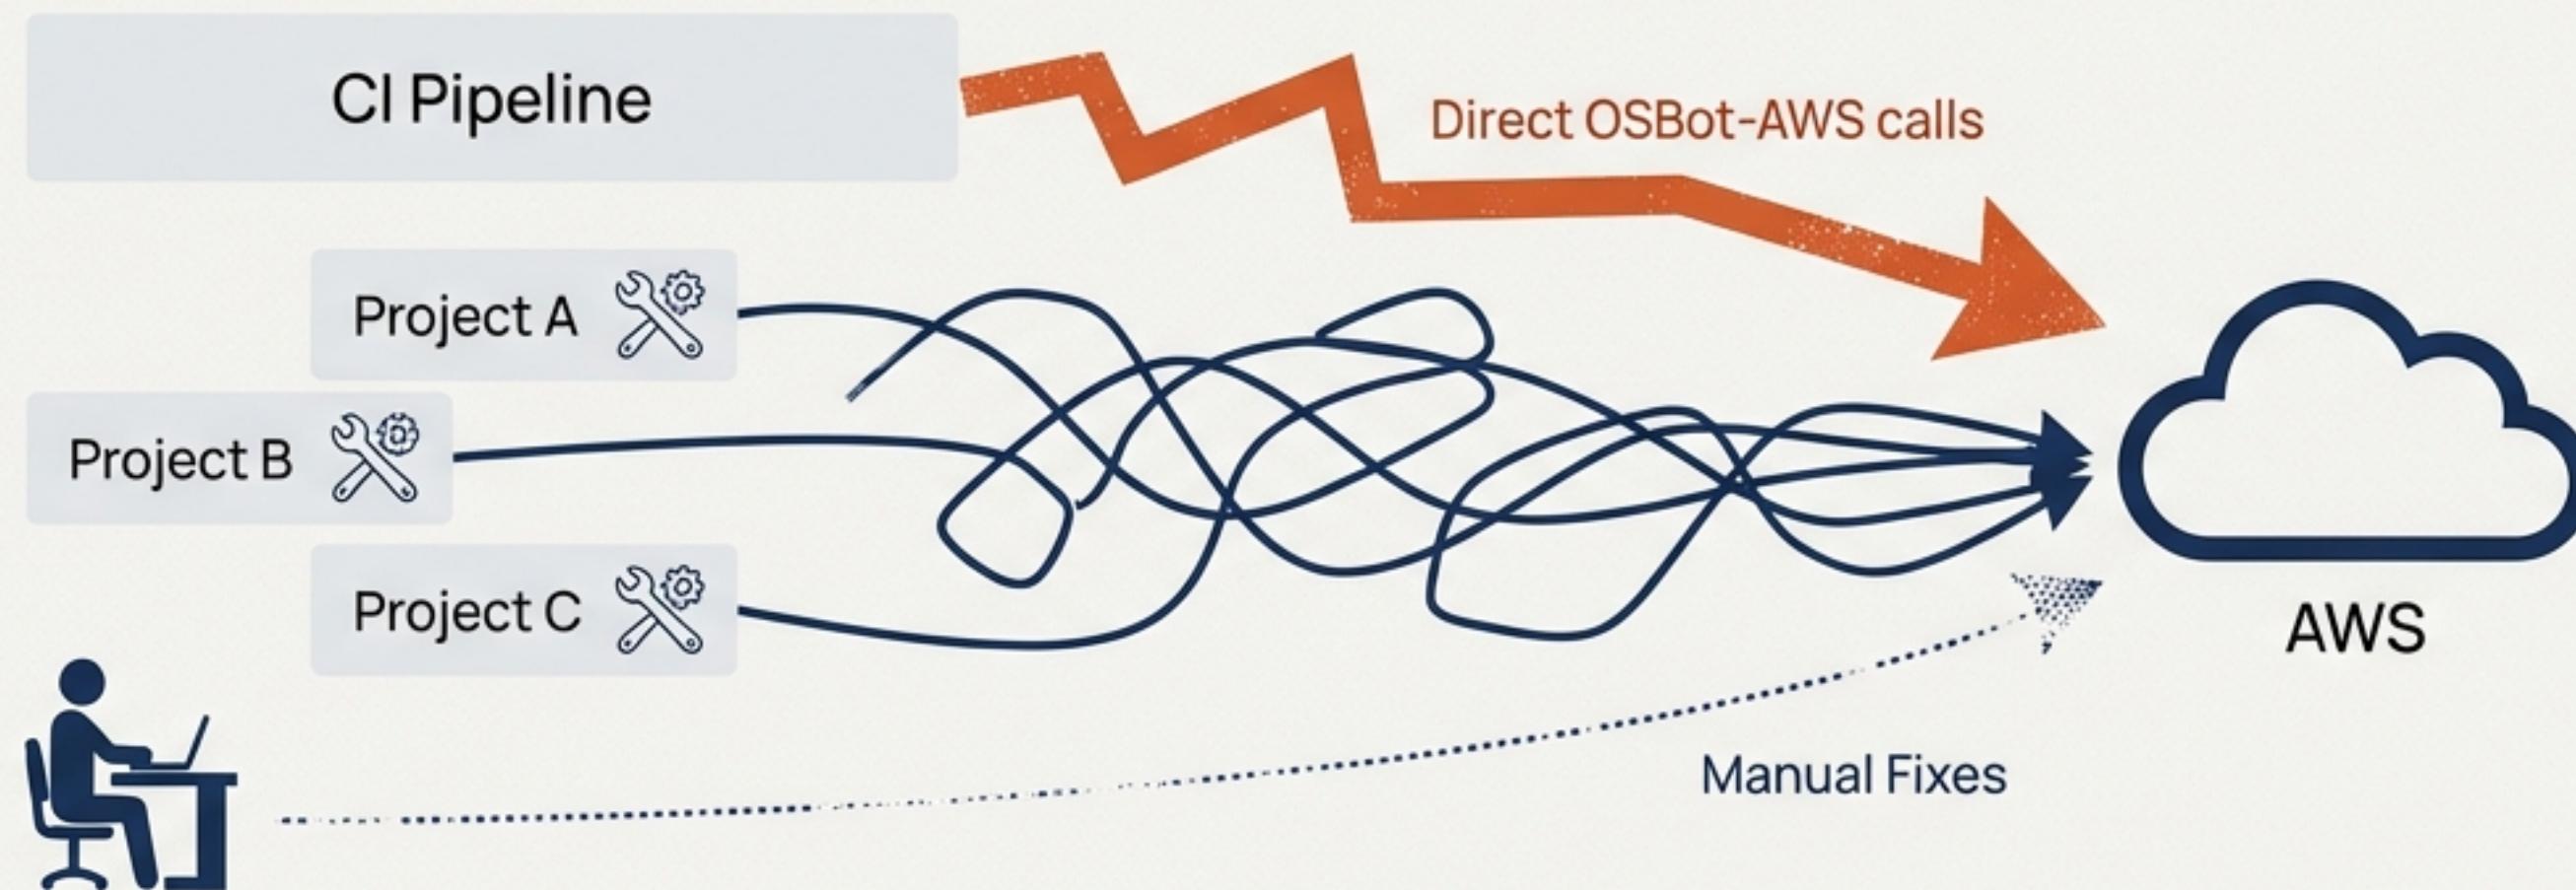
MGraph-AI Deploy Service

A Paradigm Shift in
Infrastructure Orchestration



Our current infrastructure deployment is fragmented and fragile.

We rely on a mix of CI pipeline calls, scattered helper methods, and manual interventions. This creates significant challenges.



No Unified State: Changes are untracked and often conflict.



No Verification Layer: We deploy blind, hoping for the best.



Difficult to Reproduce: Deployments are inconsistent across environments.



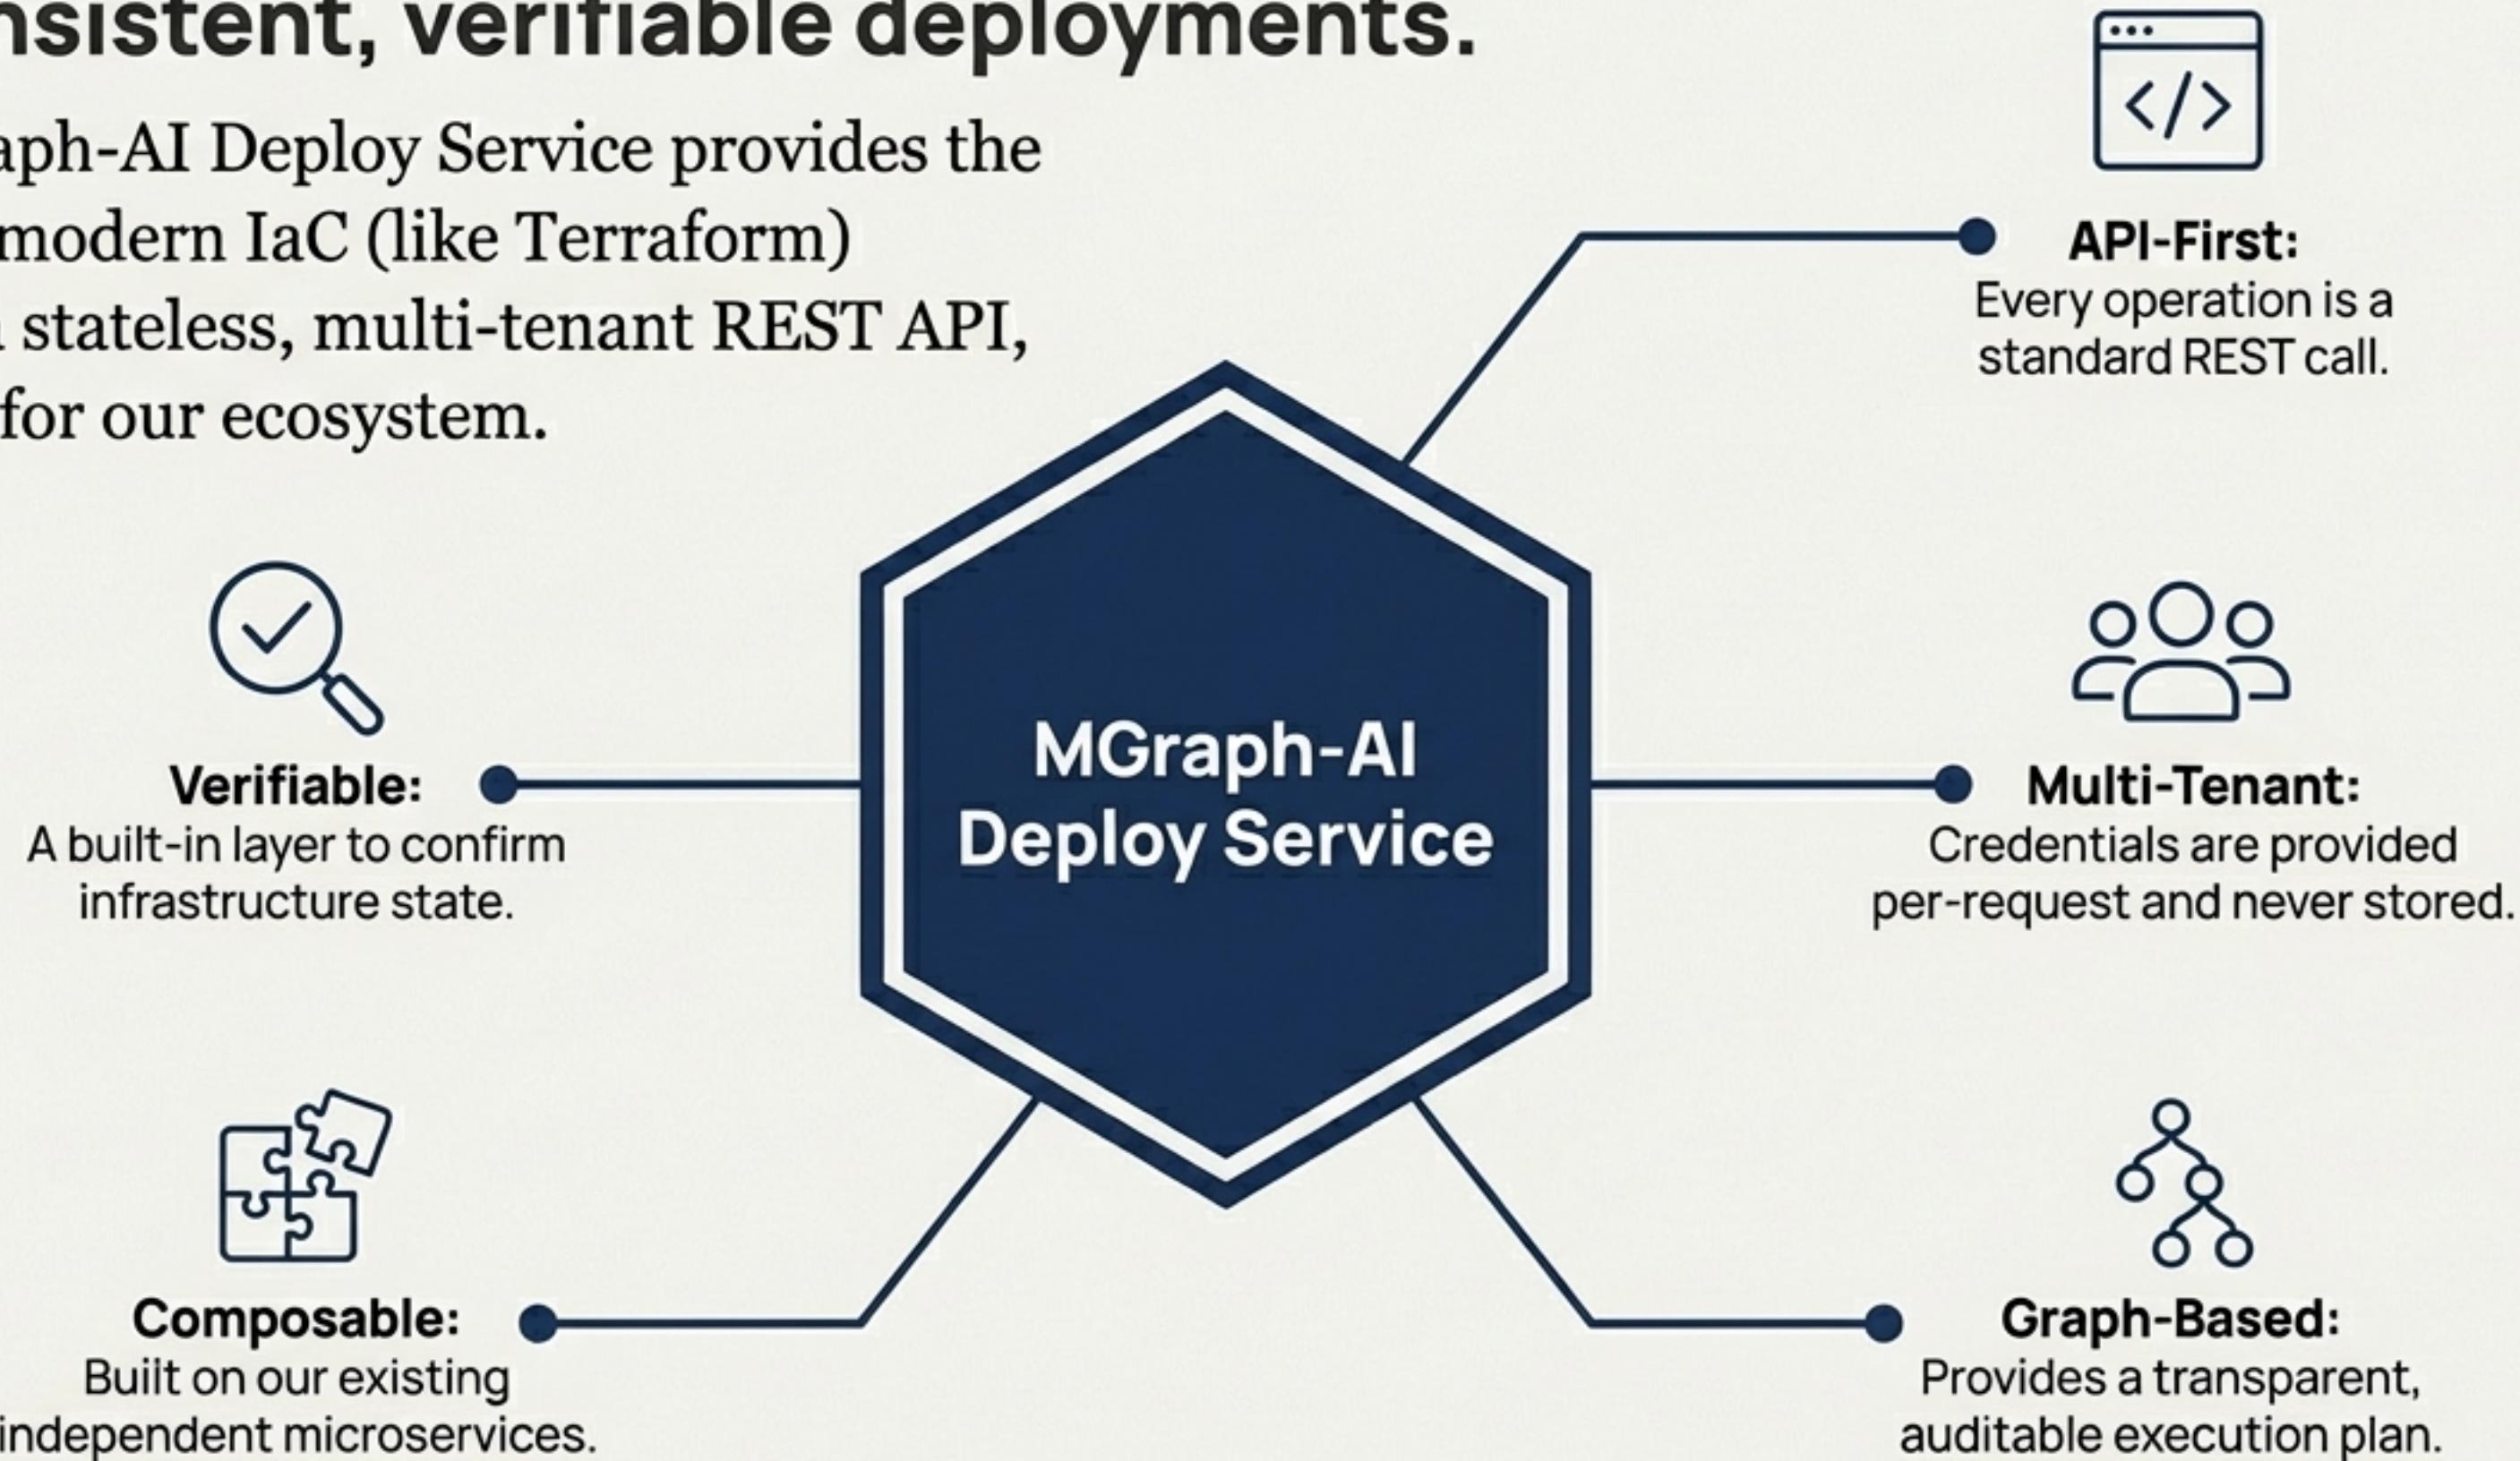
No Audit Trail: We cannot reliably answer "who deployed what, and when?"



High Migration Cost: Moving resources between regions is a manual, error-prone project.

We are building an API-driven orchestrator for consistent, verifiable deployments.

The MGraph-AI Deploy Service provides the power of modern IaC (like Terraform) through a stateless, multi-tenant REST API, designed for our ecosystem.



The service provides a complete lifecycle toolkit for infrastructure management.



Plan

Preview changes without execution (dry run)



Deploy

Create new resources across AWS and GitHub



Update

Modify existing resources efficiently



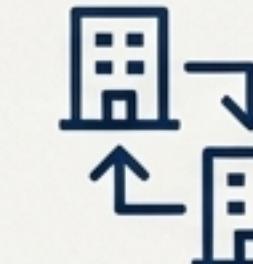
Verify

Test that infrastructure matches the expected state



Delete

Remove resources and their dependencies cleanly



Move

Migrate resources between regions or accounts



Resume

Continue a deployment from the point of failure



Step-Through

Execute one operation at a time for debugging



Multi-Region

Deploy identical resources to multiple regions simultaneously

Three core principles make this a fundamentally different approach.

Our design is not an iteration on existing tools; it's a re-imagining of how infrastructure state and operations should be managed. Understanding these three decisions is key to understanding the service.

1.

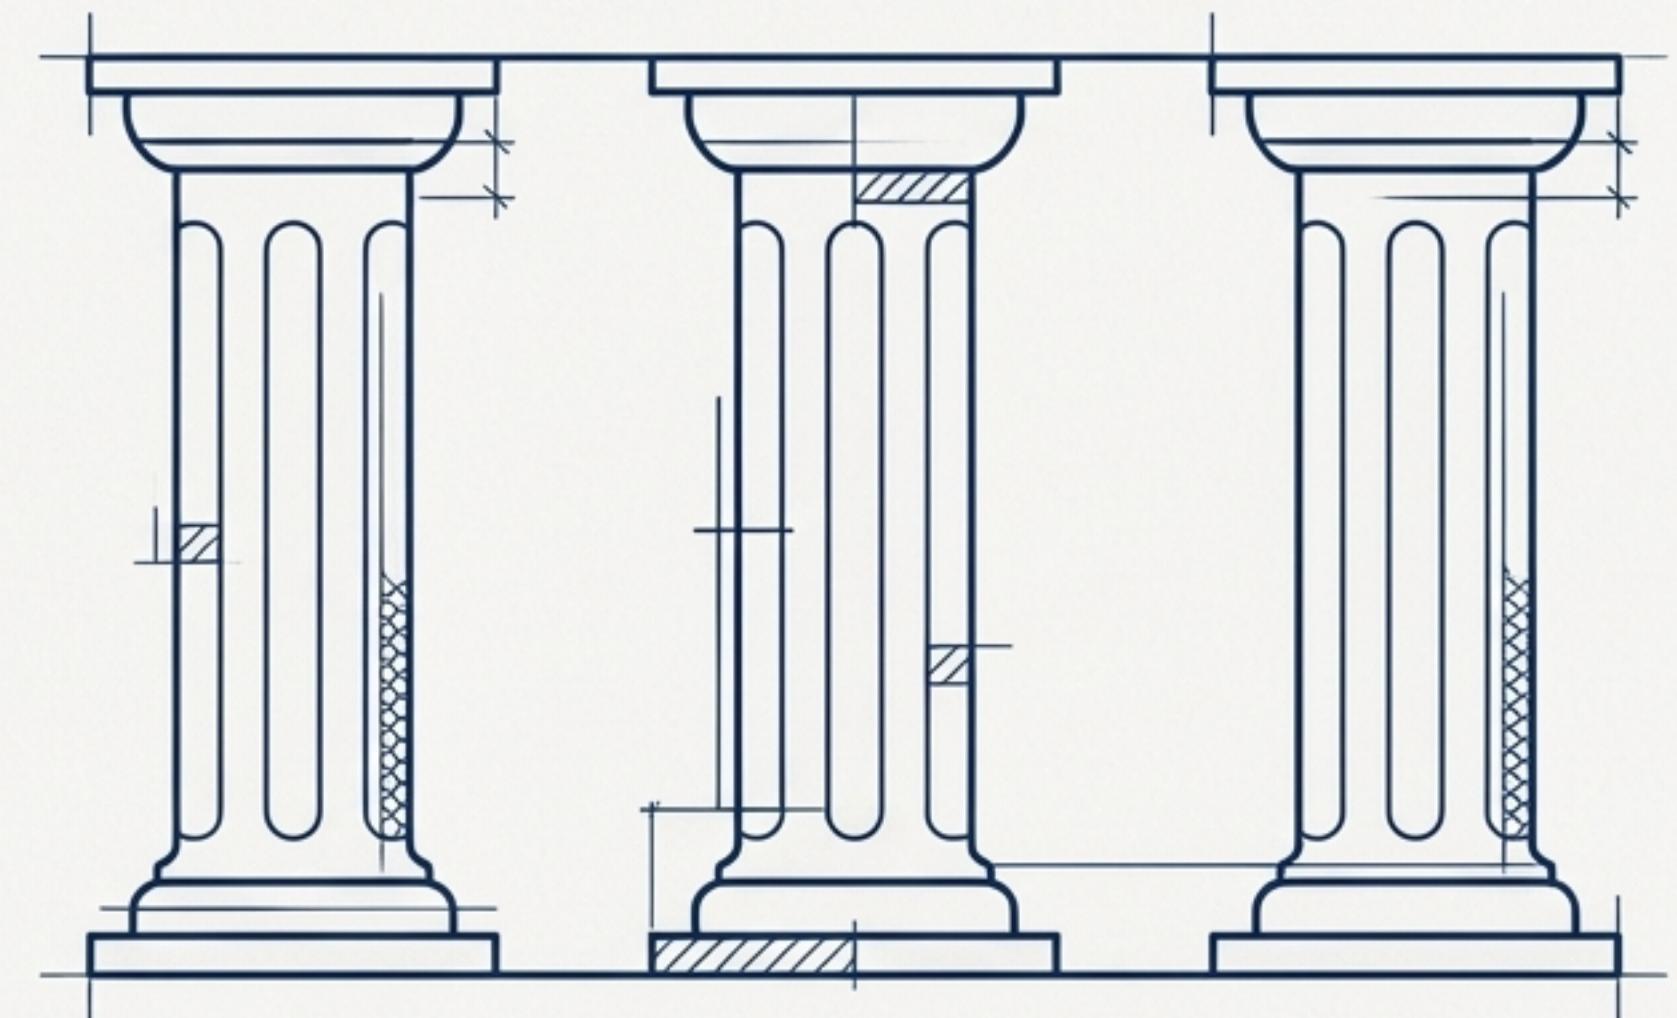
The Execution Graph IS the State

2.

Resources Discover Each Other via Tags

3.

Orchestration is Separate from Implementation



Principle 1: The Execution Graph IS the State.

We eliminate the concept of a separate, opaque ‘state file’. The **MGraph execution graph**, which is inherently auditable and versioned, becomes the single source of truth for the deployment’s history and status.



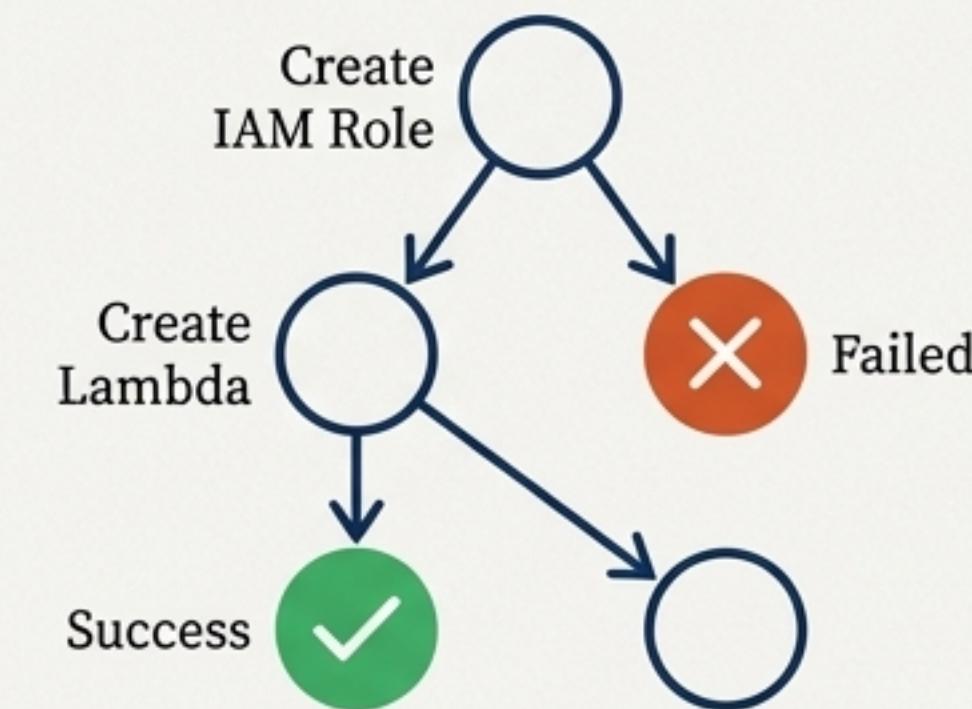
Opaque

Localised

Prone to
drift

The New Way

MGraph Execution Graph



Transparent
Audit Trail

Resumable
from Failure

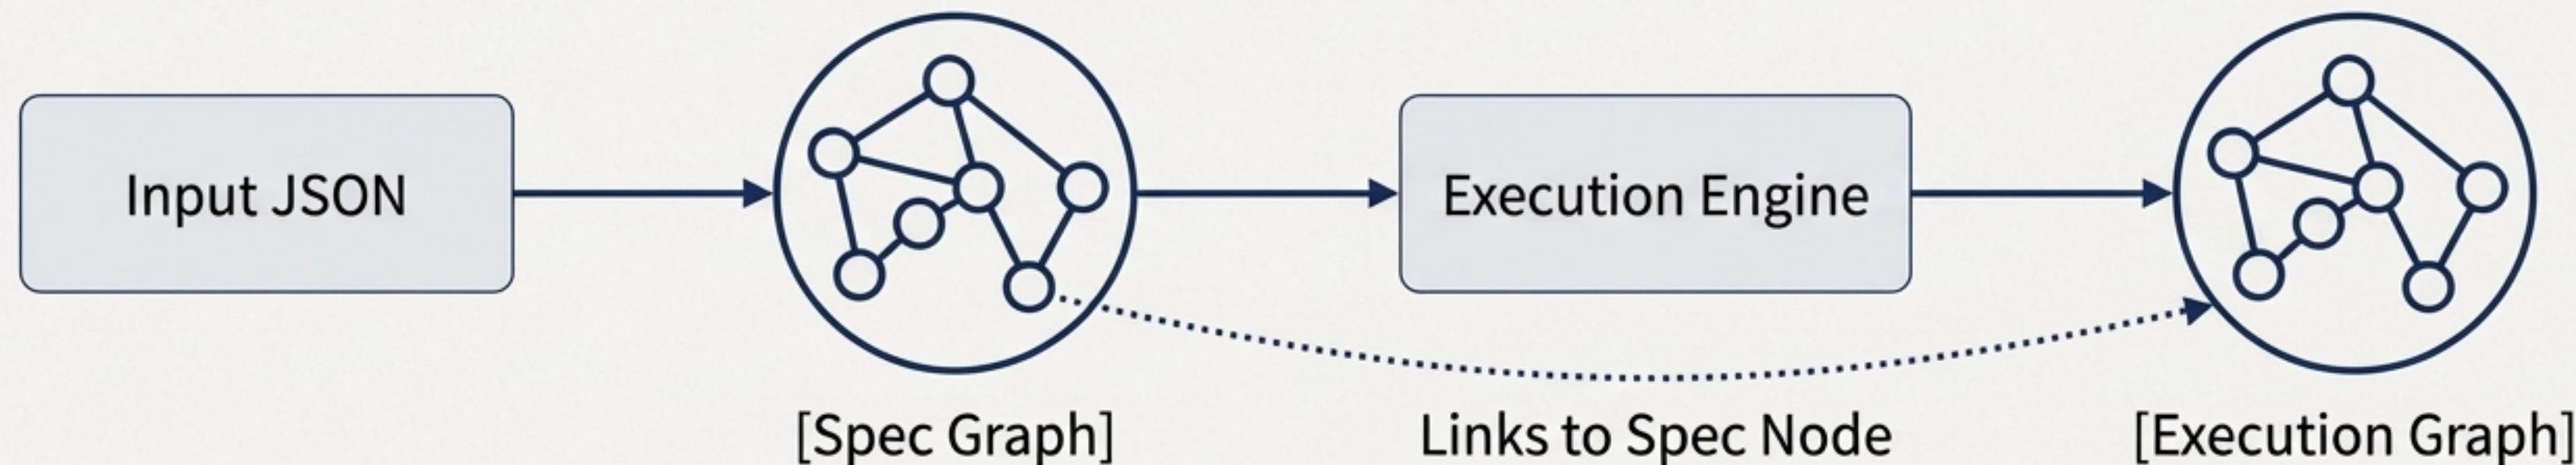
Inherent
Parallelism

Explicit
Dependencies

We use two distinct but linked graphs: one for the plan, one for the result.

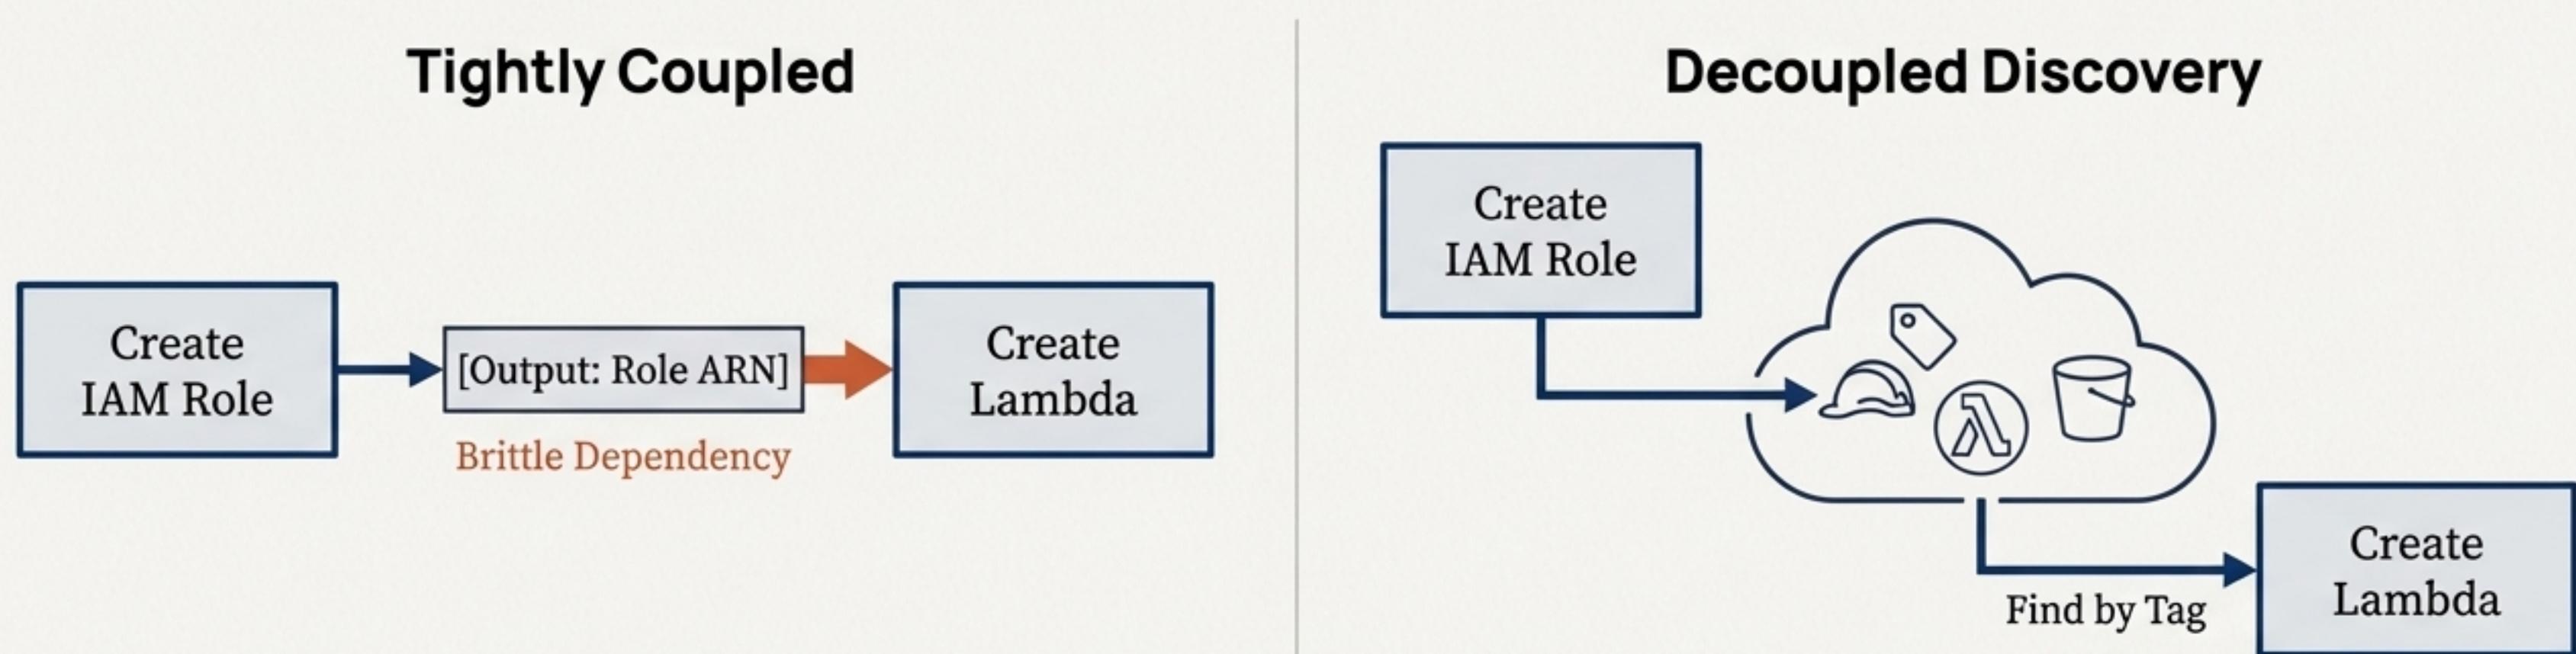
This separation provides a clear distinction between the immutable deployment specification and the mutable record of what actually occurred during execution.

Aspect	Spec Graph	Execution Graph
Purpose	What to deploy	What happened during deployment
Mutability	Immutable once created	Updated live during execution
Contents	Resource definitions, properties	Status (Success, Fail), timing, outputs, errors



Principle 2: Resources Discover Each Other via Tags.

Operations are designed to be atomic and independent. Instead of passing resource IDs (like ARNs) from one step to the next, subsequent operations find their dependencies by querying for resources with specific, predictable tags.



- ✓ Keeps operations truly atomic and independent.
- ✓ Enables state to be reconstructed by simply scanning existing infrastructure.
- ✓ Works seamlessly for multi-region and verification workflows.

We embed the graph's identity directly into the infrastructure itself.

Every resource created in AWS receives a standard set of tags linking it back to the exact graph node that created it. For services without tags, like GitHub, we use a consistent naming convention.

AWS Standard Tags

```
"Tags": {  
    "mgraph:graph-id": "65b..",  
    "mgraph:node-id": "65b..",  
    "mgraph:edge-id": "65b.."}  
  
```

This provides a complete audit trail directly on the resource.

GitHub Naming Convention

Resource:
`Repository Secret`
Pattern:
{secret_name}__{node-id}

The unique `node-id` ensures no name collisions.

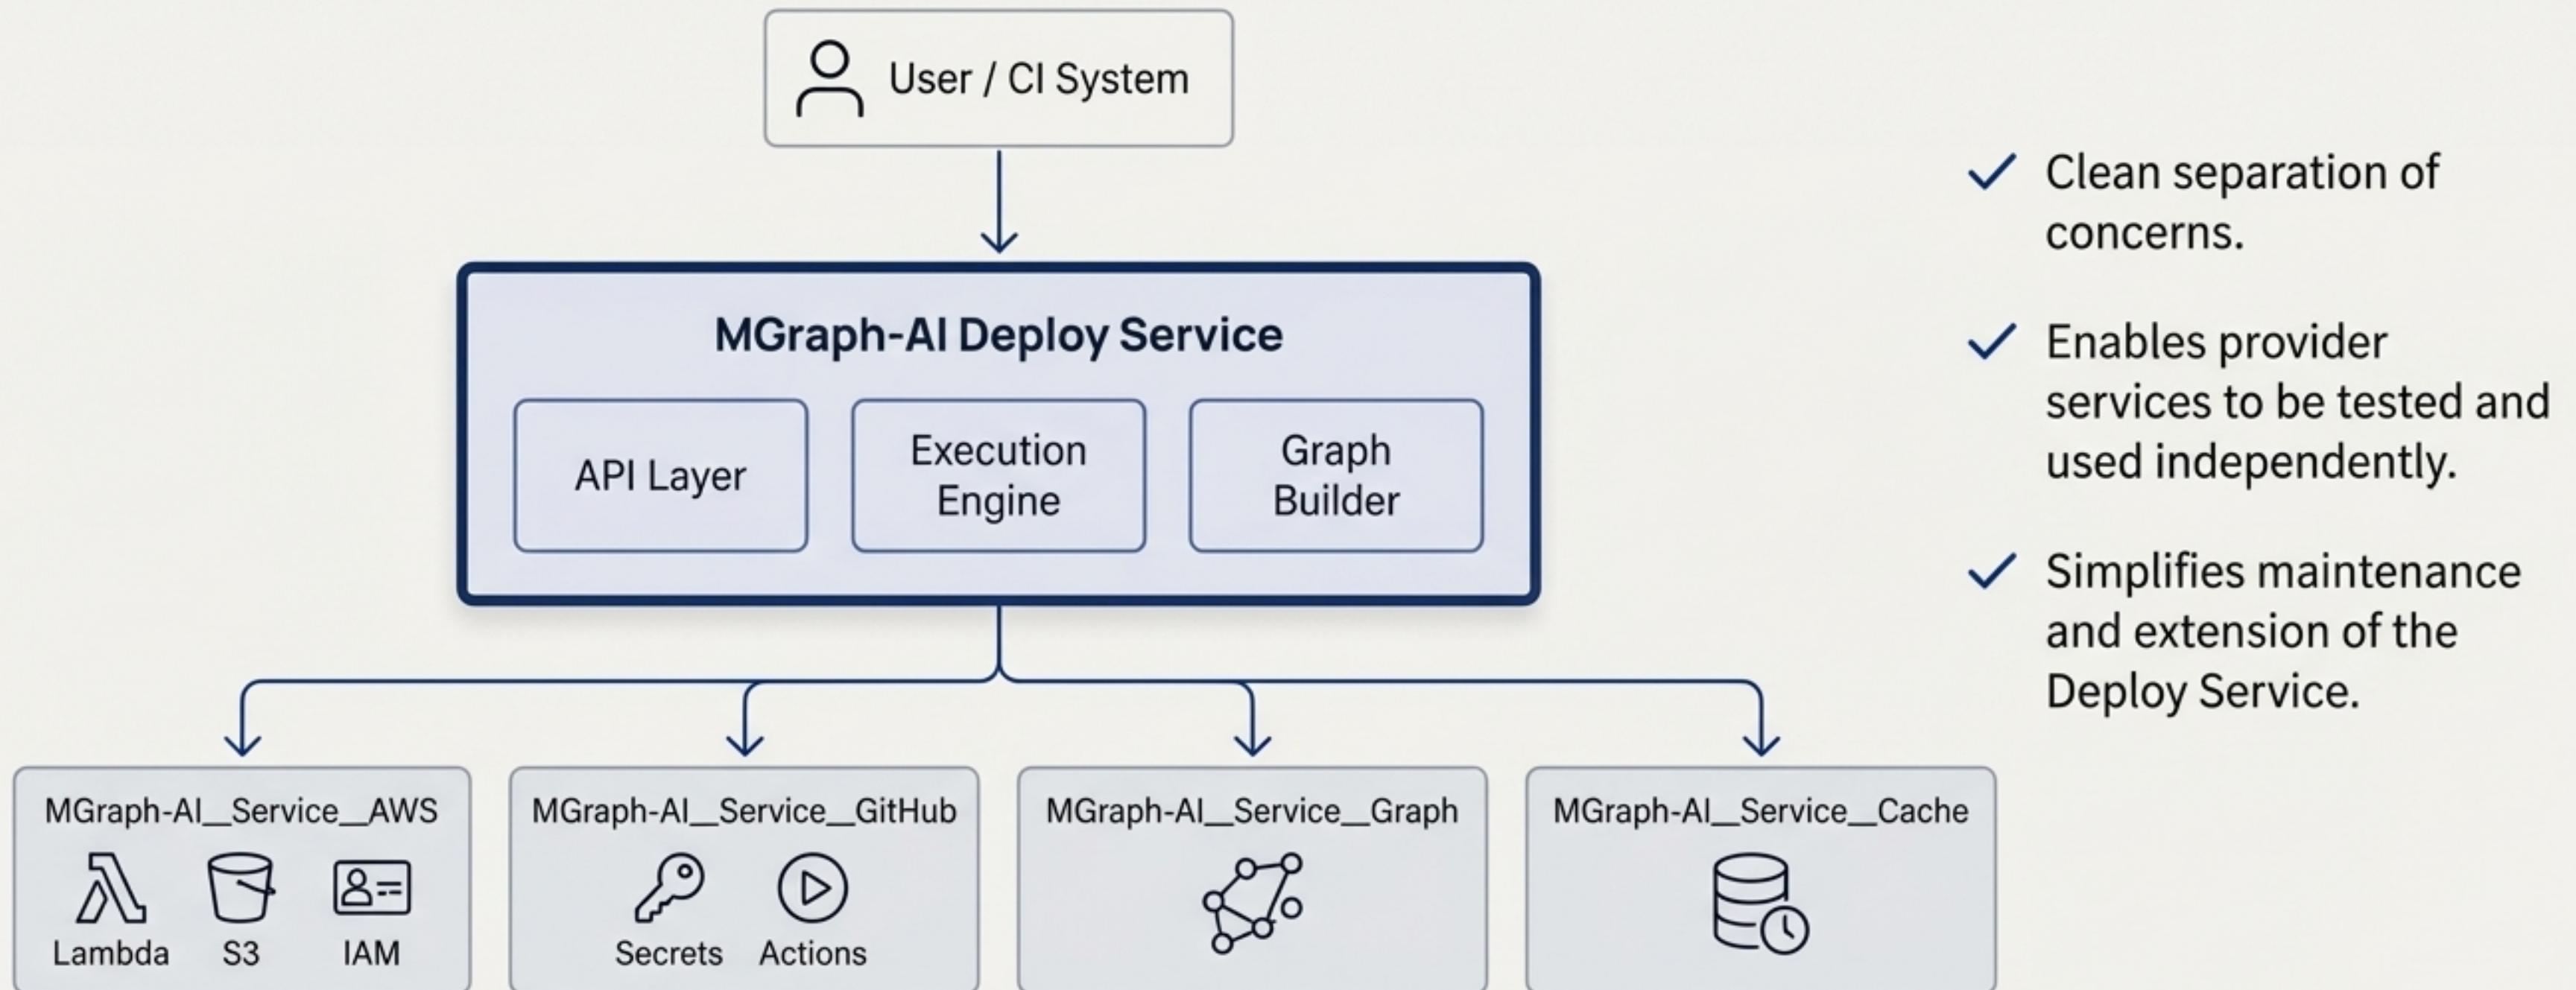
S3 Global Namespace

Resource:
`S3 Bucket`
Pattern:
{bucket-name}-{node-id}

Appending the unique `node-id` solves for the globally unique bucket name requirement.

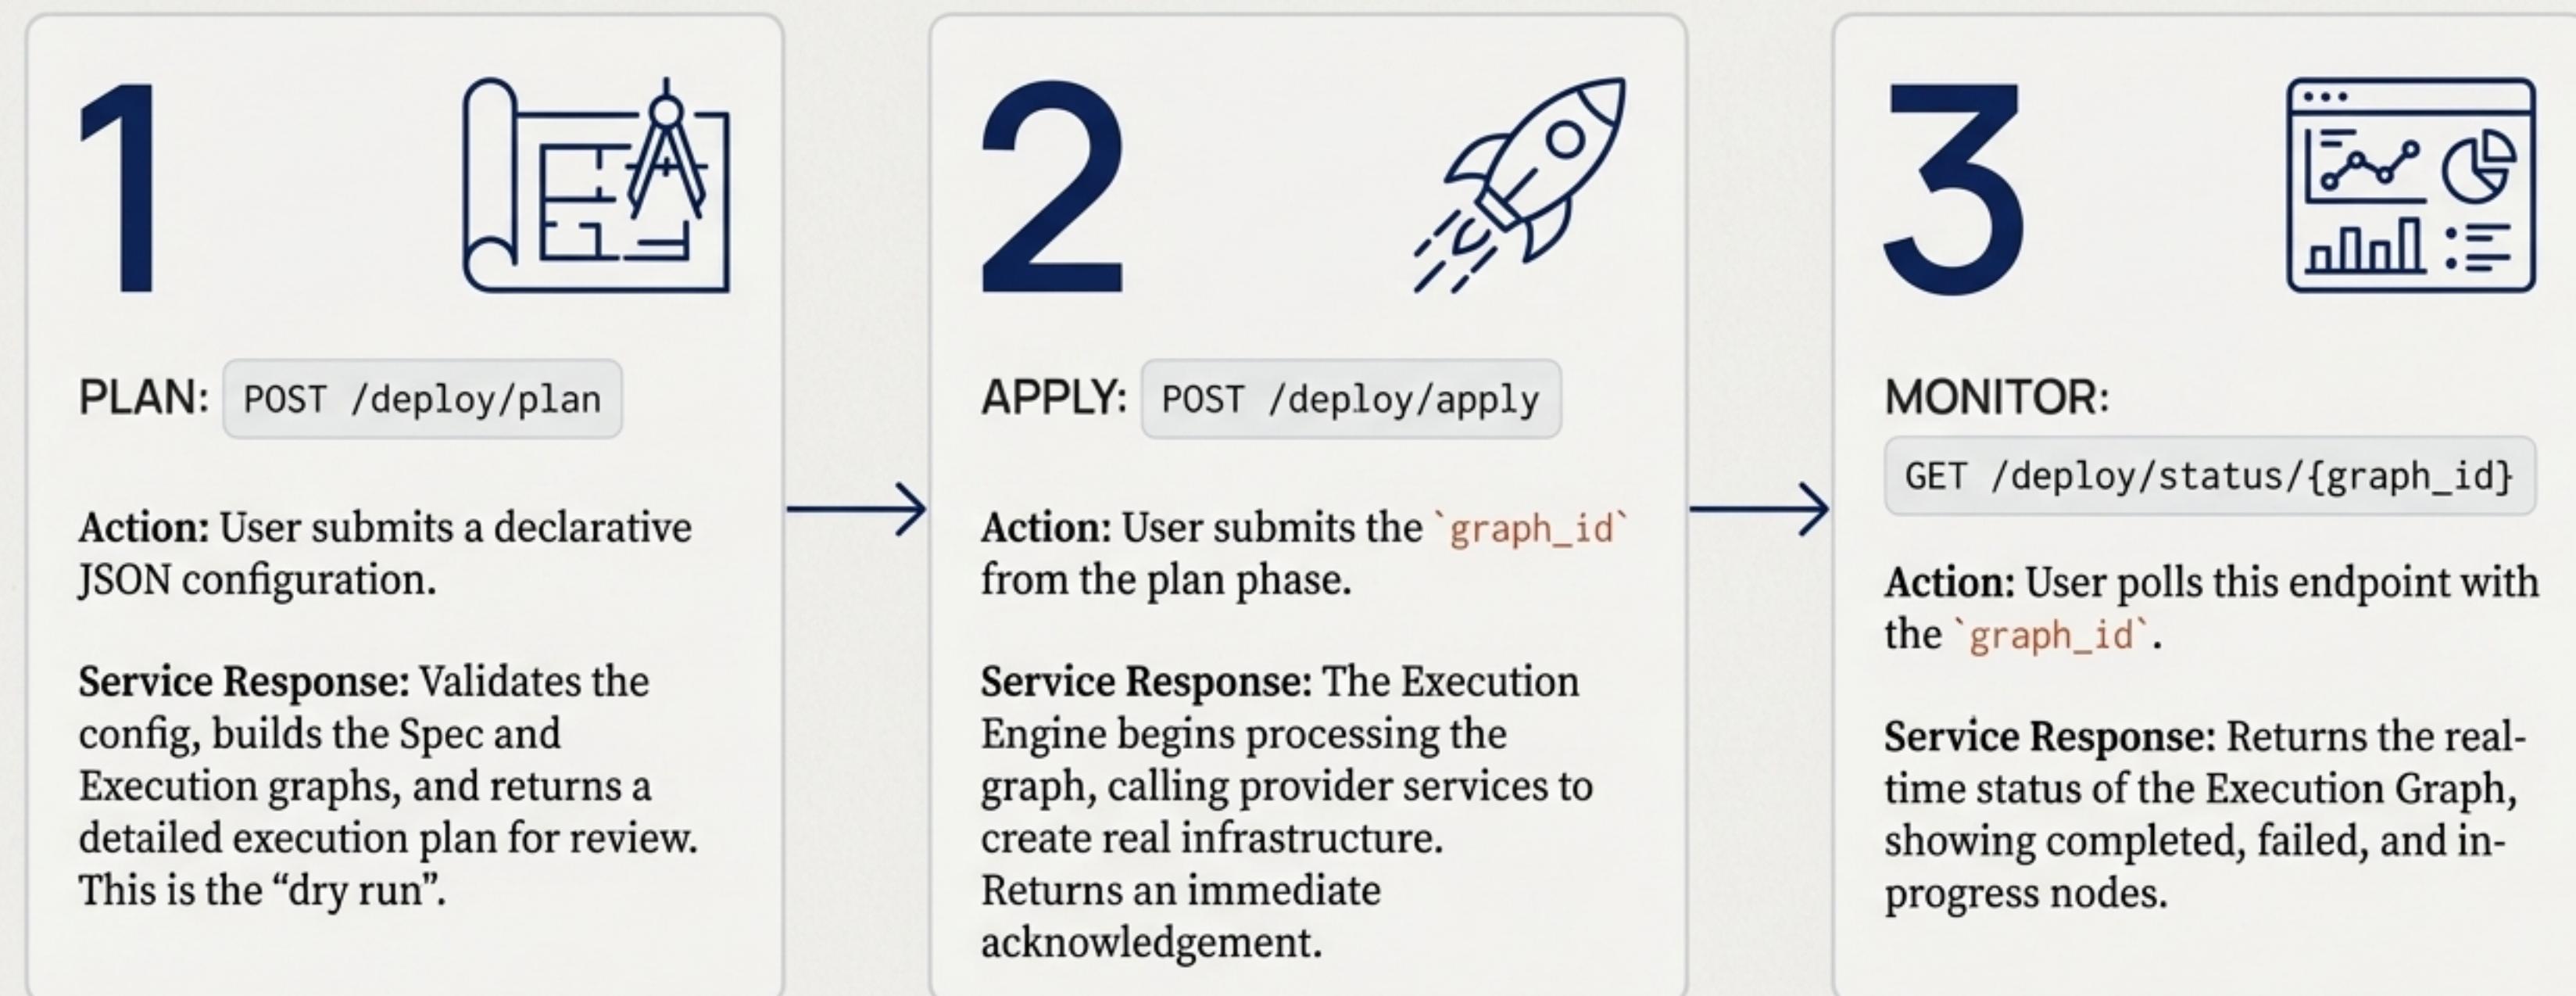
Principle 3: Orchestration is Separate from Implementation

The Deploy Service is a pure orchestrator. It focuses entirely on executing the graph logic—determining what runs when. The actual work of interacting with cloud providers is delegated to specialised “sister services”.



A typical deployment follows a simple, three-step API workflow.

The main visual is a clean, horizontal flow diagram with three numbered stages, moving from left to right.



The engine provides flexible execution modes and a clear storage strategy.

This slide details the operational flexibility and secure data handling within the engine.

Execution Modes

	Mode	Description	Use Case
	FULL	Execute all operations in parallel where possible.	Production deployments
	STEP	Execute one operation, then wait for a continue call.	Debugging & manual review
	DRY_RUN	Validate config and build plan; no execution.	Previewing changes

Storage Strategy

Data Type	Storage Service(s)	Rationale
Spec & Execution Graphs	Graph Service & Cache Service	Dual storage for different access patterns
Execution Logs & Outputs	Cache Service	Persistent, per-node audit and data
Deployment Metadata	Cache Service	Quick lookup of deployments
Credentials	In-Memory Only	Never persisted to disk
Execution State	Cache Service	Transient state for active executions

We will deliver the service incrementally over six distinct phases.

The main visual is a high-level horizontal timeline graphic.



Phase 1 (Wk 1-2): Foundation	Phase 2 (Wk 3-4): Graph Building	Phase 3 (Wk 5-6): Execution Engine	Phase 4 (Wk 7-8): Operations Library	Phase 5 (Wk 9-10): Advanced Features	Phase 6 (Wk 11-12): Testing & Docs
Project setup, schemas, and basic health routes.	Convert input JSON into a persisted Spec Graph.	Implement <code>FULL</code> , <code>STEP</code> , and <code>DRY_RUN</code> modes and the <code>/apply</code> route.	Build out the full set of AWS & GitHub operation primitives.	Implement Pause/Resume, Destroy, and Move capabilities.	Achieve production quality with comprehensive test suites and documentation.

Success is defined by these 14 measurable outcomes.

The implementation will be considered complete when all of the following criteria are met and verified.

Core API Lifecycle

- ✓ `/plan` validates and returns an execution plan.
- ✓ `/apply` creates resources across AWS and GitHub.
- ✓ `/verify` confirms infrastructure state.
- ✓ `/destroy` cleanly removes all resources.

Engine & State

- ✓ Tag-based discovery is fully functional.
- ✓ Spec & Execution graphs are correctly linked and persisted.
- ✓ `FULL`, `STEP`, `DRY_RUN` modes are working.
- ✓ Pause/Resume with `/continue` is operational.

Advanced Workflows

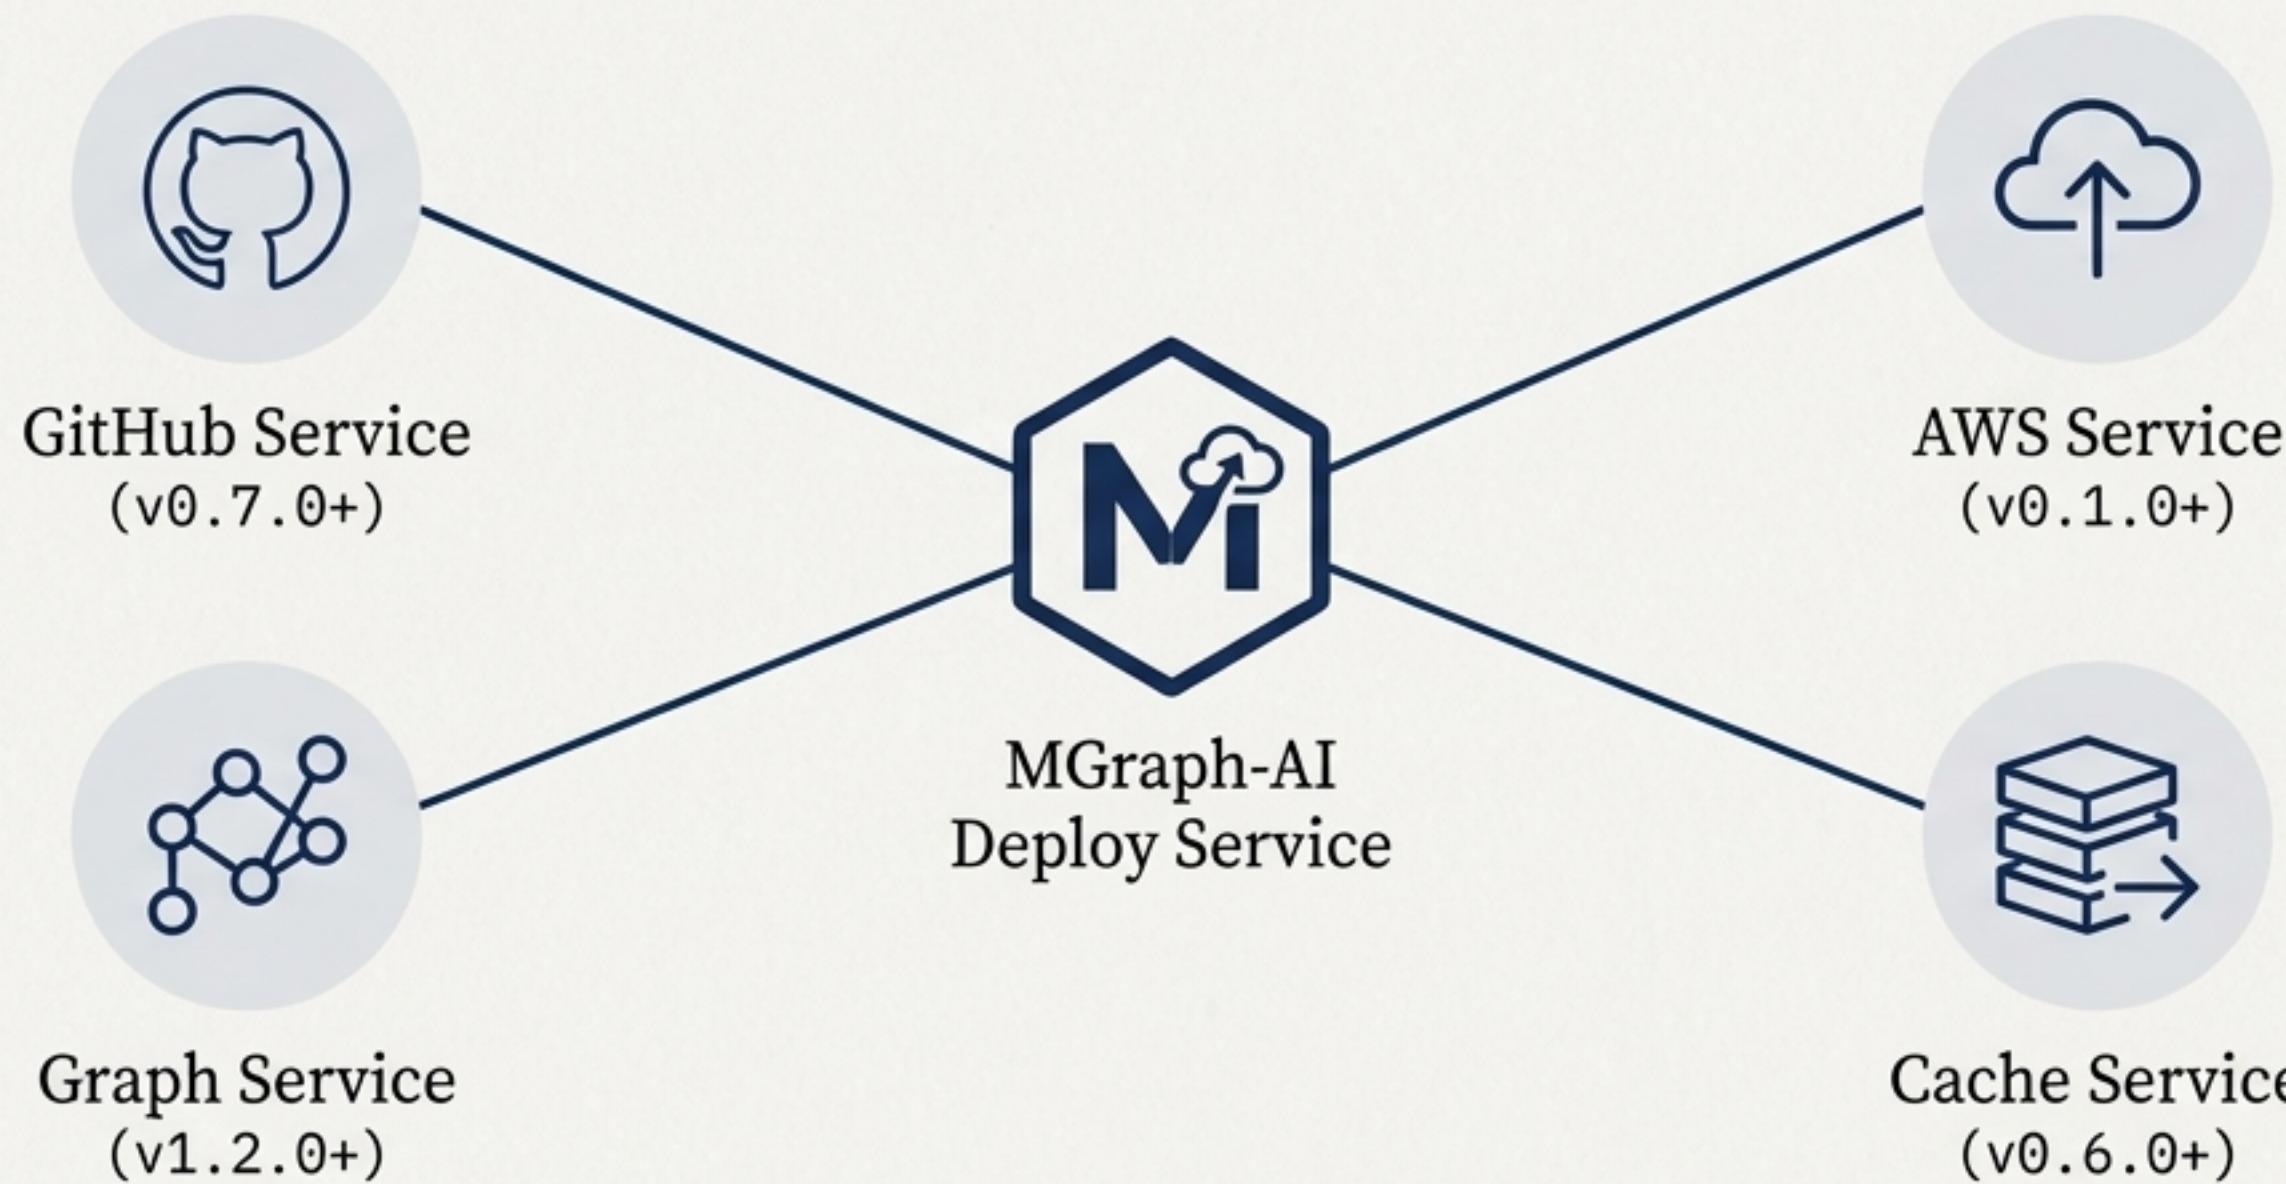
- ✓ Multi-region deployment is supported.

Quality & Readiness

- ✓ All state persisted with discoverable paths in Cache Service.
- ✓ Unit, Integration, and QA tests are all passing.
- ✓ OpenAPI documentation is complete and auto-generated.

The Deploy Service is a key orchestrator within the MGraph-AI ecosystem.

It leverages and enhances our existing service-oriented architecture, relying on these core dependencies to perform its work.



Target Repository: `the-cyber-boardroom/MGraph-AI__Service__Deploy`