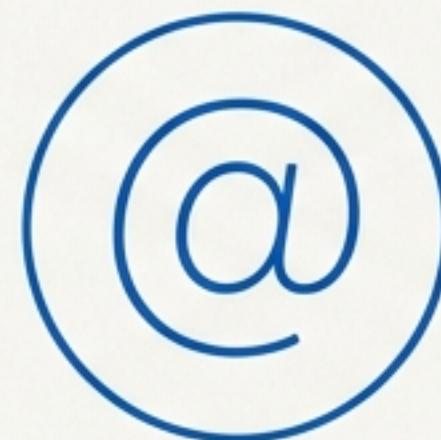


Effortless Performance Profiling for Python

An introduction to the `timestamp_capture` utility.



Simple: Add a single
decorator.



Insightful: Pinpoint exact
bottlenecks.



Production-Safe: Negligible
overhead when inactive.

From `osbot_utils` | `pip install osbot-utils`

Find a Bottleneck in Under 60 Seconds

1. Decorate

```
from osbot_utils.helpers.timestamp_capture  
import Timestamp_Collector, timestamp
```

```
@timestamp  
def slow_function():  
    # ... complex logic ...
```

```
@timestamp  
def main_process():  
    slow_function()
```

Add `@timestamp` to the methods you want to measure.

2. Capture

```
# In your main script...  
_timestamp_collector_ = Timestamp_Collector()  
main_process()  
_timestamp_collector_.print_report()
```

Create the collector with a specific variable name. That's it.

3. Analyse

```
-=[ Timing Report ]=-  
Target           | Total Time  
-----|-----  
main_process   | 150.23ms  
slow_function  | 145.12ms
```

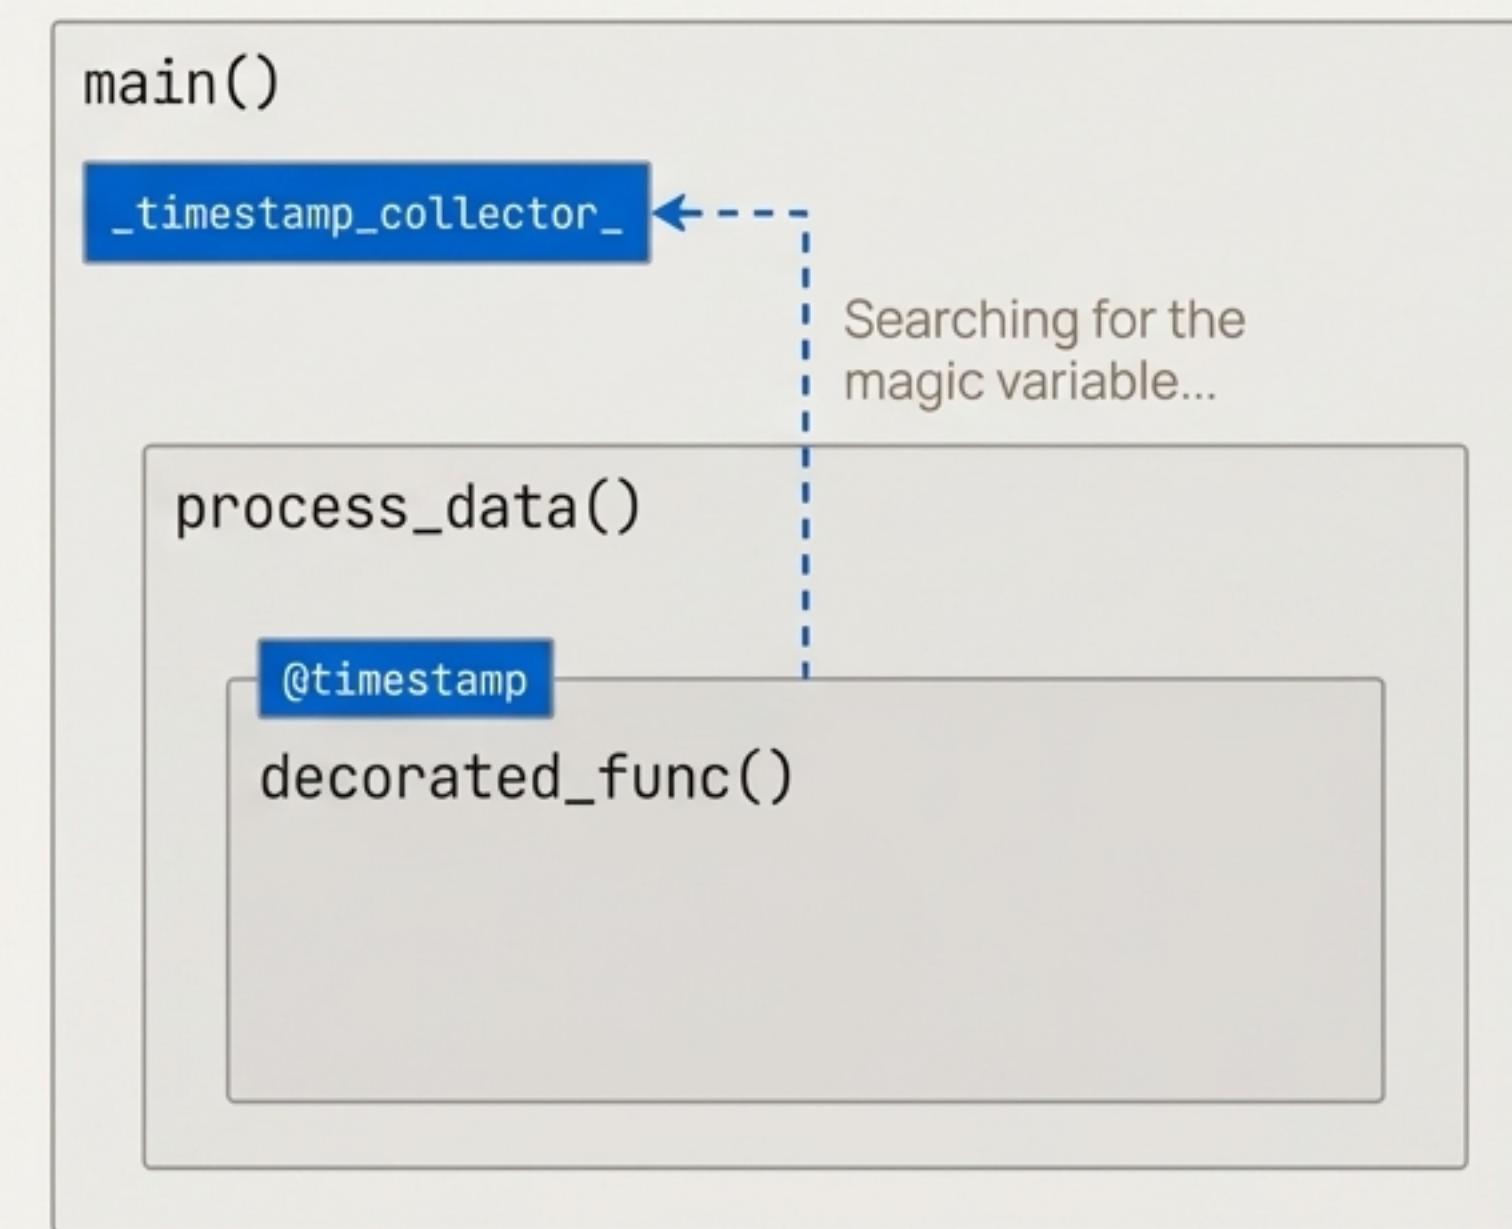
How It Works: The Magic is in the Name

The `@timestamp` decorator doesn't need to be told where the collector is. It finds it automatically.

Stack-Walking: The decorator walks up the call stack, frame by frame.

The Magic Variable: It actively looks for a local variable named *exactly* `_timestamp_collector_`.

Result: This decouples your business logic from your instrumentation code. No need to pass a 'collector' object through every function signature.



Designed to Be Left in Production Code

The decorator is dormant if no collector is found, with minimal overhead.

Scenario	Overhead per Call	Production Safe?
@timestamp (Inactive) No `_timestamp_collector_` found in stack.	~3 μ s	<input checked="" type="checkbox"/> Yes. Negligible impact.
@timestamp (Active) Collector is present and capturing data.	~8 μ s	<input checked="" type="checkbox"/> Yes. Ideal for methods >100 μ s.
Nested Decorators N levels of decorated calls.	N \times overhead	<input checked="" type="checkbox"/> Yes. Overhead scales linearly and predictably.

Rule of Thumb: If your method takes over 100 μ s to execute, the active capture overhead is less than 8%. This is acceptable for most production profiling scenarios.

Profiling More Than Just Methods

Sometimes the code you need to measure isn't a single function. Use a context manager to profile arbitrary blocks of code.

For Methods

```
# Use the decorator for entire methods
@timestamp
def data_processing_pipeline():
    # ...
```

For Code Blocks

```
# Use the block for specific phases
def data_processing_pipeline():
    # ... setup code (not timed)
    with timestamp_block("data.extraction"):
        # ... extraction logic
    with timestamp_block("data.transformation"):
        # ... transformation logic
    # ... teardown code (not timed)
```

Use Cases:

- Timing specific phases within a large function.
- Profiling code inside a loop without decorating the looped function.
- Measuring I/O or network operations that are not in their own methods.

Go Beyond Function Names for Clearer Reports

Default method names (`__qualname__`) can be long or uninformative. For complex systems, you need better organisation.

Use the `name` parameter to define a **custom, hierarchical metric name**.

****Decorator Example:****

```
# Instead of the default 'process_user_data'  
@timestamp(name="pipeline.stage1.extract")  
def process_user_data(user_id):  
    # ...
```

****Block Example:****

```
with timestamp_block(name="io.read.open_file"):  
    # ...
```

Benefit: This allows you to group related operations, version algorithms, and build reports that mirror your system's architecture, not just its code structure.

A Practical Guide to Naming Your Timestamps

A consistent naming convention is key to building insightful reports. Use a dot-separated format to create a logical hierarchy.

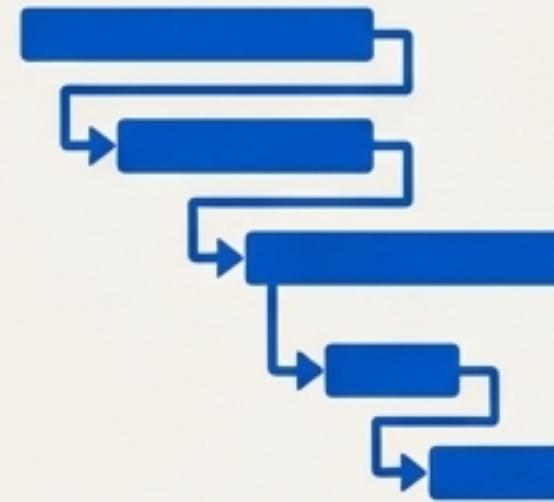
Pattern	Example	Typical Use Case
domain.operation	parser.parse_xml	Simple, effective categorisation.
component.stage.action	pipeline.stage1.extract	Tracking stages in a data pipeline.
feature.version.method	algorithm.v2.sort	A/B testing or comparing implementations.
io.category.operation	io.read.open_file	Grouping related operations like I/O.

“Key Insight: Treat your metric names as a form of documentation. They should clearly communicate the purpose and context of the code being measured.”

Turning Raw Data into Actionable Insight

`timestamp_capture` provides three distinct views into your performance data, each designed to answer a different question.

The Timeline

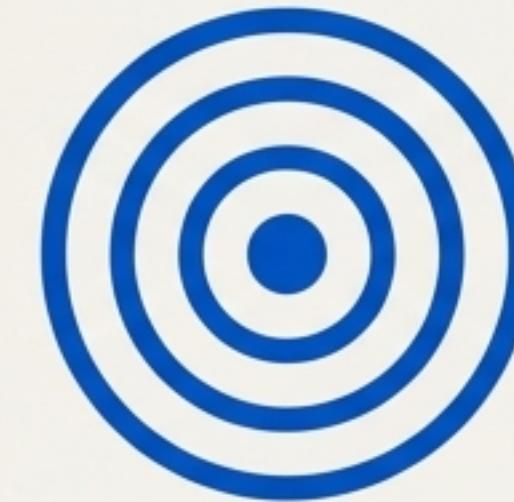


What is the sequence of events?

`print_timeline()`

Shows the chronological flow and call hierarchy.

The Hotspots



Where is the CPU actually spending its time?

`print_hotspots()`

Ranks methods by the work done **inside** them.

The Full Report

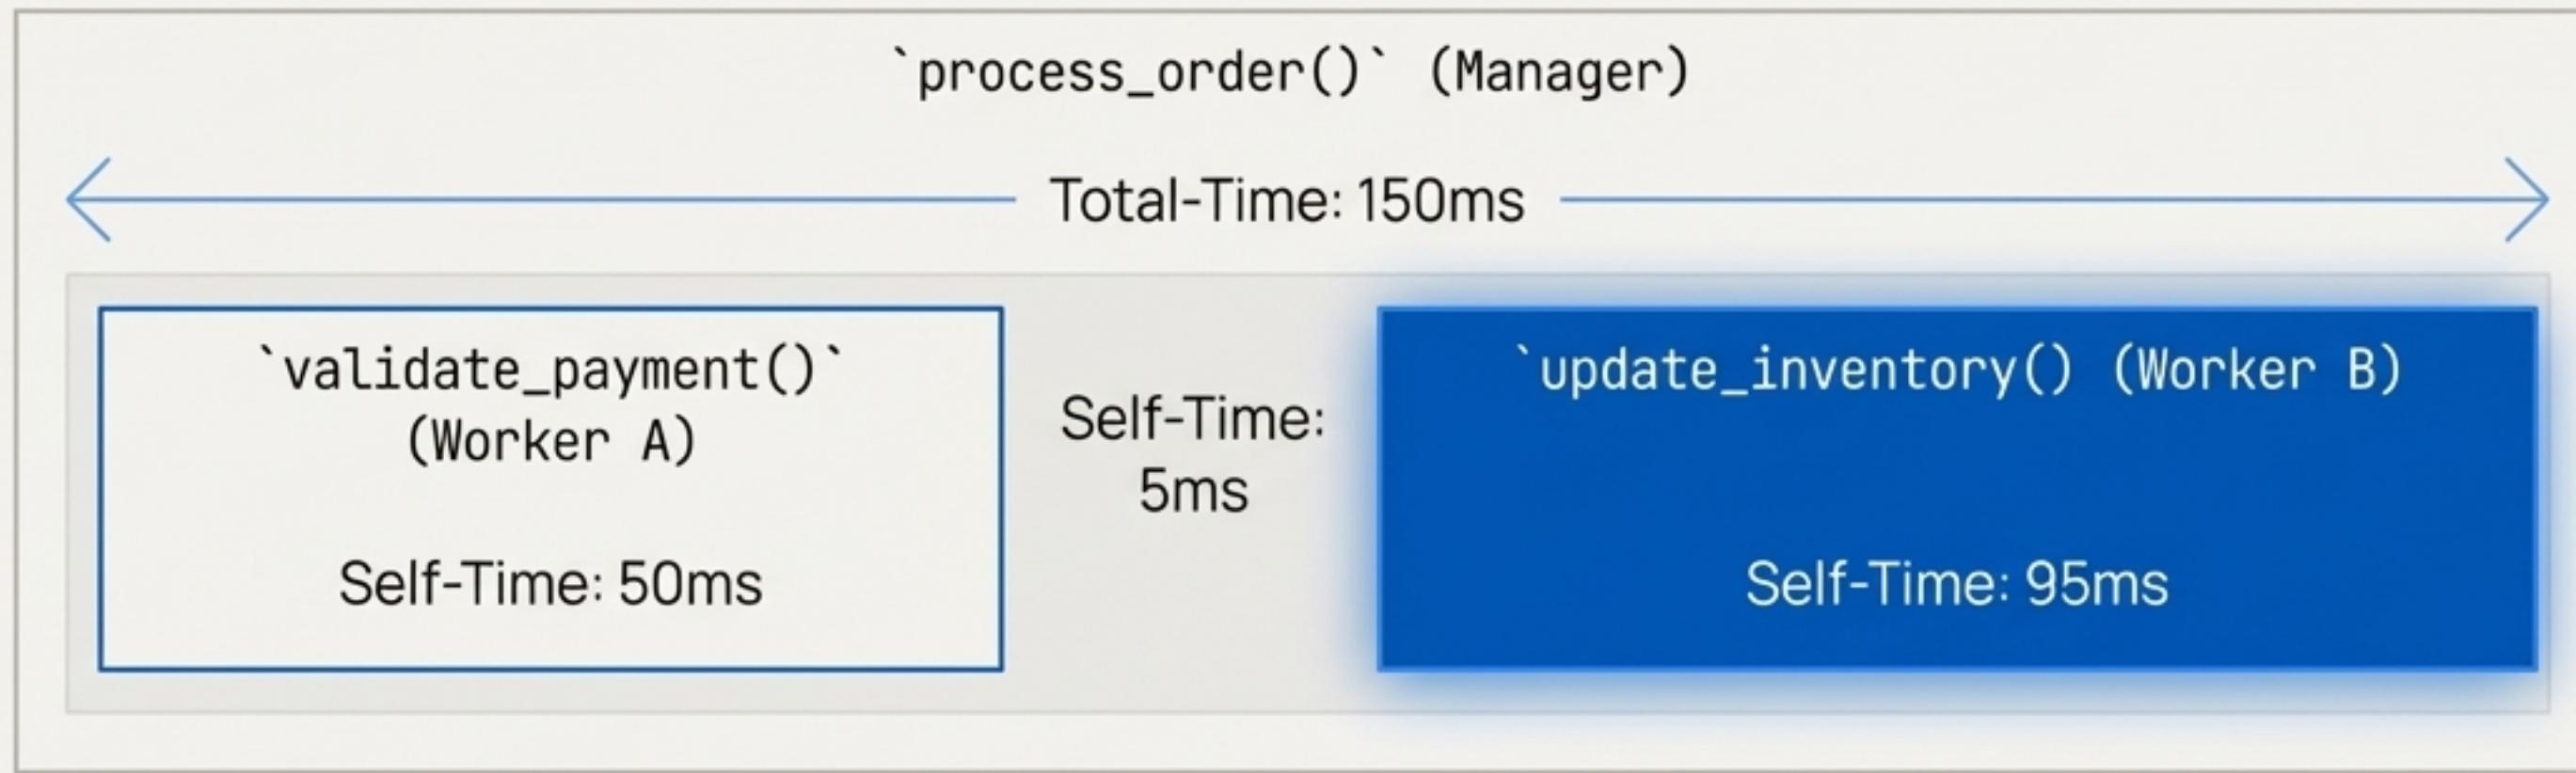


What are the complete timing details?

`print_report()`

A comprehensive table of all captured metrics.

The Most Important Concept: Self-Time vs. Total-Time



Total-Time: The entire time spent in a method, *including* all the time spent in other methods it calls (its children). A ‘busy manager’.

Self-Time: The time spent executing code *directly within* a method, *excluding* time spent in its children. This is the actual work done. **This is what you need to optimise.**

Find the Real Bottleneck with the Hotspots Report

Don't be misled by high Total-Time. The `print_hotspots()` report sorts methods by **Self-Time** to show you where your program is actually spending its clock cycles.

-[Hotspots Report (by self time)]-			
Target	Self Time	Total Time	...
update_inventory	95.0ms	95.0ms	...
validate_payment	50.0ms	50.0ms	...
process_order	5.0ms	150.0ms	...



Interpretation

In this example, `process_order` has the highest Total-Time (150ms), but it's a red herring. The real work, and the place to start optimising, is `update_inventory`, which has the highest Self-Time.

Visualise Your Code's Execution Journey

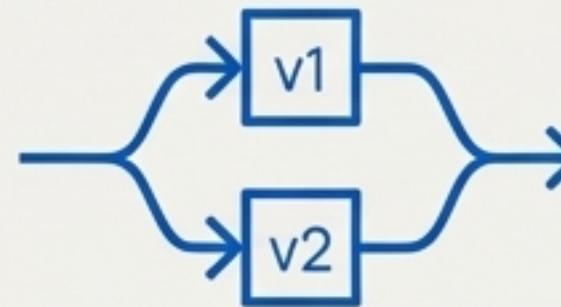
The `print_timeline()` report provides a chronological, indented view of your program's execution, making it easy to understand complex interactions.

```
▶ root.process_order      (t=0.00ms)
  ▶ a.validate_payment    (t=2.05ms)
  ▶ a.validate_payment    (t=52.15ms)
  ▶ b.update_inventory    (t=52.20ms)
    ▶ c.db_connection     (t=53.10ms)
    ▶ c.db_connection     (t=73.50ms)
  ◀ b.update_inventory    (t=147.30ms)
  ◀ root.process_order    (t=150.00ms)
```

Use this report to:

- ◆ Understand execution flow and method call hierarchy at a glance.
- ◀ Identify unexpected or recursive call sequences.
- ▶ Spot methods that are called far more times than you expect.
- ▬ Visualise where time is spent in deeply nested operations.

Putting It All Together: Common Scenarios



Profiling a Data Pipeline

- **Goal:** Find the slowest stage in a multi-step data conversion process.
- **Method:** Use `@timestamp(name="pipeline.stage.X")` on each major function. Use `print_hotspots()` to immediately identify the bottleneck stage by its high Self-Time.

Comparing Implementations

- **Goal:** A/B test a new algorithm against an old one.
- **Method:** Decorate both with `name="alg.v1.process"` and `name="alg.v2.process"`. Run both and compare their average execution times in the final `print_report()`.

Continuous Performance Monitoring

- **Goal:** Track performance over time and catch regressions.
- **Method:** Leave decorators in production code. Use `get_method_timings()` to programmatically access the data and send it to a monitoring service (e.g., StatsD, Prometheus).

Best Practices for Effective Profiling

DO

- **Decorate entry points and key methods:** Start at a high level to understand the overall flow.
- **Use descriptive collector names:**
`_timestamp_collector_ =
Timestamp_Collector("Process X").`
- **Use `timestamp_block` for phases:** Isolate and time important logic within larger methods.
- **Use `print_hotspots()` first:** It's the fastest way to find actionable bottlenecks.

DON'T

- **Decorate micro-functions in hot loops:** The overhead can become significant and create noise. Profile the loop itself instead.
- **Forget the magic variable name:** It must be exactly `_timestamp_collector_`.
- **Misinterpret Total-Time:** Don't optimise a method with high Total-Time but low Self-Time.

Troubleshooting Common Issues

Problem: Decorator isn't capturing any data.

Likely Cause: The collector variable is not named exactly `_timestamp_collector_` or is out of scope.

Solution: Check for typos. Ensure the collector is defined in a frame that is an ancestor of the decorated function call.

Problem: A method shows 0ms Self-Time.

Likely Cause: This is expected if the method only calls other decorated methods and performs no other significant work itself. It's a 'manager' method.

Solution: This isn't an error. Look at the methods it calls to find the real work.

Problem: The report is too noisy, with too many entries.

Likely Cause: You have decorated a small utility function that is called thousands of times inside a loop.

Solution: Remove the decorator from the inner function. Instead, wrap the entire loop in a `timestamp_block` to measure its total impact.

Your Checklist for Codebase Instrumentation

Use these steps to integrate `timestamp_capture` into your project today.

- Import** `@timestamp` and `Timestamp_Collector`.
- Add** `@timestamp` to key methods (entry points, major processing steps).
- Use** `name="..."` for hierarchical naming where clarity is needed.
- Use** `timestamp_block` for code phases not in dedicated methods.
- Create the collector** and name the variable exactly ` _timestamp_collector_`.
- Run your code** and generate reports.
- Start** with `print_hotspots()` to find the real bottlenecks (Self-Time).
- Use** `print_timeline()` to understand the execution flow.

Remember: The decorators are production-safe and can be committed to your codebase.