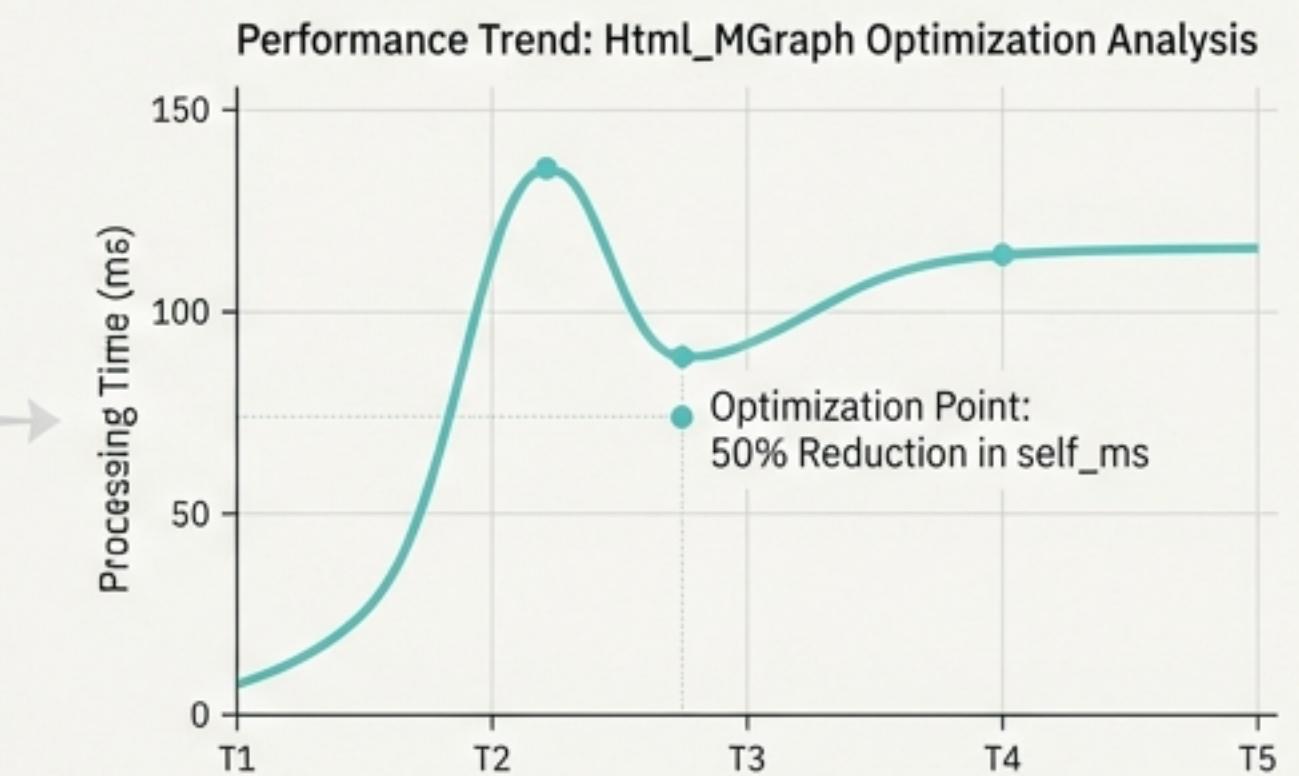
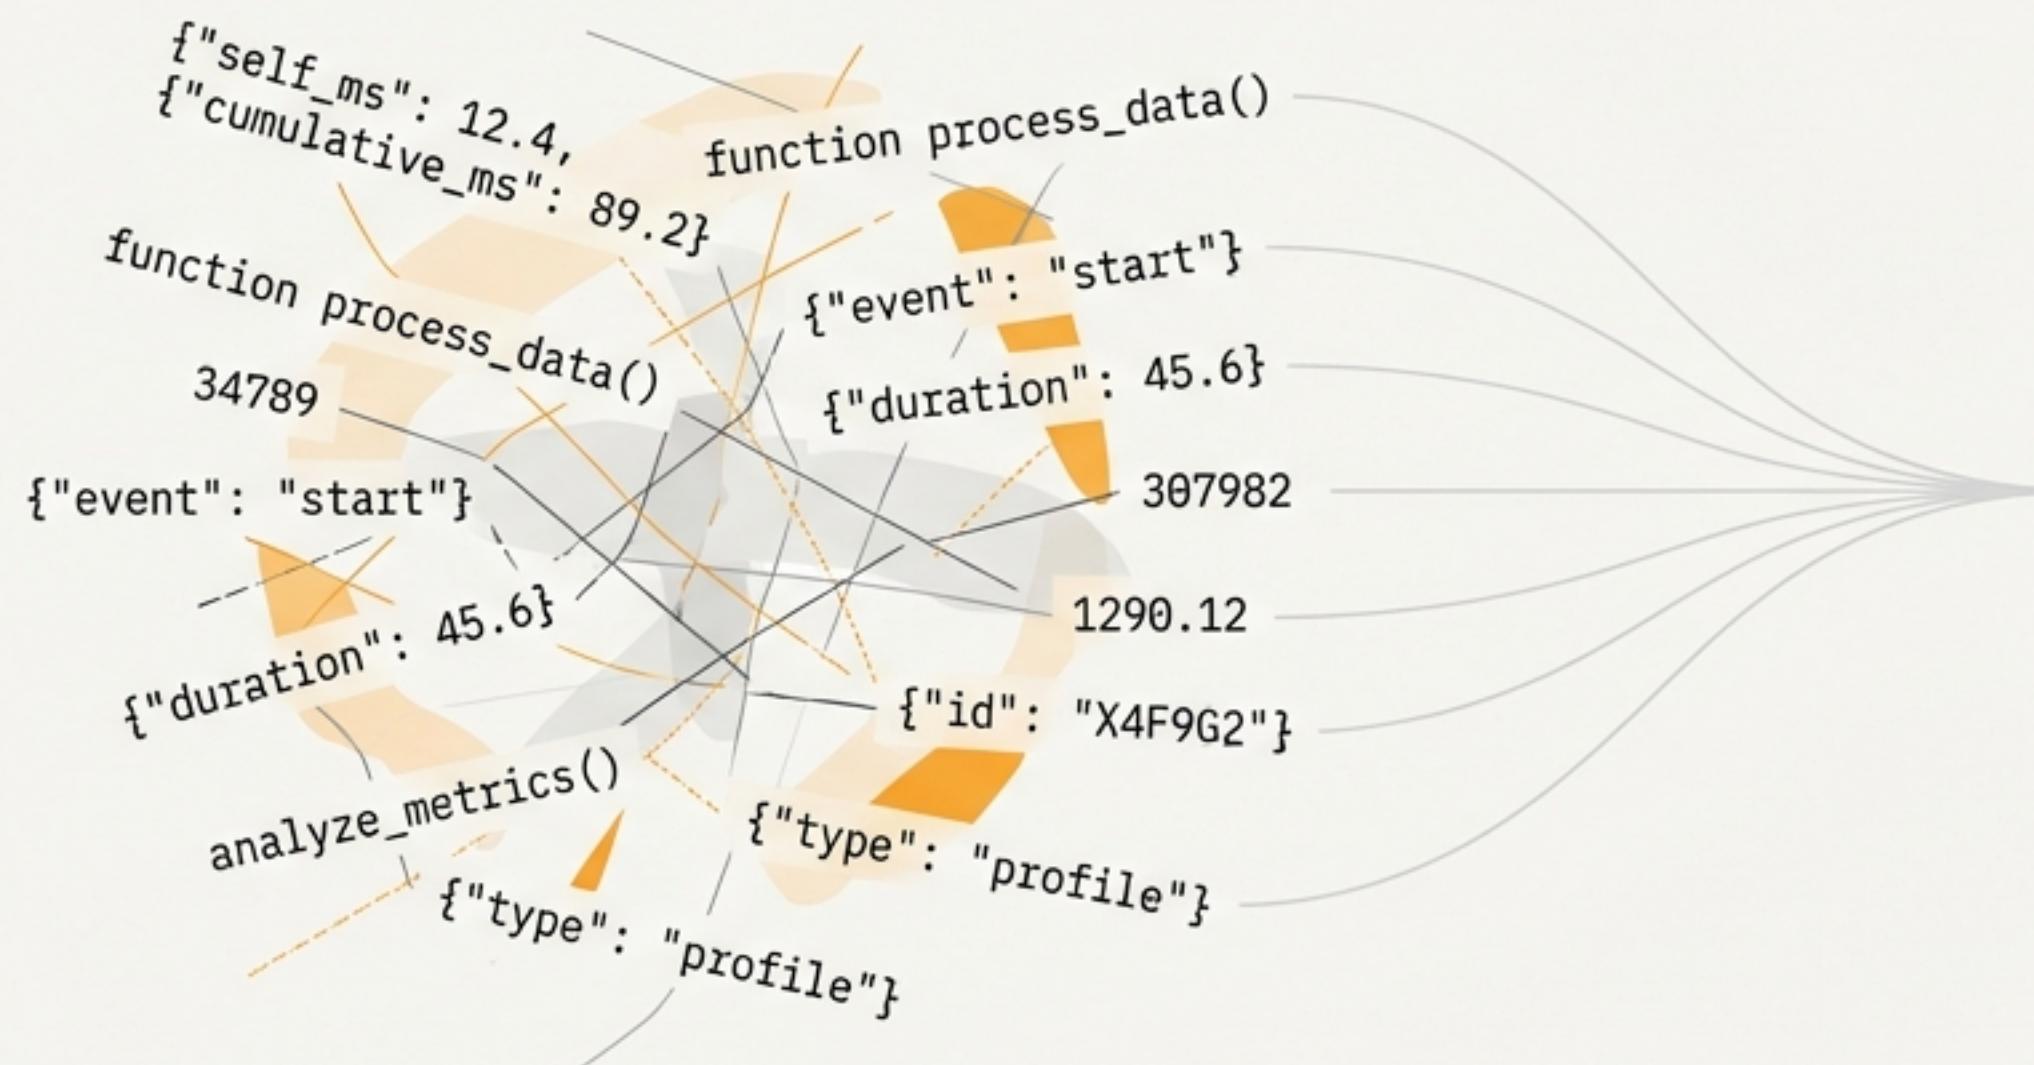


From Raw Data to Actionable Insight

A Case Study in Building a Purpose-Built Performance Analysis Workbench



Project: "Profile Analyzer"
Context: "Supporting Html_MGraph Performance Optimisation"
Stack: "A standalone, interactive React UI"

Our Starting Point: Rich Data Trapped in Raw JSON

The Profiling System

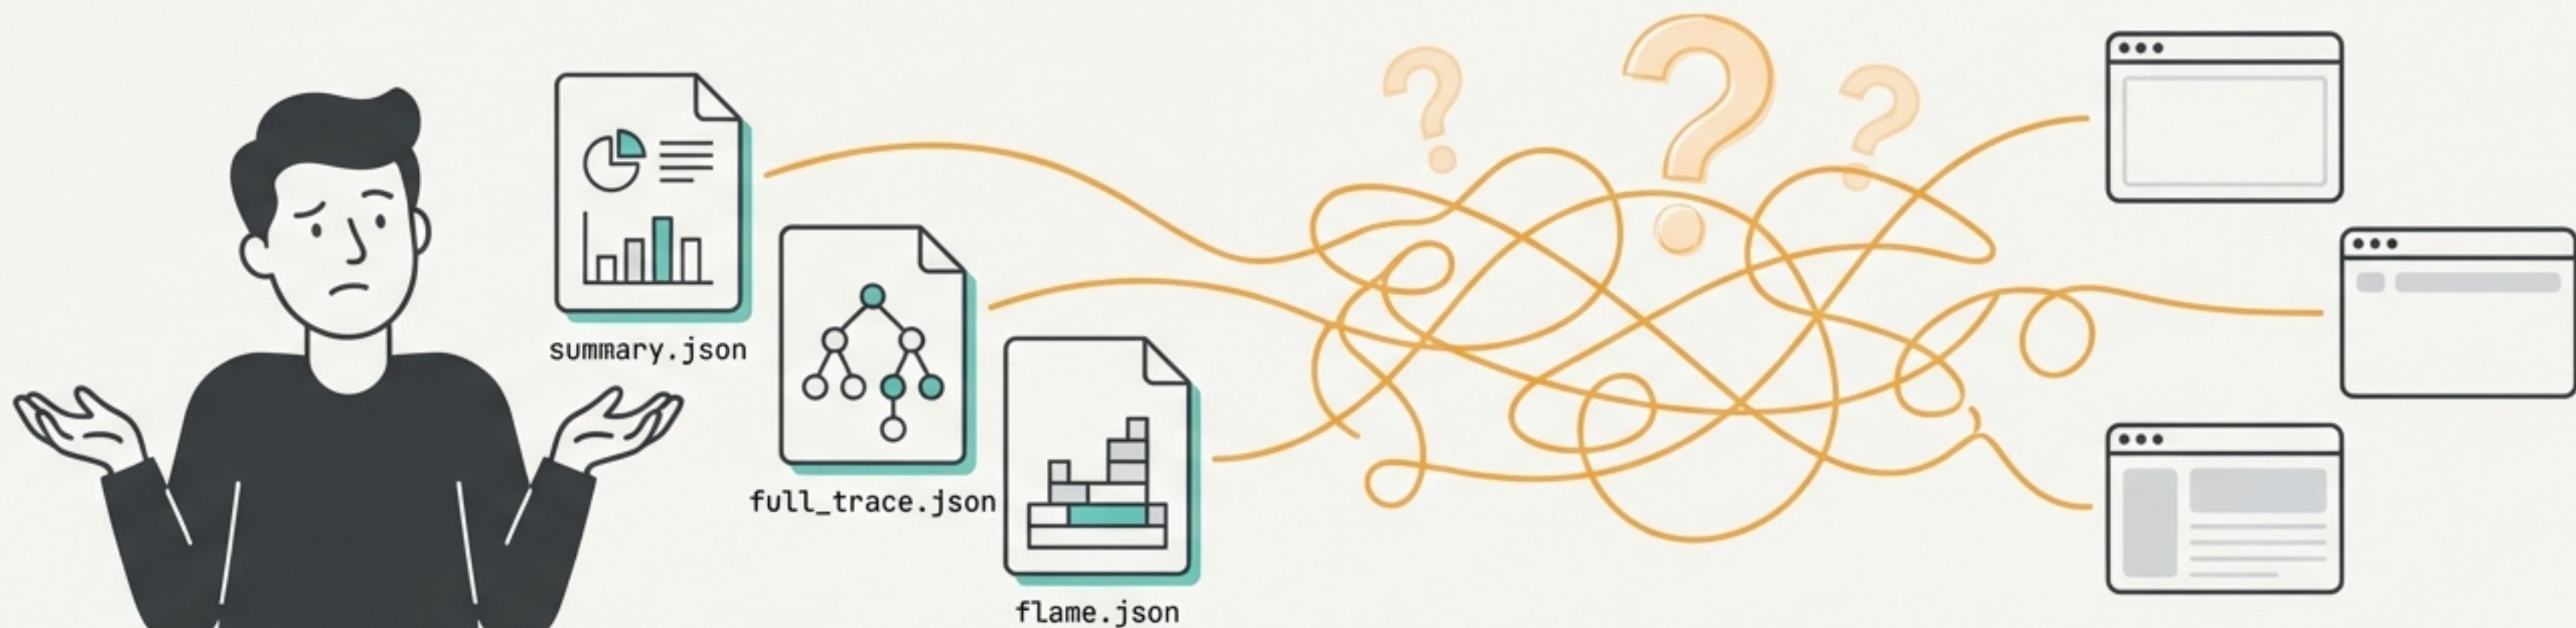
The source of the data: a sophisticated system using the `@timestamp` decorator from `osbot_utils`.

- * **Summary files:** Aggregated hotspot analysis.
- * **Full trace files:** Complete execution call trees.
- * **Speedscope files:** Flame graph format.

The Analytical Bottleneck

The existing workflow was manual, fragmented, and inefficient, creating a barrier to insight.

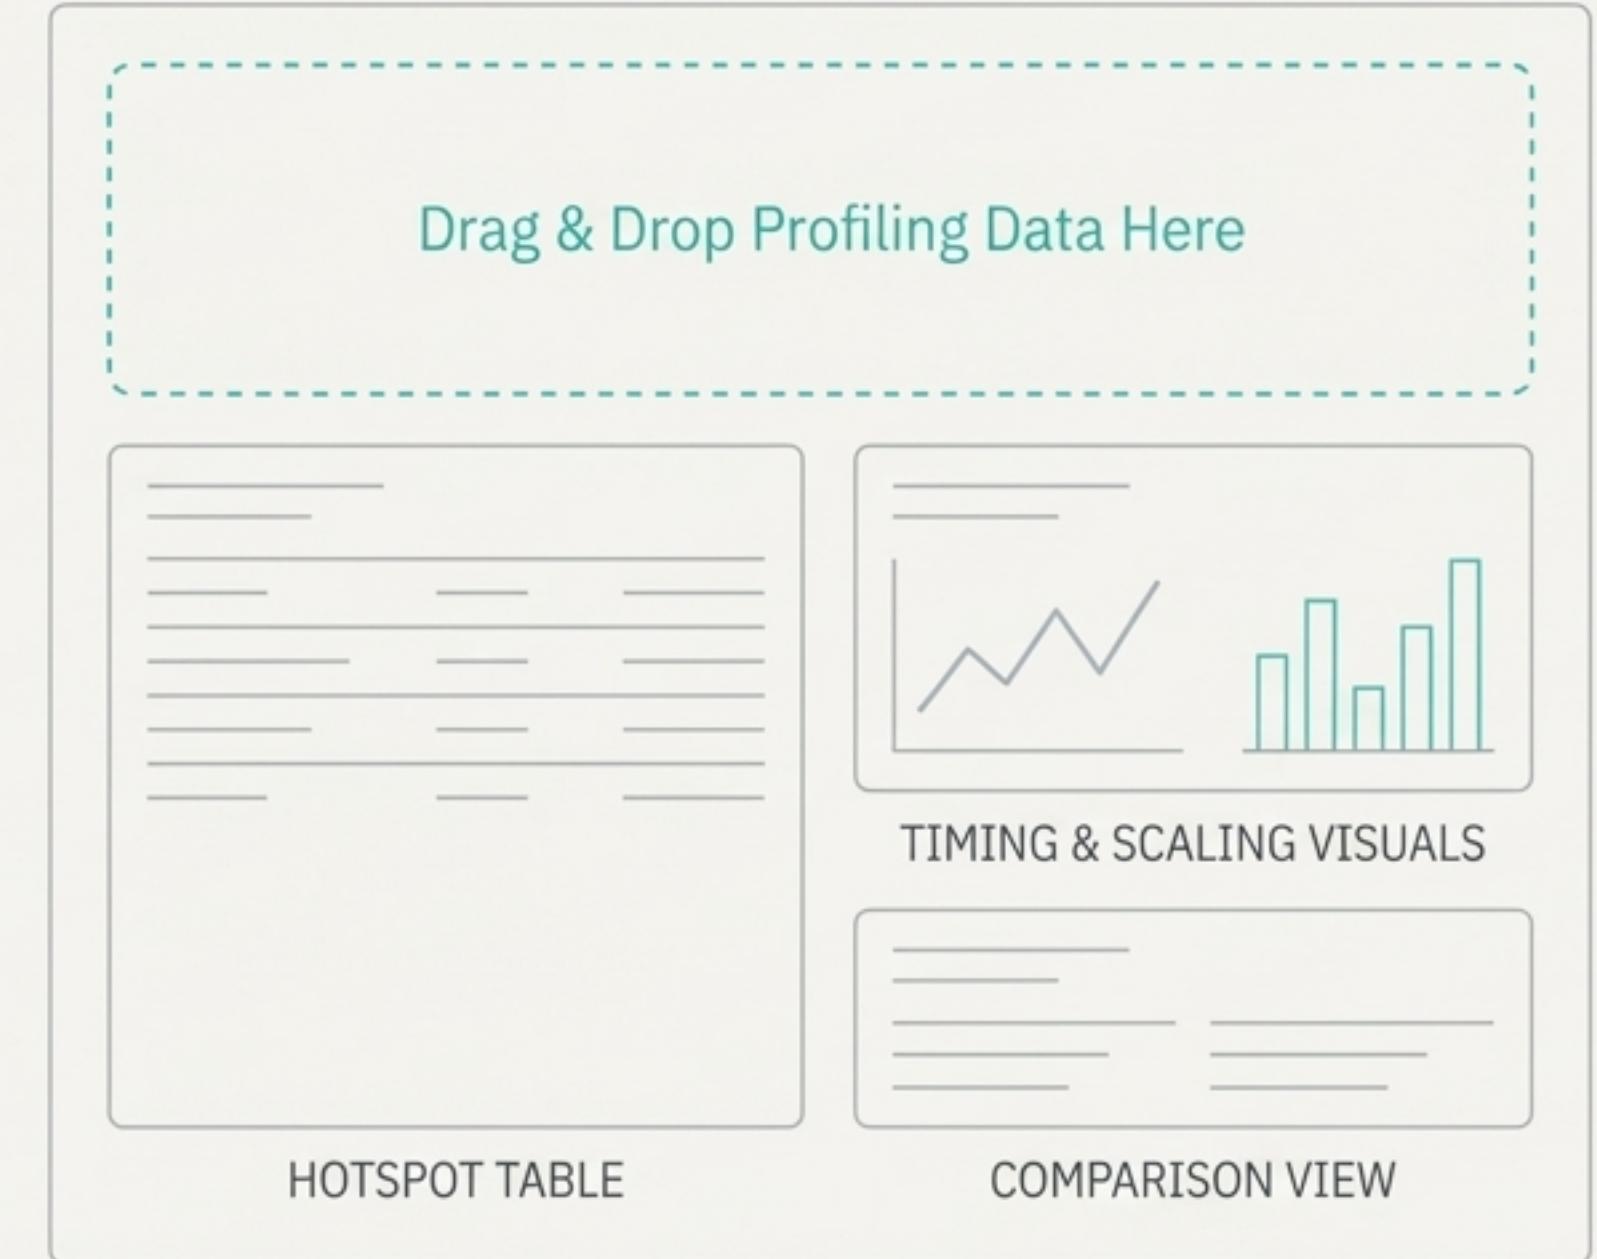
- * Manually reading complex JSON files.
- * Using generic tools that lacked specific context.
- * A clear need for a tool that understood the data's structure and the user's workflow.



The Goal: An Interactive Workbench for Performance Diagnosis

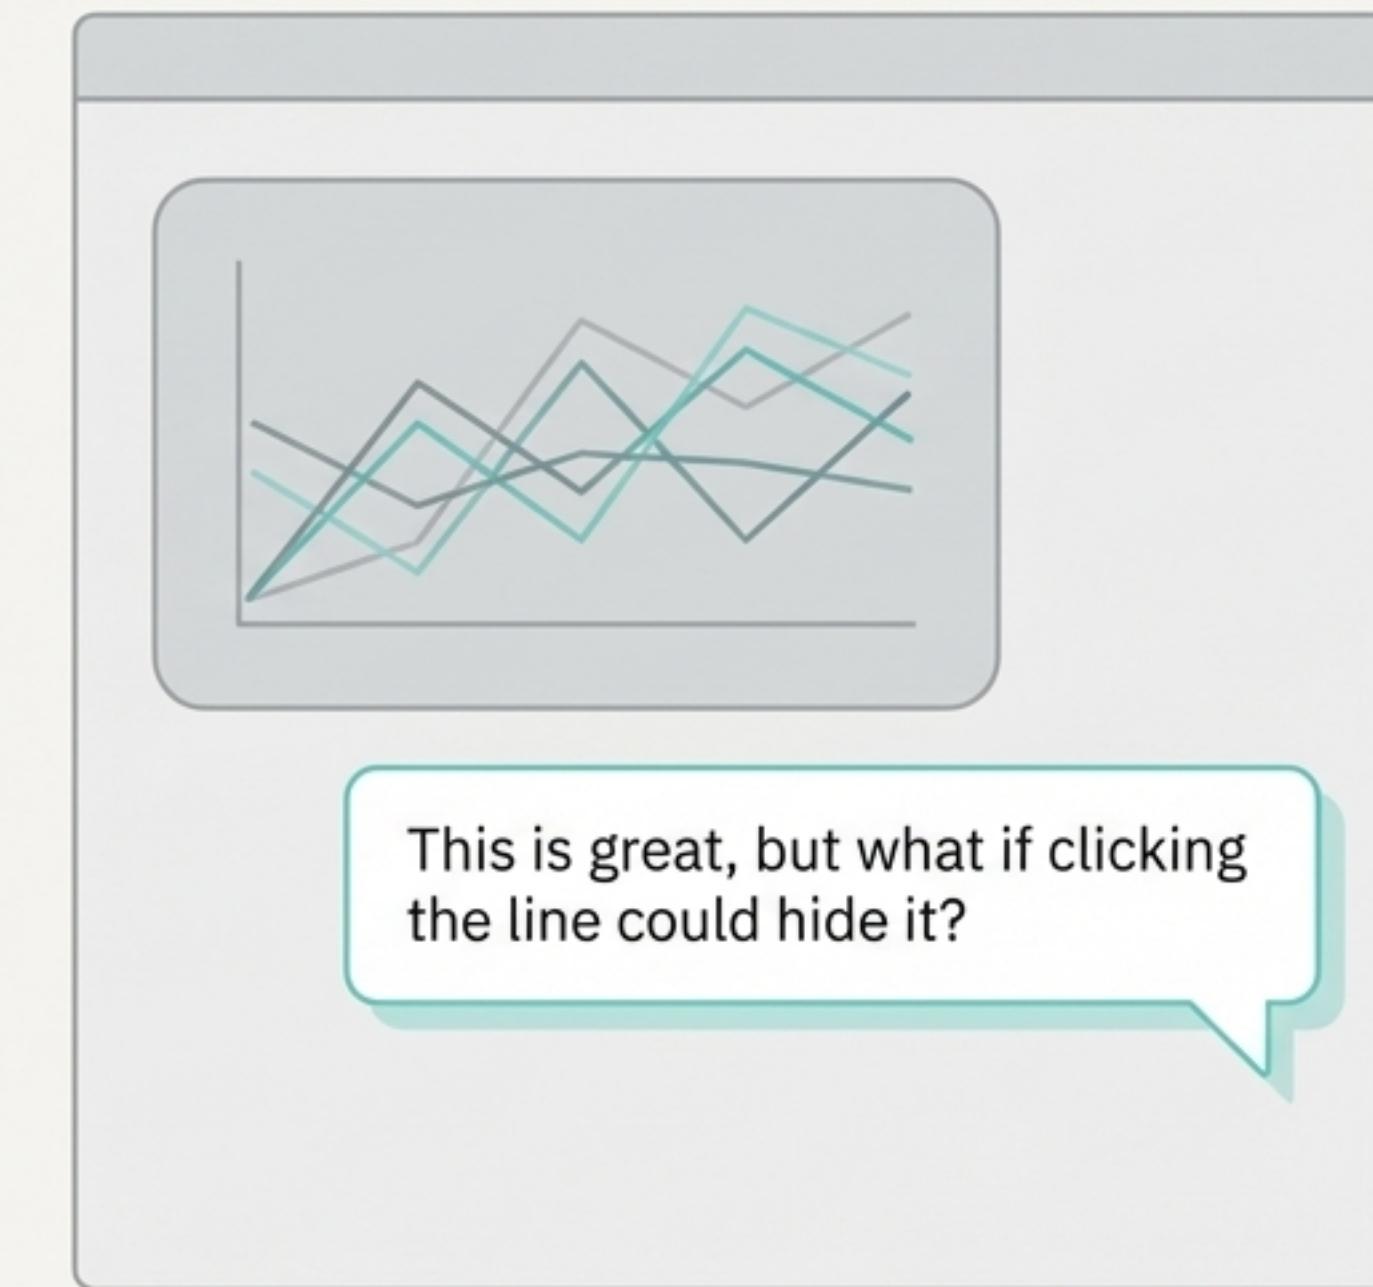
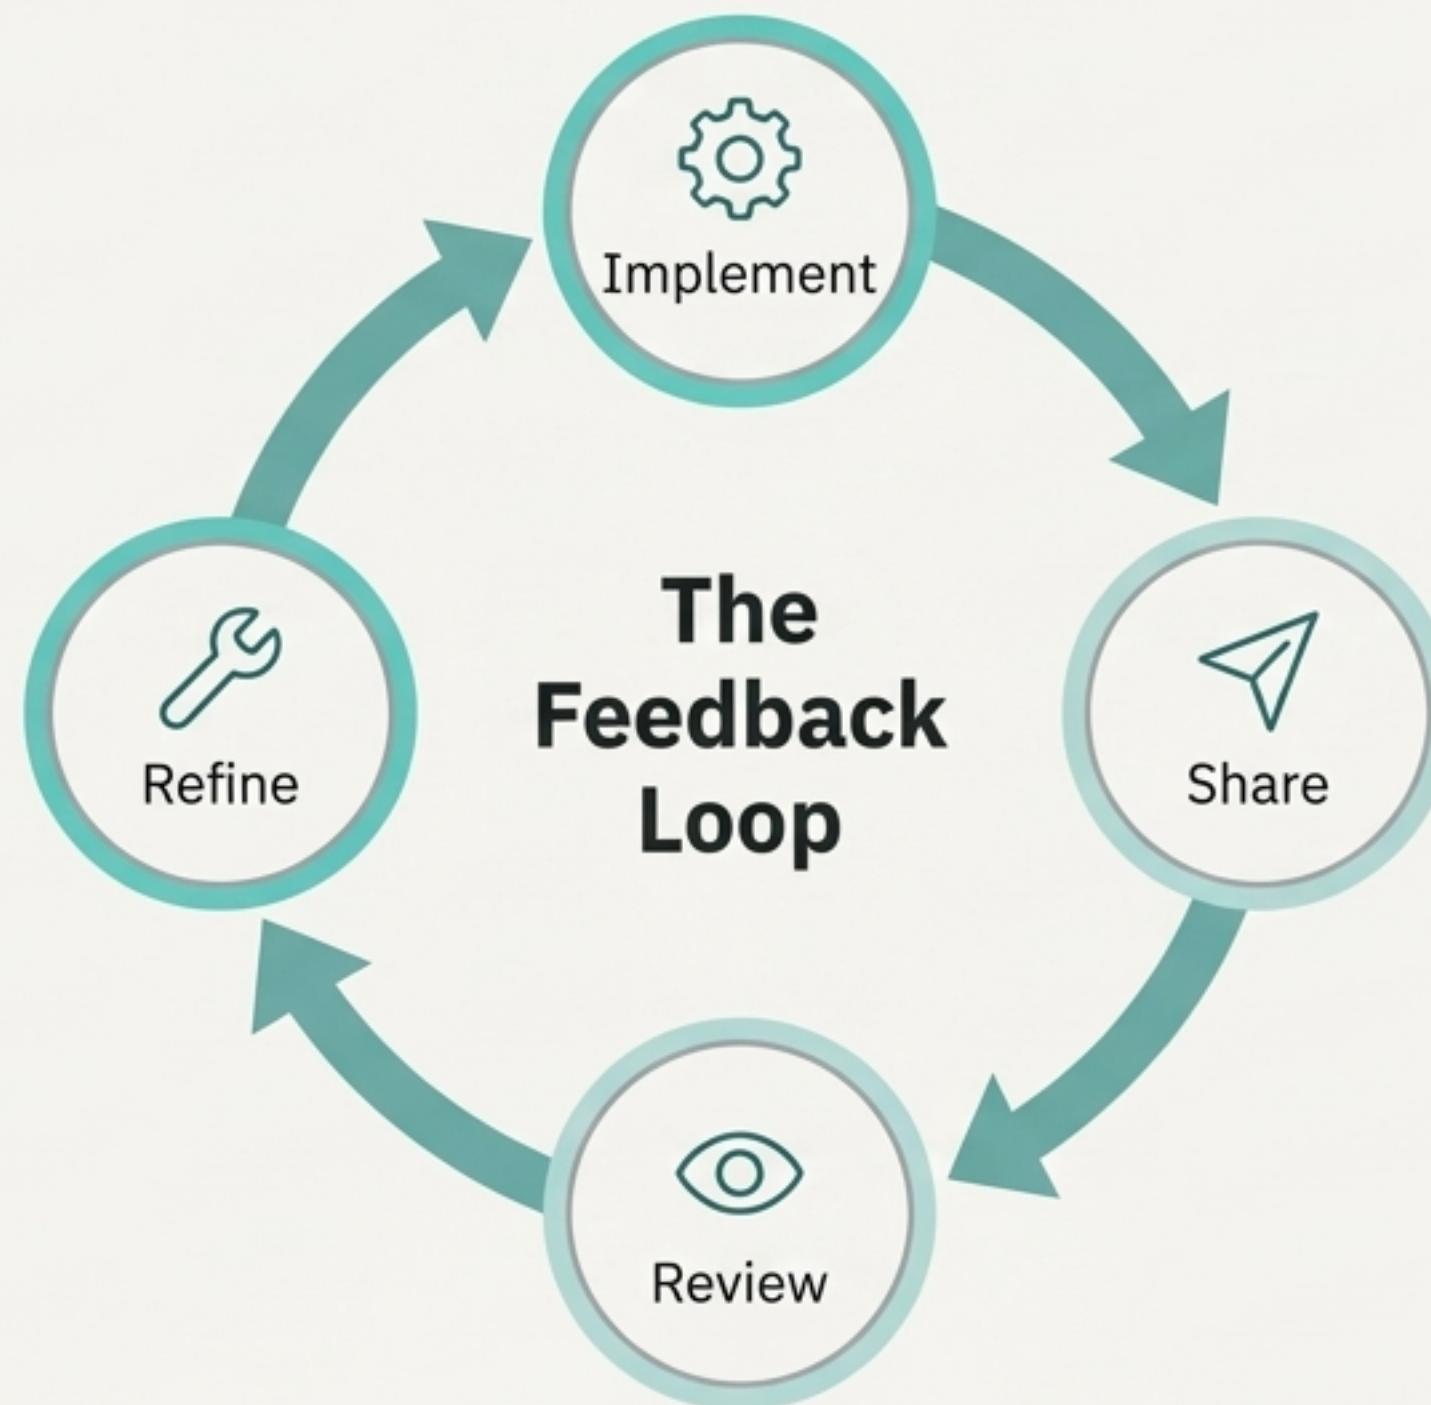
To create a web UI that could transform the analysis process from a chore into an exploration.

-  **1. Effortless Input:** Accept profiling data via drag-and-drop.
-  **2. Insightful Visuals:** Visualise hotspots and detailed timing data.
-  **3. Comparative Analysis:** Compare multiple profiling runs side-by-side.
-  **4. Scaling Behaviour:** Analyse performance across different input sizes.
-  **5. Zero Friction:** Work entirely in-browser with no backend required.



We Built the Tool Through Rapid, Feedback-Driven Iteration

What followed was a collaborative process spanning 15-20 refinement cycles.



A Key Insight Led to a Dual-Mode Architecture

We recognised that summary and full trace files serve fundamentally different analytical purposes. The UI should reflect this.



Summary Mode



- **Purpose:** Quick hotspot identification and A/B testing.
- **For Files:** summary_*.json
- **Features:** Single Profile View, Comparison View, Scaling View.



Full Mode

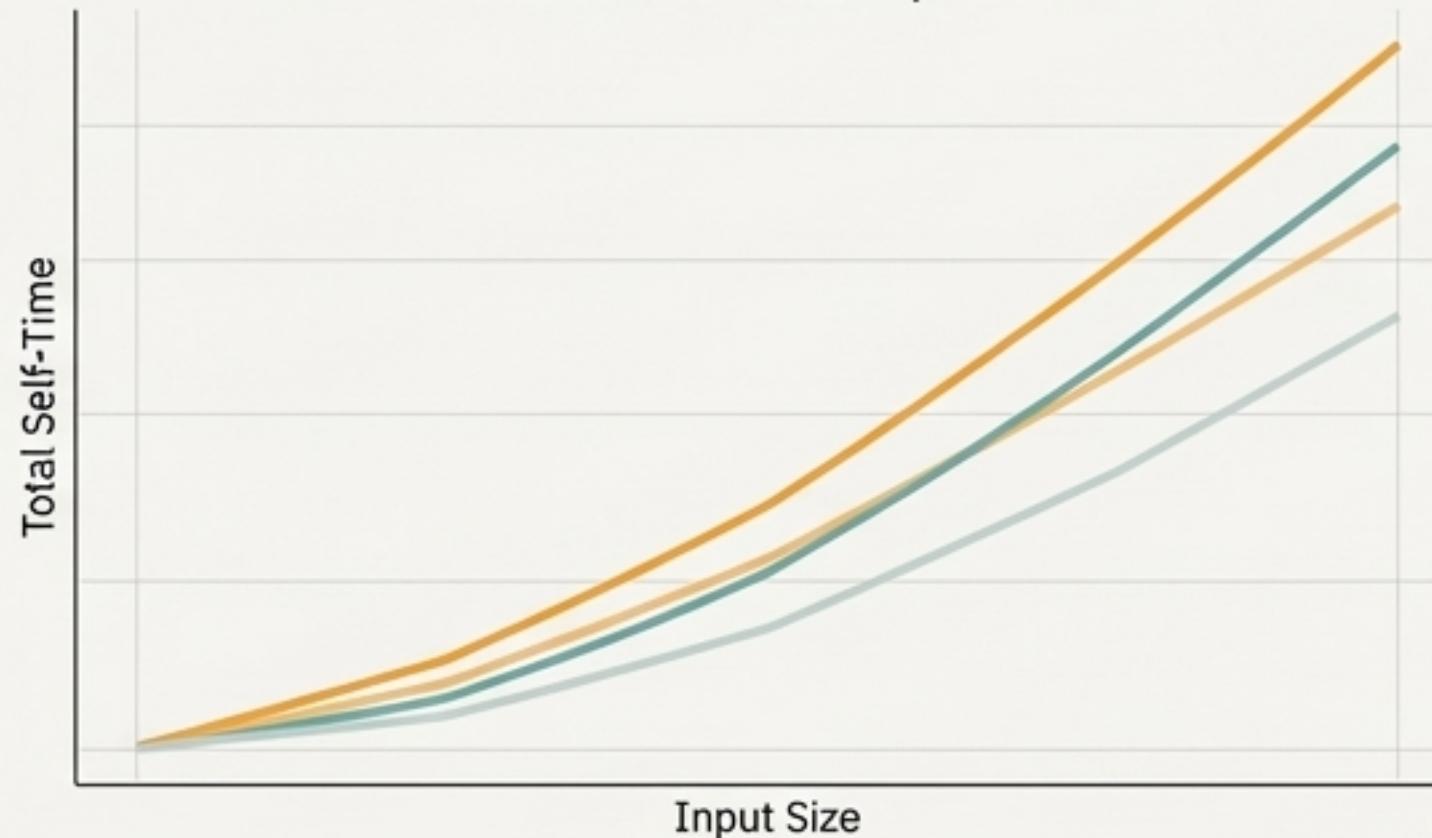


- **Purpose:** Deep algorithmic analysis and execution path exploration.
- **For Files:** full_*.json
- **Features:** Method Stats, Per-Call Scaling, Call Tree Explorer.

The Breakthrough: Realising Per-Call Averages Uncover True Algorithmic Complexity

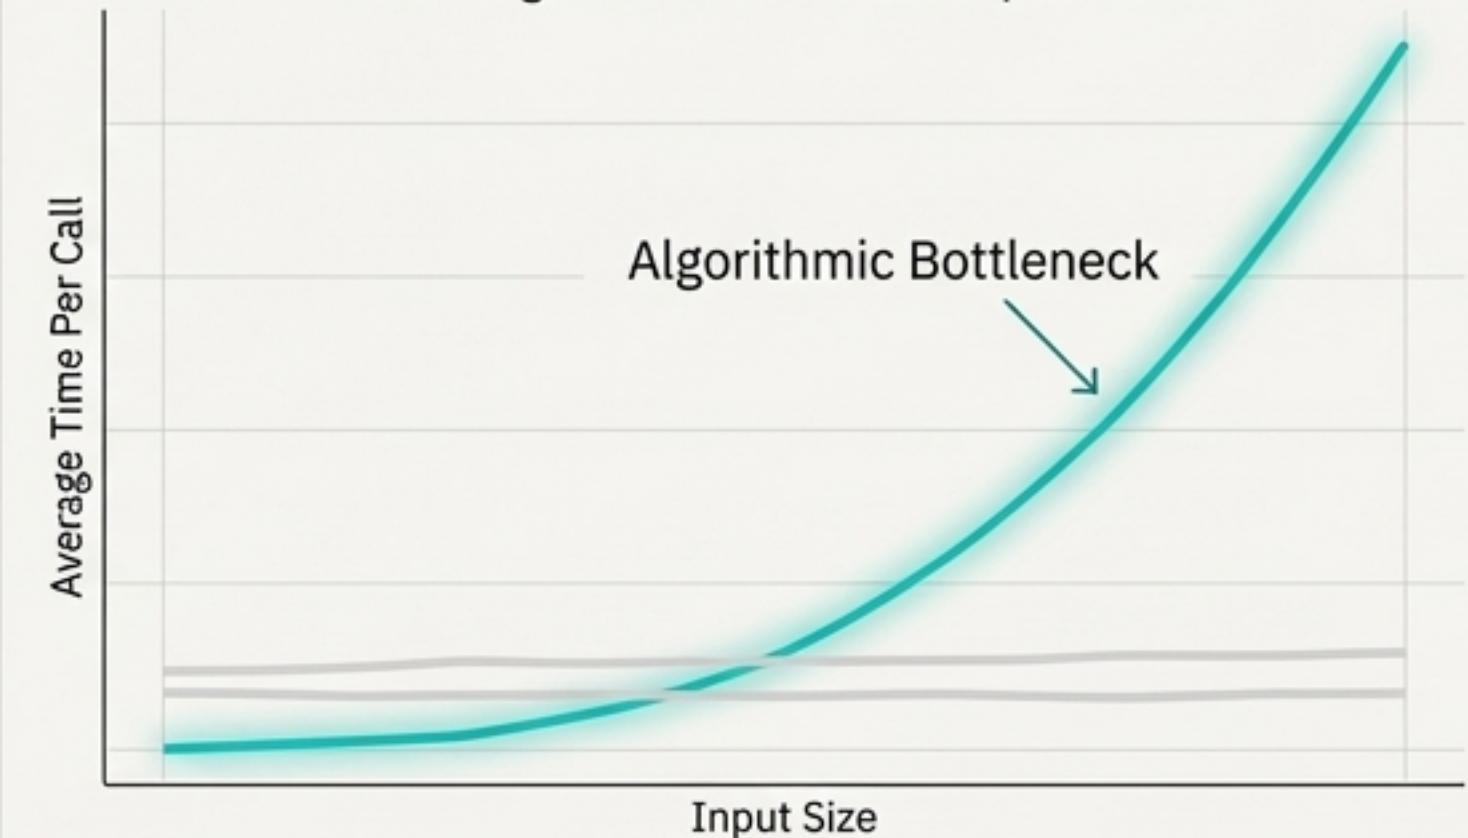
Misleading

Total Self-Time vs. Input Size



Insightful

Average Time Per Call vs. Input Size



The Flaw: Plotting **total self-time** vs. input size showed everything growing—expected but not useful.

The Insight: To find true $O(n)$ issues, we plot **average time per call**. A method with $O(1)$ complexity should show a flat line; a rising line is a red flag.

We Made Scaling Analysis More Robust by Letting the Data Describe Itself

The Problem ❌

The X-axis for scaling charts was brittle, parsing input size from filenames like `with_size_30`. This failed on real-world test files and URL-encoded names.

Before

`...with_size_30.json`



The Insight 📈

Let the data describe itself rather than imposing external naming conventions.

The Solution ⚙️

We pivoted to using the `entry_count` field from the file's own metadata, providing a universal measure of workload present in *all* profiling files.

After

```
{  
  "metadata": {  
    "entry_count": 30  
  }  
}
```

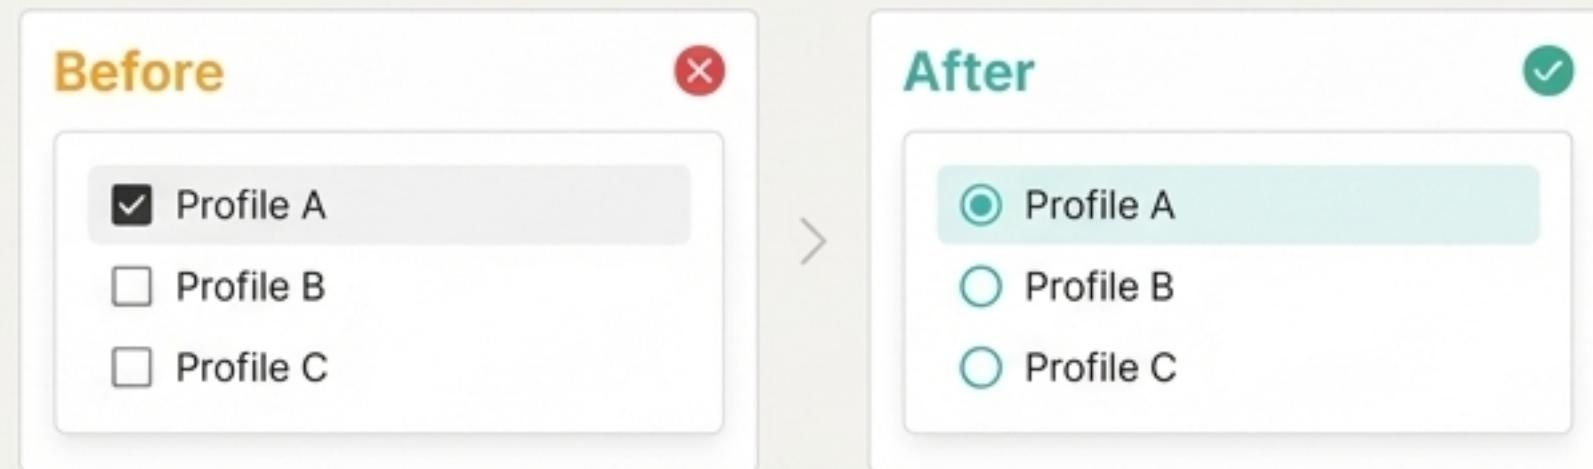


Every Interaction Was Refined to Reduce Friction and Enhance Exploration

1. Smarter Selection

Problem: Checkbox-style selection was awkward for single-profile views.

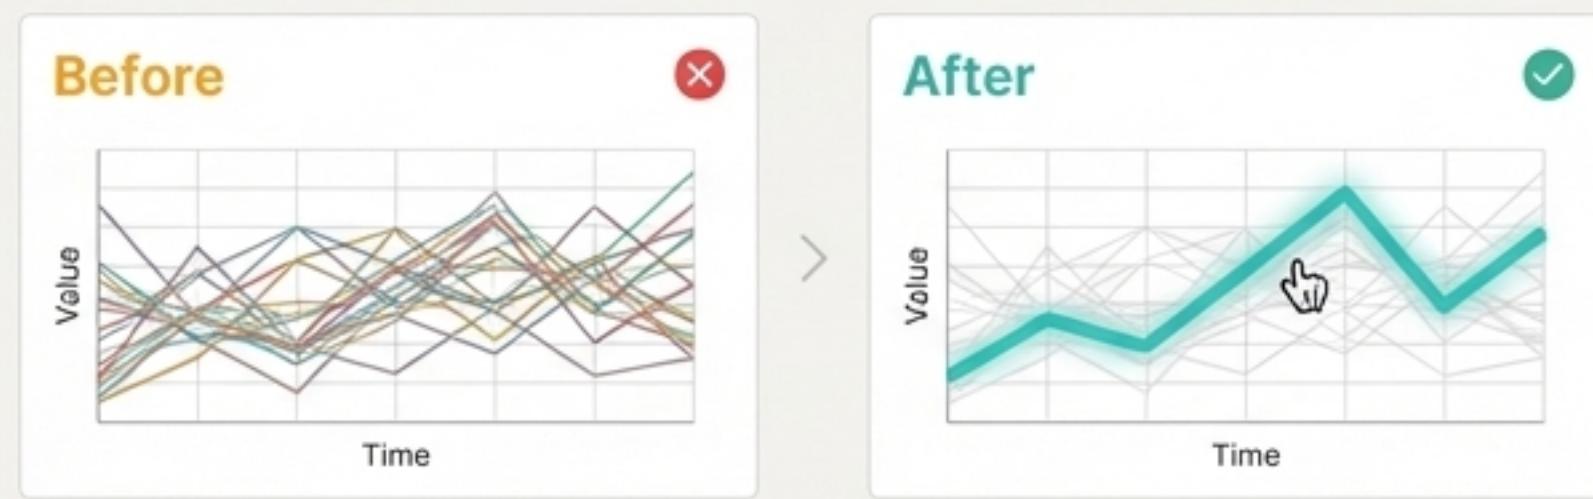
Solution: Implemented context-aware selection: radio-button behaviour in single views, checkbox behaviour in multi-select views.



2. Clearer Chart Lines

Problem: Charts with 12+ methods were cluttered and unreadable.

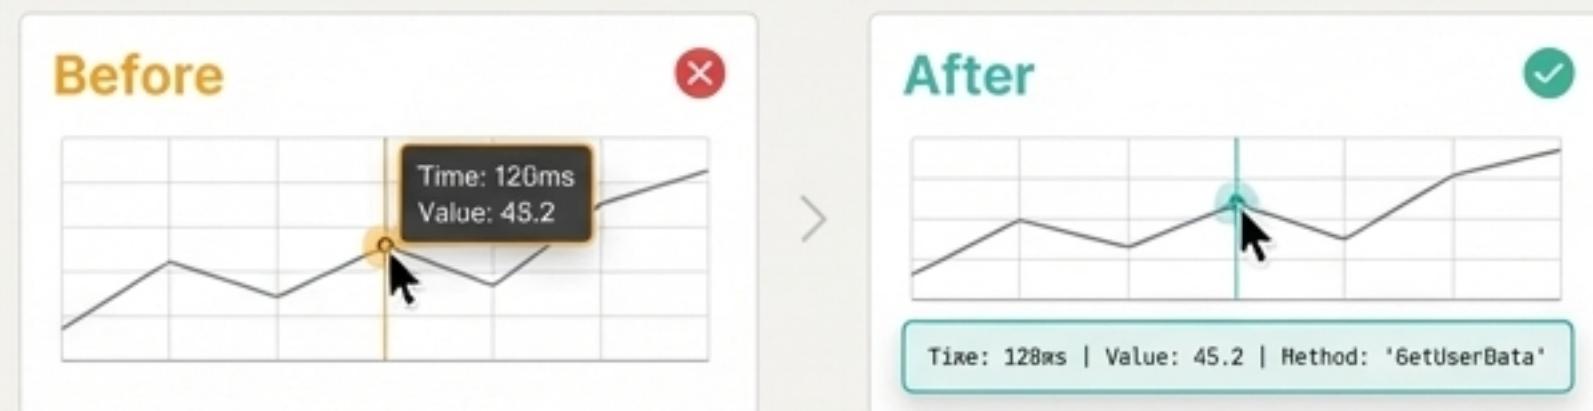
Solution: Added multiple interaction layers: table checkboxes, click-to-hide on lines, and hover highlighting where the focused line thickens and others fade.



3. Stable Tooltips

Problem: Default tooltip followed the cursor, obscuring data points.

Solution: Implemented a fixed tooltip panel below the chart that is always visible but never blocks the view.

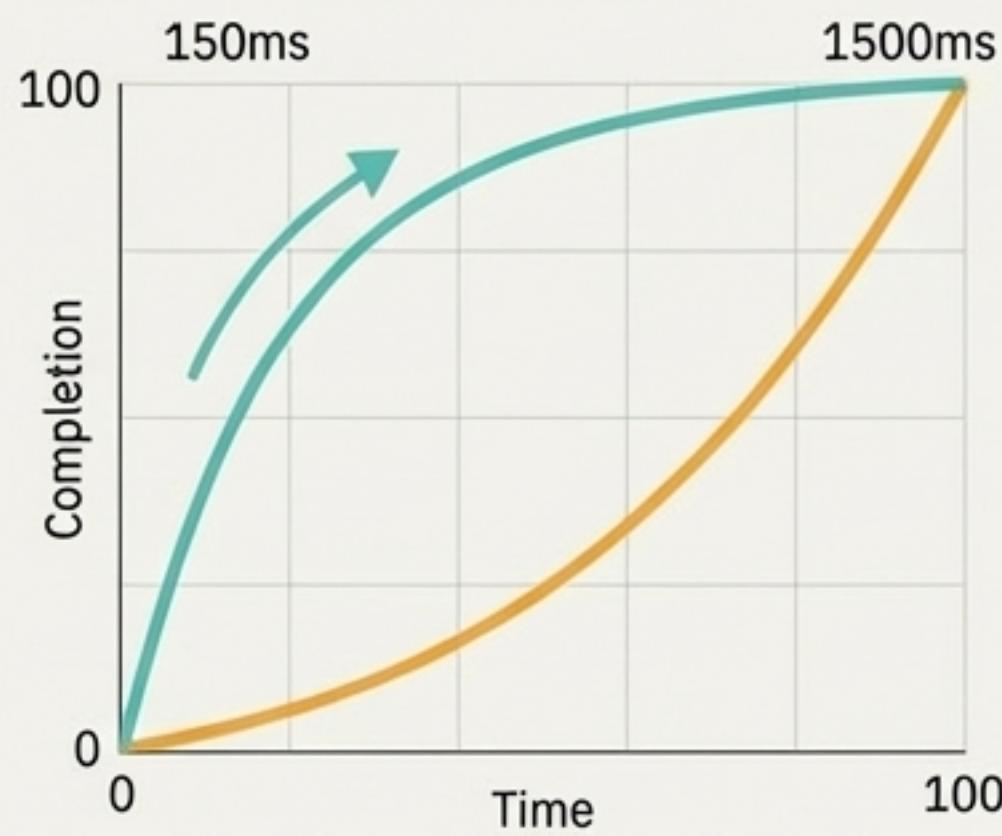


Final Polish Focused on Perceived Performance and Layout Stability

1. Fine-Tuning Animation

Problem: Default 1500ms Recharts animations felt sluggish when toggling line visibility.

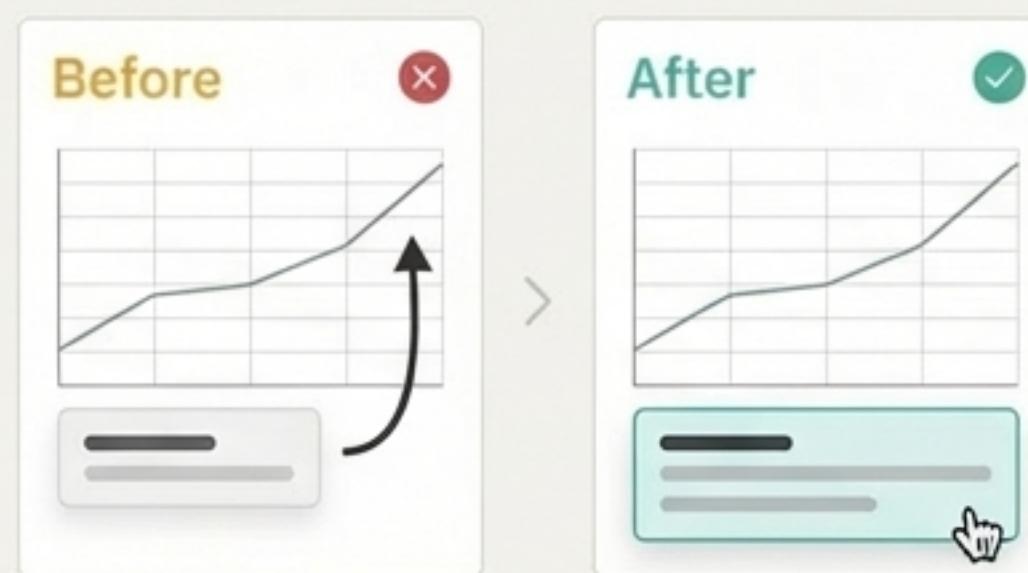
Solution: Reduced duration to 150ms (10x faster) for a responsive feel that maintains visual continuity.



2. Ensuring Layout Stability

Problem: The tooltip panel's changing height caused the entire chart layout to 'jump' during interaction.

Solution: Gave the panel a fixed height with overflow hidden to eliminate visual disruption.



Insight: Predictable layouts build user confidence; unexpected movement creates cognitive friction.

3. Maximising Focus

Feature: Added a full-screen overlay mode (+) for any chart, allowing for detailed analysis of dense visualisations without distraction.



The Result: A Summary Mode for High-Level Hotspots and A/B Testing

Single Profile Analysis

Implementation Comparison

Top-Level Scaling

Metrics Cards

Total Duration: 2.4s	CPU Time: 1.8s	Memory: 250MB
-----------------------------	-----------------------	----------------------

Hotspots

	Duration (ms)	Self Time (ms)	% of Total
DataProcessing::process	2.4s	53ms	23.2%
UI::render	1.8s	256ms	18.2%
Network::fetch	0.8s	258ms	7.8%
DataProcessing::test	0.8s	185ms	6.8%
	0.5s	185ms	3.3%
	0.5s	155ms	2.3%
	0.5s	158ms	0.2%
	0.5s	38ms	6.8%

Instantly identify hotspots

Profile A vs. Profile B

Method	Profile A (ms)	Profile B (ms)
Method A	~1.7	~1.2
Method B	~2.5	~2.1
Method C	~1.4	~1.8
Method D	~1.6	~1.5
Method E	~0.3	~0.4

Delta Table

Method	Profile A (ms)	Profile B (ms)	Delta (ms)
Method A	72ms	85ms	+120ms
Method B	20ms	40ms	-80ms
Method C	12ms	40ms	-28ms

Quantify performance improvements

Performance Trend

Input Size	Total Duration (s)
1x	~0.5
2x	~1.0
4x	~2.0
8x	~4.0

Quick overview of performance trends

NotebookLM

The Result: A Full Mode for Deep Algorithmic Investigation

Per-Call Scaling (The O(n) Detector)

Time (s)

Input Size

Trend (→) Trend (→) Trend (↑)

Pinpoint algorithmic complexity issues at a glance.

Method	Min (ms)	Max (ms)	Avg (ms)
DataProcessing::process	1.2	4.5	2.8
...			

Method Statistics

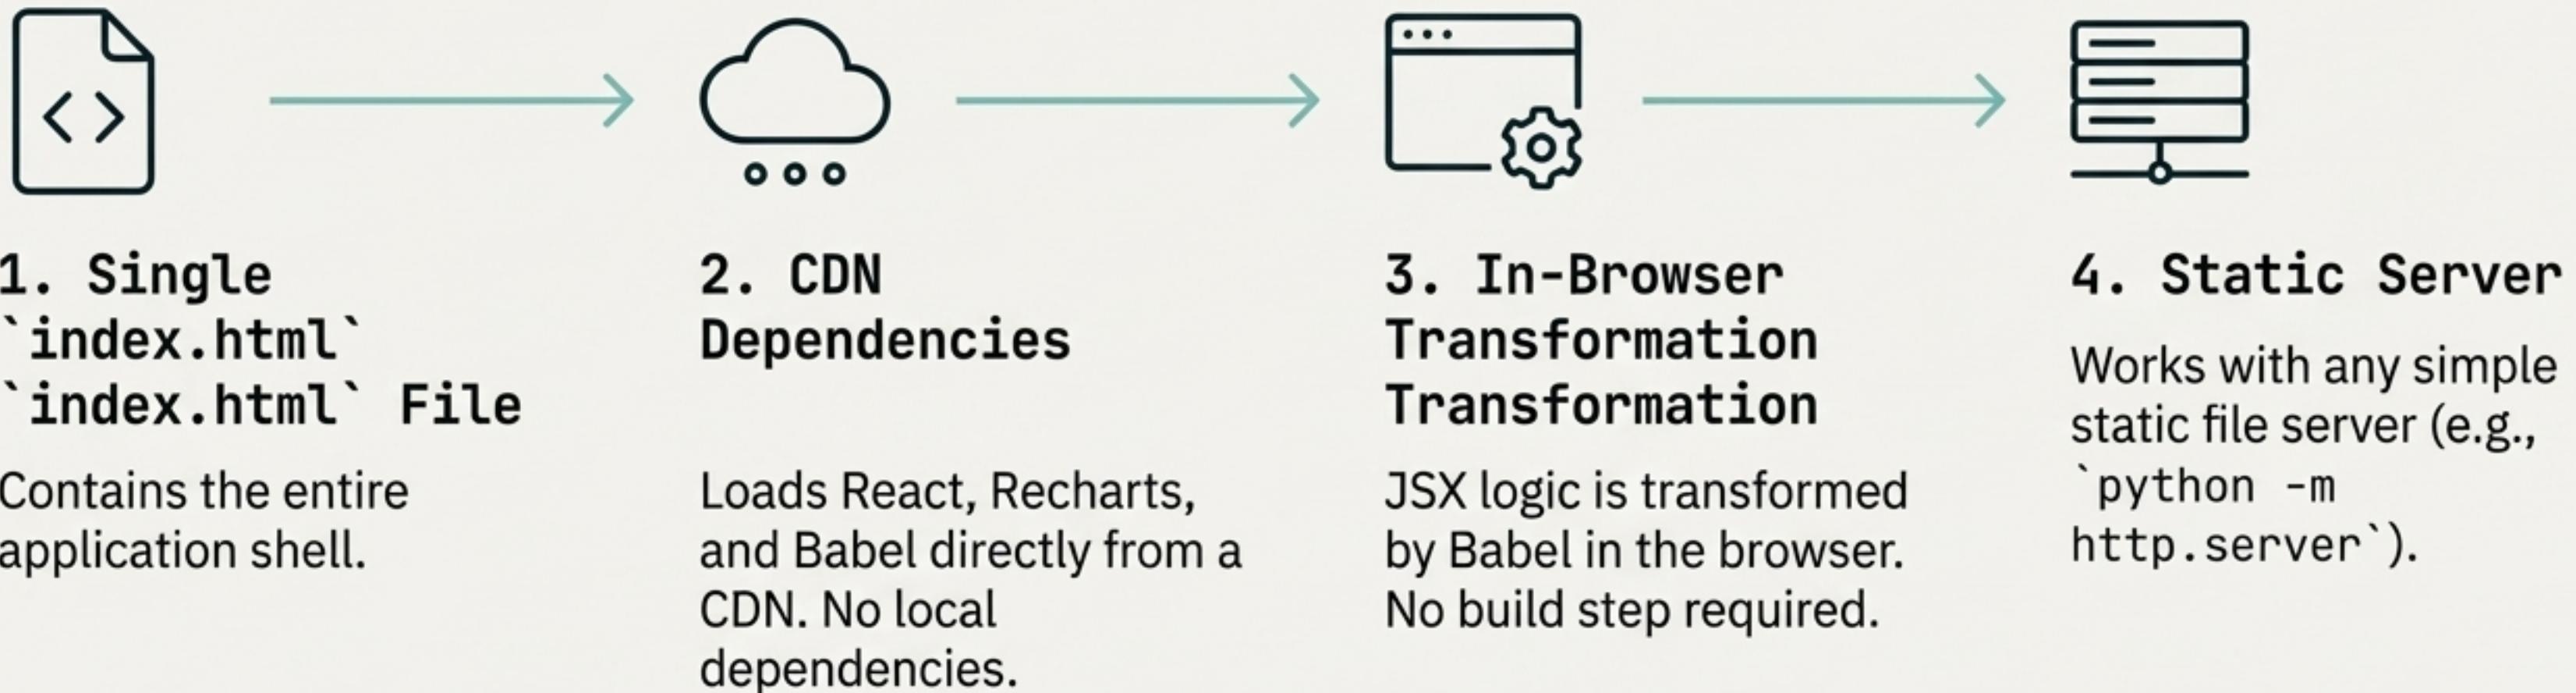
Call Tree Explorer

- ▼ Main
 - ▶ DataProcessing::process (2.8s) [diagonal lines]
 - ▶ UI::render (1.5s) [diagonal lines]
 - ▼ Network::fetch (0.5s) [diagonal lines]
 - ▶ fetch::data (0.2s) [diagonal lines]

NotebookLM

The Final Artefact is a Single, Portable HTML File

The project was packaged for maximum portability and zero installation friction.



Key Benefit: An incredibly simple distribution model that allows anyone to run the tool instantly.

Key Lessons from the Development Journey

-  • **Start with the data:** Its structure should drive every design decision.
-  • **Iterate with real feedback:** Rapid feedback loops are essential for refining interactions.
-  • **Separate concerns by use case:** Distinct interfaces for distinct analytical tasks (Summary vs. Full).
-  • **Make everything interactive:** Enable hypothesis testing, don't just display static visuals.
-  • **Prioritise layout stability:** Unexpected visual movement creates cognitive friction.
-  • **Use the data's own metadata:** It is always more robust than external conventions like filename parsing.

A Purpose-Built Workbench is More Than a Visualiser; It's a Workflow Enabler

What began as a need to “visualise some JSON” evolved into a specialised performance analysis workbench. The iterative process was essential in transforming a generic viewer into a diagnostic tool aligned with a real-world workflow.

Looking Ahead: Future Opportunities

-  Speedscope format integration for flame graphs.
-  Persistent storage of user preferences.
-  Automated anomaly detection and regression tracking.

The tool is immediately useful for the `Html_MGraph` project and generalisable to any codebase instrumented with the `@timestamp` system.