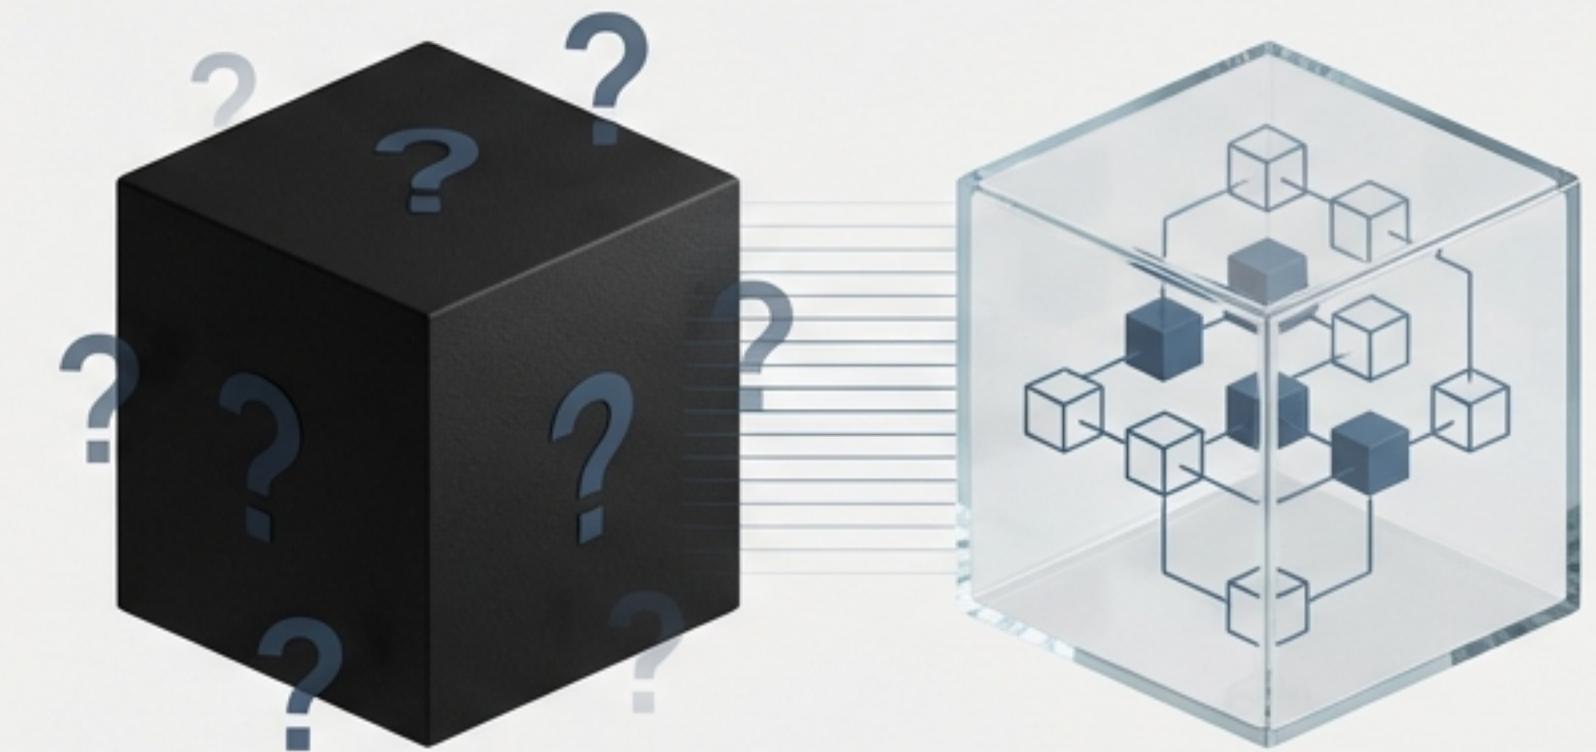


# From Black Box to Glass Box

A Proposal for Non-Invasive,  
On-Demand Performance Tracing



# When production services are slow, we need answers

Every developer has faced these critical questions when a transformation pipeline behaves unexpectedly:

-  **Where is the time being spent?** Which specific methods are the performance hotspots?
-  **What is the call flow?** How do the methods nest and in what sequence do they execute?
-  **How does this run compare?** Is this specific execution slower than previous ones, and if so, why?

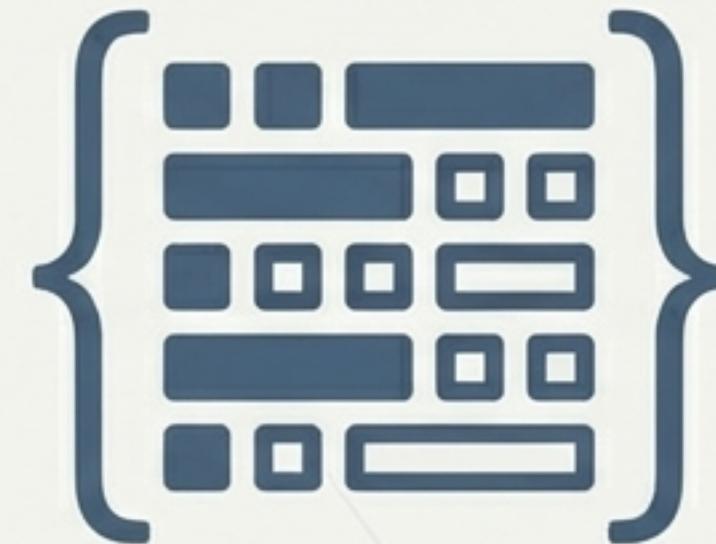
# We need a solution that is non-invasive and on-demand

Traditional profiling often requires attaching debuggers, modifying application code, or restarting services. Our approach is designed to be fundamentally better.



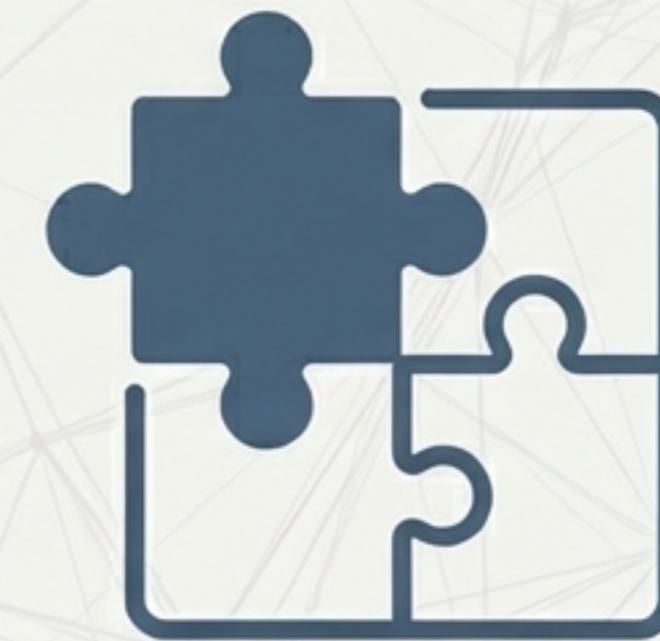
## Enable Per-Request

Activate detailed tracing for specific API calls without affecting the entire service.



## Return Structured Data

Output traces in a serialisable, machine-readable format for analysis.

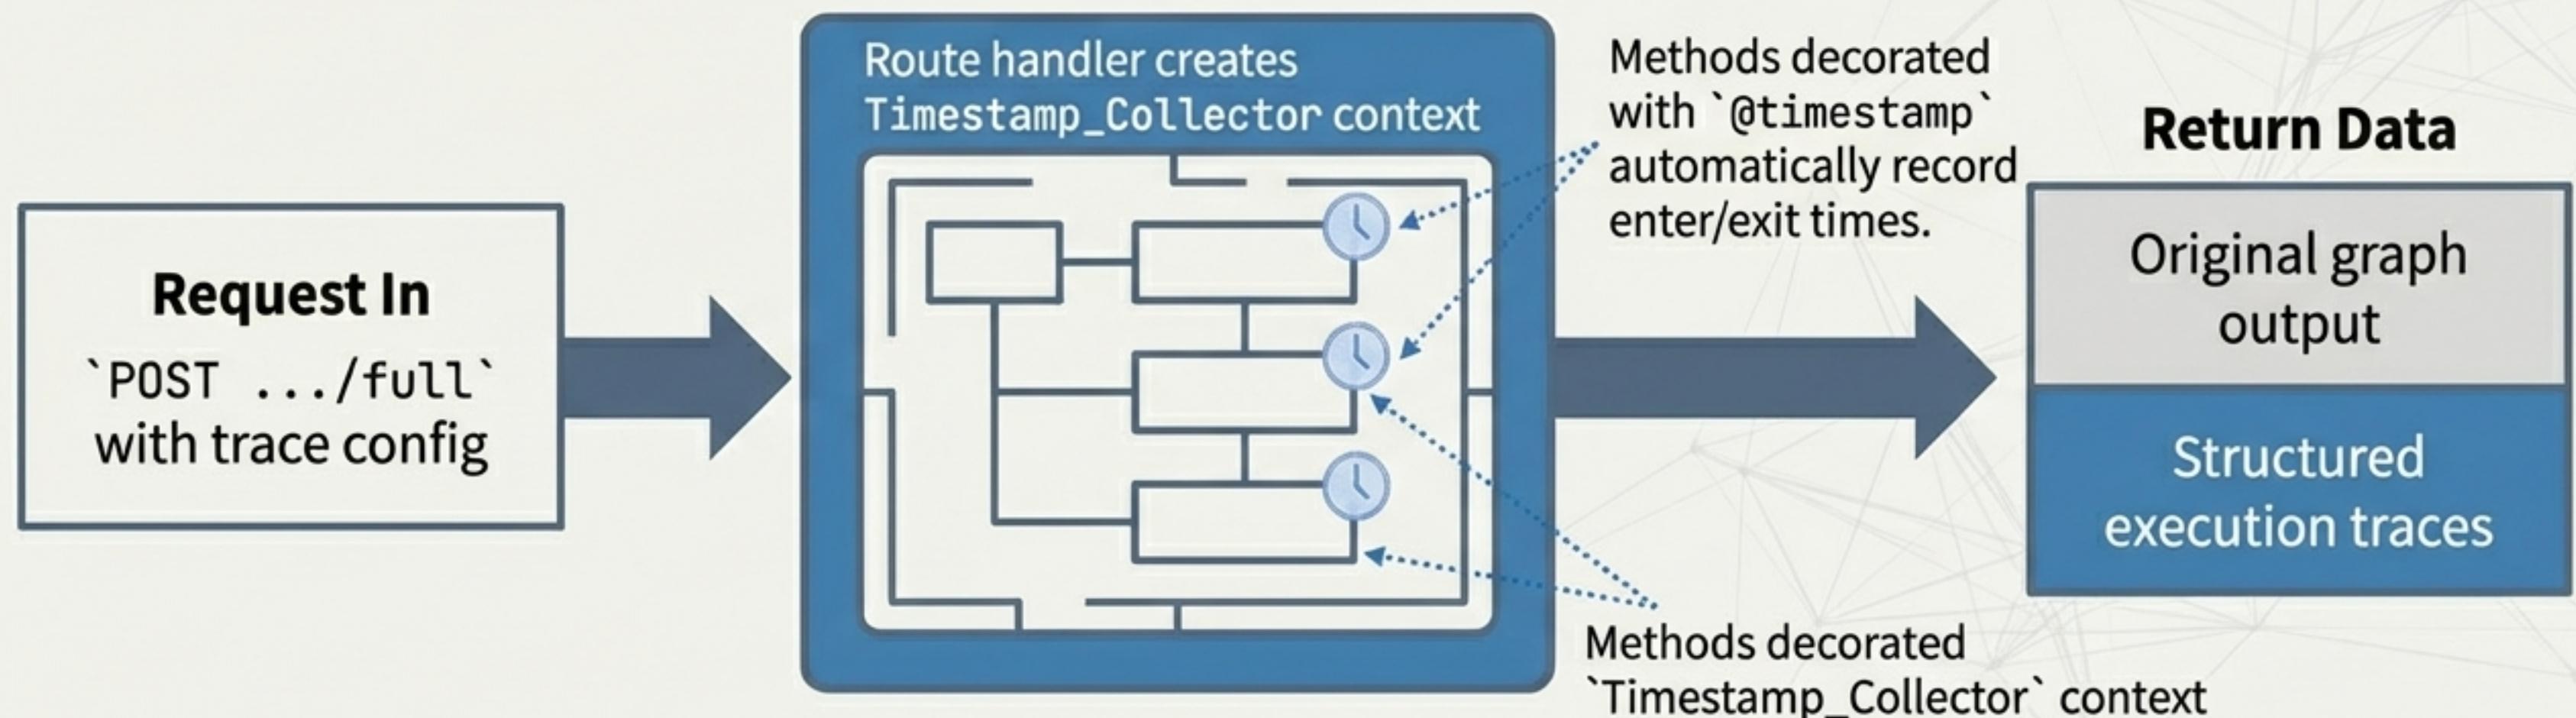


## Integrate with Tooling

Natively support visualisation tools like speedscope and flame graphs.

# The solution separates trace capture from core execution

The transformation pipeline remains completely unaware it is being traced. The route handler orchestrates the entire process, wrapping the execution and exporting the results.



**Separation of concerns:** The client chooses the format, the route handles the capture, and the core logic remains unchanged.

# Three new routes provide flexible access to execution traces

We will introduce three new routes that mirror the existing graph endpoints, but are enhanced with the ability to capture and return timestamp data.

Route Suffix	Response Type	Primary Use Case
.../full	Response__Full	Complete trace data for storage and deep analysis.
.../summary	Response__Summary	A quick, high-level overview of performance hotspots.
.../speedscope	Response__Speedscope	An interactive format for immediate visualisation in speedscope.app.

The API provides the graph output alongside the trace data

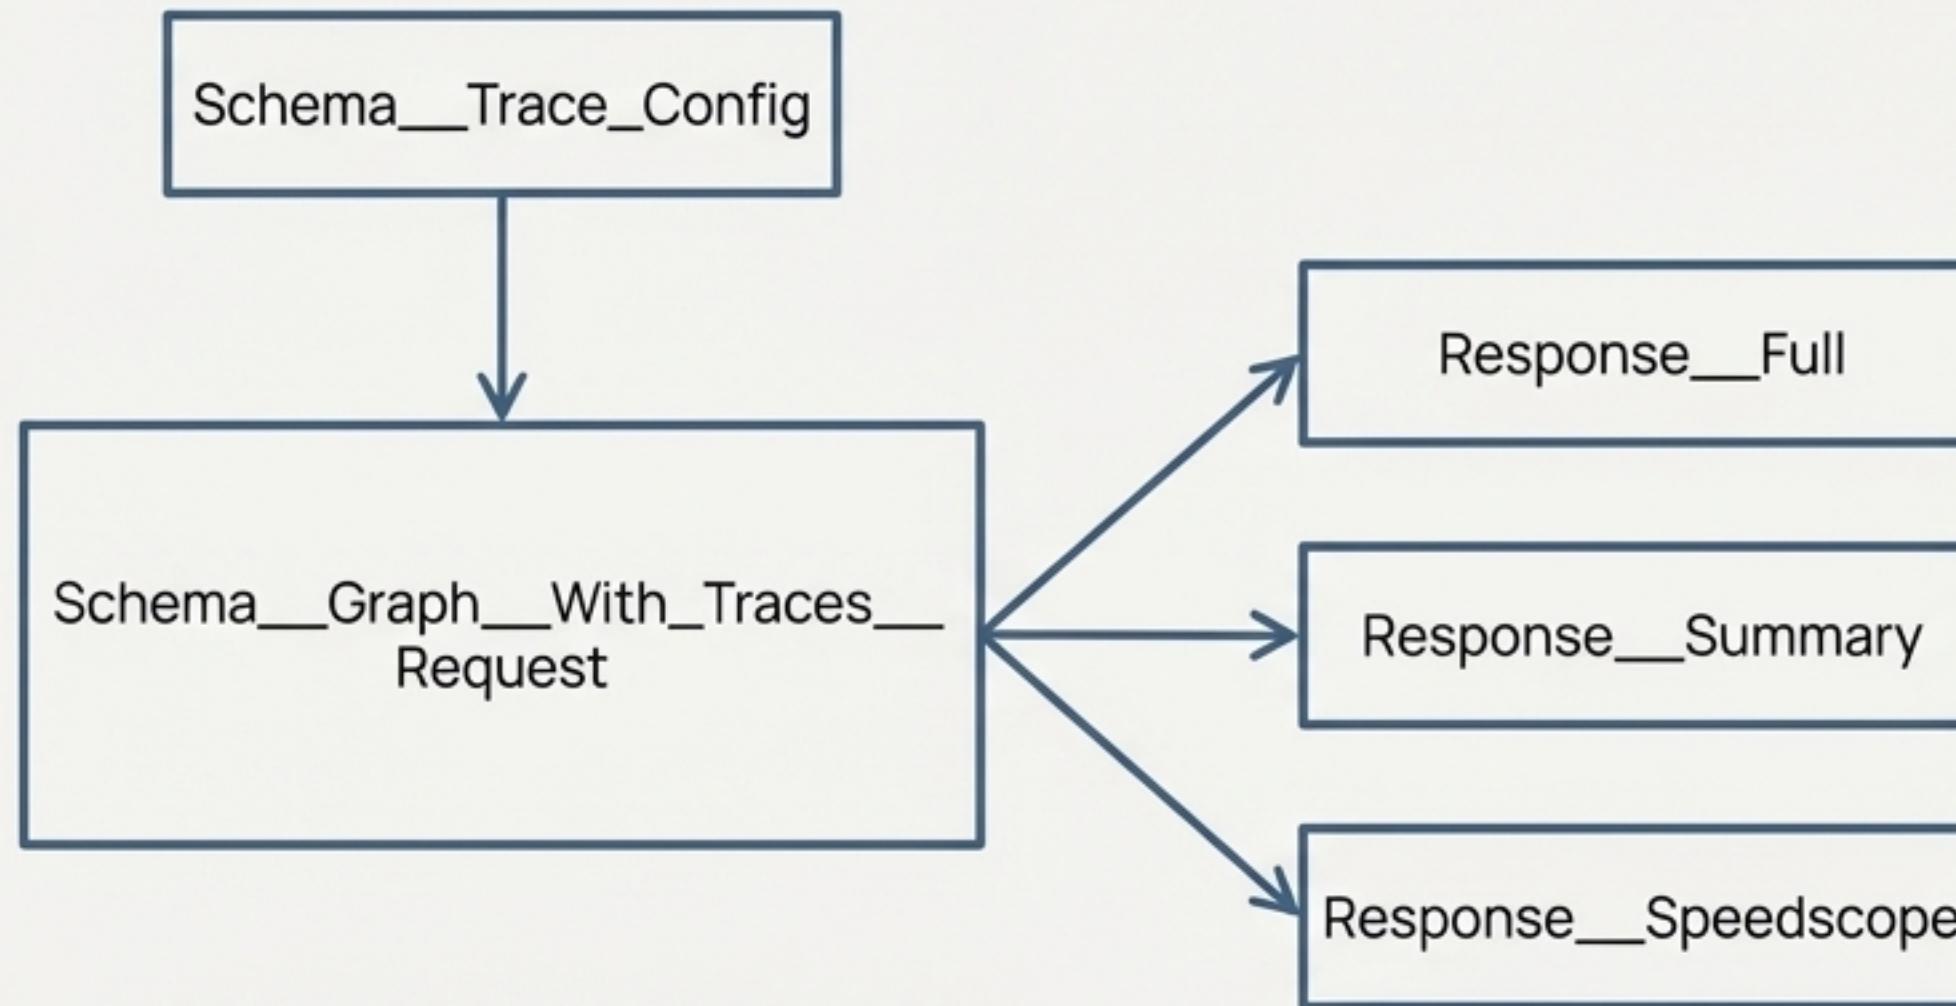
**'POST .../speedscope'**

```
{  
  "graph_data": { ... },  
  "transform_config": { ... },  
  "trace_config": {  
    "capture_traces": true,  
    "trace_format": "speedscope"  
  }  
}
```

**Response (speedscope format)**

```
{  
  "graph_output": { ... },  
  "traces": {  
    "$schema": "https://www.speedscope.  
    app/file-format-schema.json",  
    "shared": { ... },  
    "profiles": [ ... ]  
  }  
}
```

# Type-safe schemas ensure data integrity and clear contracts



## Key Schemas

**Enums\***: `Enum__Trace_Format`, `Enum__Trace_Output`

**Configuration\***: `Schema__Trace_Config`

**Request\***:  
`Schema__Graph__With_Traces__Request`

**Responses\***: Full, Summary, and Speedscope schemas.

All data structures are defined and validated using Type\_Safe schemas, eliminating ambiguity.

# Interacting with the new endpoints is straightforward

## cURL

```
curl -X POST "http://127.0.0.1:8000/.../  
summary" \  
-H "Content-Type: application/json" \  
-d '{  
    "graph_data": { ... },  
    "trace_config": { "capture_traces": true }  
}'
```

## Python Client

```
import requests  
  
url = "http://127.0.0.1:8000/.../speedscope"  
payload = {  
    "graph_data": { ... },  
    "trace_config": {  
        "capture_traces": true,  
        "trace_format": "speedscope"  
    }  
}  
response = requests.post(url, json=payload)  
trace_data = response.json().get("traces")  
  
# Now visualise trace_data in speedscope.app
```

# The result is a powerful, integrated, and low-friction profiling system



## Non-Invasive

- Core logic is untouched.
- No need for code changes or restarts to enable tracing.
- Leverages existing @timestamp decorators.



## On-Demand

- Profiling is enabled via a simple API flag per-request.
- Zero performance overhead when not in use.
- Developers control when and what to trace.

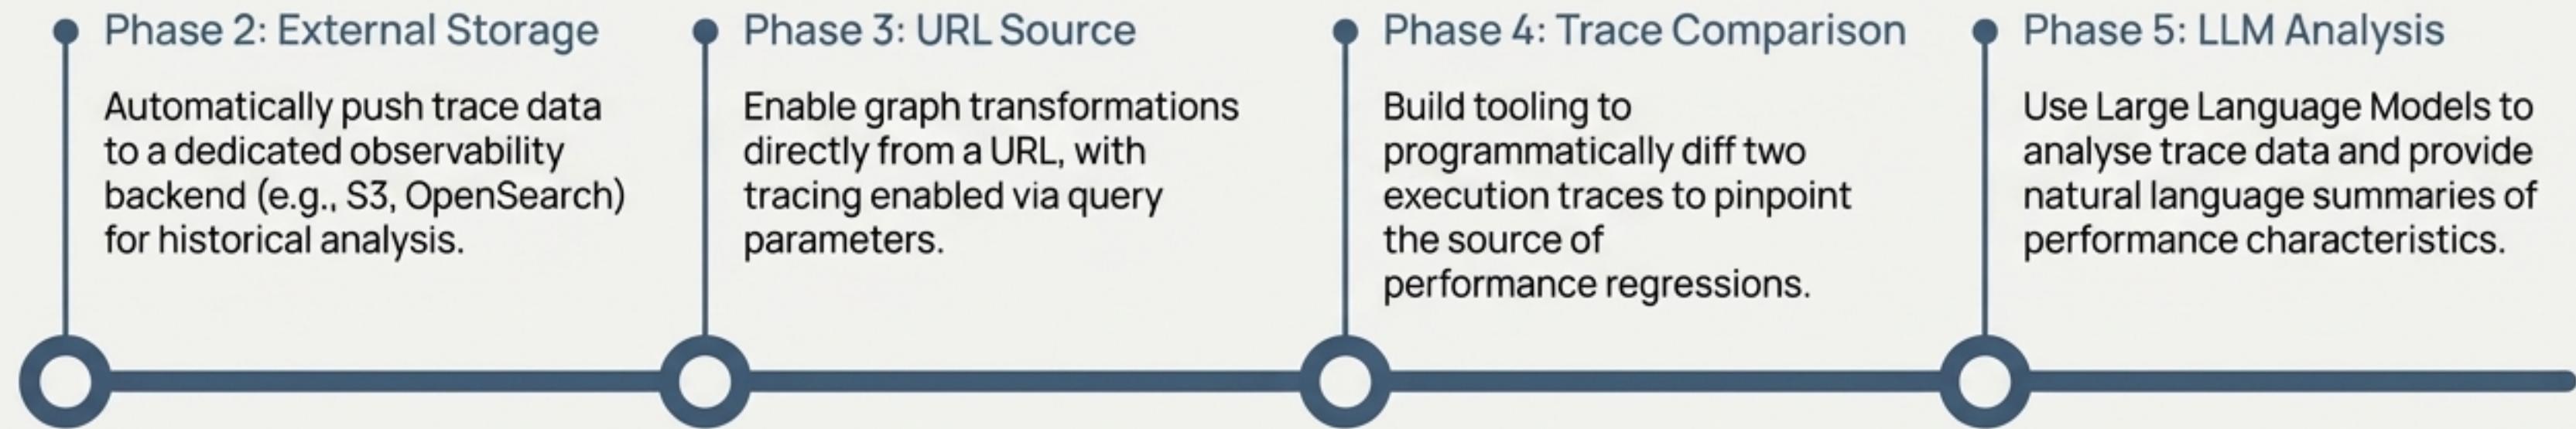


## Integrated

- Returns structured, serialisable data.
- Works out-of-the-box with tools like speedscope.app.
- Designed for distributed systems: capture here, visualise there.

# This is the foundation for a comprehensive observability platform

The initial implementation unlocks several powerful future extensions.  
We have a clear, phased plan to build upon this work.



# The implementation plan is well-defined and ready for execution

## Schema & Enum Definition

- `Enum__Trace_Format`, `Enum__Trace_Output`
- `Schema__Trace_Config`
- all Request/Response schemas

## Route Implementation

- Create `Routes__Timestamps` class
- Register routes in the FastAPI application

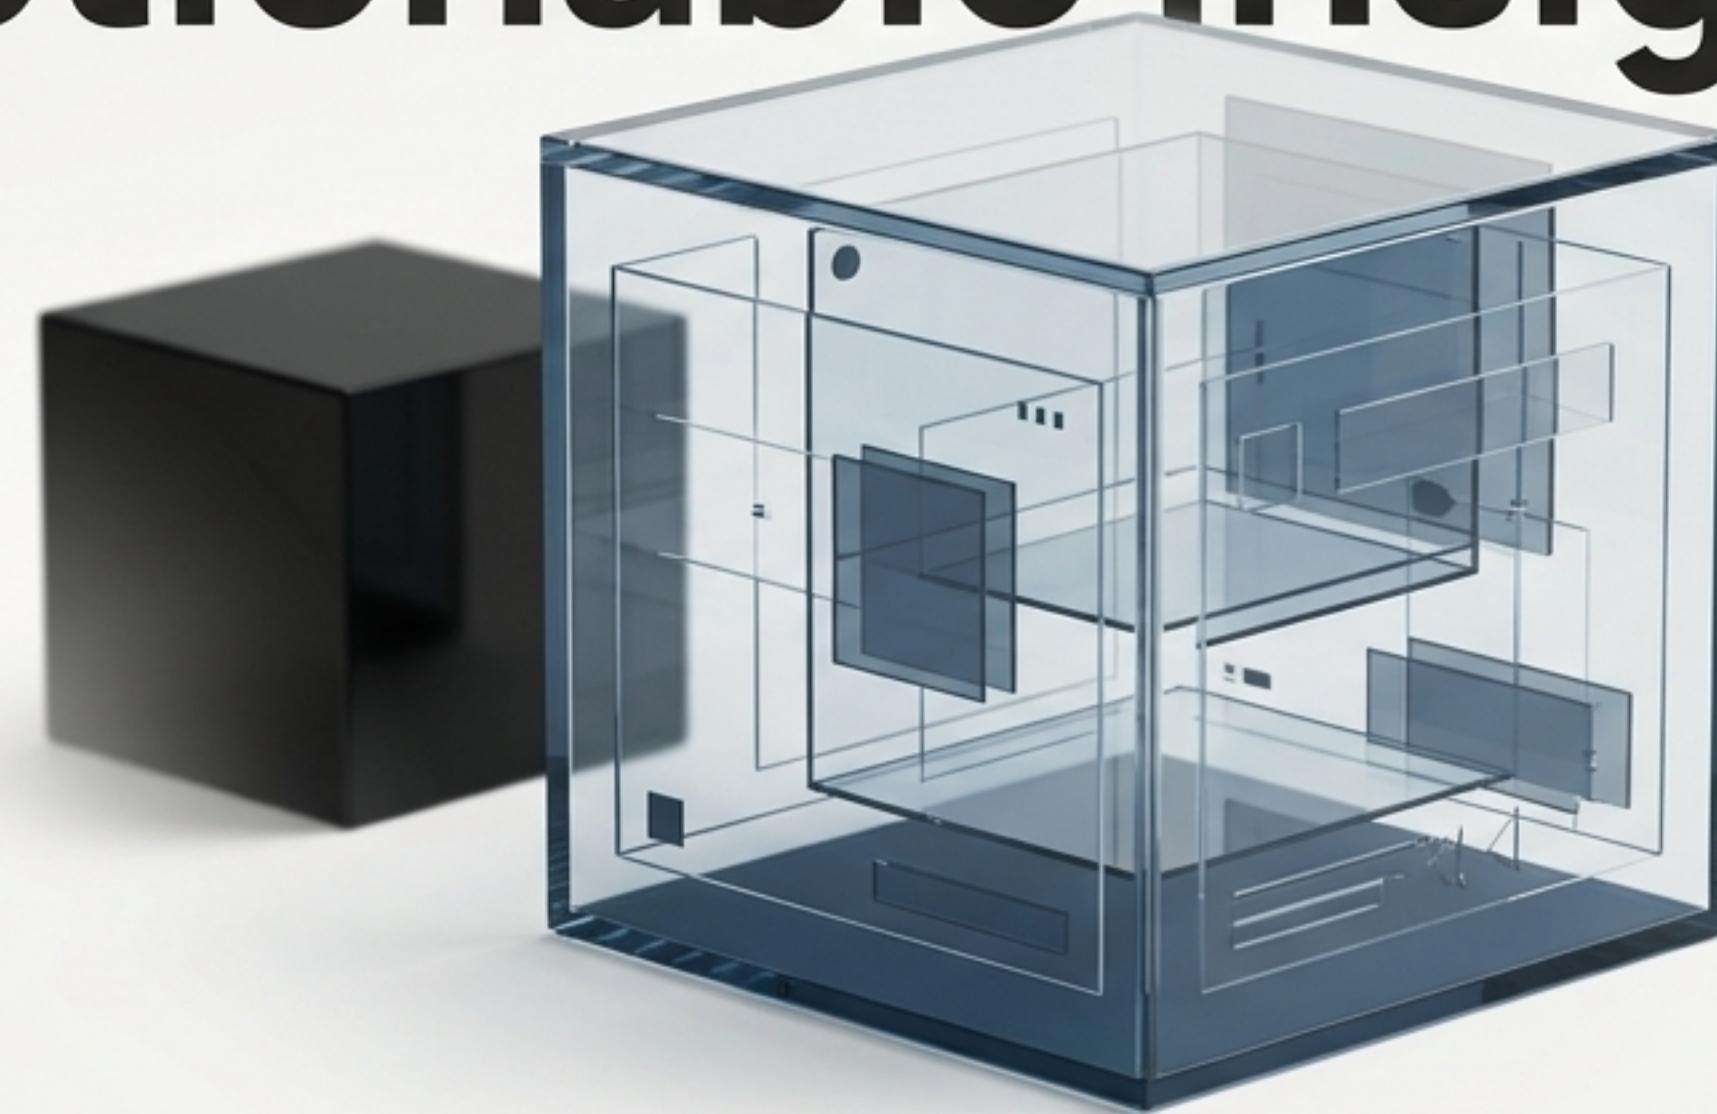
## Instrumentation

- Add `@timestamp` decorators to key service methods

## Validation

- Write comprehensive integration tests
- Test end-to-end with `speedscope.app`

# Turning opaque behaviour into actionable insight



By providing non-invasive, on-demand tracing, we empower developers to understand, debug, and optimise our services more effectively than ever before.