

# **Project M-Graph: Slaying the Silent Performance Killer**

How a shift to lazy initialisation cut  
HTML processing time by 31%

# A 31% Reduction in Total Execution Time

---

**Total Processing  
Time (Size=30  
Doc)**

**-31%**

Tiempos Text Regular  
Execution time reduced from  
441.79ms to 304.21ms.

**Single Element  
Processing  
(Size=1 Doc)**

**-70%**

Tiempos Text Regular  
Execution time reduced from  
121.5ms to 36.3ms.

**Index Object  
Construction**

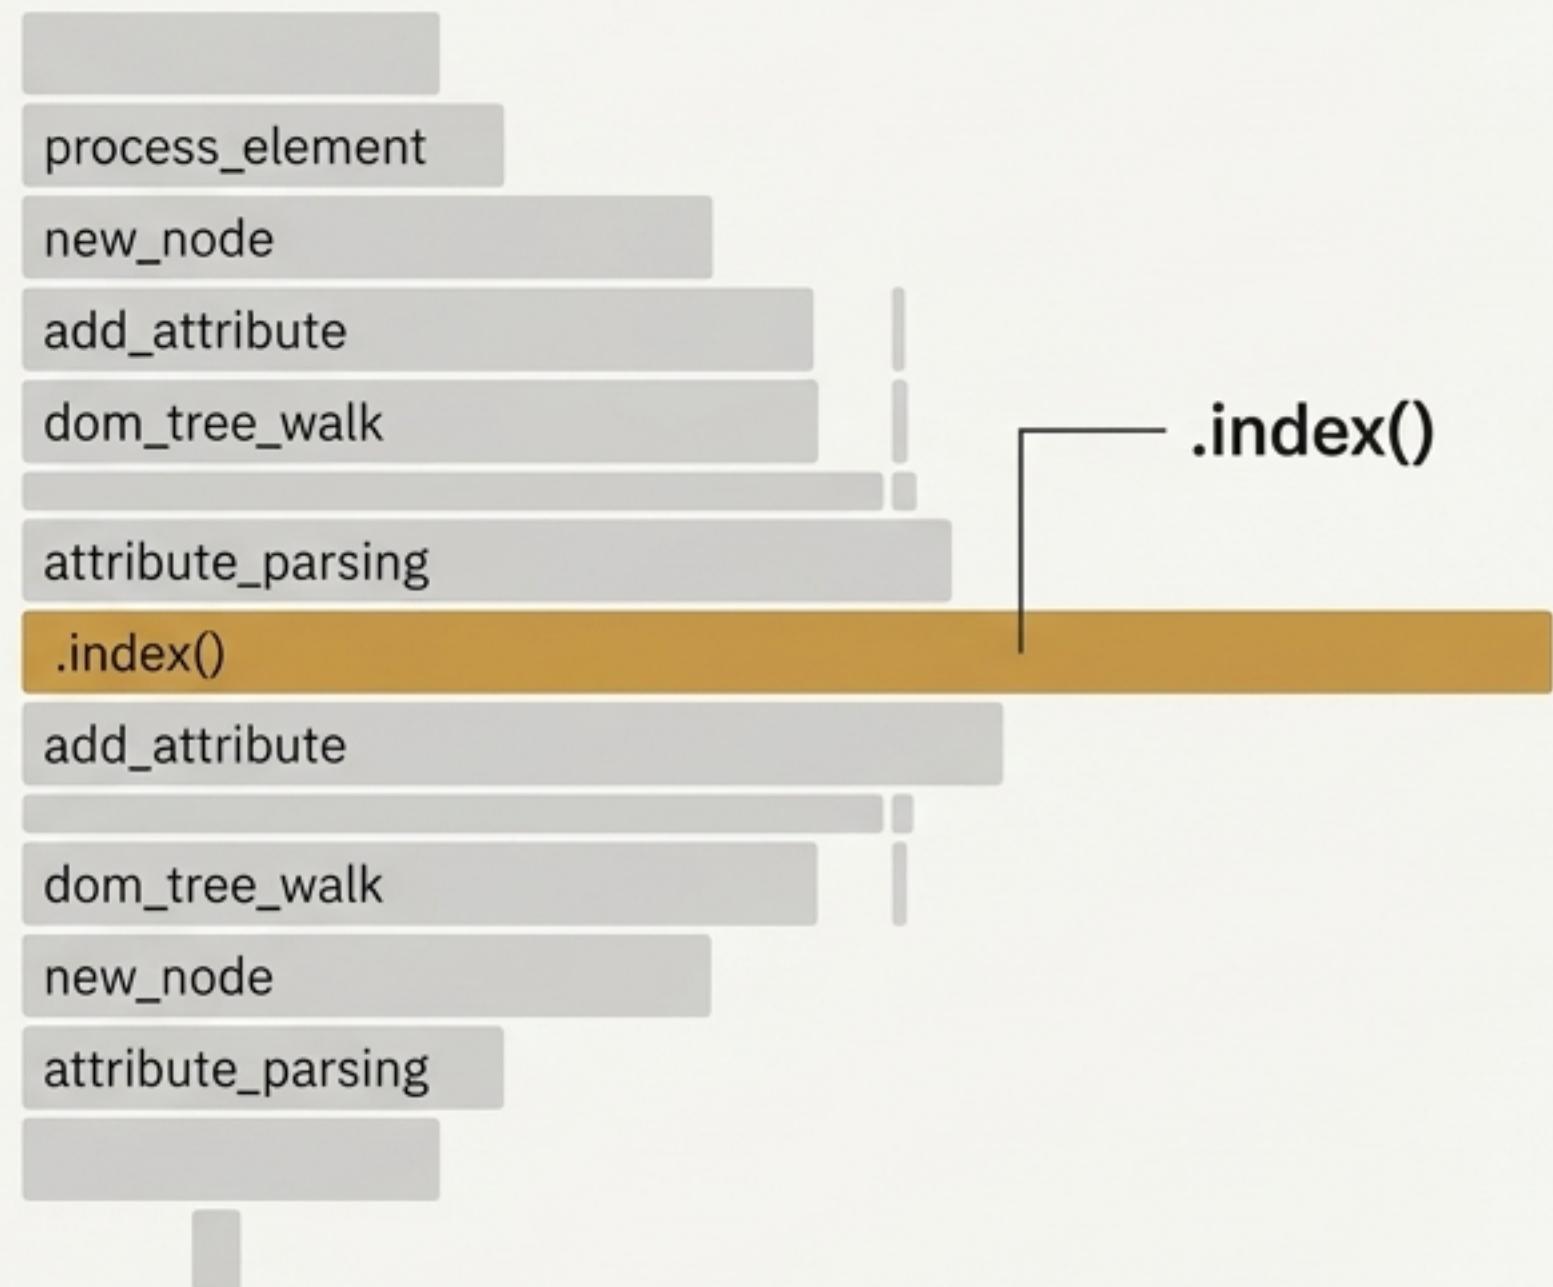
**-97%**

Tiempos Text Regular  
Construction time reduced  
from 1.9ms to just 50μs.

# The Investigation Began with a Clear Bottleneck: Indexing

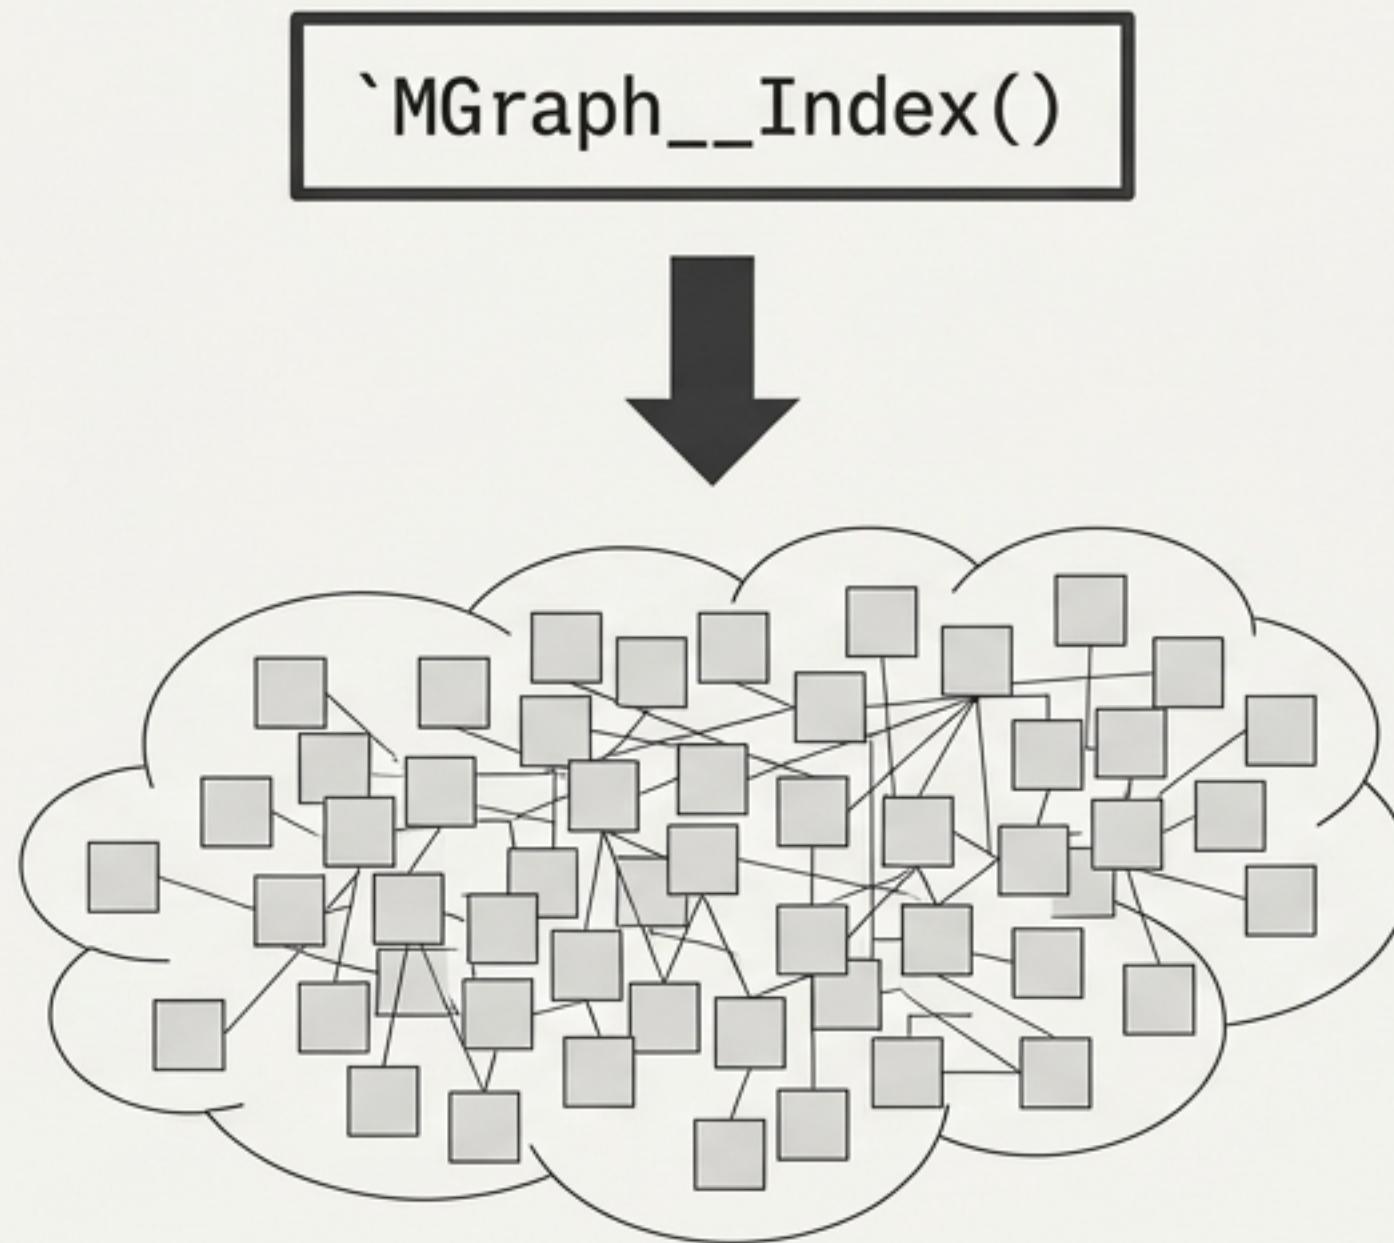
Profiling revealed that calls to `.index()` were disproportionately expensive, consuming a significant portion of the total processing time.

For a simple HTML document, `.index()` calls alone were responsible for **31%** of total processing time (14.2ms of 45.3ms).



# The Root Cause: Wasteful, Eager Object Initialisation

The `Type\_Safe` base class initialised all nested attributes recursively on construction. For a single `MGraph\_\_Index` object, this meant creating a cascade of **over 100 nested objects**, regardless of whether they were needed.



**“Most of these 100+ objects were immediately discarded as references were rewired.”**

# The True Cost of Construction Was Staggering

73:1

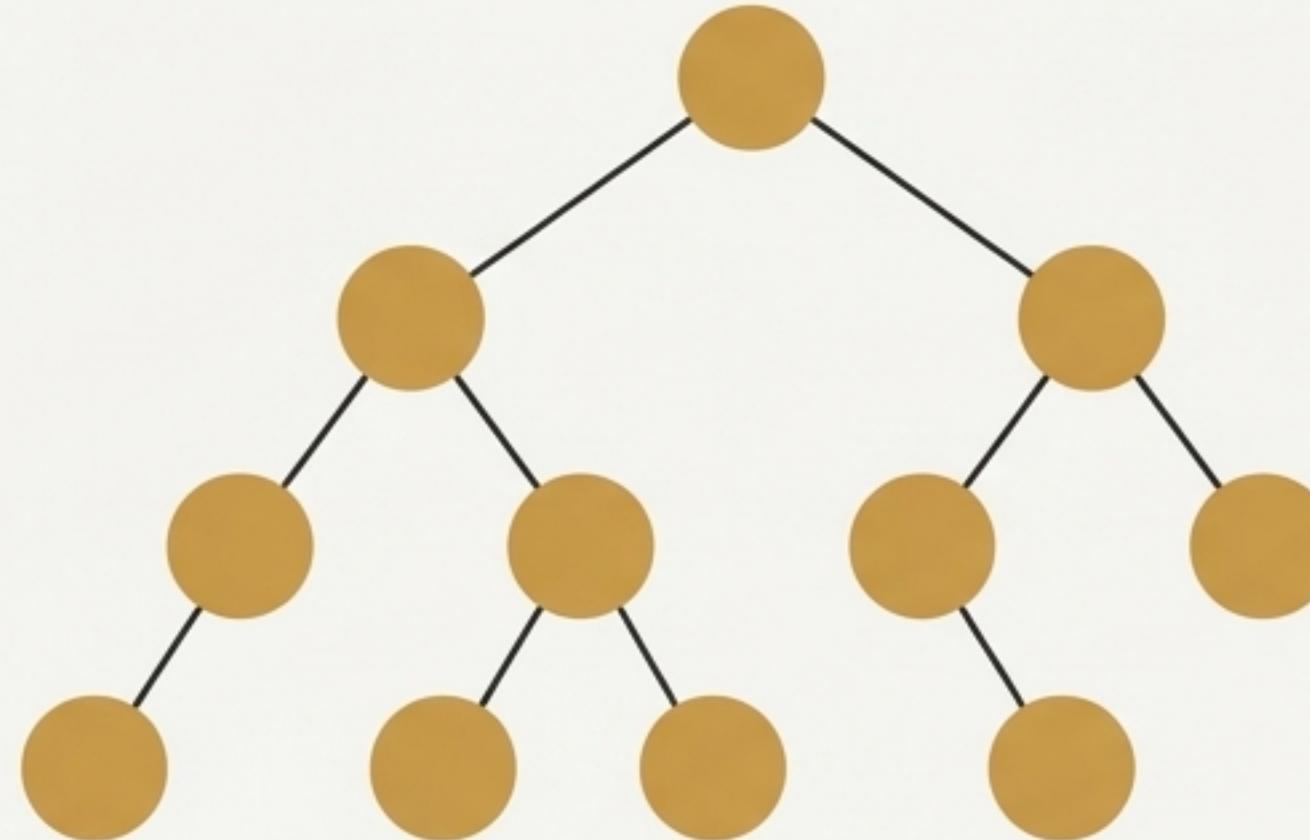
The ratio of object construction time to actual,  
useful work for every index operation.

Additional Detail: After optimisation,  
this ratio dropped to approximately **2:1**.

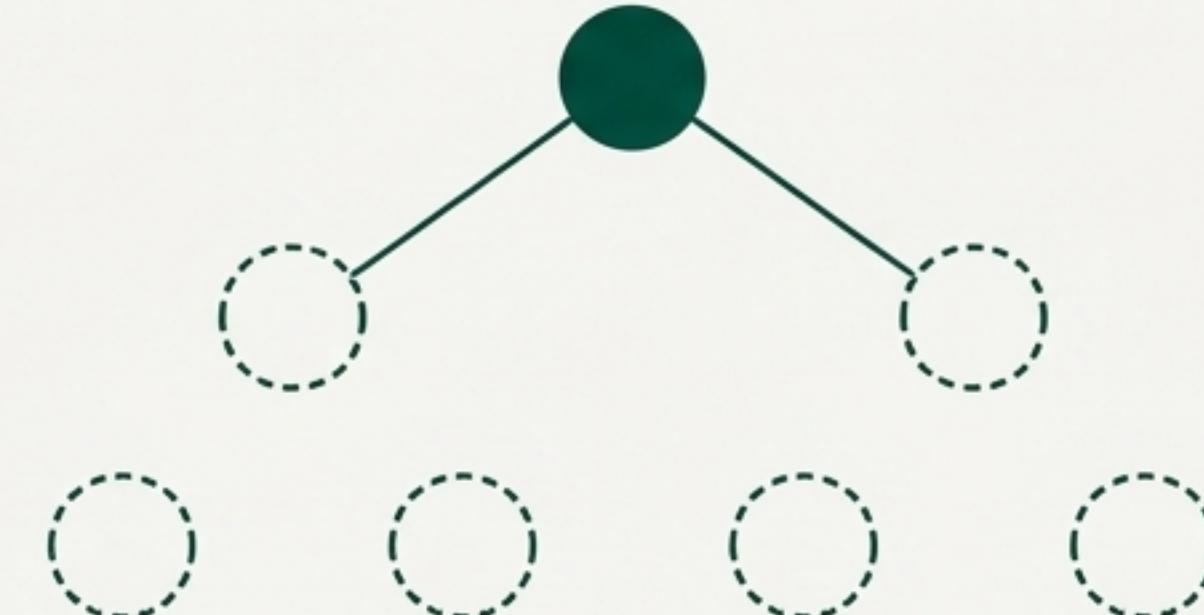
# The Solution: Build Only What You Need, When You Need It

We introduced `Type_Safe__On_Demand`, a new base class that defers attribute initialisation until the moment of first access. This eliminates the upfront construction overhead entirely.

BEFORE: `Type\_Safe`: Eager Initialisation



AFTER: `Type\_Safe\_\_On\_Demand`: Lazy Initialisation



# A Targeted Upgrade Across the Indexing Layer

The shift to lazy initialisation was applied to all classes involved in the M-Graph indexing mechanism.

## Index Layer

MGraph\_\_Index

Schema\_\_MGraph\_\_Index\_\_Data

Schema\_\_MGraph\_\_Index\_\_Data\_\_Edges

Schema\_\_MGraph\_\_Index\_\_Data\_\_Labels

Schema\_\_MGraph\_\_Index\_\_Data\_\_Paths

Schema\_\_MGraph\_\_Index\_\_Data\_\_Types

## Supporting Indexes

MGraph\_\_Index\_\_Edges

MGraph\_\_Index\_\_Labels

MGraph\_\_Index\_\_Paths

MGraph\_\_Index\_\_Types

MGraph\_\_Index\_\_Values

# Head-to-Head: A 31% Speedup on the Same Workload

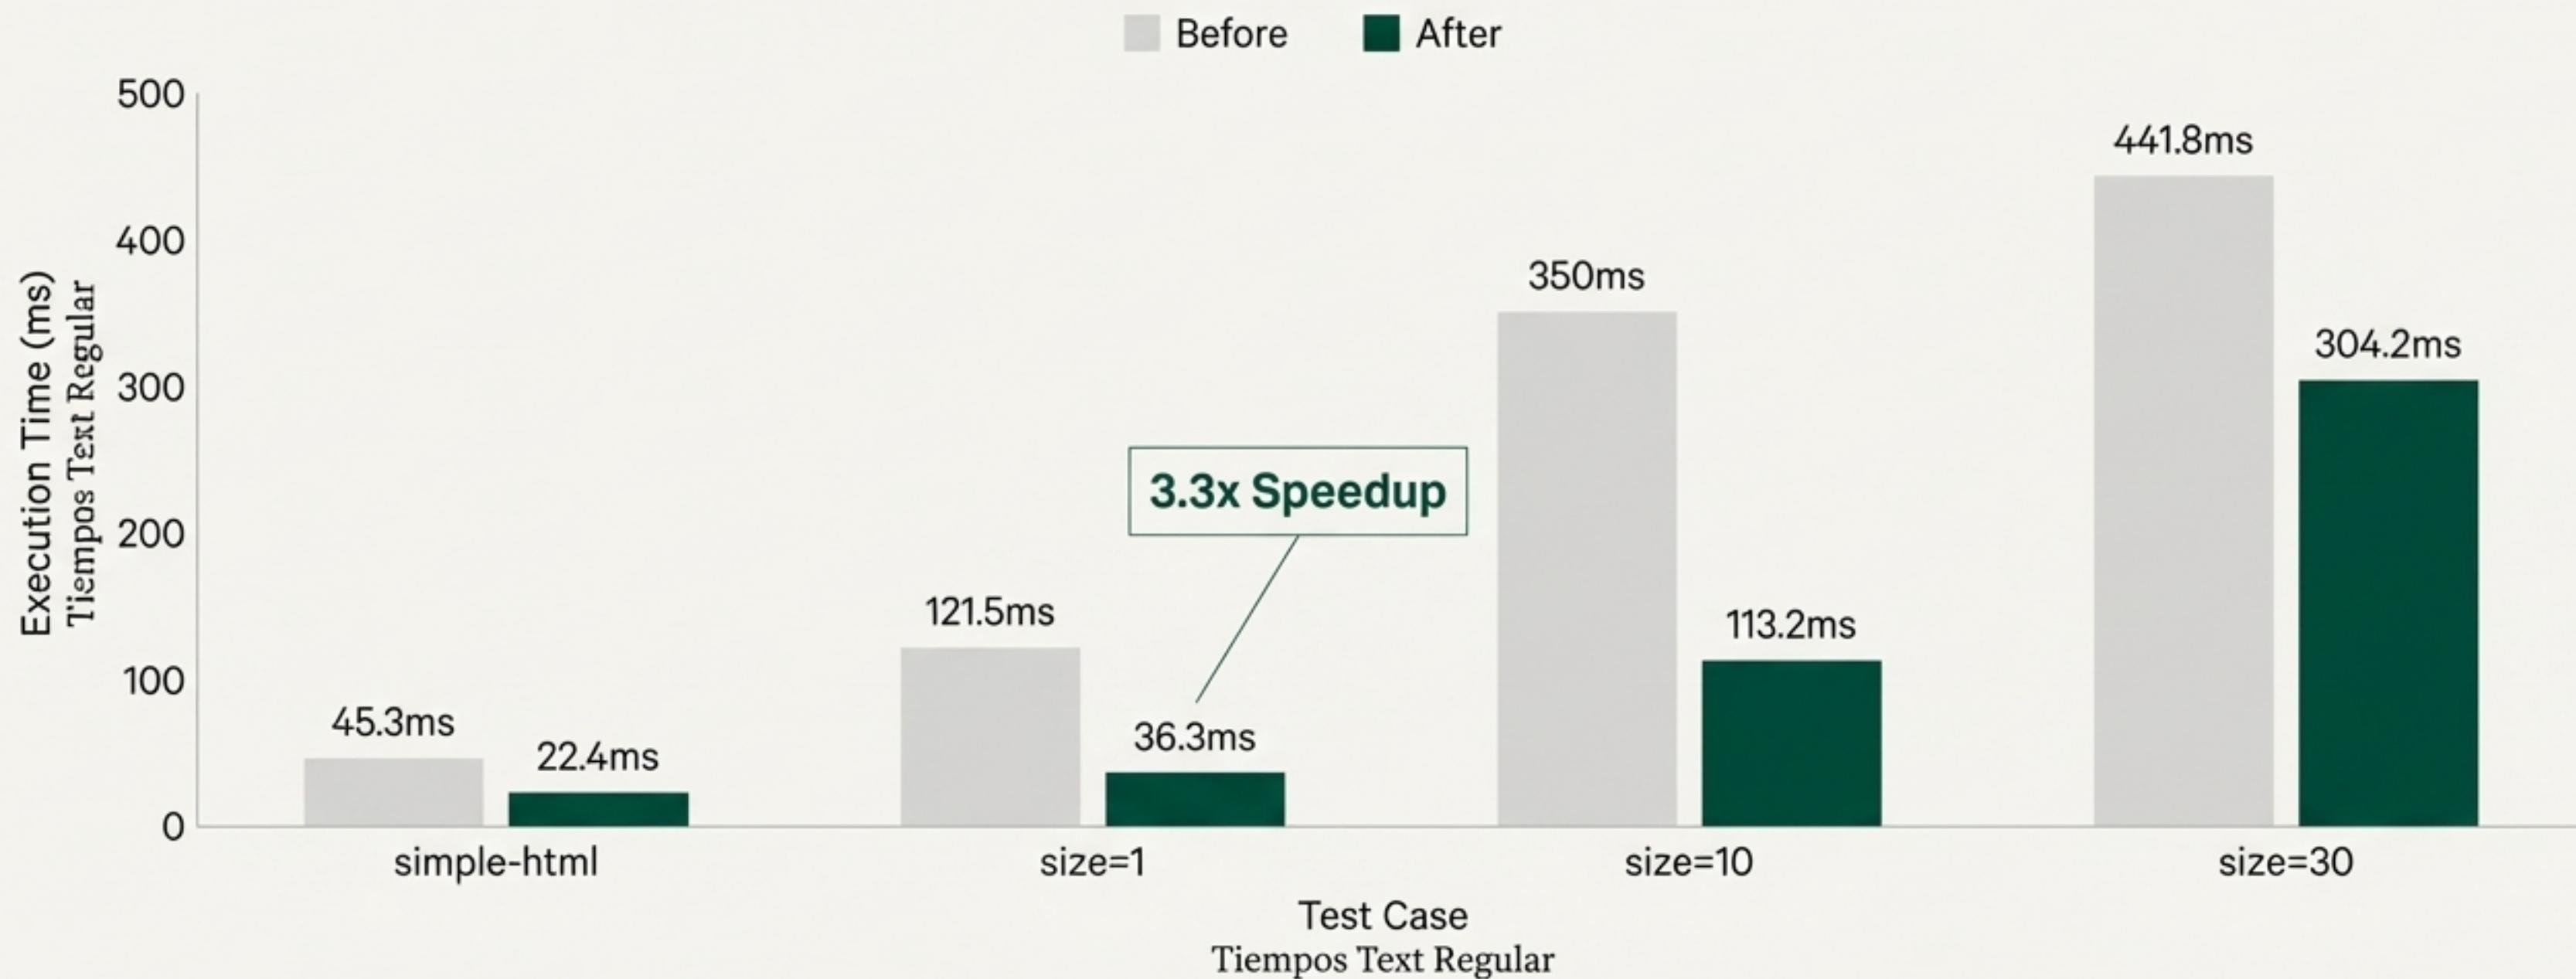
Comparison for a "Size=30" document, ensuring an apples-to-apples test.  
Both runs executed exactly **4,322 calls**.

Frame	Before (27 Nov)	After(29 Nov)	Improvement
process_body_children	904.2ms	636.6ms	-30%
_process_body_element	873.3ms	614.3ms	-30%
<b>execute_pipeline</b>	<b>441.8ms</b>	<b>304.2ms</b>	<b>-31%</b>
add_attribute	252.8ms	177.5ms	-30%
new_node (mgraph_edit)	148.2ms	100.8ms	-32%
mgraph_node	46.5ms	29.9ms	-36%

**-31%**  
Total Time

# The Performance Gains Hold Up at Every Scale

Execution Time Before vs. After Optimisation



# Microseconds Matter: Transforming Core Operations

``MGraph__Index()` Construction`

**-97%**

From 1.9ms down to 50 $\mu$ s

``add_attribute` call`

**-67%**

From 4.2ms down to 1.4ms

`First`.index()` call`

**-76%**

From 2.5ms down to 0.6ms

``new_node (model_mgraph)``

**-65%**

From 0.31ms down to 0.11ms

# Key Architectural Insights From This Optimisation

## 1. Construction is Not Free

The hidden cost of object creation, even when implicit, can easily dominate total processing time. Our initial 73:1 construction-to-work ratio is a stark reminder.

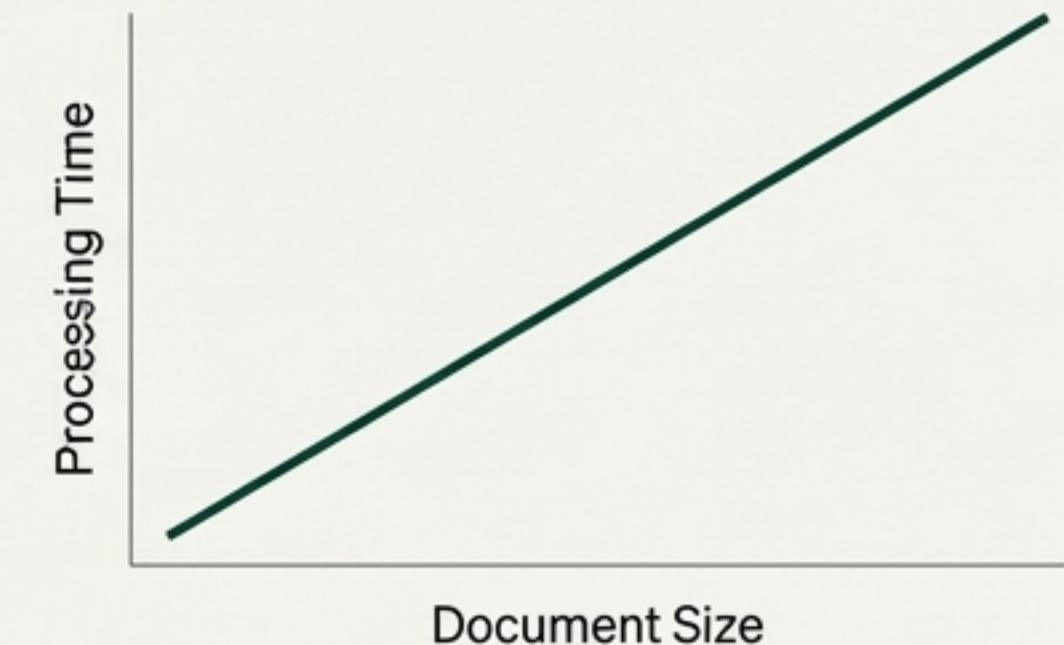
## 2. Fixed Costs Now Amortise Correctly

The setup cost is now truly fixed, ensuring predictable performance. The `document.setup` phase is constant regardless of document size.

Size	document.setup
1	54.4ms
10	53.5ms
30	52.7ms

## 3. Linear Scaling is Now a Reality

With the construction bottleneck removed, processing time for `process\_body\_element` now scales linearly with document size, matching its  $O(n)$  design.



# Our Next Performance Frontiers

This optimisation has cleared the way for us to tackle the next layer of performance bottlenecks.

-  **Flatten Delegation Chains:** The `add_attribute` call chain is 6 levels deep. Flattening it could save significant function call overhead.
-  **Batch Edge Creation:** Each attribute currently creates 3 separate edges. A batch API could reduce this overhead.
-  **Reduce Index Wiring Overhead:** Further lazy wiring of `_sync_index_data()` can eliminate more unnecessary object creation.
-  **Optimise Node Schema:** Simplify or cache complex conditional logic within the `new_node(model_mgraph)` method.

# A 31% Improvement with Zero API Changes

A complex HTML document that took **442ms** to process now completes in **304ms.**

*"This success demonstrates the immense value of systematic profiling and micro-benchmarking."*

# Appendix: Profiling Methodology

## Tools Used

-  Custom @timestamp decorator for method-level timing
-  Speedscope for flame graph visualization
-  Profile Analyzer UI for comparison views
-  Perf utility for micro-benchmarks

## Test Cases

- **simple-html:**

```
<html><body><div class="main">  
id="content">Hello World</div>  
</body></html>
```

- **with\_size\_N:** Generated HTML containing N repeated div elements.